

Simple and Efficient Concurrent Programming via Synchronization Synthesis

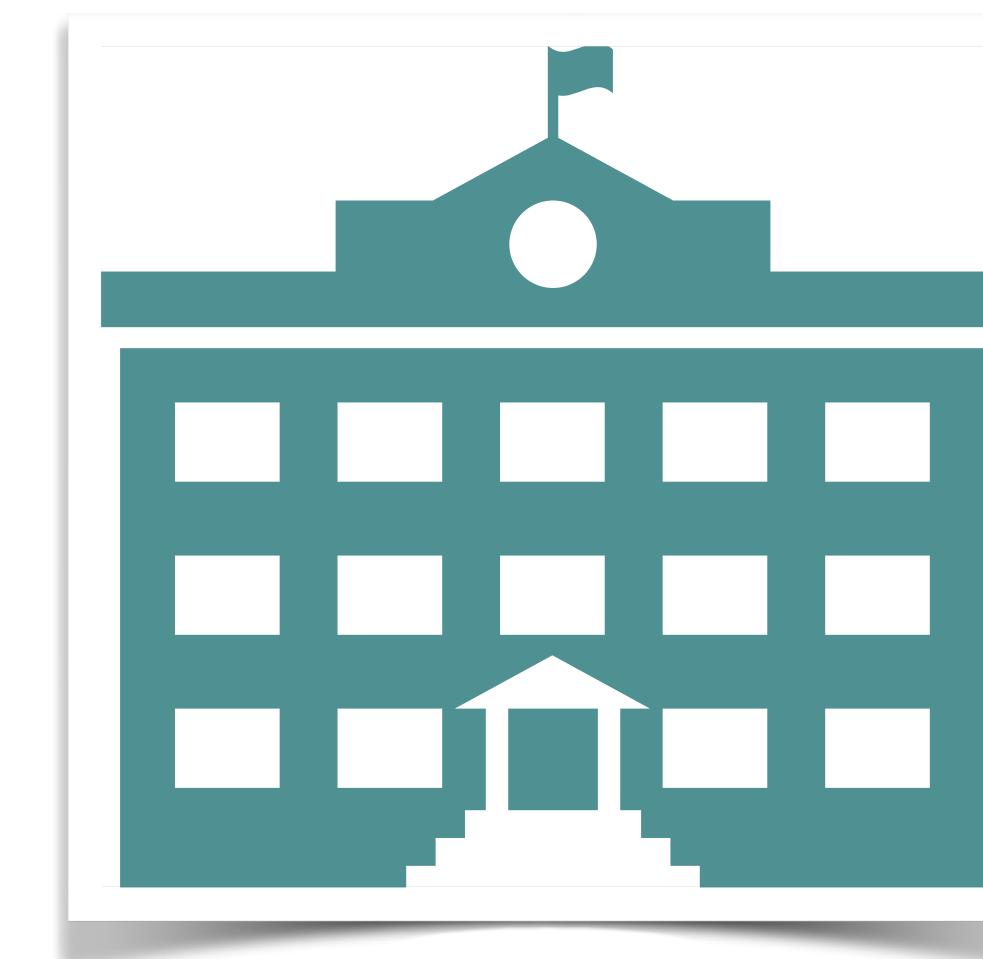
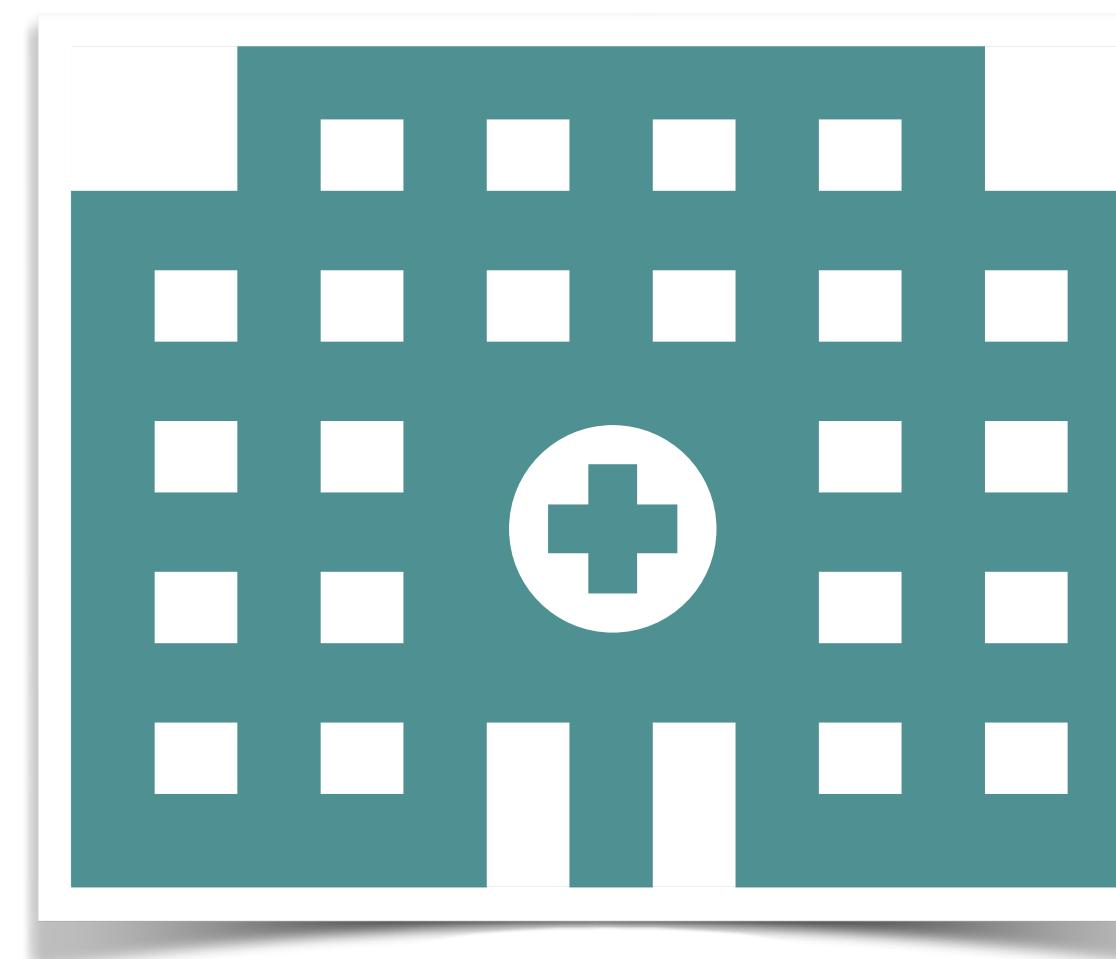
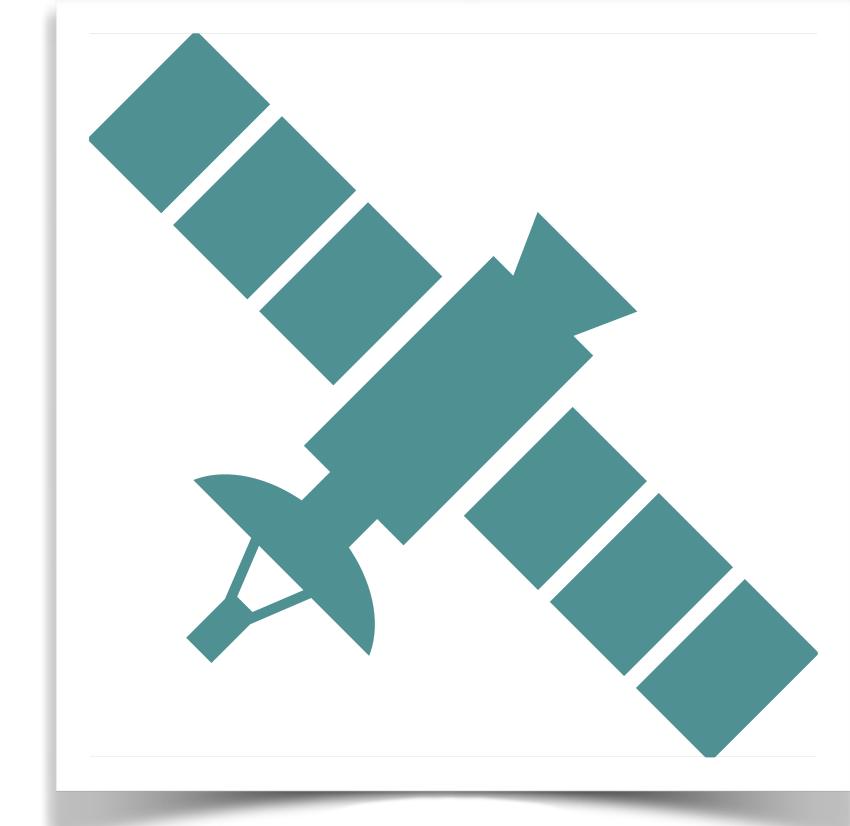
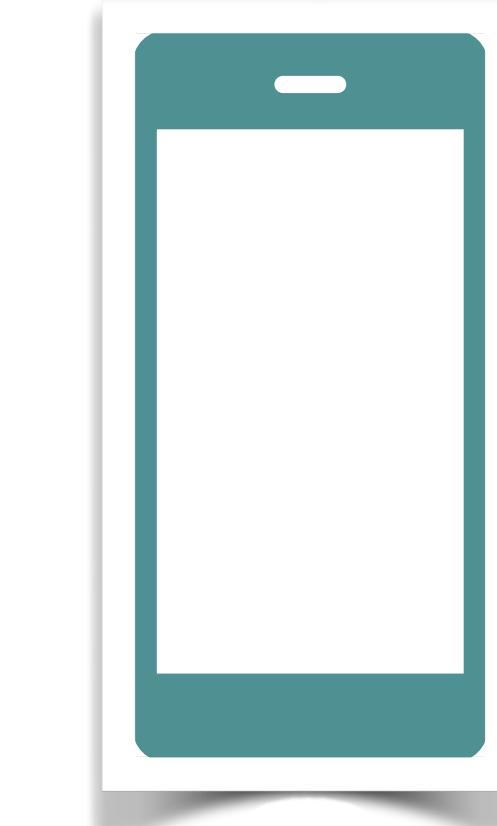
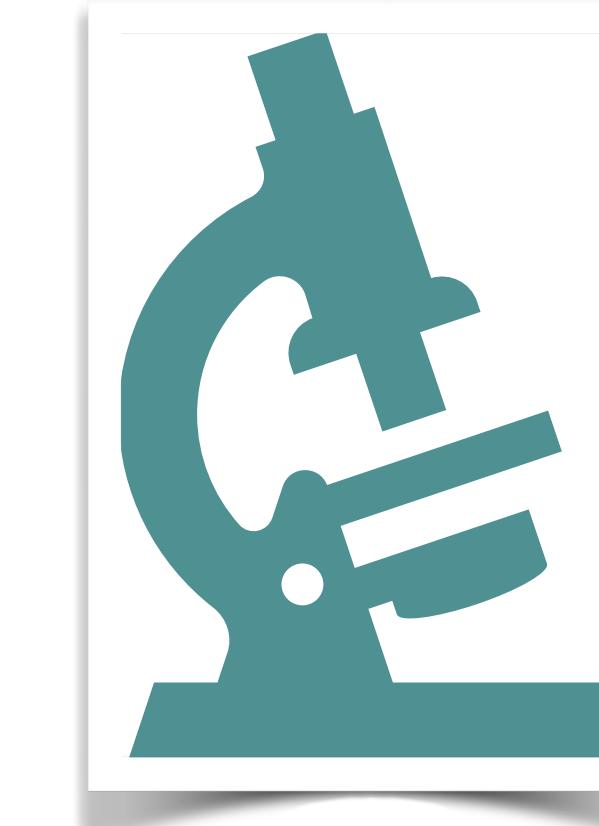
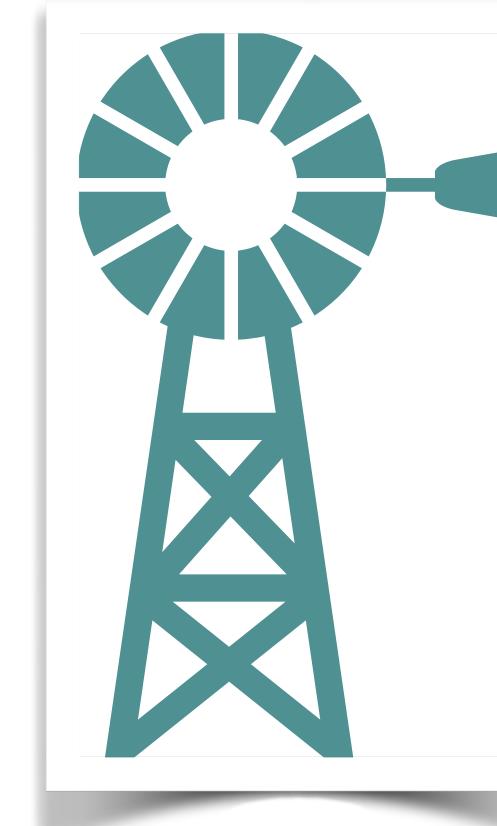
Kostas Ferles

UToPIA



The University of Texas at Austin
Computer Science

Software-Powered Society



Bugs Everywhere!

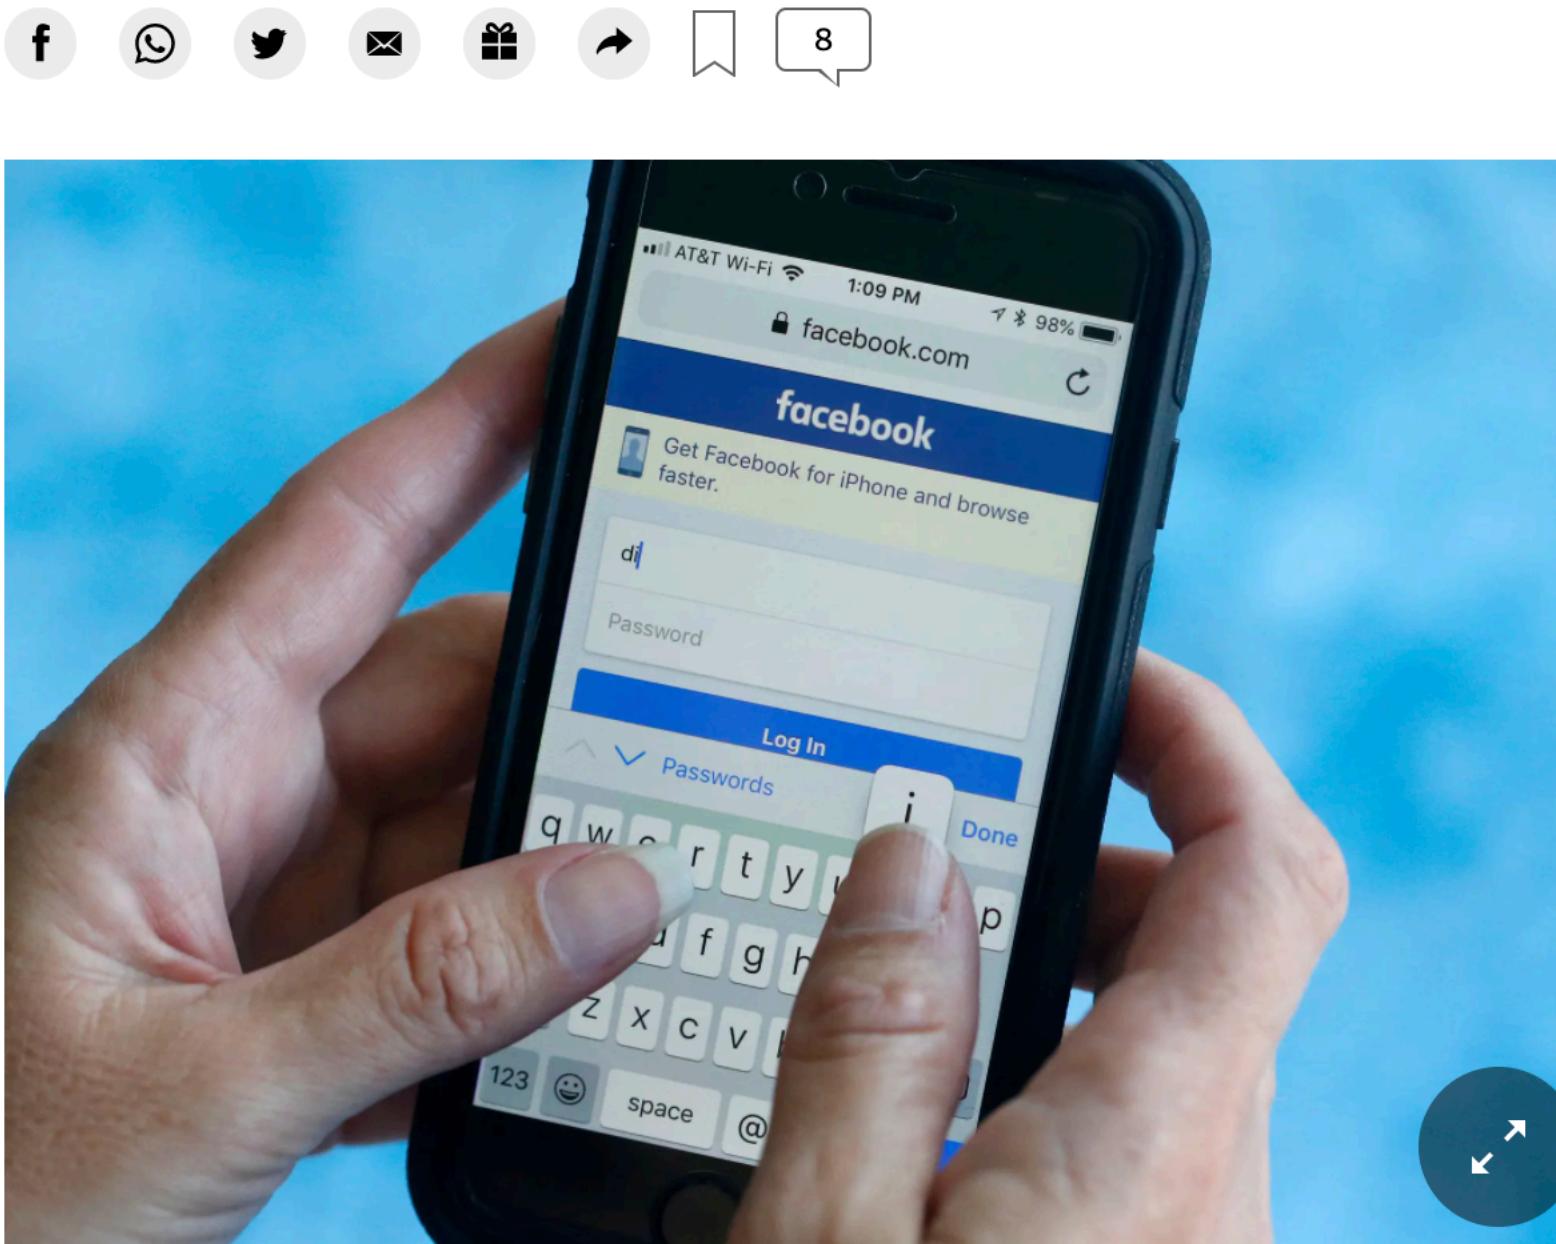
Bugs Everywhere!

Irish Hospitals Are Latest to Be Hit by Ransomware Attacks

Hospitals in Ireland, New York and San Diego are reeling from a ransomware attack.

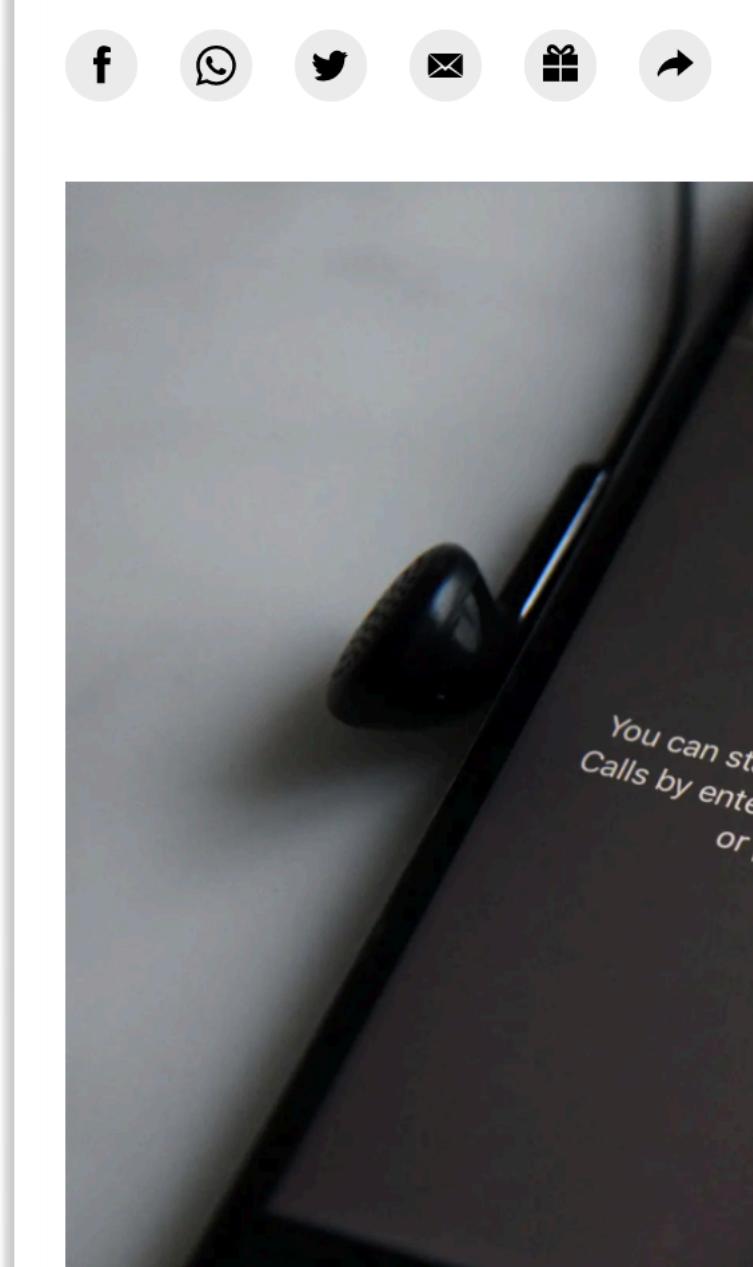


Facebook Hack Included Search History and Location Data of Millions



Facebook said Friday that a security breach had affected 30 million users, 20 million fewer than originally thought. Wilfredo Lee/Associated Press

In a Stumble for Apple, a FaceTime Bug Lets iPhone Users Eavesdrop



The FaceTime bug could also give a caller access to the camera. Roslan Rahman/Agence France-Presse

Cisco to Pay \$8.6 Million to Settle Government Claims of Flawed Tech



A vulnerability in Cisco surveillance software was identified in 2008 by a whistleblower, but the company continued to sell the software until July 2013. Albert Gea/Reuters

Bugs Everywhere!

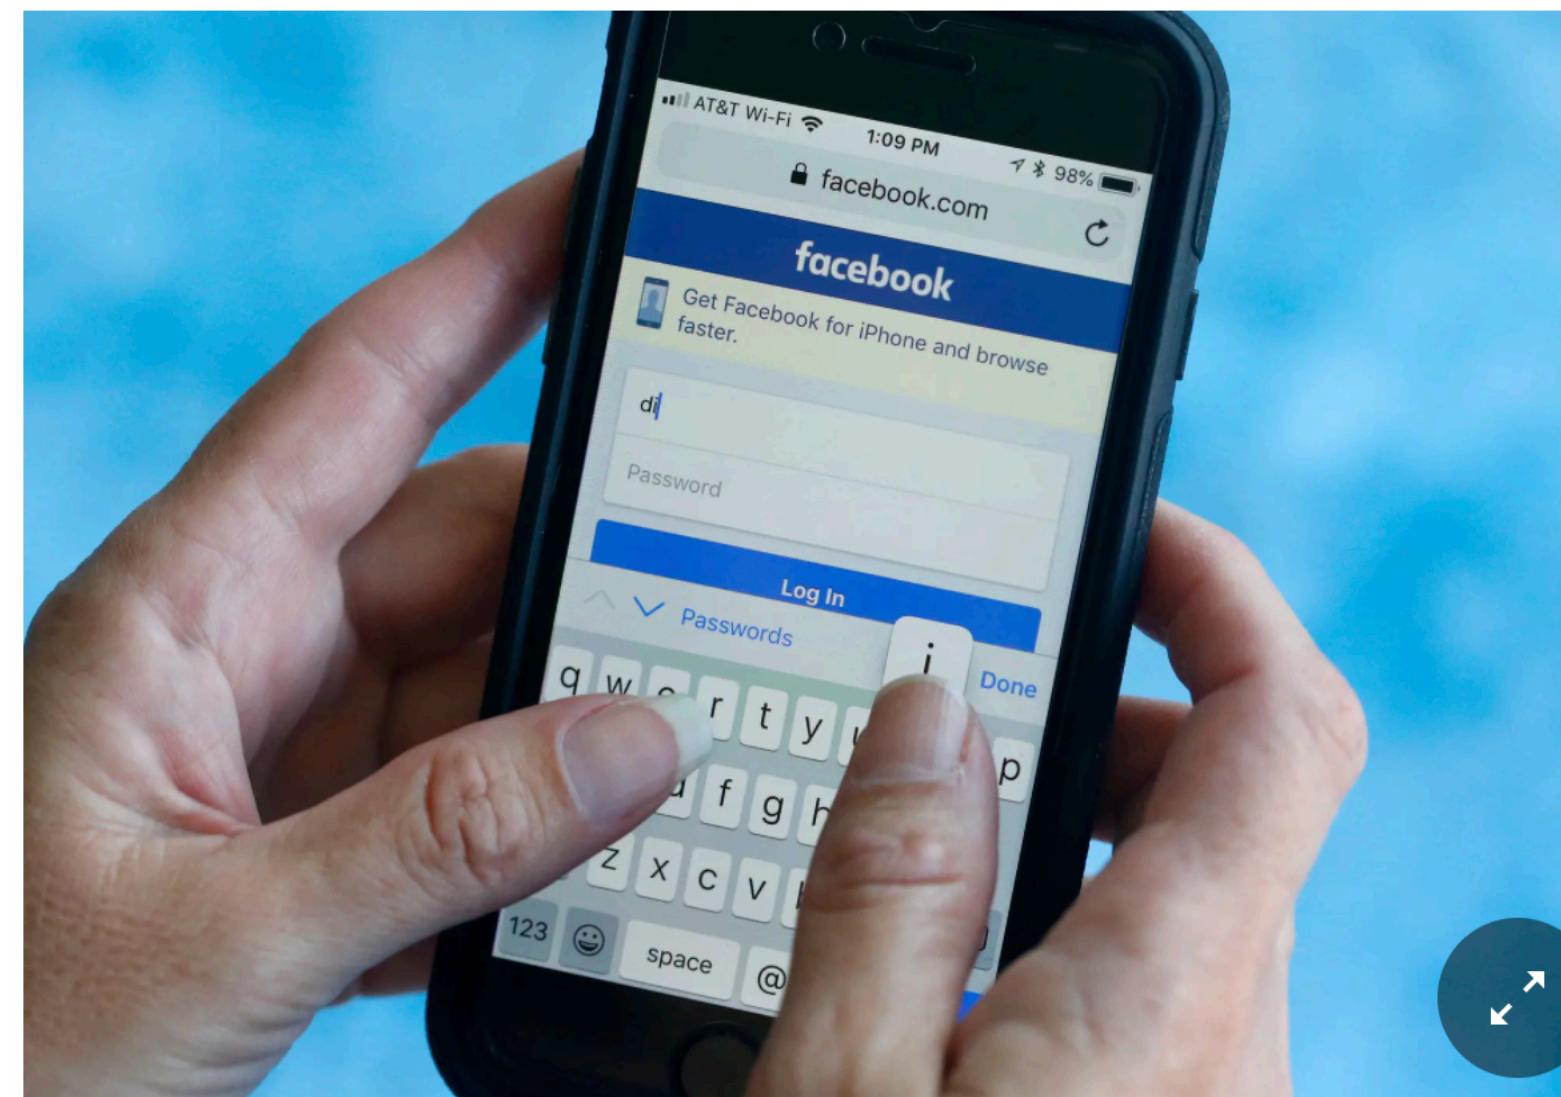
Formal Methods to The Rescue

Irish Hospitals Are Latest to Be Hit by Ransomware Attacks

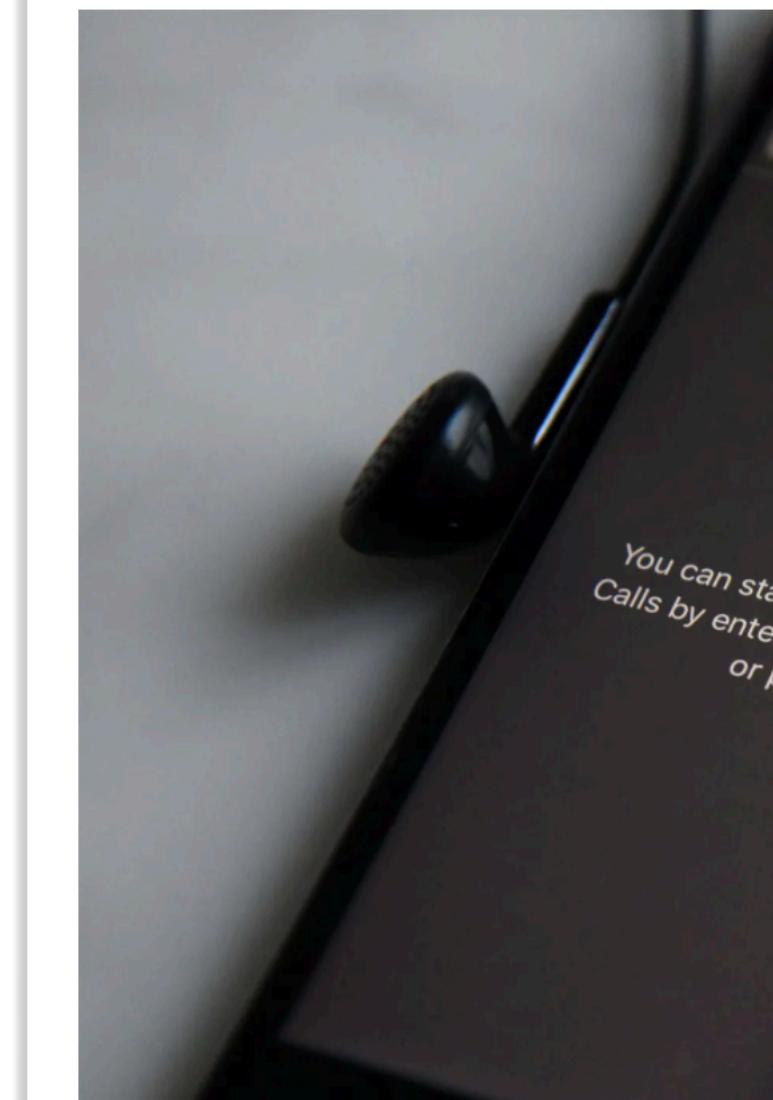
Hospitals in Ireland, New
Diego are reeling from c



Facebook Hack Included Search History and Location Data of Millions



In a Stumble for Apple, a FaceTime Bug Lets iPhone Users Eavesdrop



The FaceTime bug could also give a caller access to the camera. Roslan Rahman/Agence France-Presse

4

Cisco to Pay \$8.6 Million to Settle Government Claims of Flawed Tech



A vulnerability in Cisco surveillance software was identified in 2008 by a whistle-

Bugs ~~Everywhere!~~ Nowhere

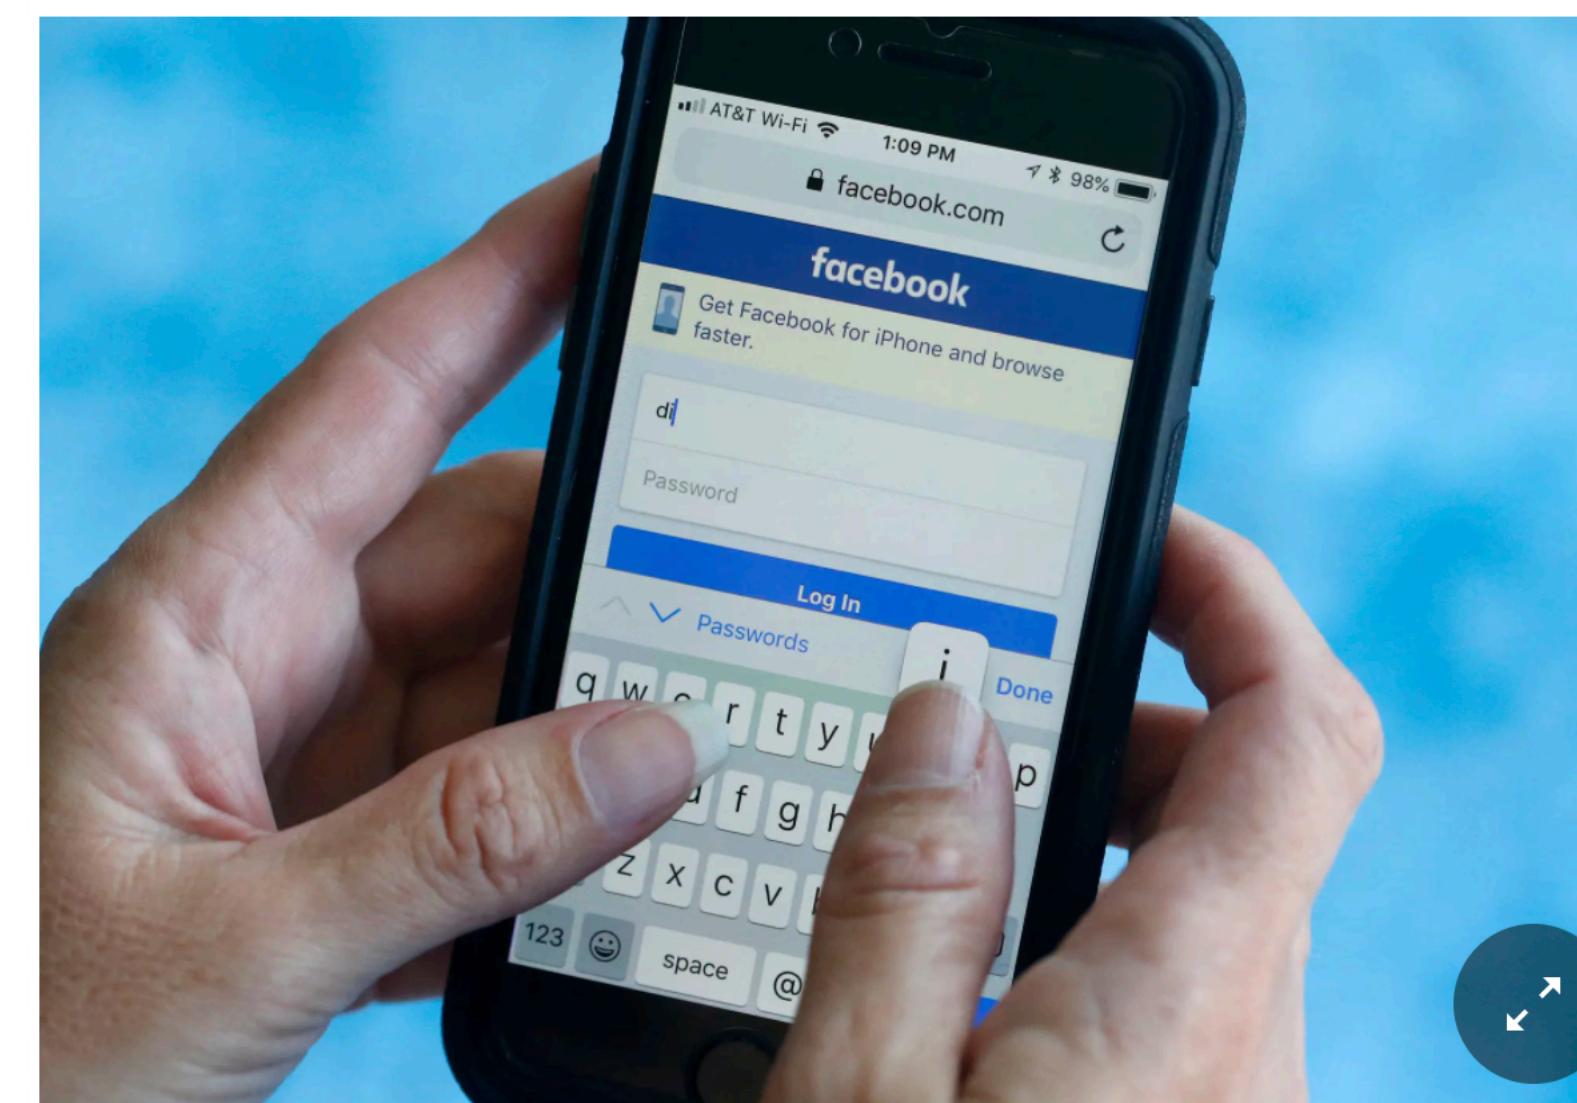
Formal Methods to The Rescue

Irish Hospitals Are Latest to Be Hit by Ransomware Attacks

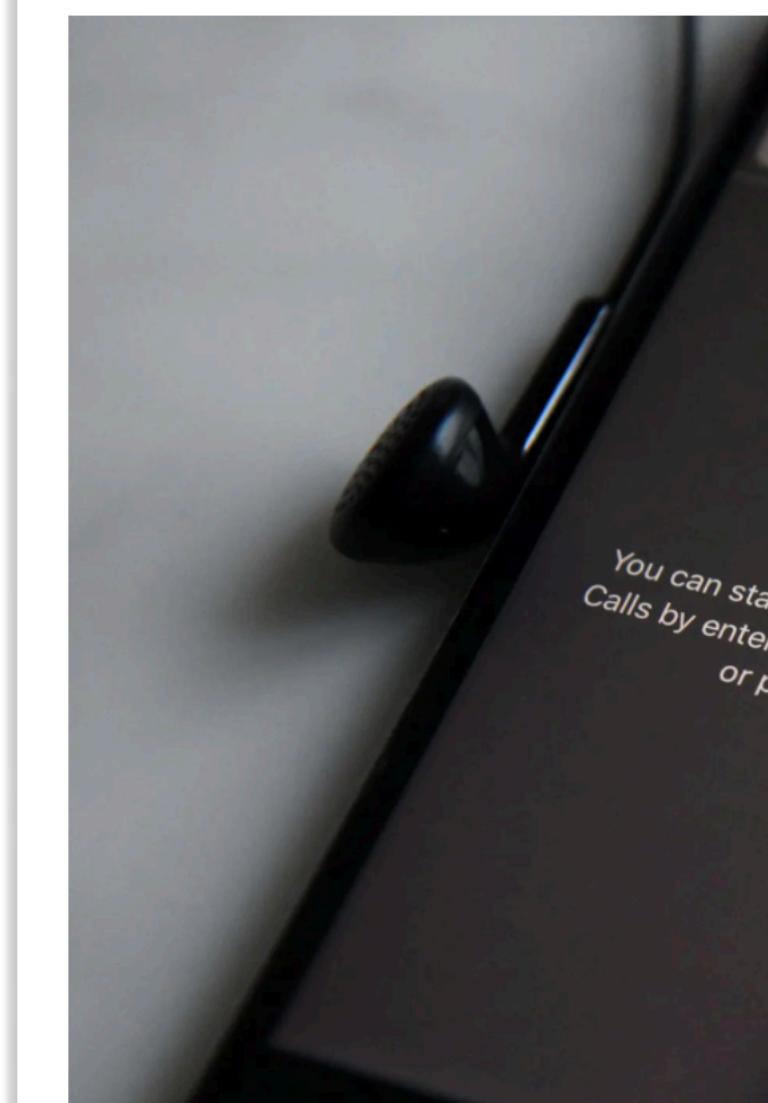
Hospitals in Ireland, New
Diego are reeling from c



Facebook Hack Included Search History and Location Data of Millions



In a Stumble for Apple, a FaceTime Bug Lets iPhone Users Eavesdrop



The FaceTime bug could also give a caller camera. Roslan Rahman/Agence France-Pres

4

Cisco to Pay \$8.6 Million to Settle Government Claims of Flawed Tech

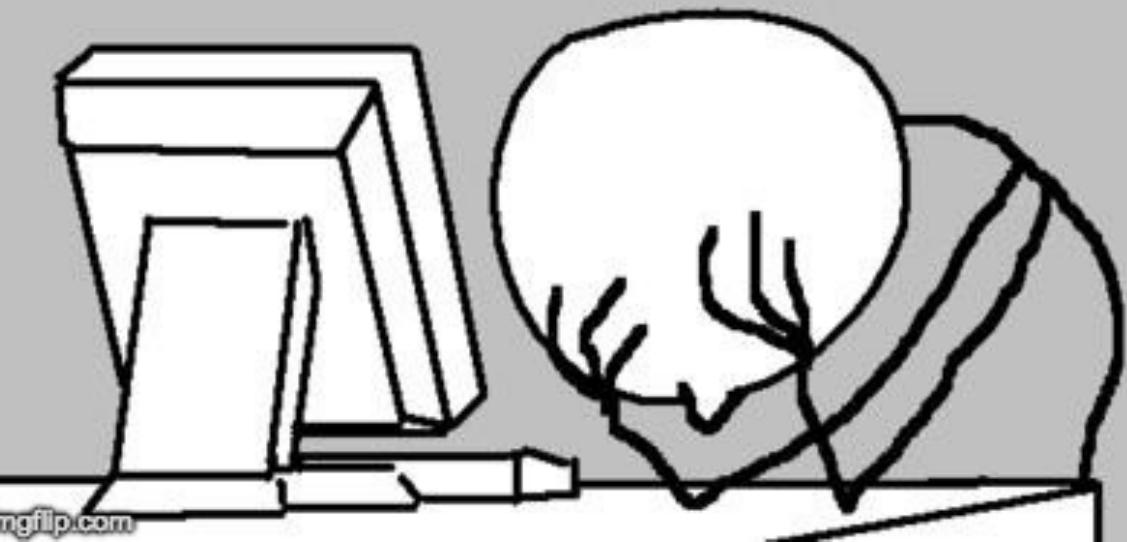


A vulnerability in Cisco surveillance software was identified in 2008 by a whistle-

Motivation

Concurrent programming is hard!

SO MANY THREAD
INTERLEAVINGS



Concurrent code is a perpetual source of painful bugs:

Deadlocks, atomicity violations, starvation, data races, ...

Motivation

If buggy concurrent code is deployed, things can go terribly wrong!

Infamous concurrency bugs



"All the News
That's Fit to Print"

The New York Times

Late Edition
New York: Today, mostly sunny and hot. High 81. Tonight, clear and warm. Low 75. Tomorrow, partly sunny, in mid, a late storm. High 87. Yesterday high 88, low 75. Details, Page B11

VOL CLIX . No. 52,576 Copyright © 2003 The New York Times NEW YORK, FRIDAY, AUGUST 15, 2003 ONE DOLLAR

**POWER SURGE BLACKS OUT NORTHEAST,
HITTING CITIES IN 8 STATES AND CANADA;
MIDDAY SHUTDOWNS DISRUPT MILLIONS**

OVERLOADED GRID
Buildings Are Evacuated and Hospitals Fill—
Bush Plans Review
By JAMES BARRON

A surge of electricity to western New York and Canada touched off a series of power failures and entire blackouts yesterday that left parts of at least eight states in the Northeast and the Midwest without electricity. The widespread failures provoked the evacuation of office buildings stranded thousands of common and flooded some hospitals with patients suffering in the stifling heat.

In an instant that one utility official called "a blink-of-the-eye sort" shortly after 4 p.m., the grid that distributes electricity to the eastern United States became overloaded. As circuit breakers tripped at generating stations from New York to Michigan and into Canada millions of people were instantly caught up in the largest blackout in American history.

In New York City, power was shut off by officials struggling in head of a wider blackout. Cleveland and De

James L. Barron/The New York Times

The view looking across Manhattan at sunset yesterday. New York City was left without traffic lights at rush hour, and officials struggling to head off a wider blackout shut off power.

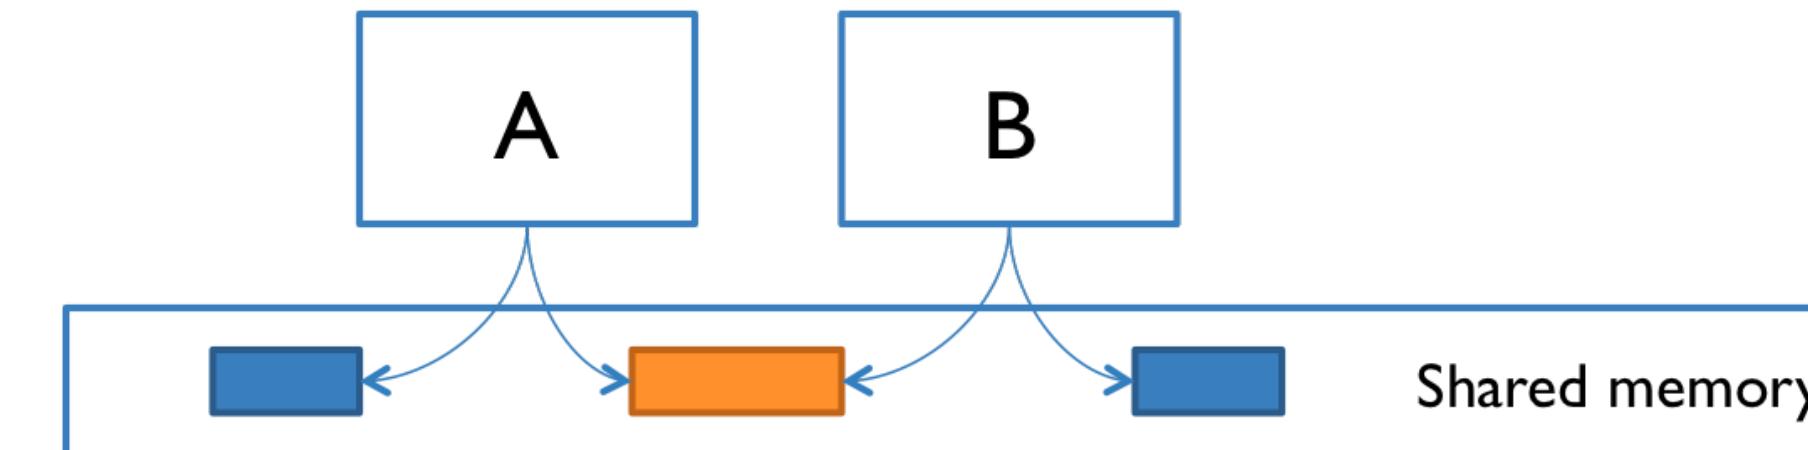
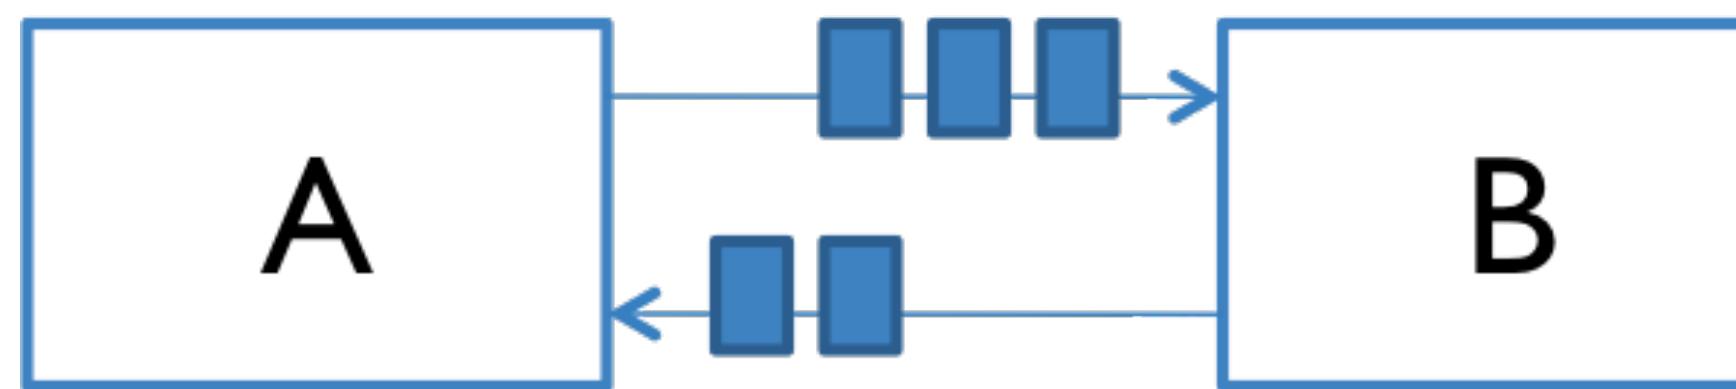
Research Question



Can we employ formal method techniques to simplify concurrent programming without sacrificing performance?

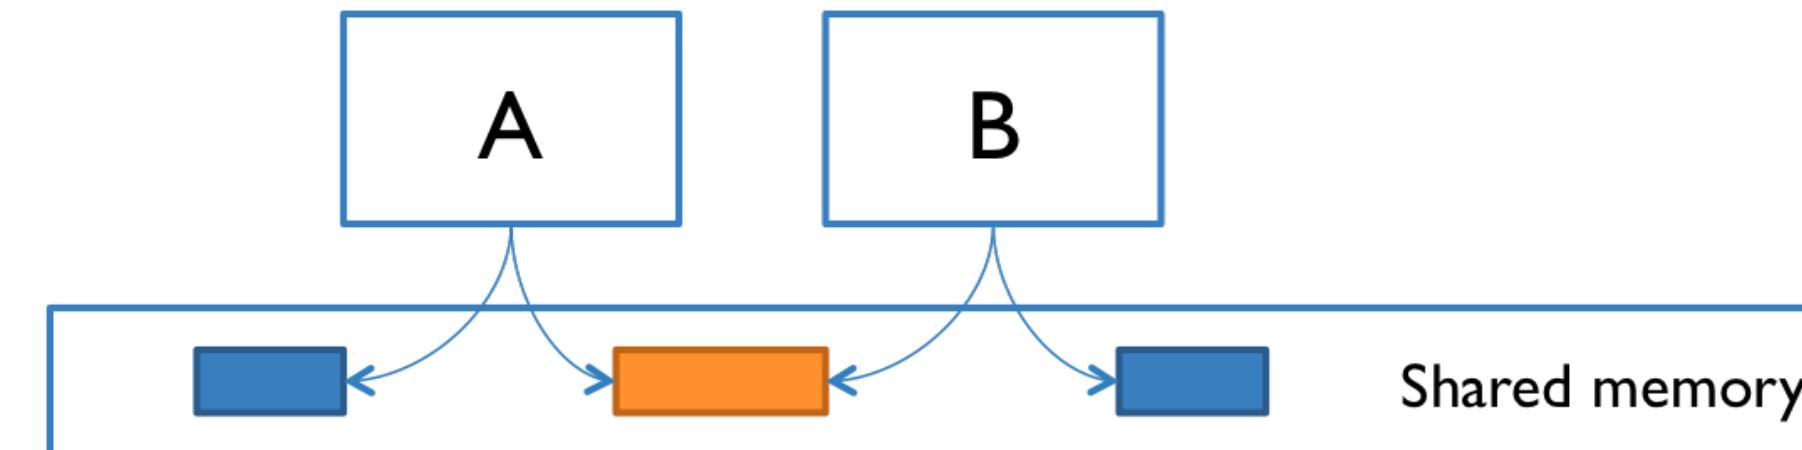
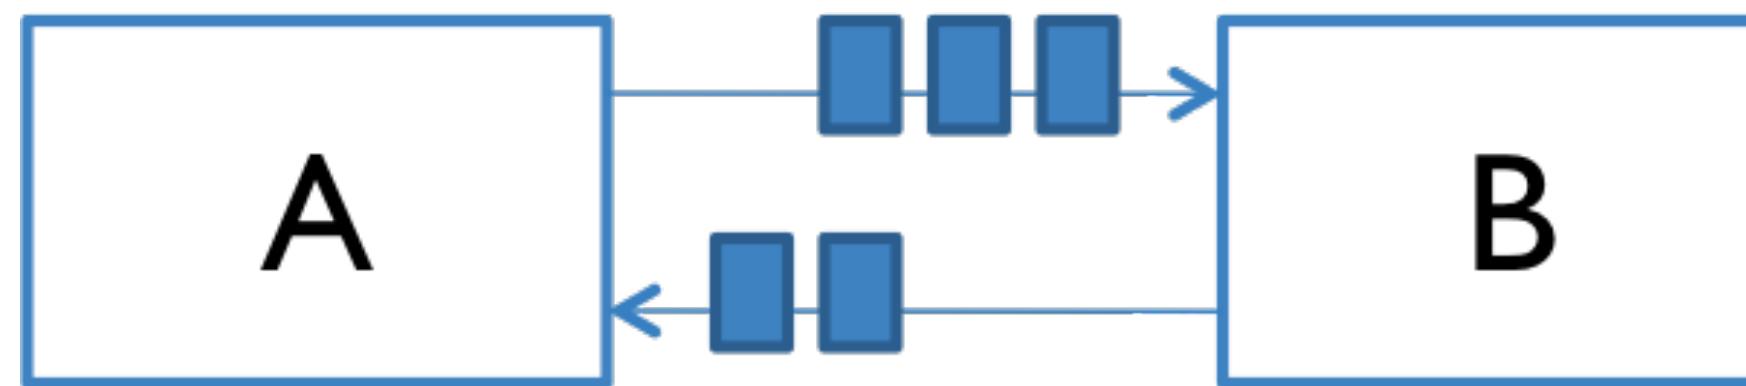
Shared Memory vs Message Passing

Two main models of concurrency:
message-passing and *shared memory*



Shared Memory vs Message Passing

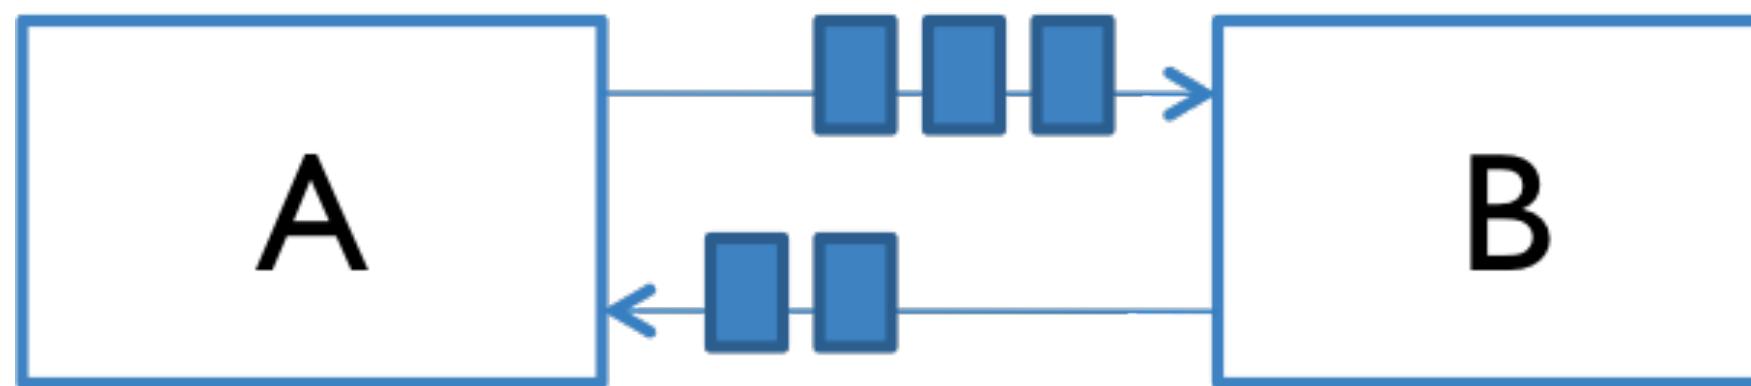
Two main models of concurrency:
message-passing and *shared memory*



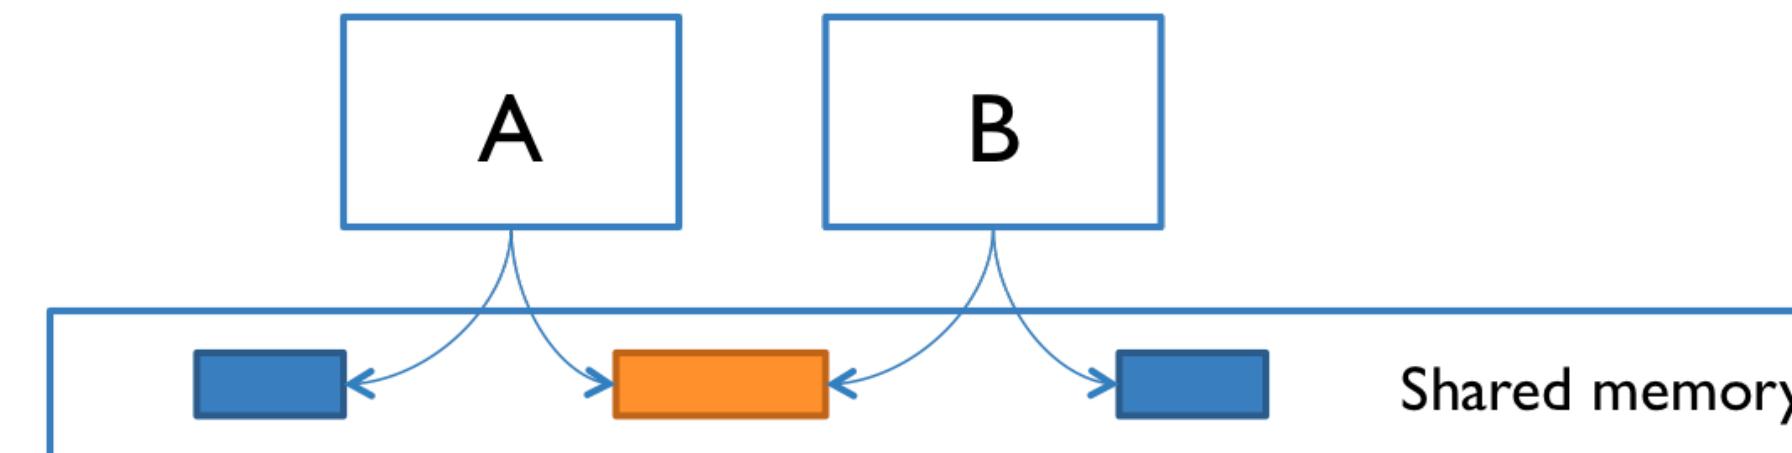
Threads pass
messages; no
shared state

Shared Memory vs Message Passing

Two main models of concurrency:
message-passing and *shared memory*



Threads pass
messages; no
shared state

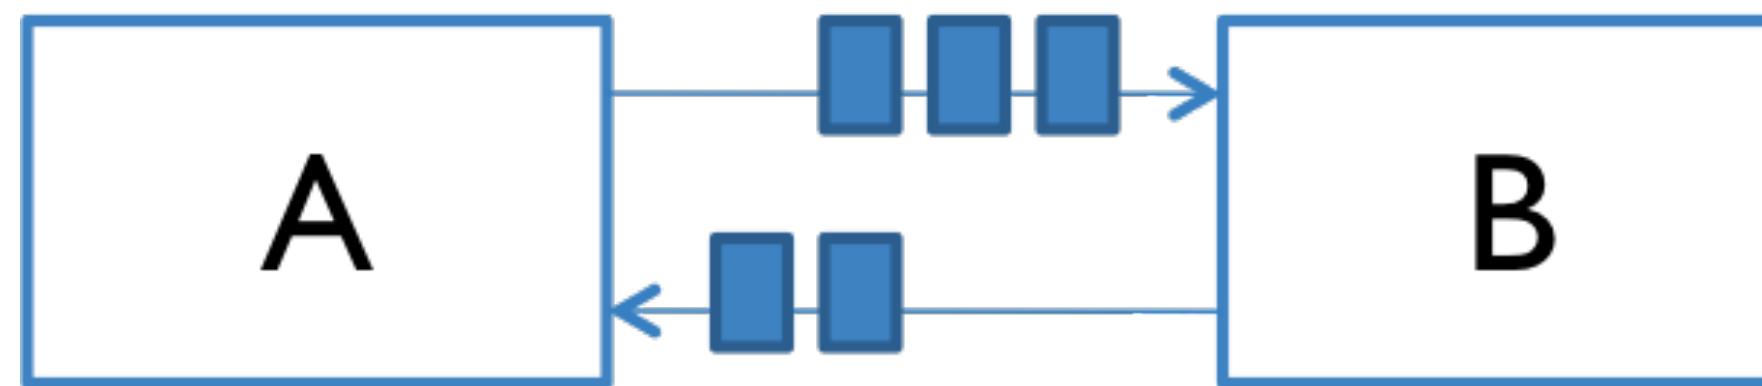


Threads exchange
data through shared
memory

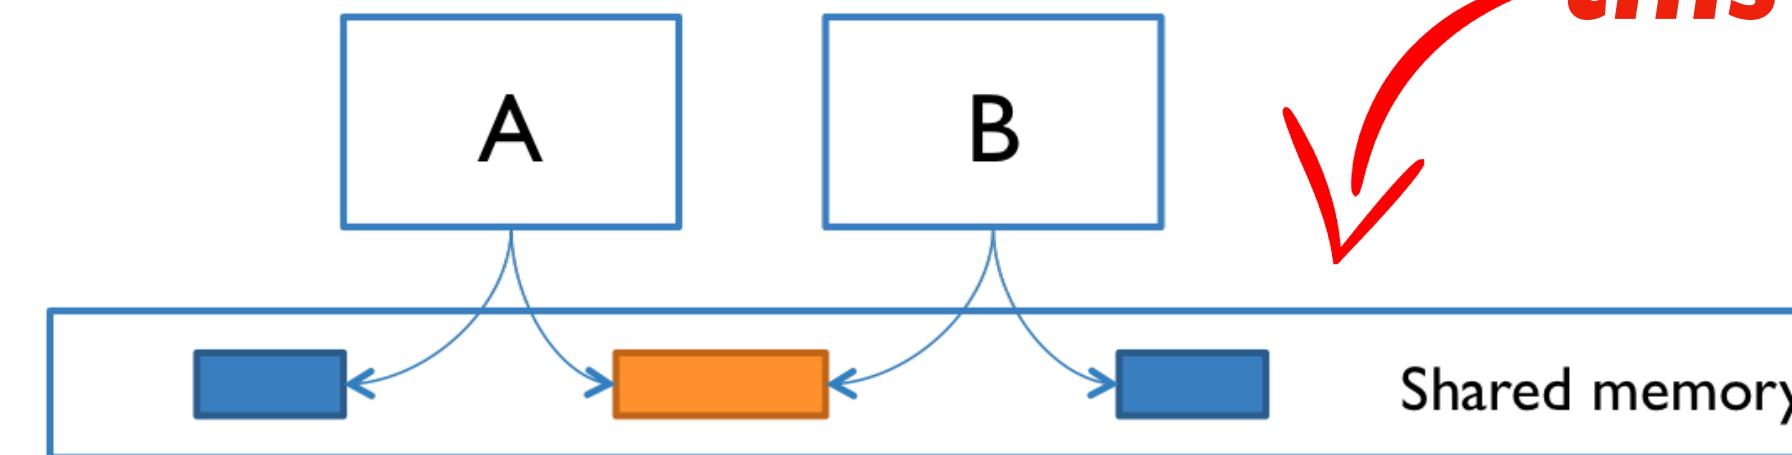
Shared Memory vs Message Passing

Two main models of concurrency:
message-passing and shared memory

**Focus of
this talk**



Threads pass
messages; no
shared state



Threads exchange
data through shared
memory

Shared Memory Primitives



Locks

Used to provide
mutual exclusion



Condition
variables

Signal threads when shared
resource becomes available

Shared Memory Primitives



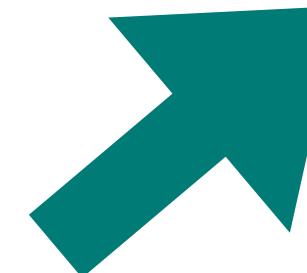
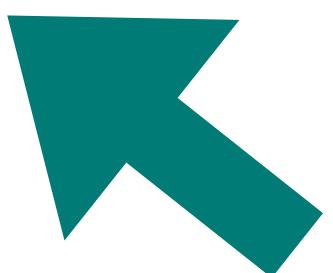
Locks

Used to provide mutual exclusion



Condition variables

Signal threads when shared resource becomes available



Typically used in conjunction with each other to implement the **monitor** programming abstraction

The Monitor Abstraction

```
class Mutex {  
    private boolean h;  
  
    private Lock l = new ...;  
    private Condition heldC  
        = l.newCond();  
  
    public void acq() {  
        l.lock();  
        while (h)  
            heldC.await();  
  
        h = true;  
        l.unlock();  
    }  
  
    public void rel() {  
        l.lock();  
        h = false;  
        heldC.signal();  
        l.unlock();  
    }  
}
```

- Class that mediates access to shared resources

The Monitor Abstraction

```
class Mutex {  
    private boolean h;  
  
    private Lock l = new ...;  
    private Condition heldC  
        = l.newCond();  
  
    public void acq() {  
        l.lock();  
        while (h)  
            heldC.await();  
  
        h = true;  
        l.unlock();  
    }  
  
    public void rel() {  
        l.lock();  
        h = false;  
        heldC.signal();  
        l.unlock();  
    }  
}
```

- Class that mediates access to shared resources
- Threads access shared resources only by calling monitor methods
- Monitor implementation should guarantee **deadlock freedom** and **atomicity**

Implementing Monitors

```
class Mutex {  
    private boolean h;  
  
    private Lock l = new ...;  
    private Condition heldC  
        = l.newCond();  
  
    public void acq() {  
        l.lock();  
        while (h)  
            heldC.await();  
  
        h = true;  
        l.unlock();  
    }  
  
    public void rel() {  
        l.lock();  
        h = false;  
        heldC.signal();  
        l.unlock();  
    }  
}
```

- Monitors typically implemented using **locks** and **condition variables**

Implementing Monitors

```
class Mutex {  
    private boolean h;  
  
    private Lock l = new ...;  
    private Condition heldC  
        = l.newCond();  
  
    public void acq() {  
        l.lock();  
        while (h)  
            heldC.await();  
  
        h = true;  
        l.unlock();  
    }  
  
    public void rel() {  
        l.lock();  
        h = false;  
        heldC.signal();  
        l.unlock();  
    }  
}
```

- Monitors typically implemented using **locks** and **condition variables**

Implementing Monitors

```
class Mutex {  
    private boolean h;  
  
    private Lock l = new ...;  
    private Condition heldC  
        = l.newCond();  
  
    public void acq() {  
        l.lock();  
        while (h)  
            heldC.await();  
  
        h = true;  
        l.unlock();  
    }  
  
    public void rel() {  
        l.lock();  
        h = false;  
        heldC.signal();  
        l.unlock();  
    }  
}
```

- Monitors typically implemented using **locks** and **condition variables**

Implementing Monitors

```
class Mutex {  
    private boolean h;  
  
    private Lock l = new ...;  
    private Condition heldC  
        = l.newCond();  
  
    public void acq() {  
        l.lock();  
        while (h)  
            heldC.await();  
  
        h = true;  
        l.unlock();  
    }  
  
    public void rel() {  
        l.lock();  
        h = false;  
        heldC.signal();  
        l.unlock();  
    }  
}
```

- Monitors typically implemented using **locks** and **condition variables**

Implementing Monitors

```
class Mutex {  
    private boolean h;  
  
    private Lock l = new ...;  
    private Condition heldC  
        = l.newCond();  
  
    public void acq() {  
        l.lock();  
        while (h)  
            heldC.await();  
  
        h = true;  
        l.unlock();  
    }  
  
    public void rel() {  
        l.lock();  
        h = false;  
        heldC.signal();  
        l.unlock();  
    }  
}
```

- Monitors typically implemented using **locks** and **condition variables**

Implementing Monitors

```
class Mutex {  
    private boolean h;  
  
    private Lock l = new ...;  
    private Condition heldC  
        = l.newCond();  
  
    public void acq() {  
        l.lock();  
        while (h)  
            heldC.await();  
  
        h = true;  
        l.unlock();  
    }  
  
    public void rel() {  
        l.lock();  
        h = false;  
        heldC.signal();  
        l.unlock();  
    }  
}
```

- Monitors typically implemented using **locks** and **condition variables**
- But hard to implement both correctly and efficiently
- Improper locking and signaling cause **atomicity violations** or **deadlocks**

Our Goal



Simplify concurrent programming by making it much easier to implement **correct and efficient** monitors

Our Goal



Simplify concurrent programming by making it much easier to implement **correct and efficient** monitors

Approach:

- (I) Use **Conditional Critical Regions (CCRs)** instead of locks and condition variables
 - Provides higher level of abstraction — hence easier
- (2) Synthesize efficient & correct low-level code with locks and condition variables

Conditional Critical Region (CCRs)

```
waituntil(P) {S}
```

- Blocks when P is false and atomically executes S when P becomes true

Conditional Critical Region (CCRs)

```
waituntil(P) {S}
```

- Blocks when P is false and atomically executes S when P becomes true
- Provides a higher level of abstraction:
 - No need to use locks to ensure atomicity of S
 - No need to explicitly signal using condition variables

Conditional Critical Region (CCRs)

```
waituntil(P) {S}
```

- Blocks when P is false and atomically executes S when P becomes true
- Provides a higher level of abstraction:
 - No need to use locks to ensure atomicity of S
 - No need to explicitly signal using condition variables

Called *implicit monitor*

Assumption: Each monitor method is a sequence of CCRs

Back to Monitors

```
class Mutex {  
    private boolean h;  
  
    private Lock l = new ...;  
    private Condition heldC  
        = l.newCond();  
  
    public void acq() {  
        l.lock();  
        while (h)  
            bCond.await();  
  
        h = true;  
        l.unlock();  
    }  
  
    public void rel() {  
        l.lock();  
        h = false;  
        bCond.signal();  
        l.unlock();  
    }  
}
```

Explicit monitor

```
class Mutex {  
    private boolean h;  
  
    public void atomic acq() {  
        waituntil (!h) {  
            h = true;  
        }  
    }  
  
    public void atomic rel() {  
        h = false;  
    }  
}
```

Implicit monitor

Back to Monitors

```
class Mutex {  
    private boolean h;  
  
    private Lock l = new ...;  
    private Condition heldC  
        = l.newCond();  
  
    public void acq() {  
        l.lock();  
        while (h)  
            bCond.await();  
  
        h = true;  
        l.unlock();  
    }  
  
    public void rel() {  
        l.lock();  
        h = false;  
        bCond.signal();  
        l.unlock();  
    }  
}
```

Explicit monitor

```
class Mutex {  
    private boolean h;  
  
    public void atomic acq() {  
        waituntil (!h) {  
            h = true;  
        }  
    }  
  
    public void atomic rel() {  
        h = false;  
    }  
}
```

Omit waituntil if predicate is true.

Implicit monitor

Back to Monitors

```
class Mutex {  
    private boolean h;  
  
    private Lock l = new ...;  
    private Condition heldC  
        = l.newCond();  
  
    public void acq() {  
        l.lock();  
        while (h)  
            bCond.await();  
  
        h = true;  
        l.unlock();  
    }  
  
    public void rel() {  
        l.lock();  
        h = false;  
        bCond.signal();  
        l.unlock();  
    }  
}
```

Explicit monitor

Preds(M)

```
class Mutex {  
    private boolean h;  
  
    public void atomic acq() {  
        waituntil (!h) {  
            h = true;  
        }  
    }  
  
    public void atomic rel() {  
        h = false;  
    }  
}
```

Implicit monitor

Back to Monitors

```
class Mutex {  
    private boolean h;  
  
    private Lock l = new ...;  
    private Condition heldC  
        = l.newCond();  
  
    public void acq() {  
        l.lock();  
        while (h)  
            bCond.await();  
  
        h = true;  
        l.unlock();  
    }  
  
    public void rel() {  
        l.lock();  
        h = false;  
        bCond.signal();  
        l.unlock();  
    }  
}
```

Explicit monitor

Implicit Monitor = Explicit Monitor - Synchronization Code

```
class Mutex {  
    private boolean h;  
  
    public void atomic acq() {  
        waituntil (!h) {  
            h = true;  
        }  
    }  
  
    public void atomic rel() {  
        h = false;  
    }  
}
```

Implicit monitor

Explicit vs Implicit Trade-offs

Explicit

Pro: Efficient

Con: Hard to get right



Implicit

Pro: Easy to get right

Con: Inefficient

Explicit vs Implicit Trade-offs

Explicit

Pro: Efficient

Con: Hard to get right



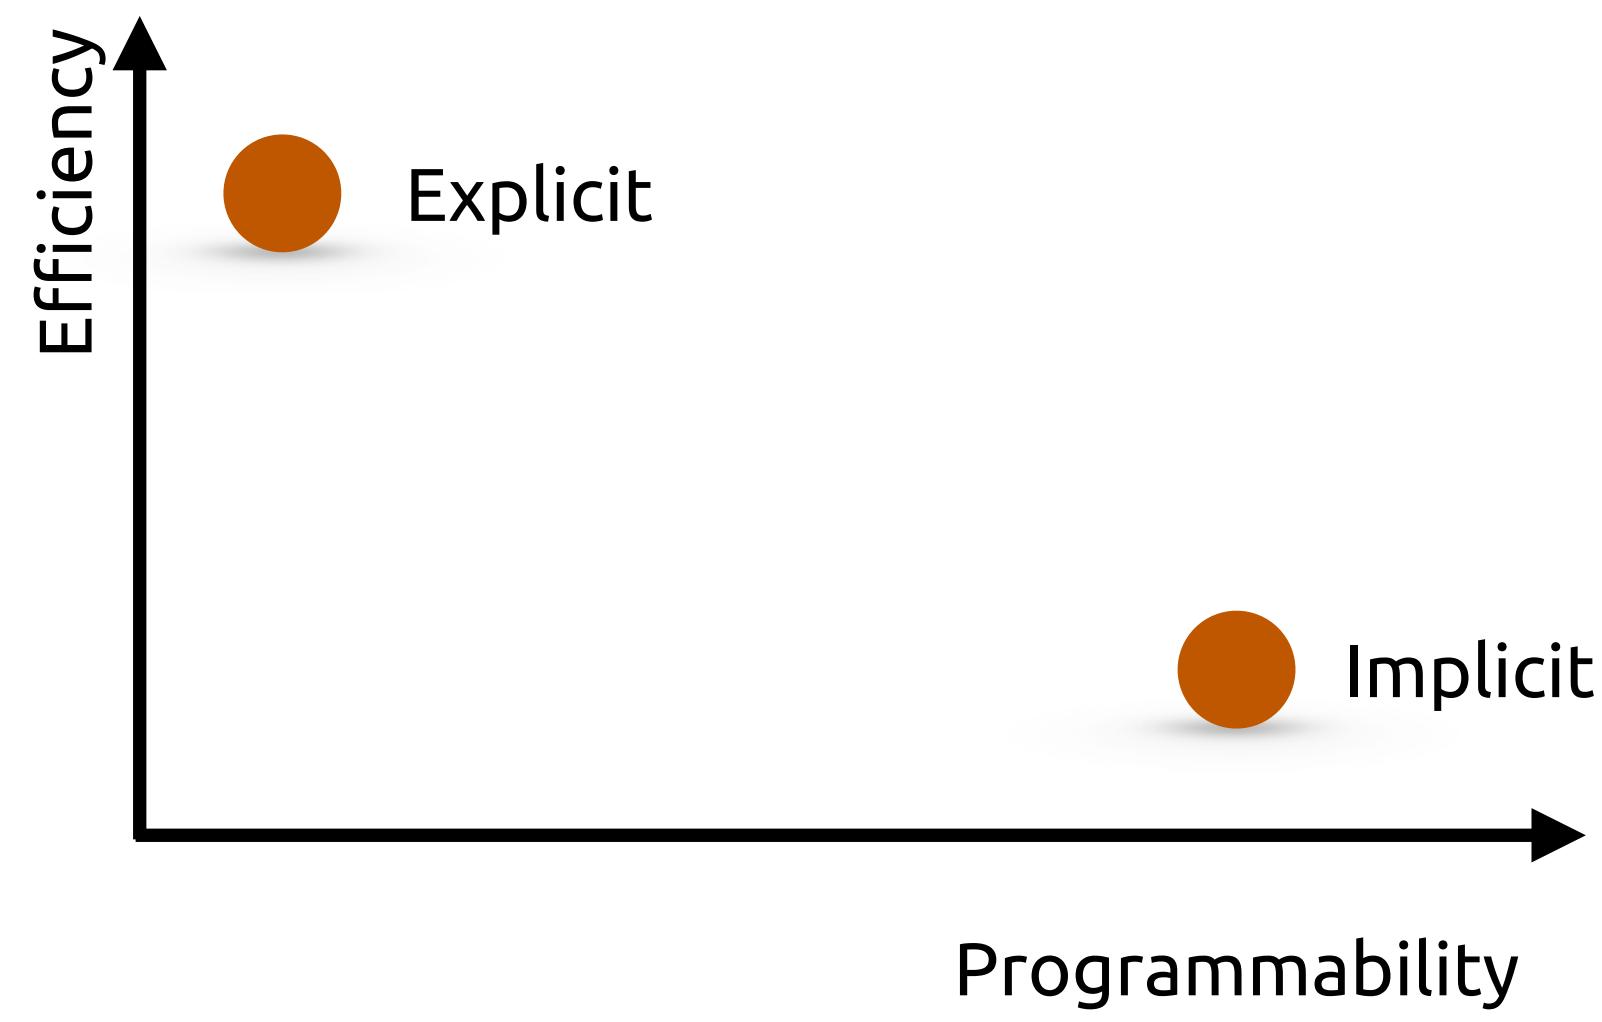
Implicit

Pro: Easy to get right

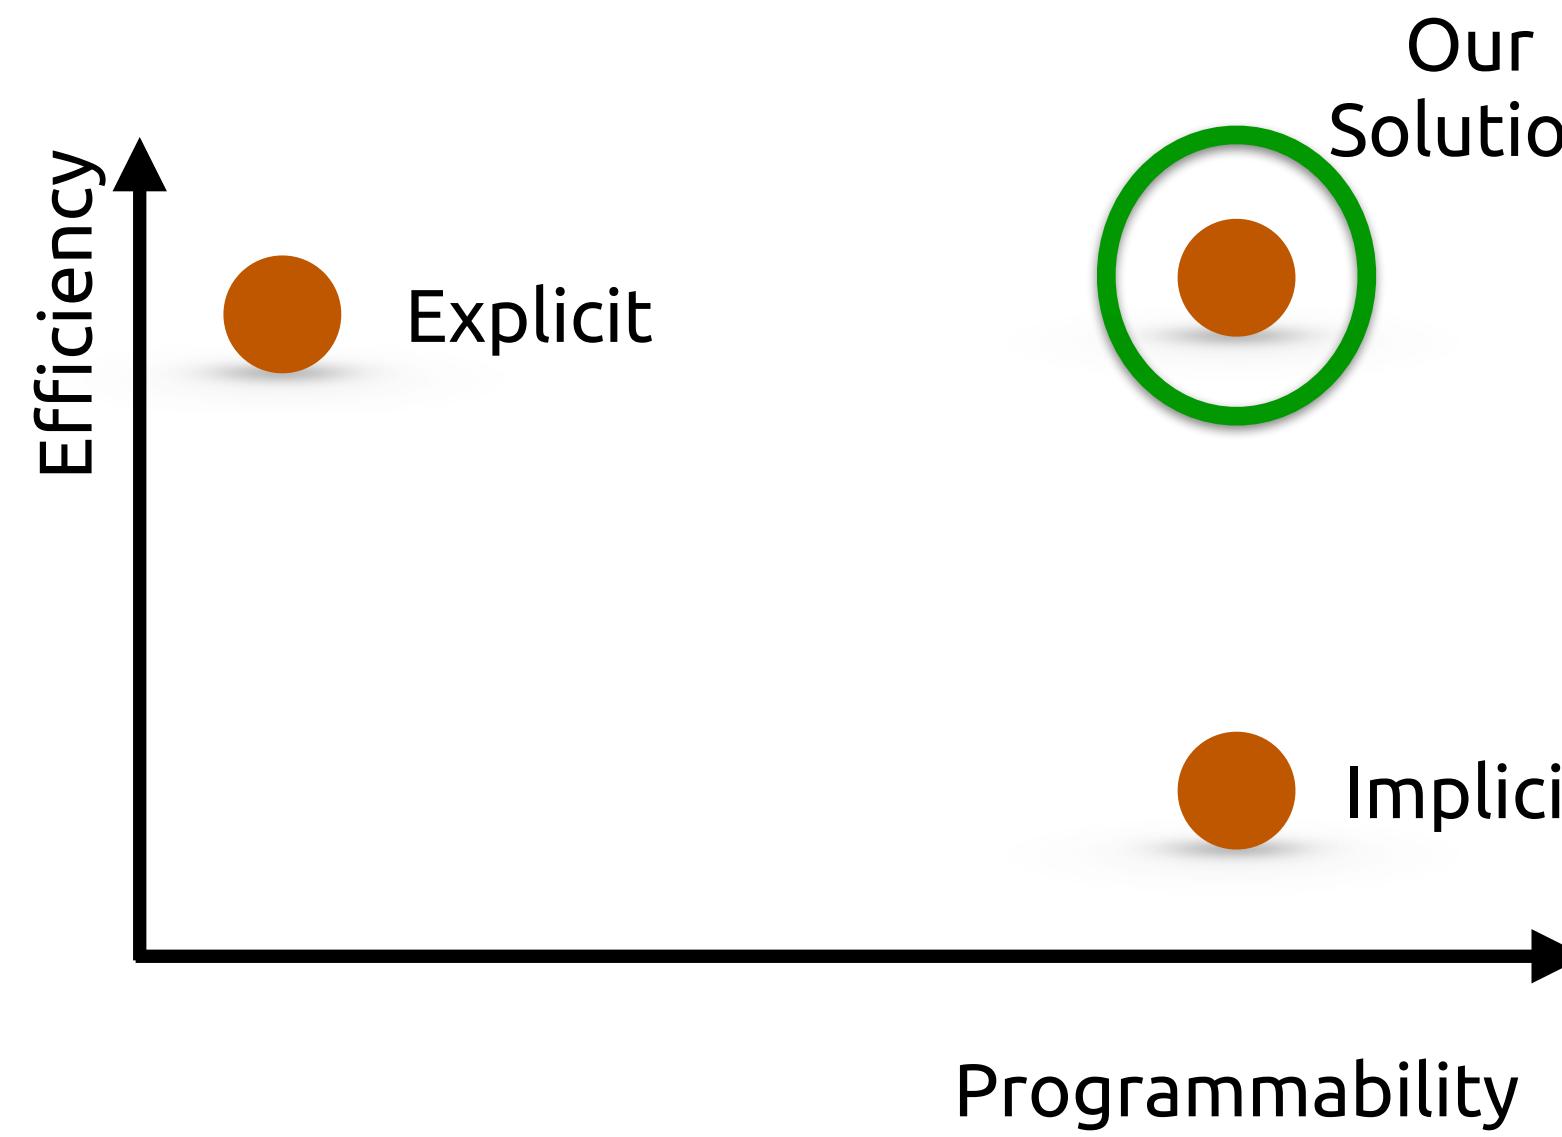
Con: Inefficient

All existing solutions require run-time support

Synthesis-Based Solution



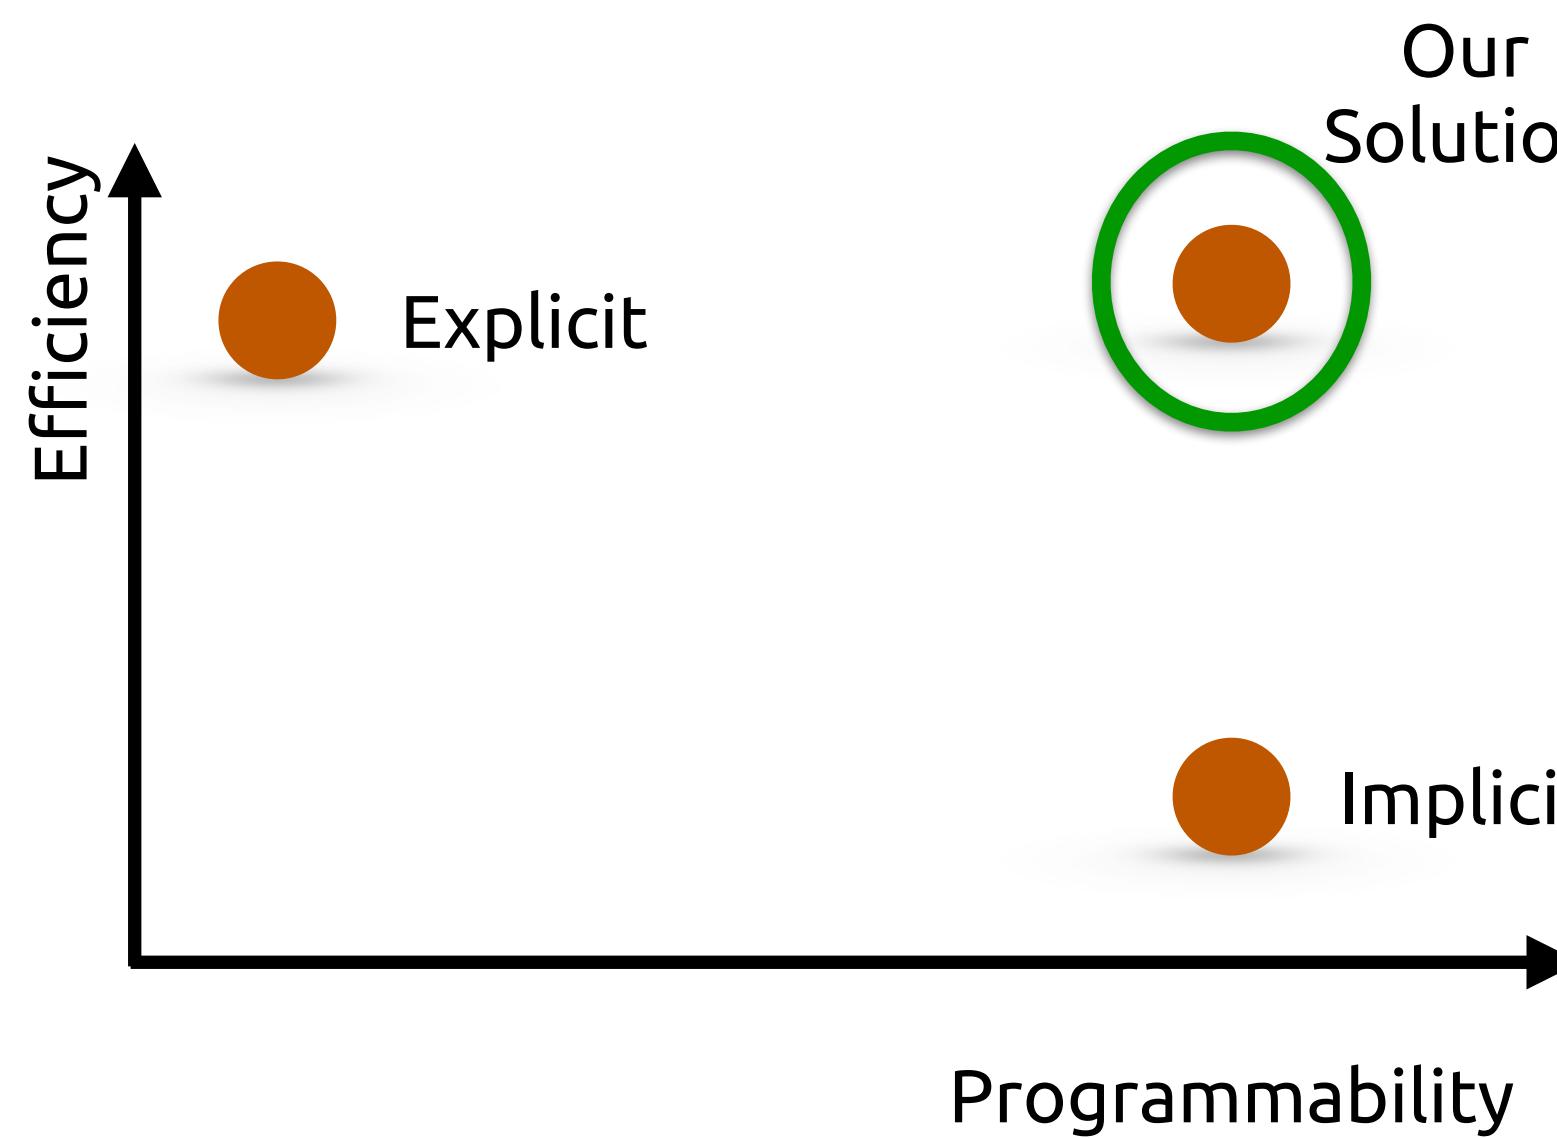
Synthesis-Based Solution



Goal: Synthesize explicit monitor from implicit

- Correct-by-construction
- As efficient as possible

Synthesis-Based Solution



Goal: Synthesize explicit monitor from implicit

- Correct-by-construction
- As efficient as possible

Provides best of both worlds in terms of efficiency & programmability!

Road Map

Part I: How to automatically synthesize necessary **signaling** code



Symbolic Reasoning for Automatic Signal Placement.
Ferles et al, [PLDI'18](#).



Part II: How to automatically synthesize efficient **locking**



Synthesizing Fine-grained Synchronization Protocols for Implicit Monitors.
Ferles et al, [OOPSLA'22](#).

Part I: Synthesizing Signaling Code



The task:

Given an implicit monitor:

- instrument code with condition vars
- add all necessary wait/signal operations on these variables

Part I: Synthesizing Signaling Code

The task:

Given an implicit monitor:

- instrument code with condition vars
- add all necessary wait/signal operations on these variables



Key difficulty: Where and how to signal so that each thread can make progress and there are no spurious wake-ups?

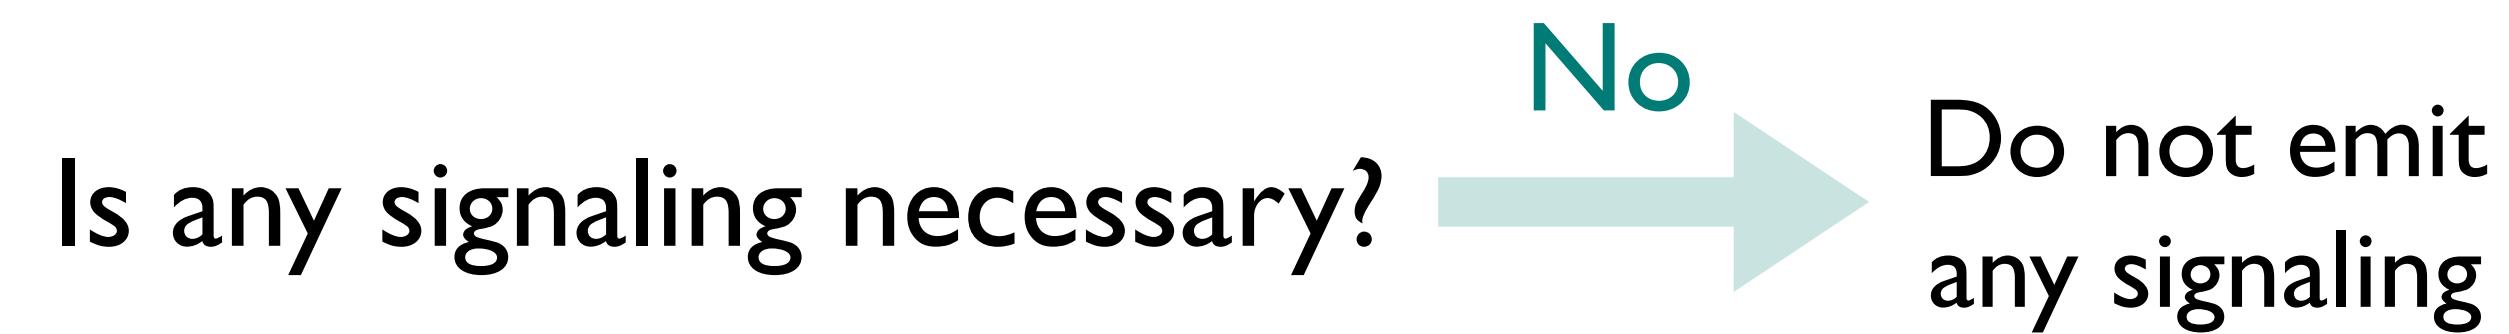
Signaling Decisions

For each CCR C and predicate $p \in Preds(M)$, decide whether to signal or not, and if so, how?

Is any signaling necessary?

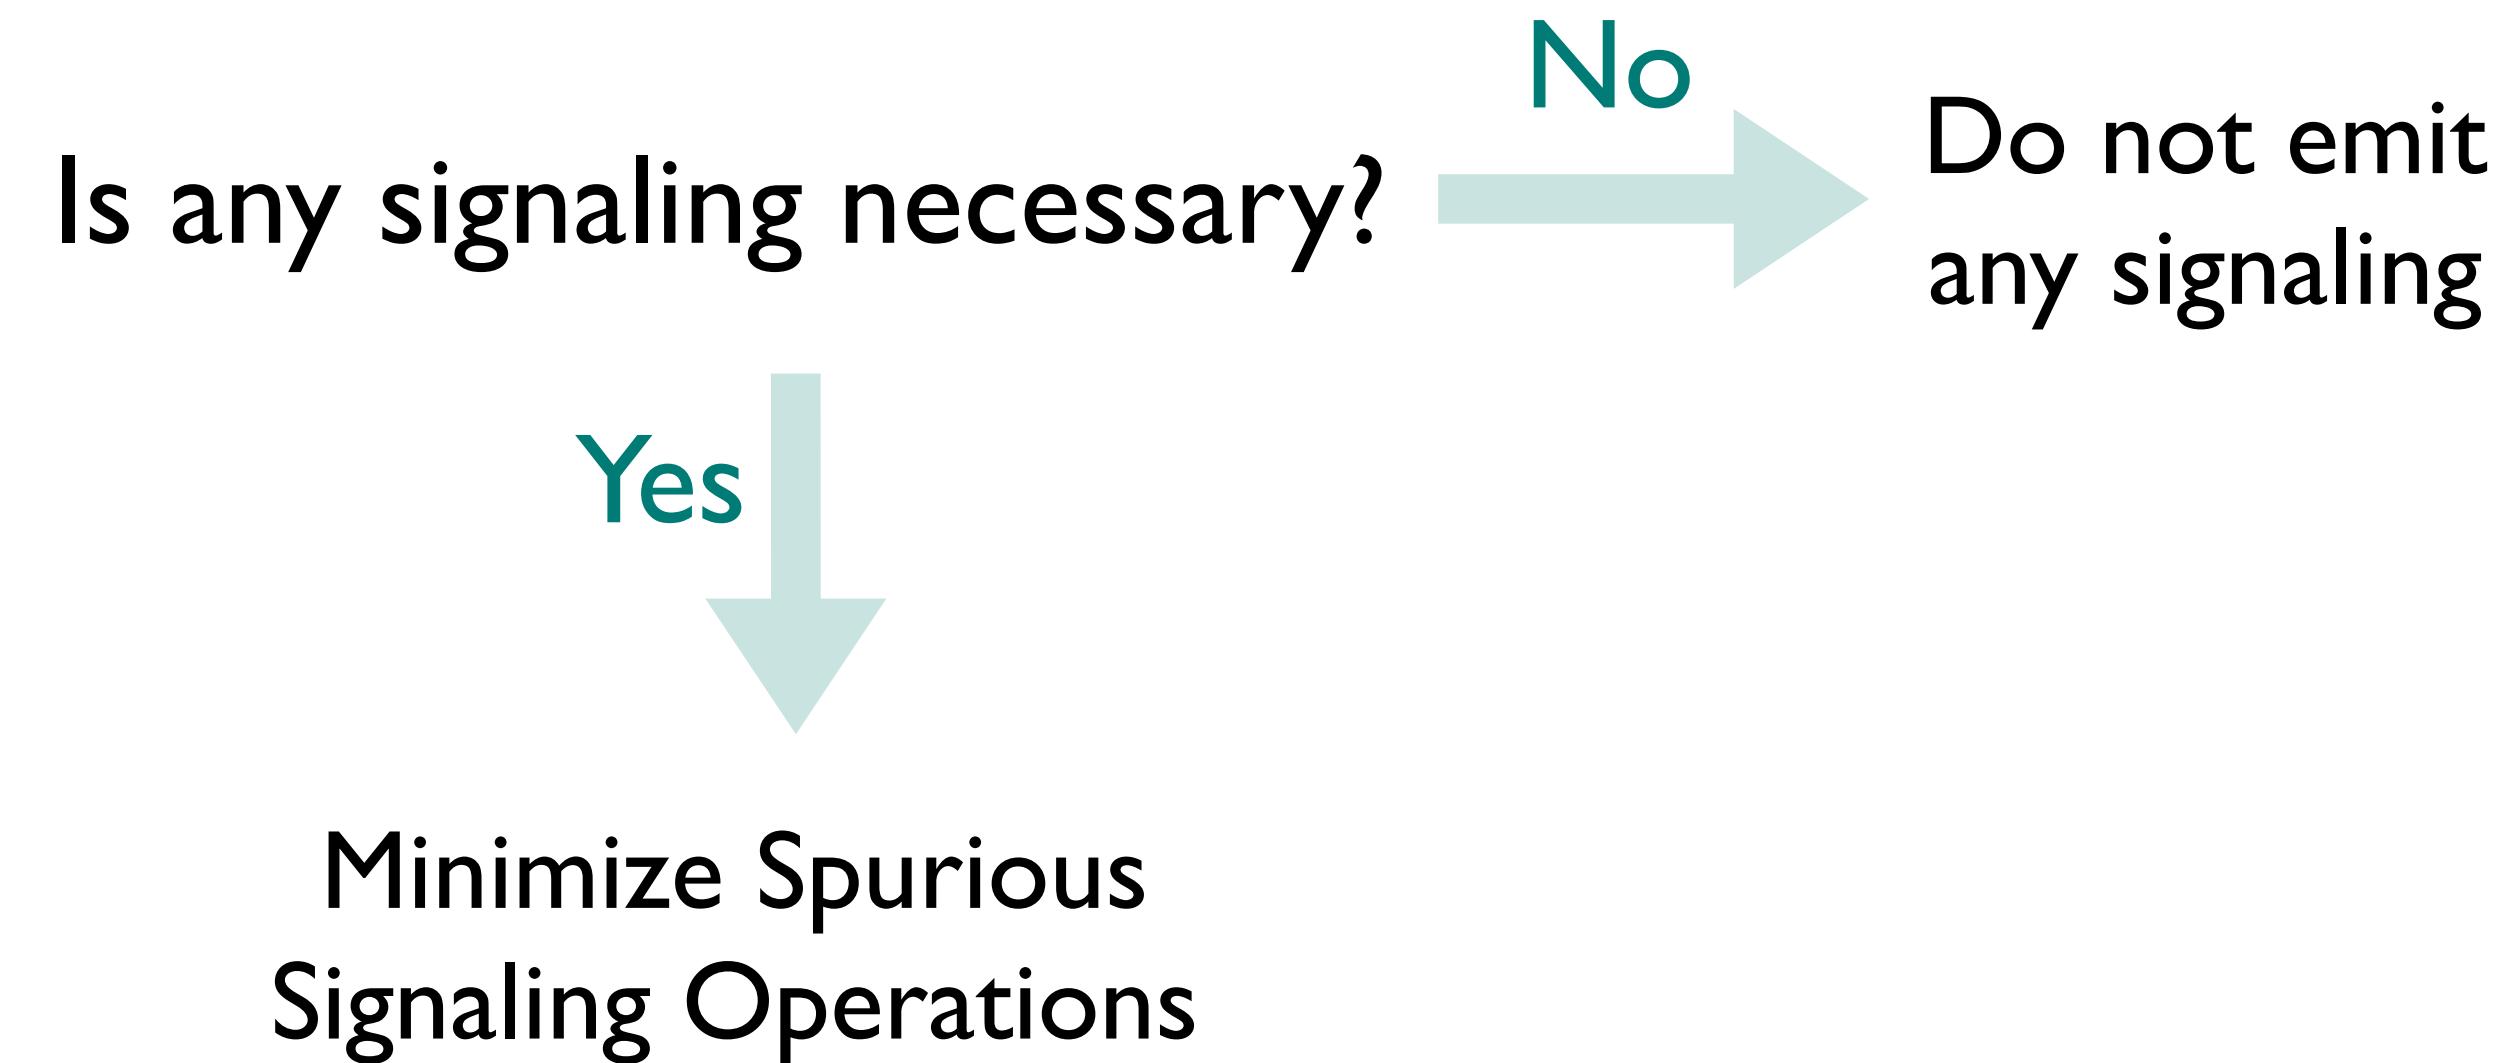
Signaling Decisions

For each CCR C and predicate $p \in Preds(M)$, decide whether to signal or not, and if so, how?



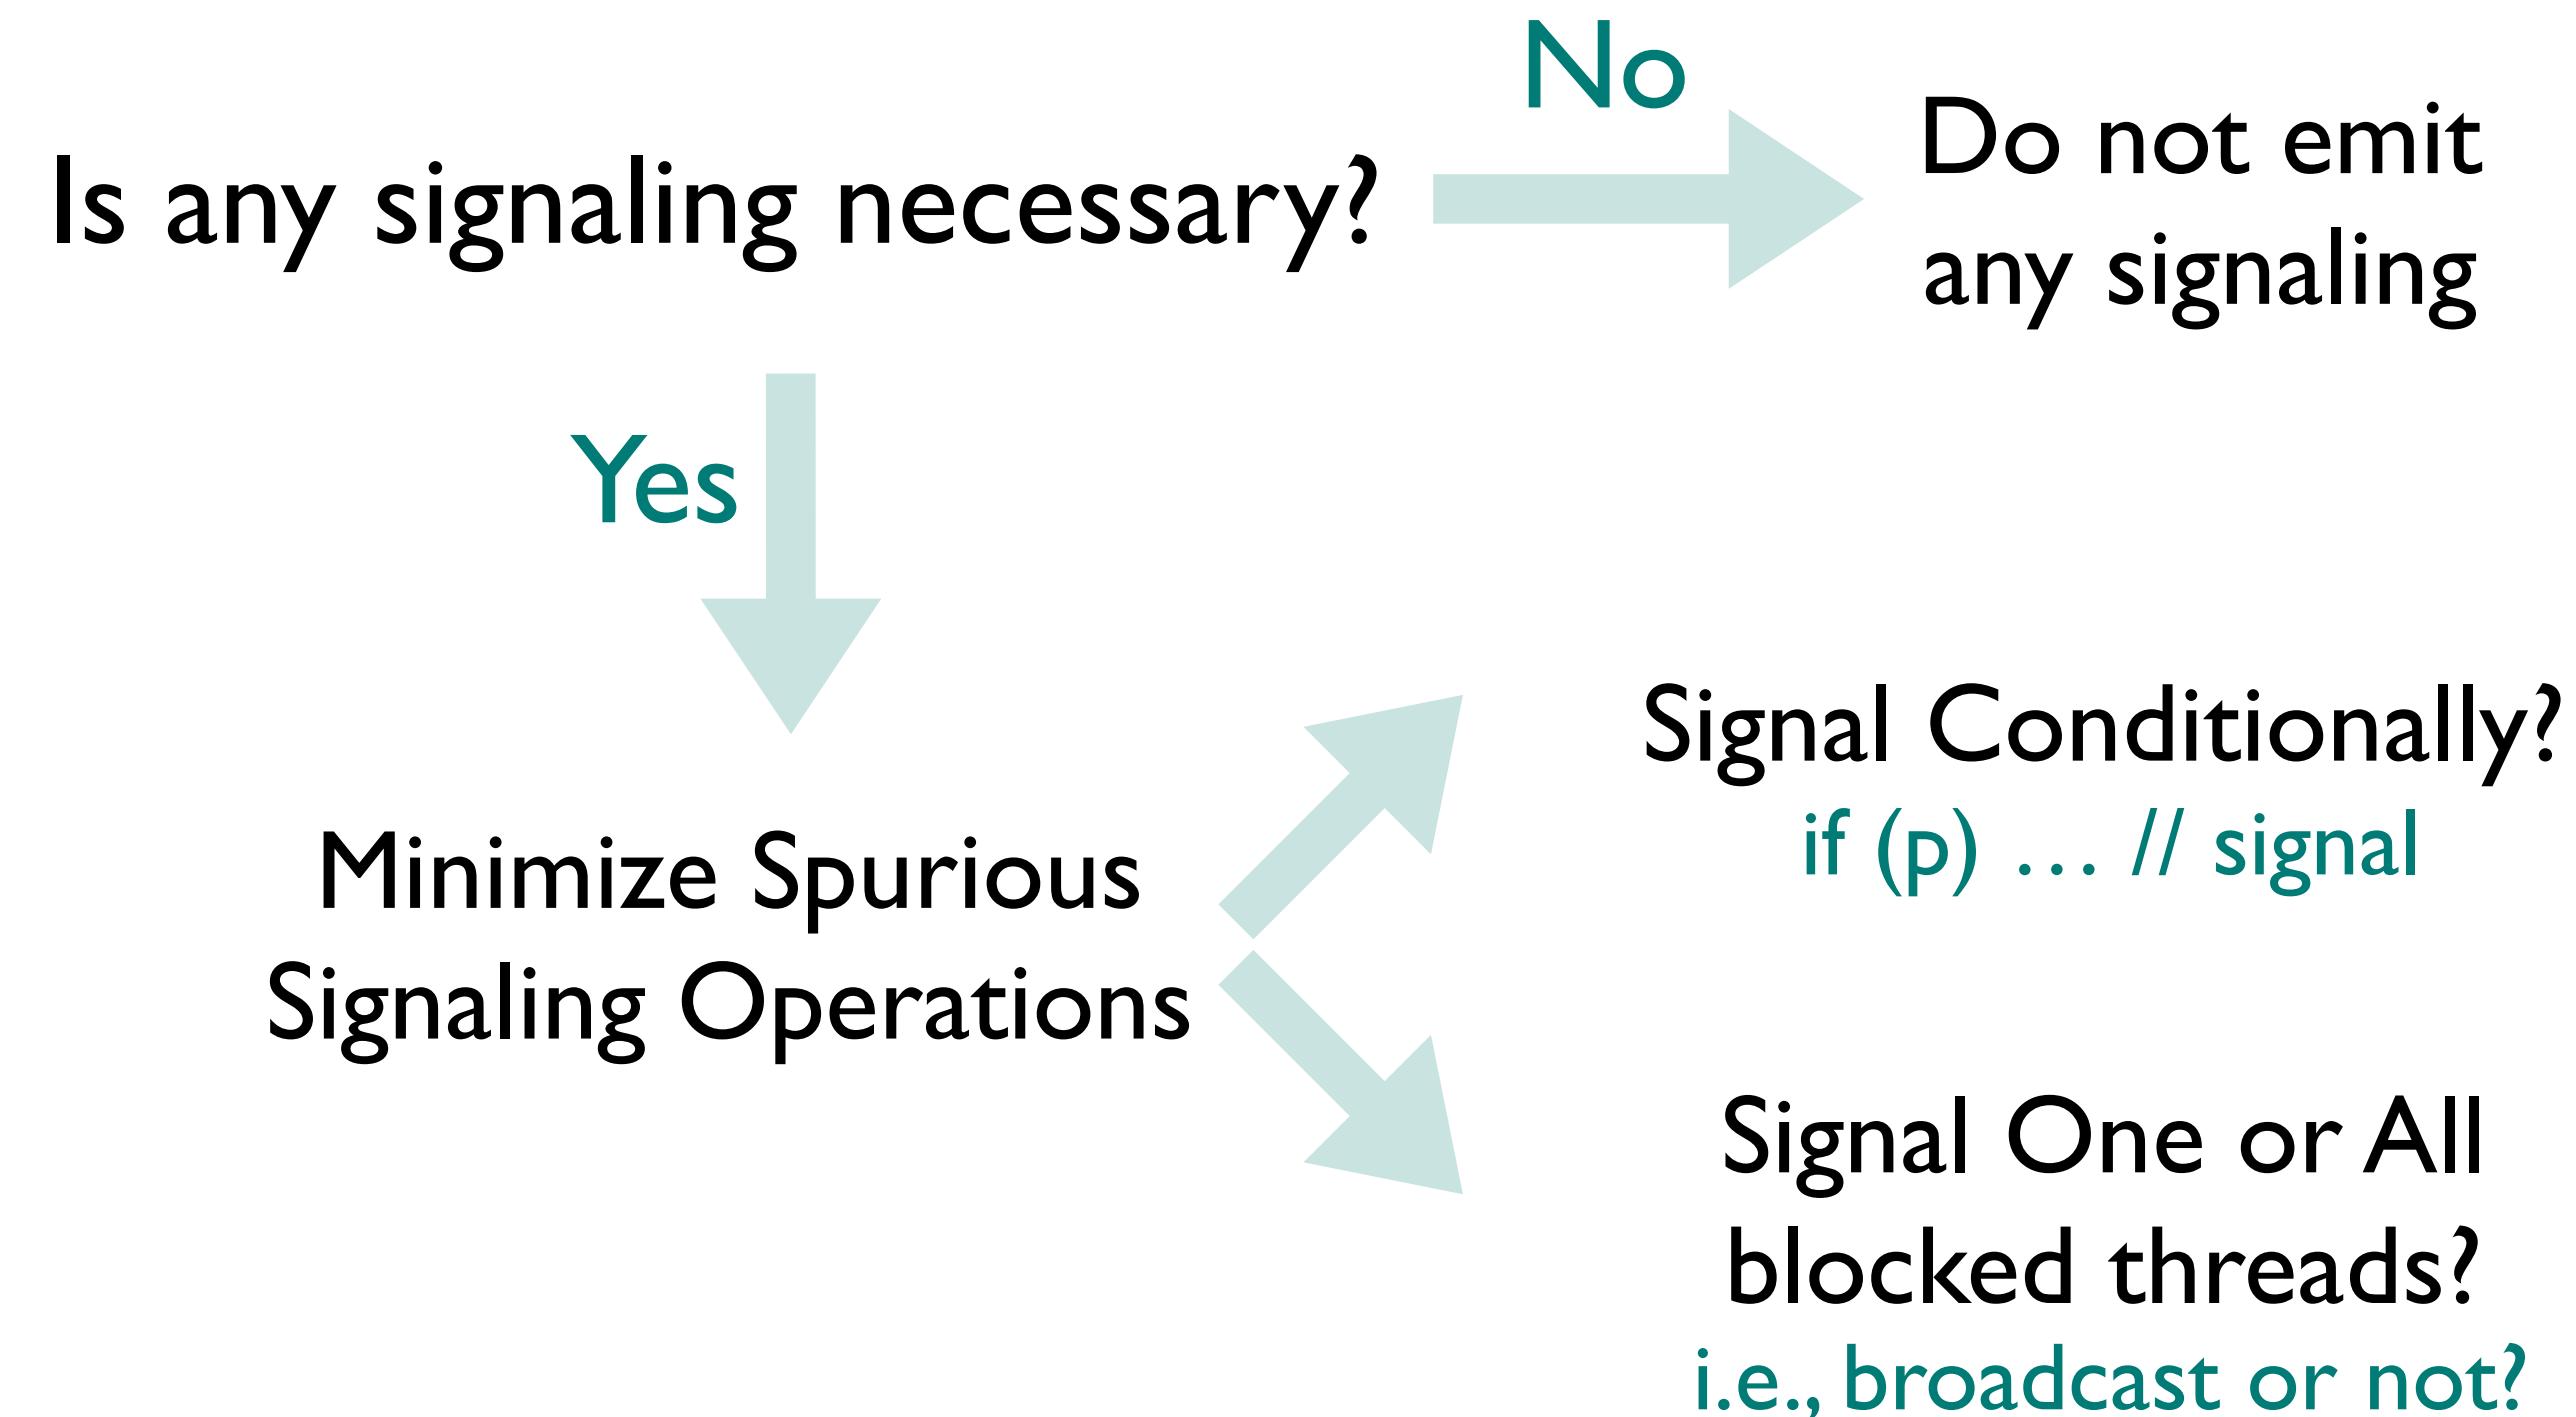
Signaling Decisions

For each CCR C and predicate $p \in Preds(M)$, decide whether to signal or not, and if so, how?



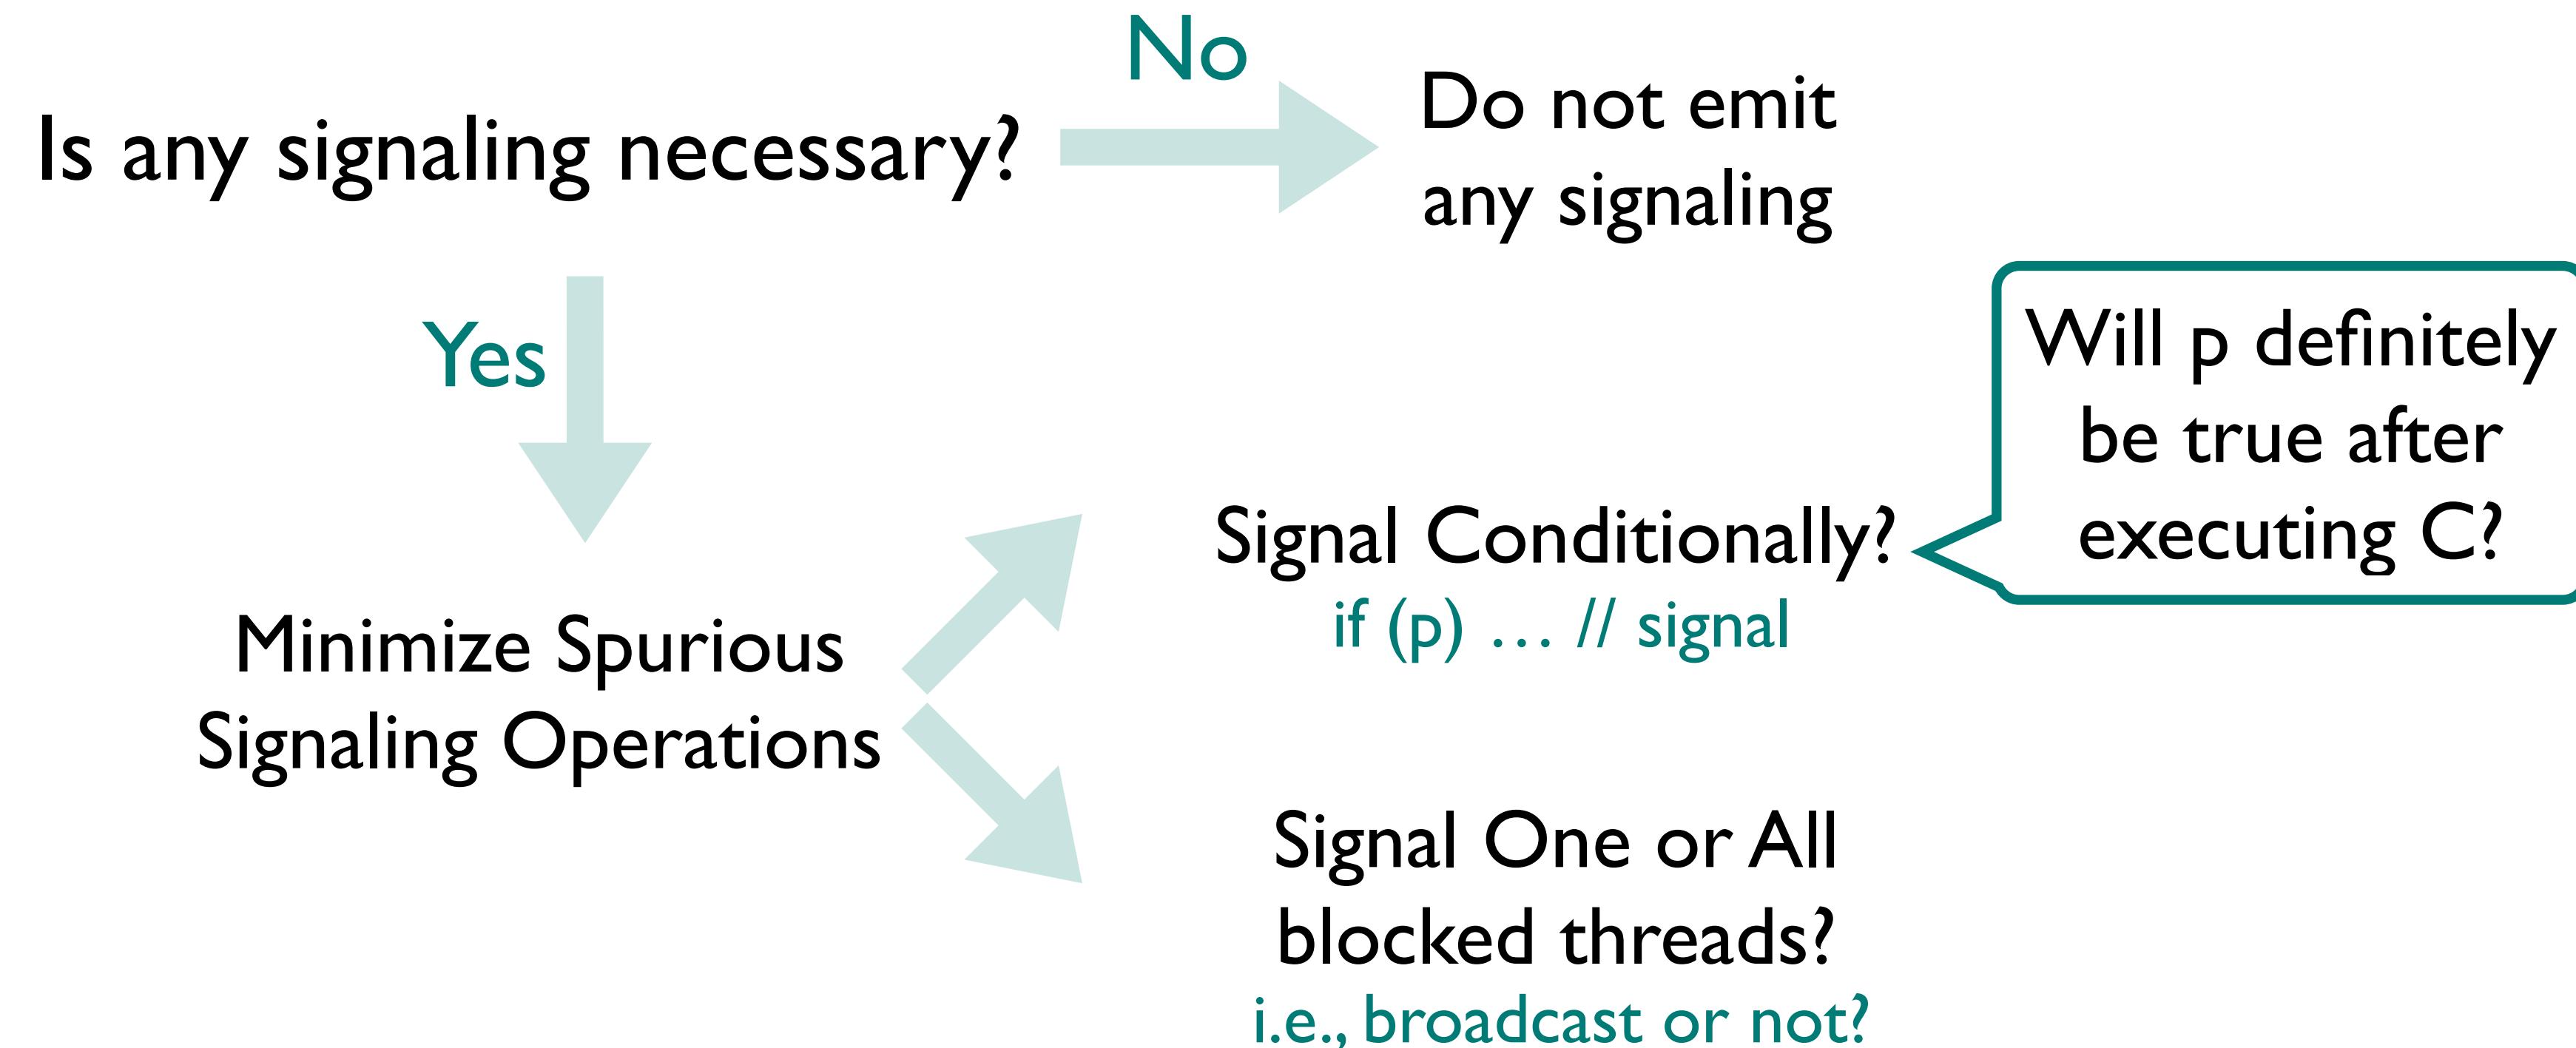
Signaling Decisions

For each CCR C and predicate $p \in Preds(M)$, decide whether to signal or not, and if so, how?



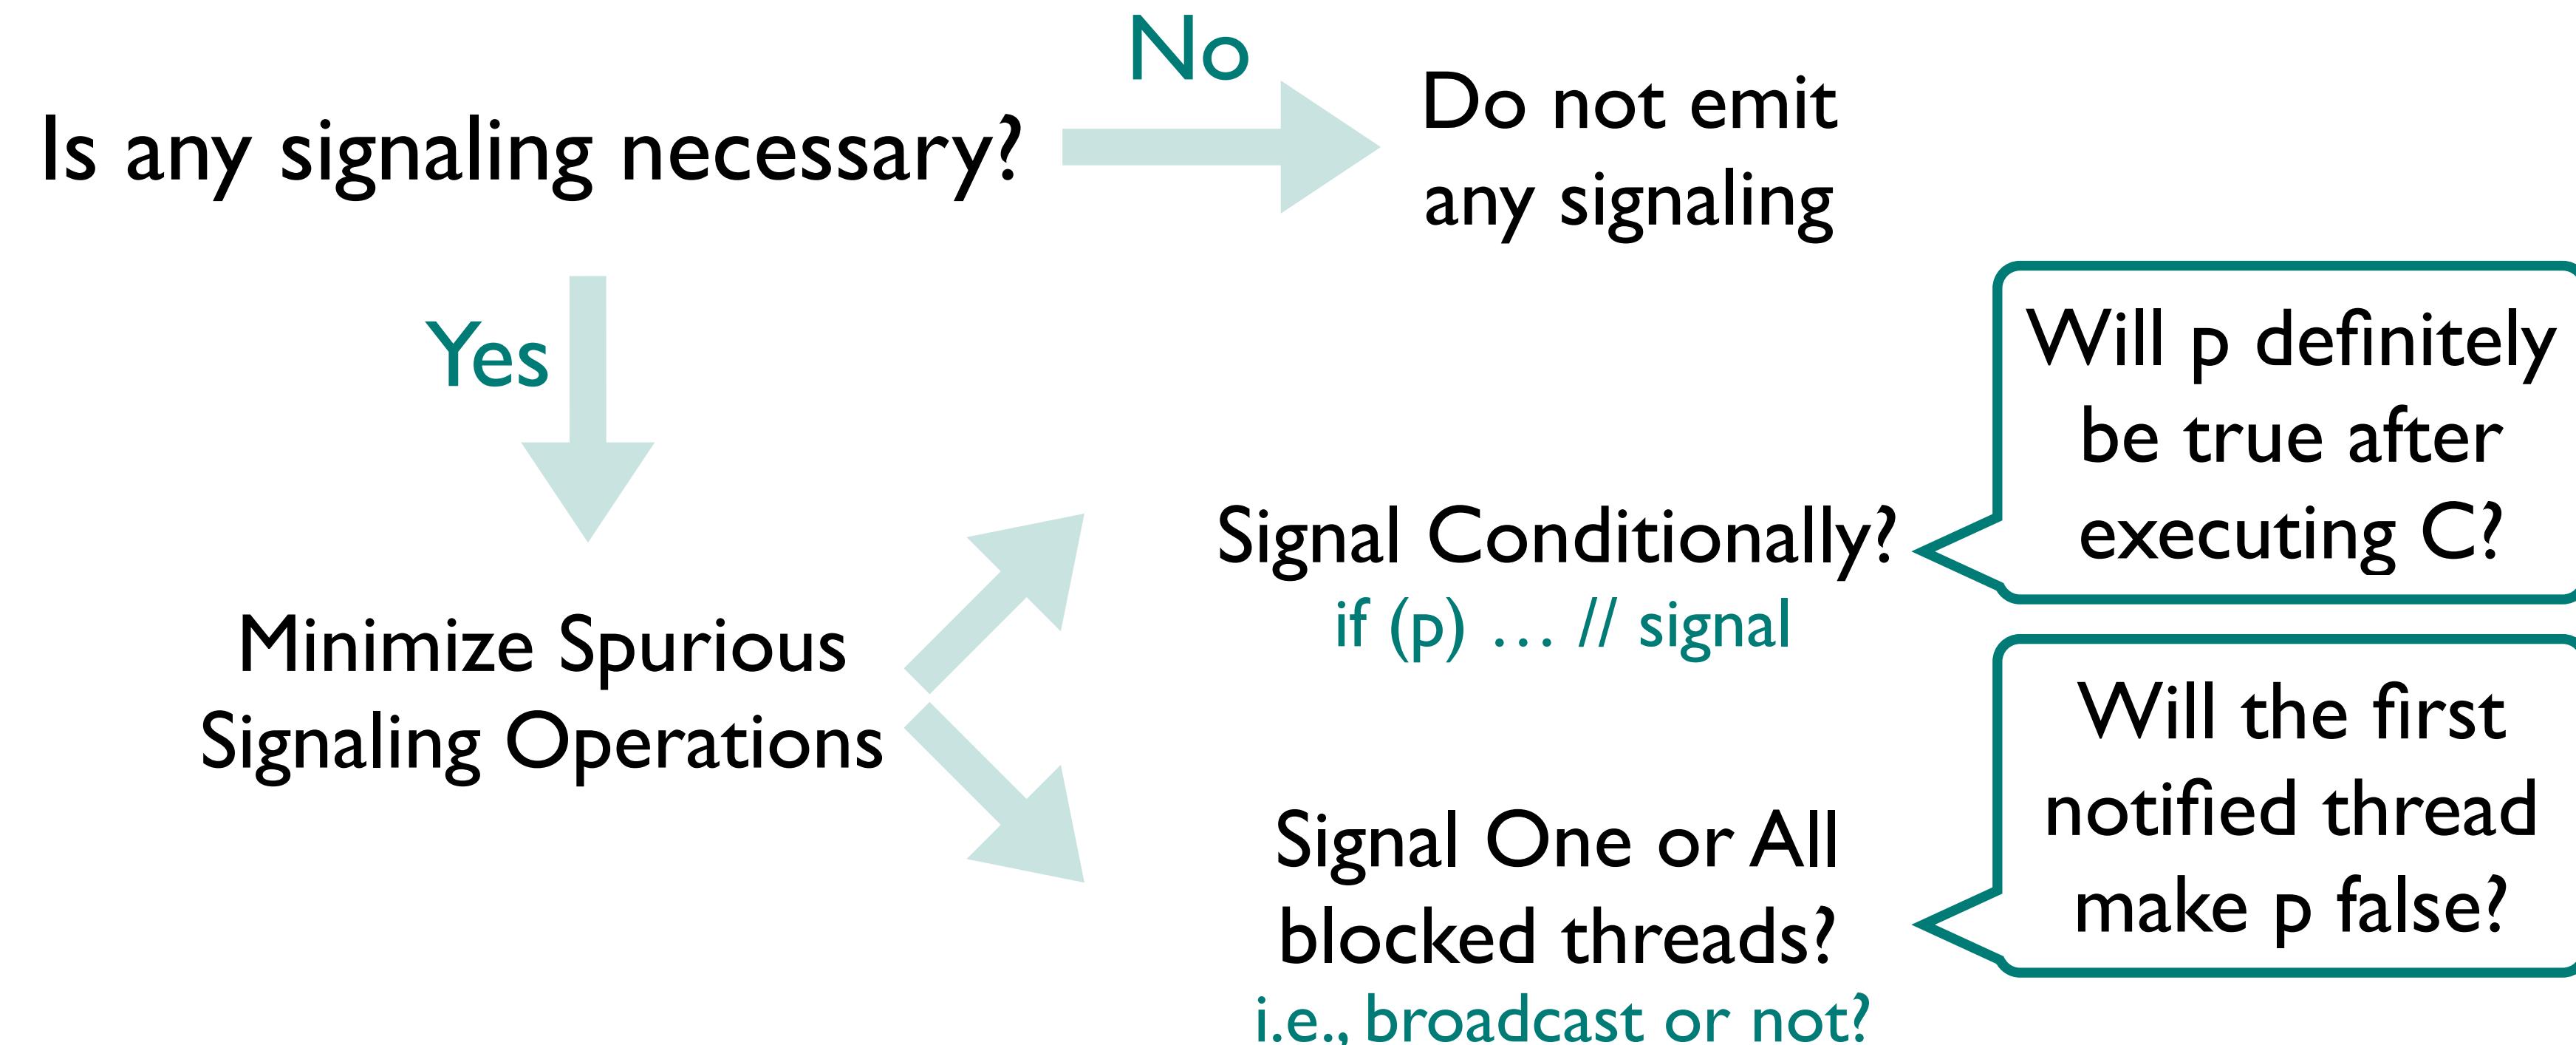
Signaling Decisions

For each CCR C and predicate $p \in Preds(M)$, decide whether to signal or not, and if so, how?



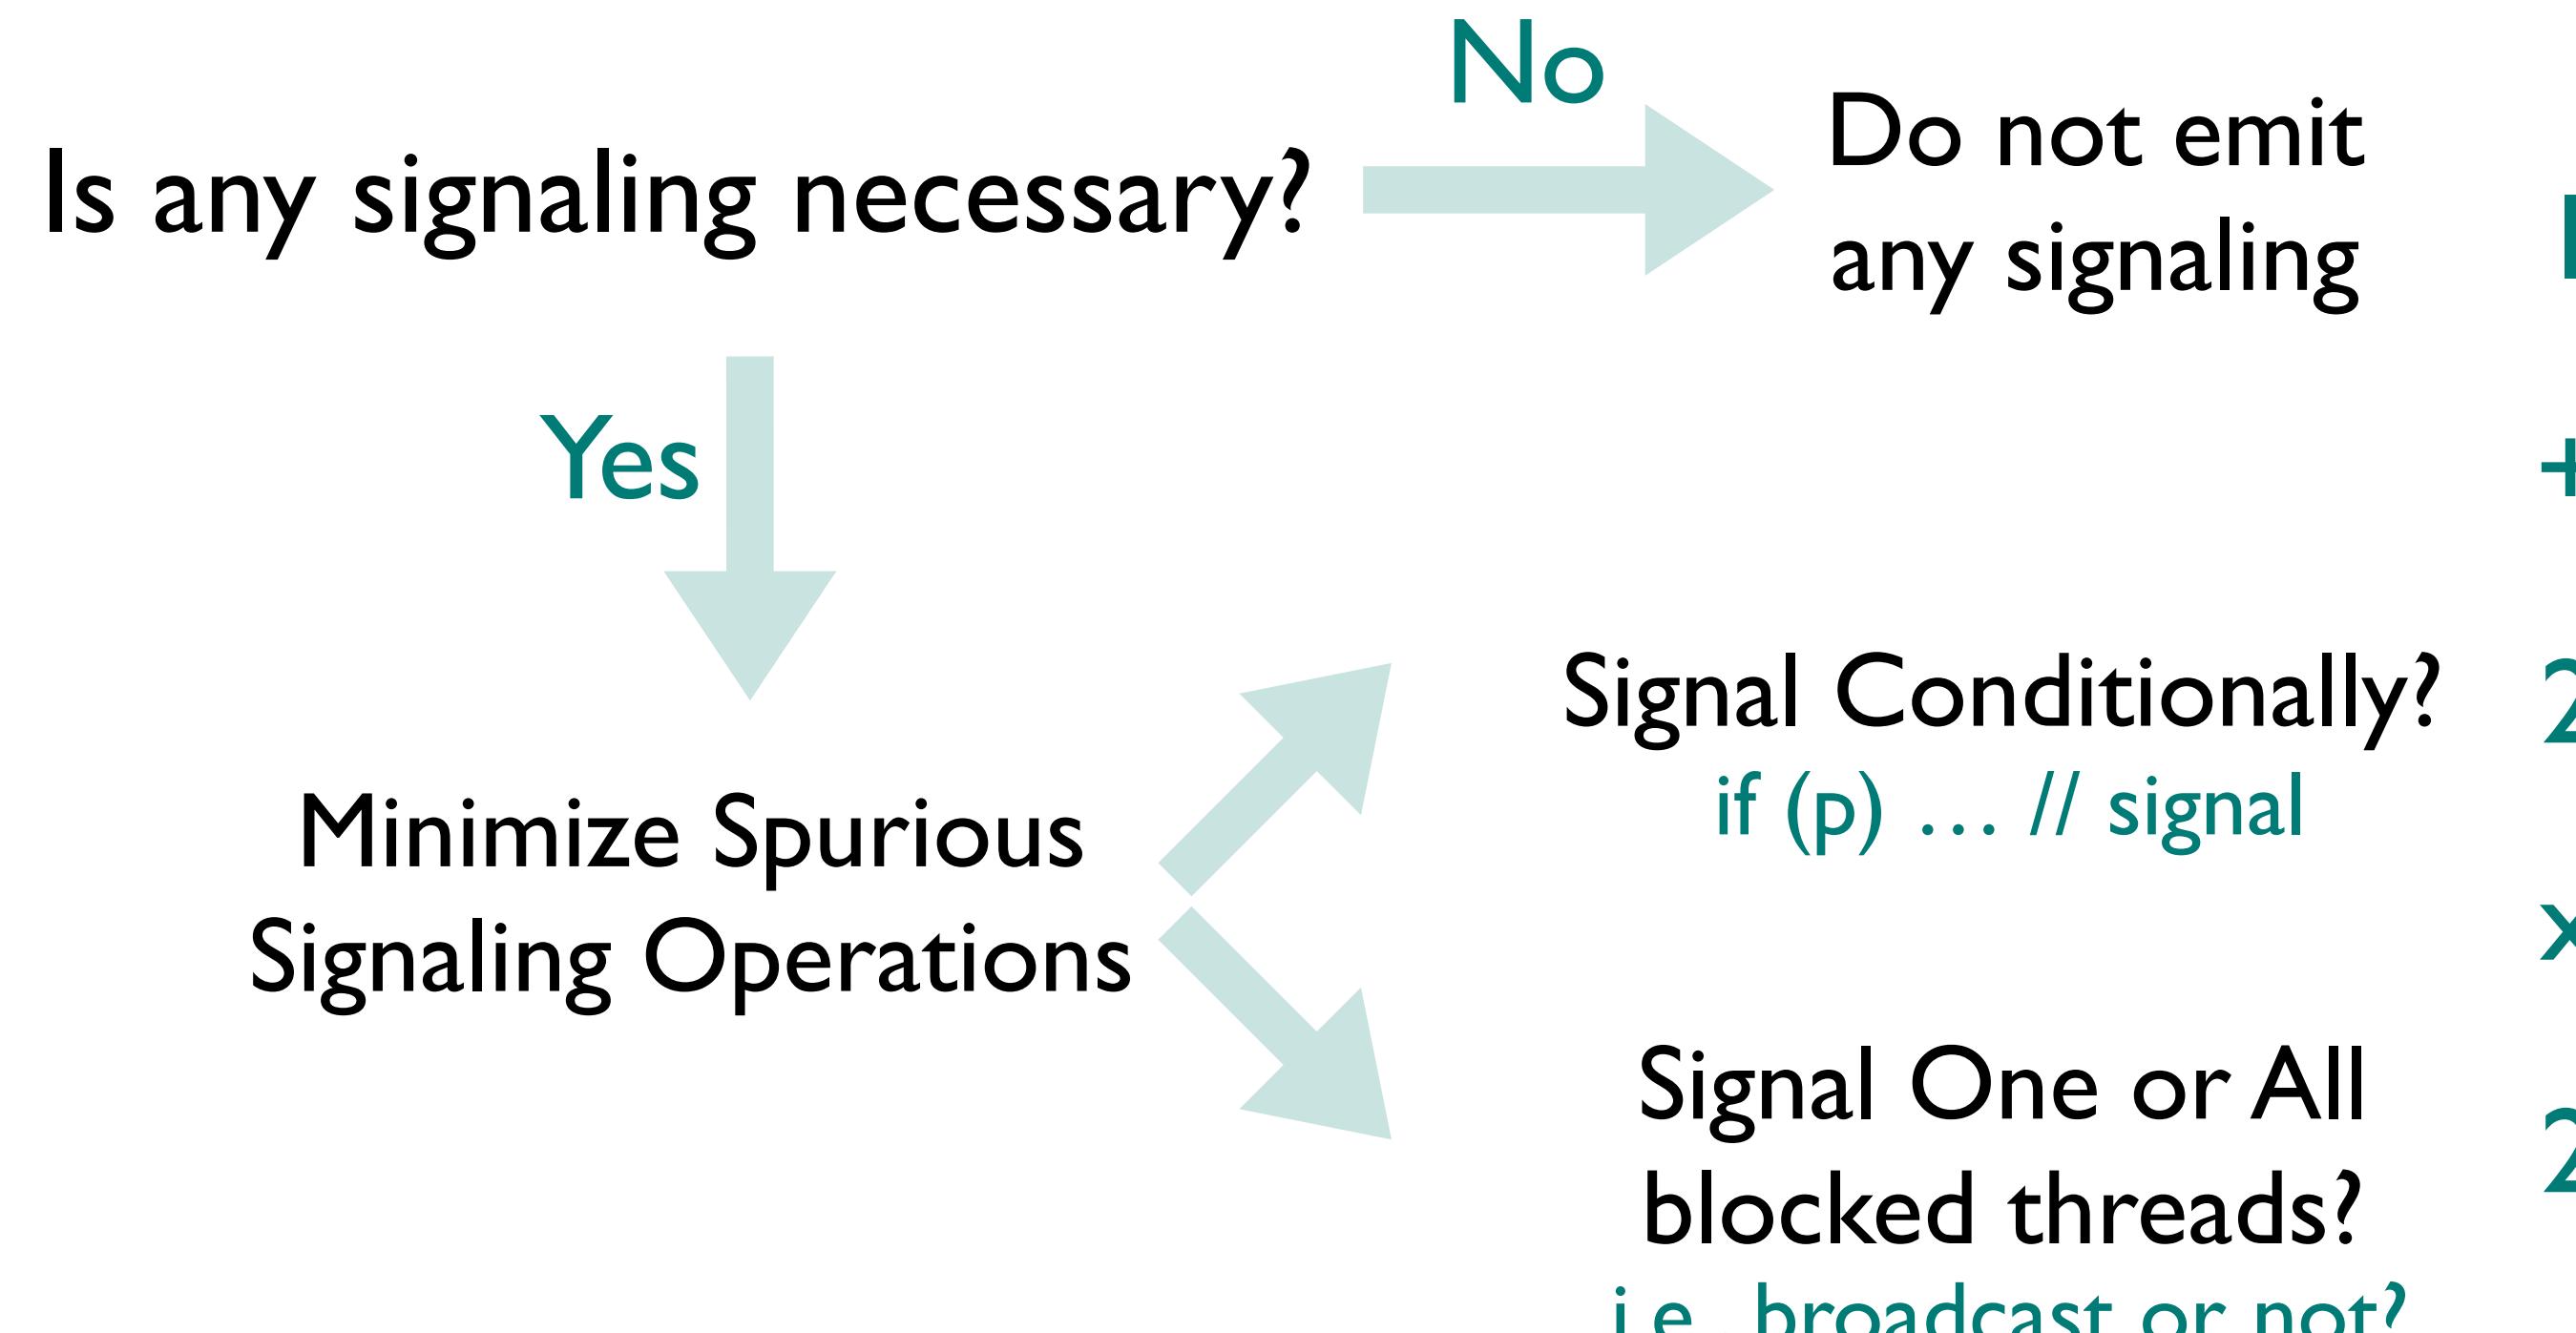
Signaling Decisions

For each CCR C and predicate $p \in Preds(M)$, decide whether to signal or not, and if so, how?



Signaling Decisions

For each CCR C and predicate $p \in Preds(M)$, decide whether to signal or not, and if so, how?



For each CCR and predicate, 5 possible outcomes

Running Example

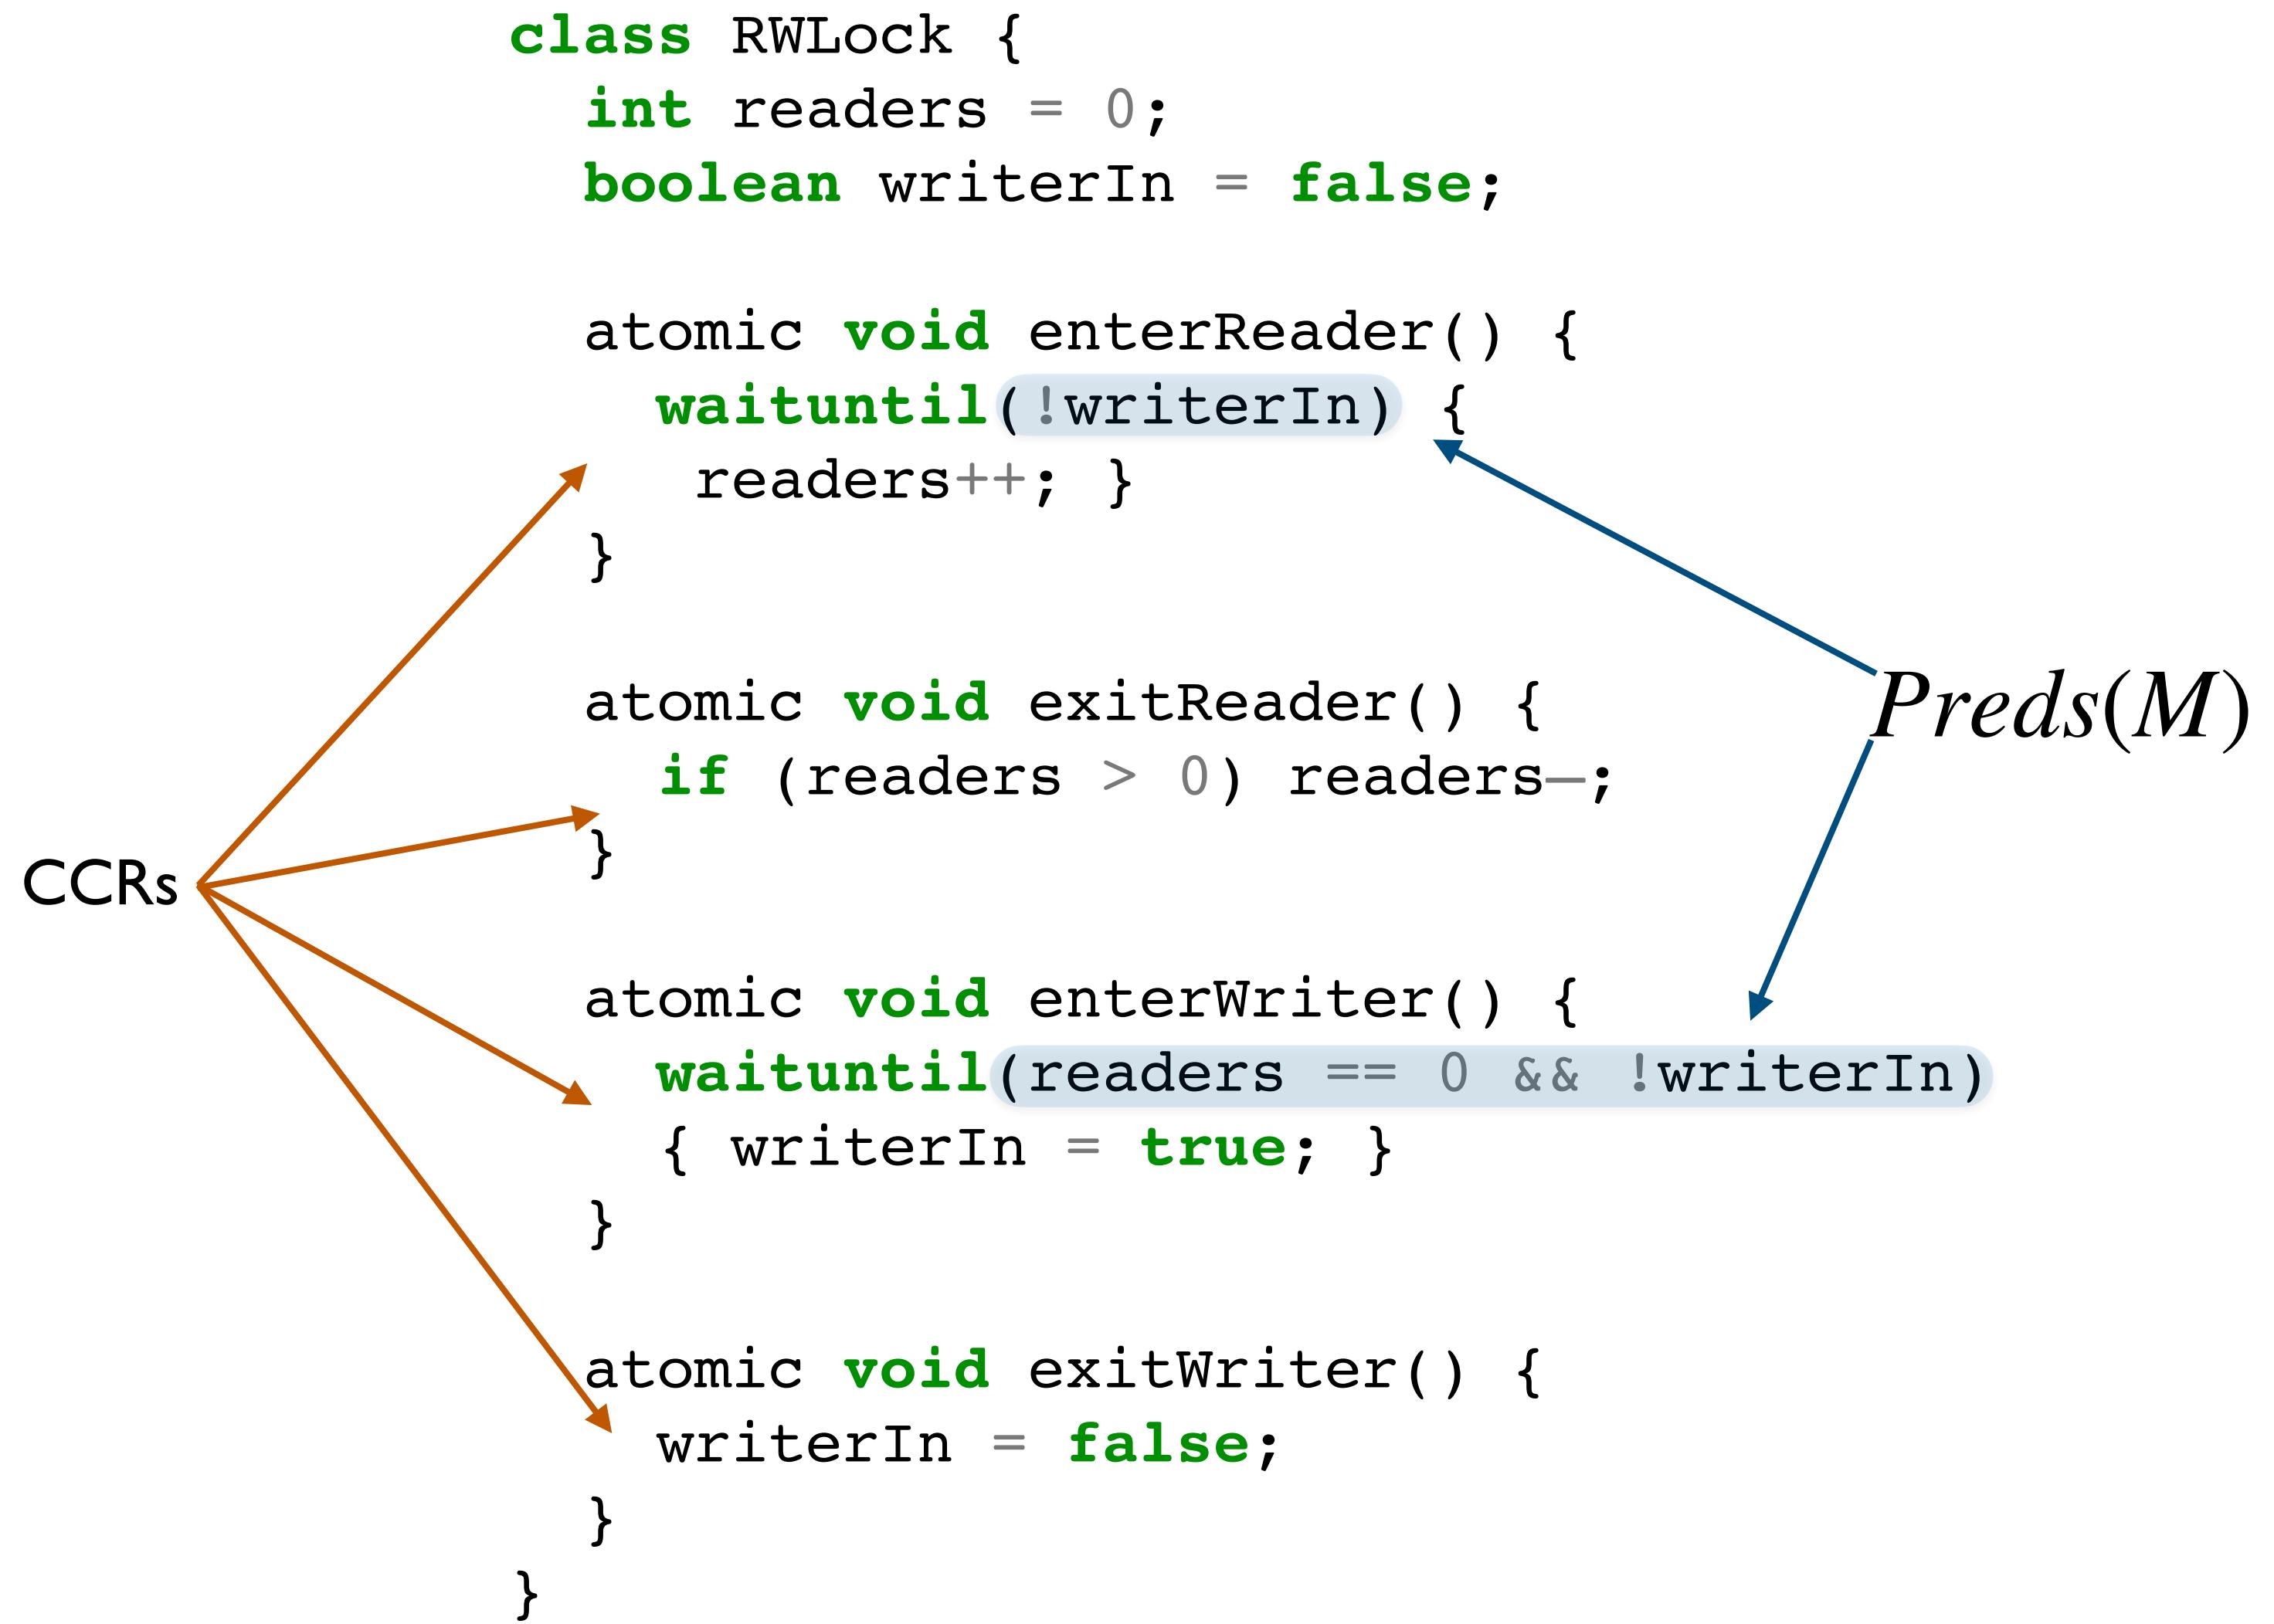
```
class RWLock {  
    int readers = 0;  
    boolean writerIn = false;  
  
    atomic void enterReader() {  
        waituntil(!writerIn) {  
            readers++; }  
    }  
  
    atomic void exitReader() {  
        if (readers > 0) readers--;  
    }  
  
    atomic void enterWriter() {  
        waituntil(readers == 0 && !writerIn)  
        { writerIn = true; }  
    }  
  
    atomic void exitWriter() {  
        writerIn = false;  
    }  
}
```

Running Example

```
class RWLock {  
    int readers = 0;  
    boolean writerIn = false;  
  
    atomic void enterReader() {  
        waituntil(!writerIn) {  
            readers++; }  
    }  
  
    atomic void exitReader() {  
        if (readers > 0) readers--;  
    }  
  
    atomic void enterWriter() {  
        waituntil(readers == 0 && !writerIn)  
        { writerIn = true; }  
    }  
  
    atomic void exitWriter() {  
        writerIn = false;  
    }  
}
```

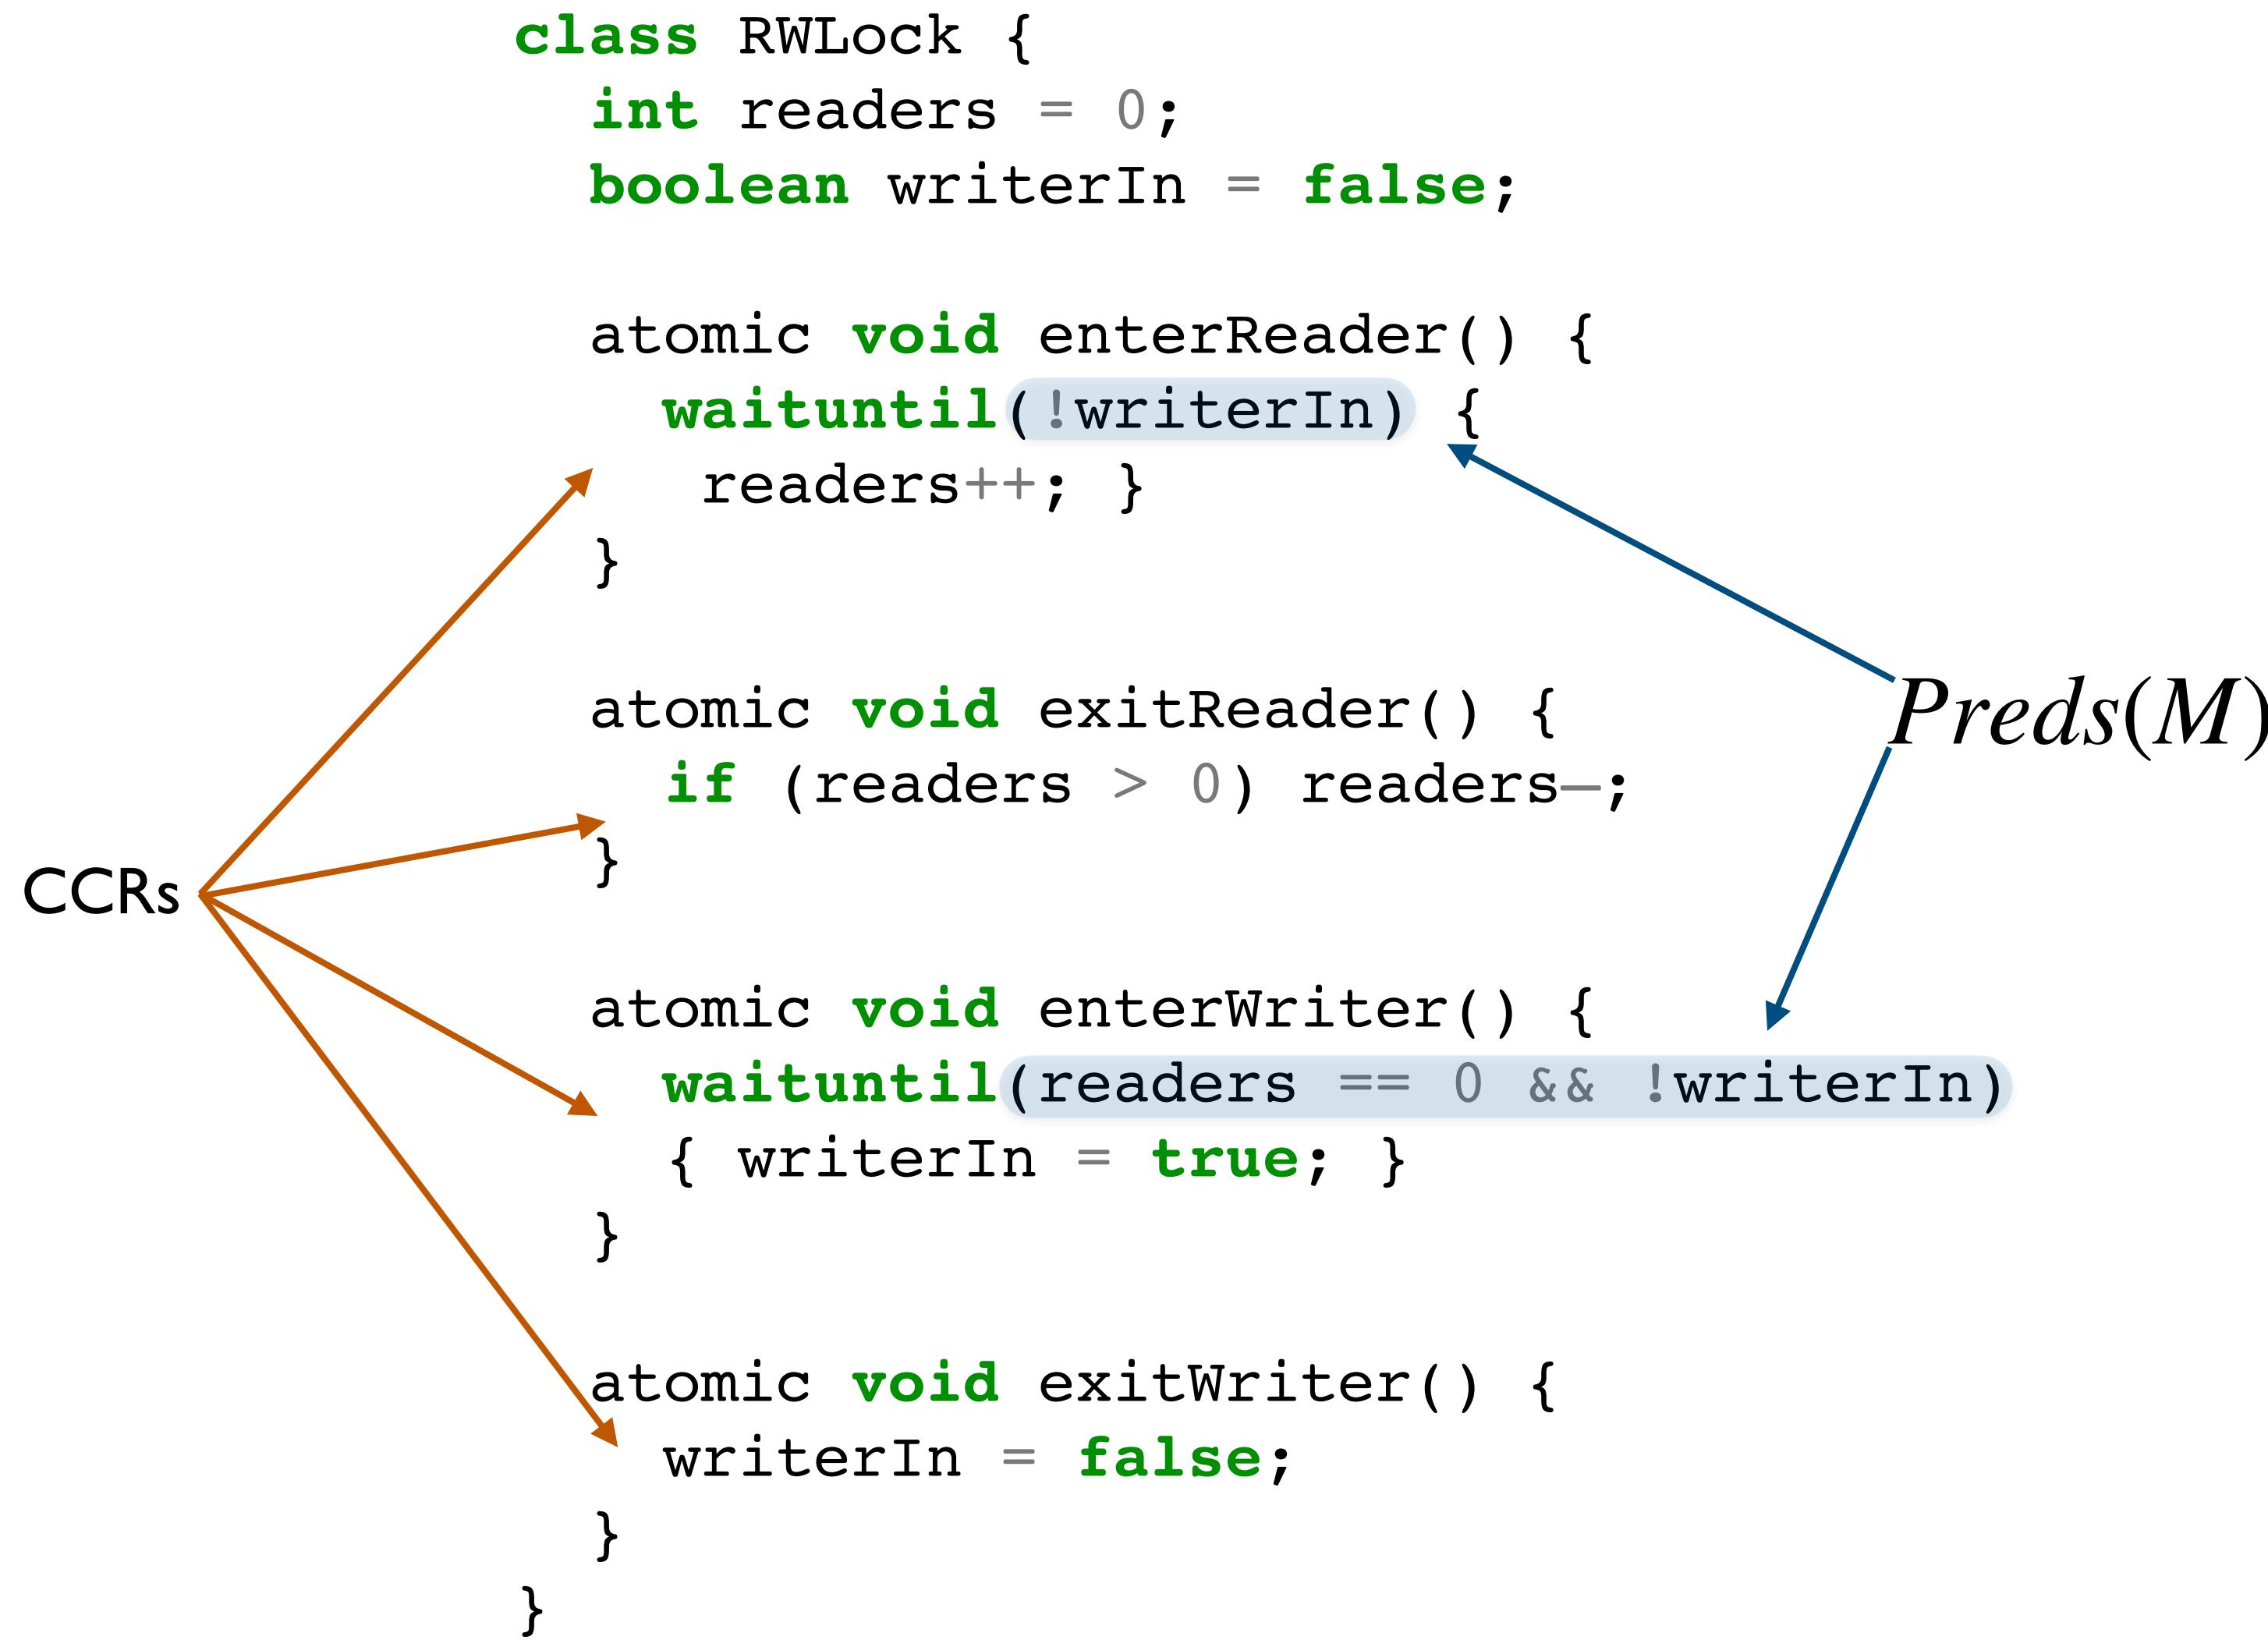
```
RWLock l = new ...();  
Readers:  
l.enterReader();  
... // Access Resource  
l.exitReader();  
  
Writers:  
l.enterWriter();  
... // Access Resource  
l.exitWriter();
```

Lots of Possible Implementations



Need to make signaling decisions for each combination of CCRs and predicates (8 pairs)

Lots of Possible Implementations



Need to make signaling decisions for each combination of CCRs and predicates (8 pairs)

So 5^8 different ways to implement explicit monitor! (ignoring locking)

Optimal Signaling for this Example

```
class RLock {
    int readers = 0;
    boolean writerIn = false;

    atomic void enterReader() {
        waituntil(!writerIn) {
            readers++;
        }
    }

    atomic void exitReader() {
        if (readers > 0) readers--;
    }

    atomic void enterWriter() {
        waituntil(readers == 0 && !writerIn) ← No signaling necessary
        { writerIn = true; }
    }

    atomic void exitWriter() {
        writerIn = false;
    }
}
```

Optimal Signaling for this Example

```
class RLock {
    int readers = 0;
    boolean writerIn = false;

    atomic void enterReader() {
        waituntil(!writerIn) {
            readers++;
        }
    }

    atomic void exitReader() {
        if (readers > 0) readers--;
    }

    atomic void enterWriter() {
        waituntil(readers == 0 && !writerIn) ← No signaling necessary
        { writerIn = true; }
    }

    atomic void exitWriter() {
        writerIn = false;
    }
}
```

Condition var representing `!writerIn`

`readC.signalAll();`

Unconditional broadcast to readers

Optimal Signaling for this Example

```
class RLock {
    int readers = 0;
    boolean writerIn = false;

    atomic void enterReader() {
        waituntil(!writerIn) {
            readers++;
        }
    }

    atomic void exitReader() {
        if (readers > 0) readers--;
    }

    atomic void enterWriter() {
        waituntil(readers == 0 && !writerIn)
        { writerIn = true; }
    }

    atomic void exitWriter() {
        writerIn = false;
    }
}
```

Conditionally signal a single writer:

```
if(readers == 0) writeC.signal();
```

No signaling necessary

Condition var representing !writerIn

```
readC.signalAll();
```

Unconditional broadcast to readers

Optimal Explicit Monitor

```
class RWLock {
    int readers = 0;
    boolean writerIn = false;

    Lock l = new ReentrantLock();
    Condition readC = l.newCondition(), writeC = l.newCondition();

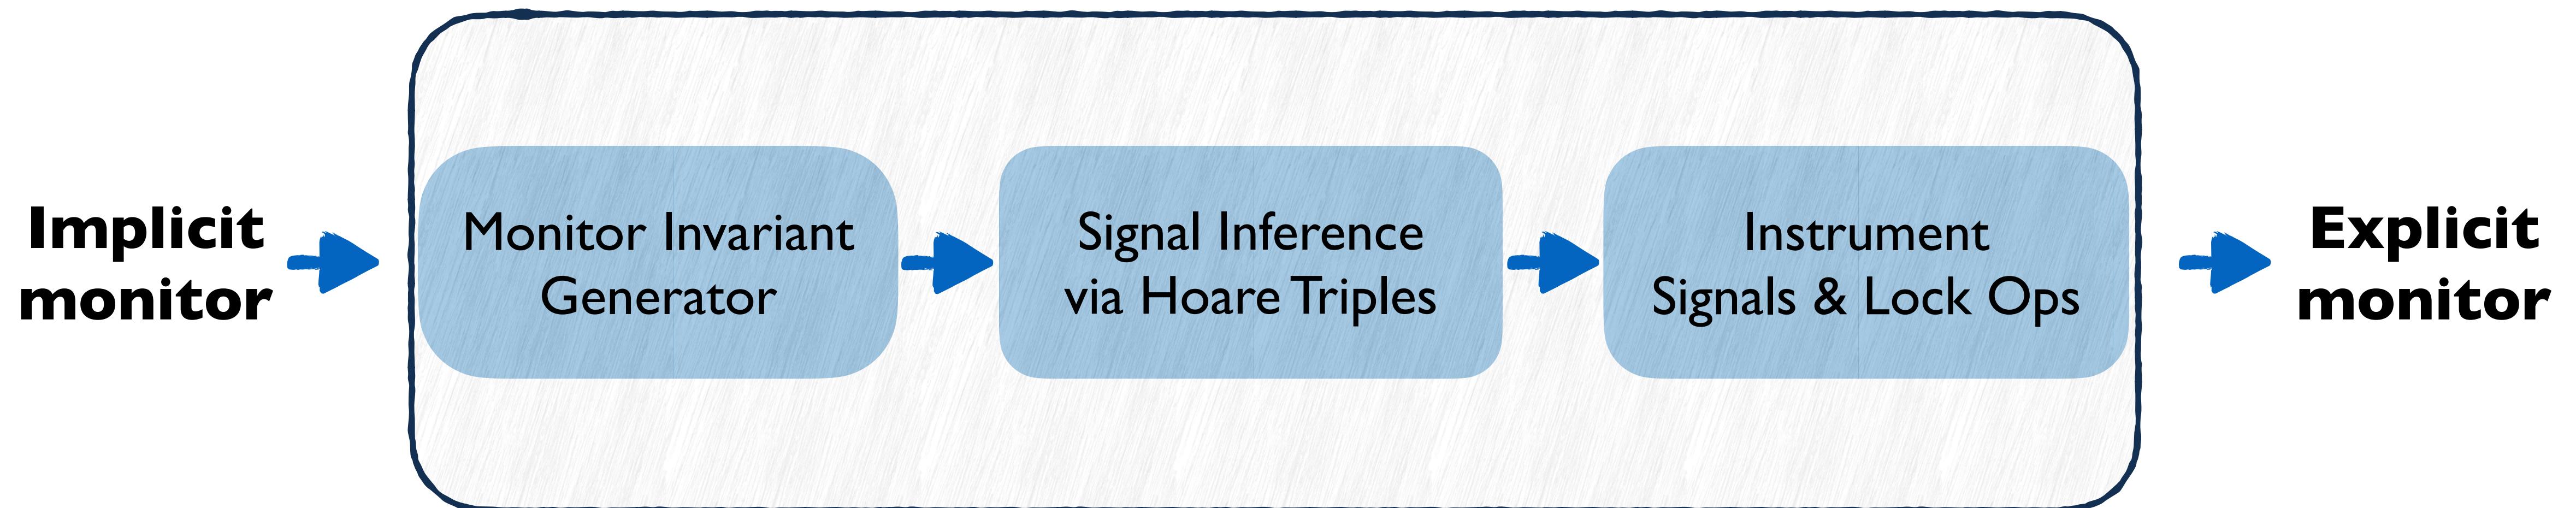
    void enterReader() {
        l.lock();
        while(writerIn) readC.await();
        readers++;
        l.unlock();
    }

    void exitReader() {
        l.lock();
        if (readers > 0) readers--;
        if (readers == 0 && !writerIn) writeC.signal();
        l.unlock();
    }

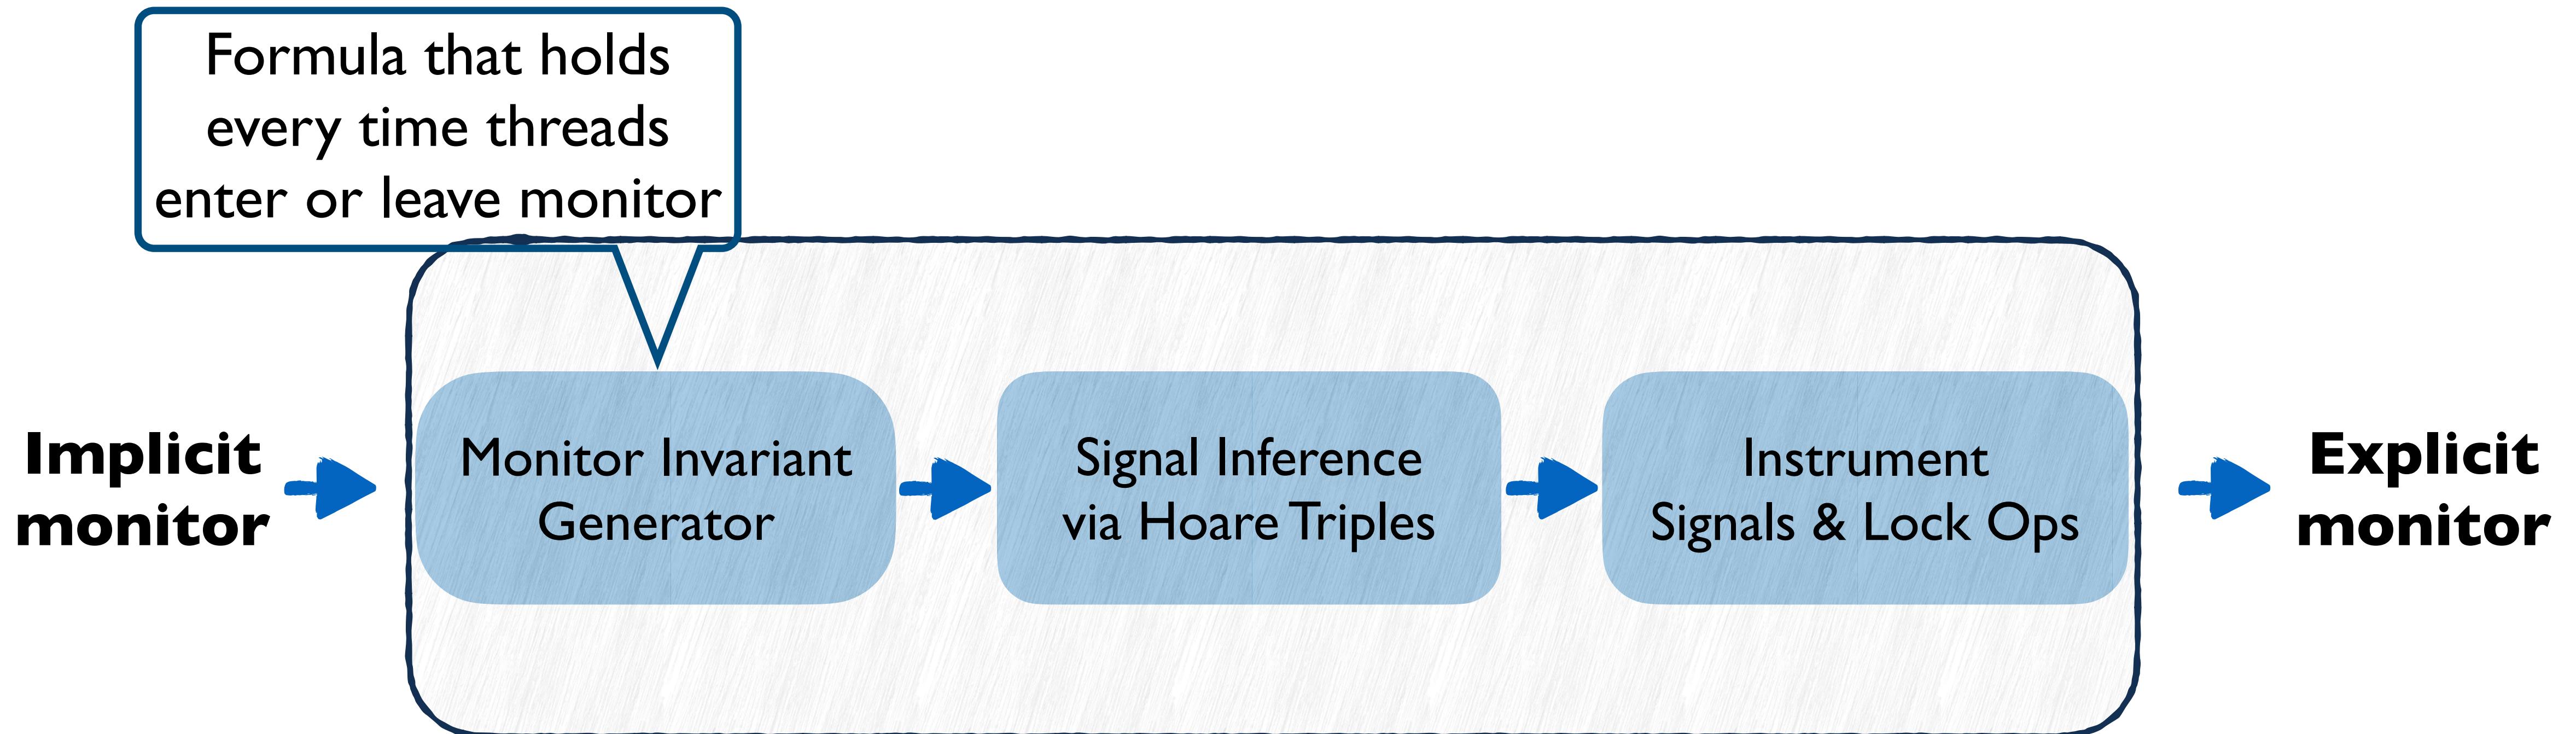
    void enterWriter() {
        l.lock();
        while(readers != 0 || writerIn) writeC.await();
        writerIn = true;
        l.unlock();
    }

    void exitWriter() {
        l.lock();
        writerIn = false;
        if (readers == 0 && !writerIn) writeC.signal();
        readC.signalAll();
        l.unlock();
    }
}
```

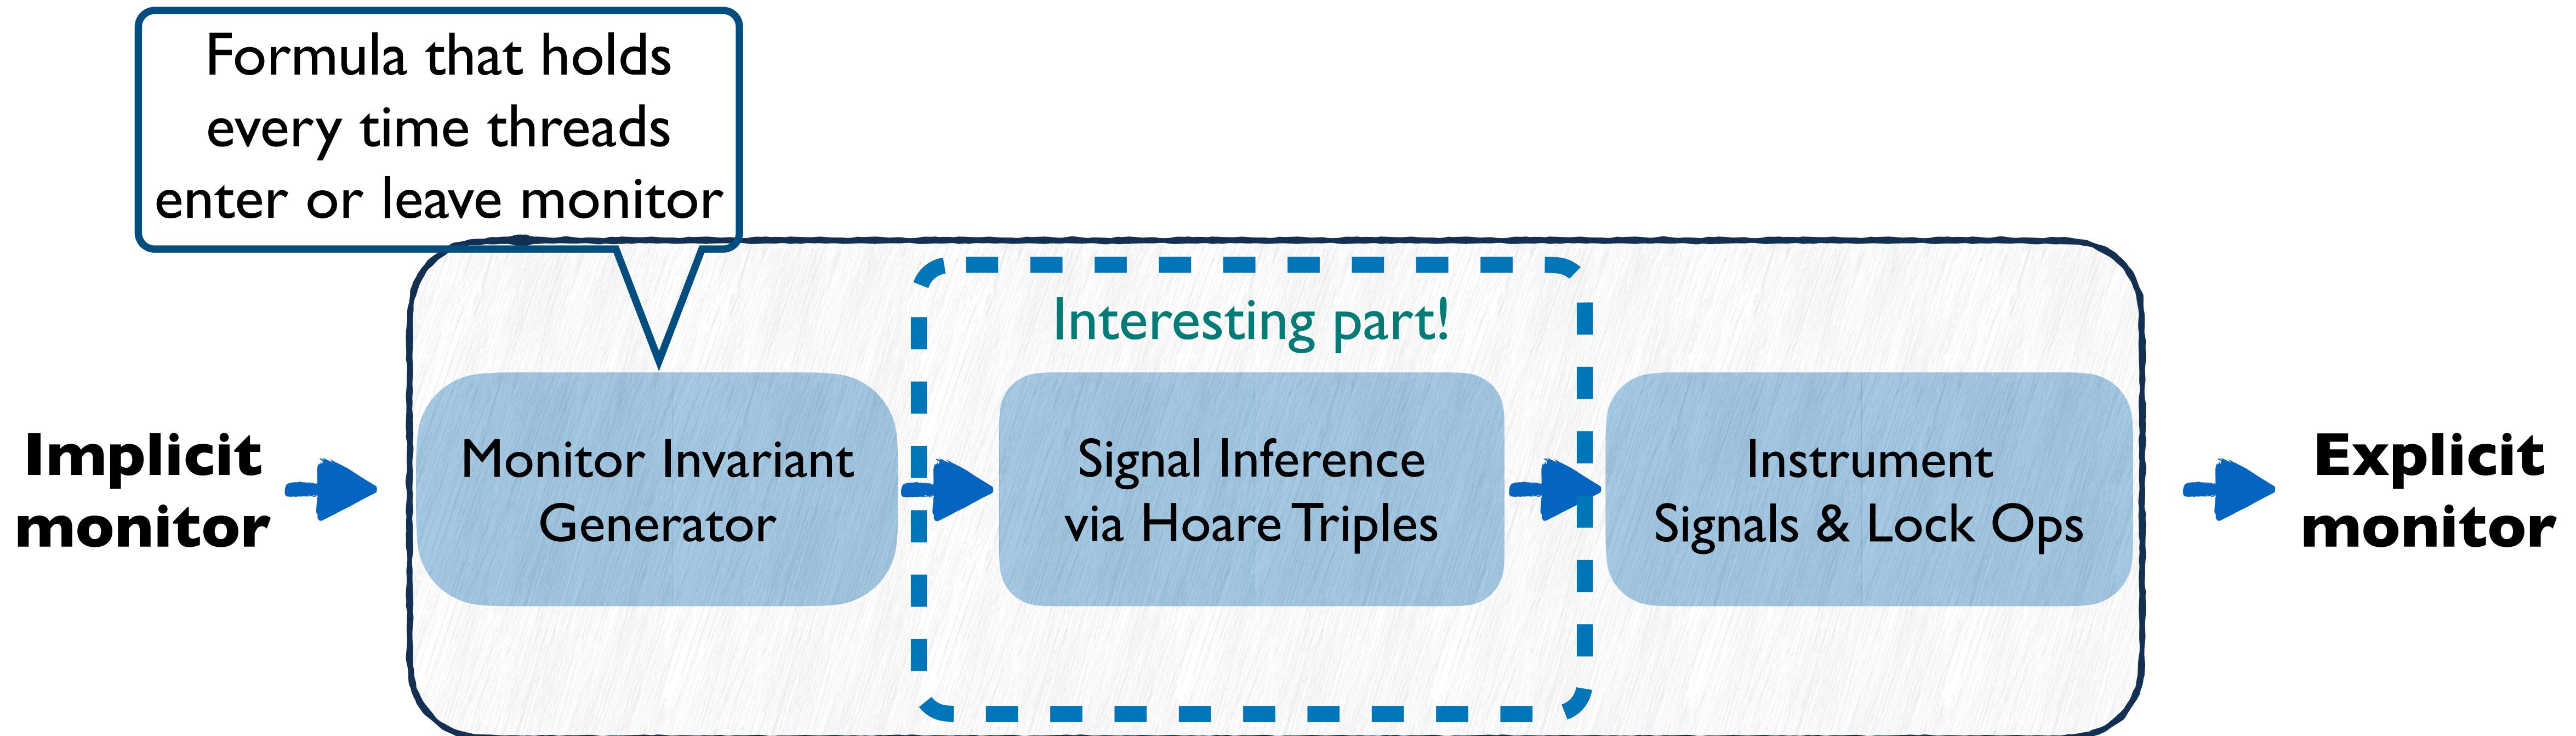
Synthesis Workflow



Synthesis Workflow



Synthesis Workflow



Primer on Hoare Triples

Hoare Triple:

$$\{ P \} S \{ Q \}$$

Primer on Hoare Triples

Hoare Triple:

$$\{ P \} S \{ Q \}$$

Is Q guaranteed to hold
after executing S ?

Primer on Hoare Triples

Starting in a state
where P holds

Hoare Triple:

$$\{ P \} S \{ Q \}$$

Is Q guaranteed to hold
after executing S?

Primer on Hoare Triples

Starting in a state
where P holds

Hoare Triple:

$$\{ P \} S \{ Q \}$$

Is Q guaranteed to hold
after executing S?

If yes, we call a triple valid

Primer on Hoare Triples

Starting in a state
where P holds

Hoare Triple:

$$\{ P \} S \{ Q \}$$

Is Q guaranteed to hold
after executing S?

If yes, we call a triple valid

Example of valid triple

Primer on Hoare Triples

Starting in a state
where P holds

Hoare Triple:

$$\{ P \} S \{ Q \}$$

Is Q guaranteed to hold
after executing S?

If yes, we call a triple valid

Example of valid triple

$$\{ count \geq 0 \}$$

Primer on Hoare Triples

Starting in a state
where P holds

Hoare Triple:

$$\{ P \} S \{ Q \}$$

Is Q guaranteed to hold
after executing S?

If yes, we call a triple valid

Example of valid triple

$$\{ count \geq 0 \} \quad count++;$$

Primer on Hoare Triples

Starting in a state
where P holds

Hoare Triple:

$$\{ P \} S \{ Q \}$$

Is Q guaranteed to hold
after executing S?

If yes, we call a triple valid

Example of valid triple

$$\{ count \geq 0 \} \quad \text{count++;} \quad \{ count \geq 0 \}$$

Signaling Decisions as Hoare Triples



To Signal or Not:

Given CCR C , do we need any signaling for threads blocked on predicate $P \in \text{Preds}(M)$?

$$\{Inv \wedge \neg P\} \subset \{\neg P\}$$

Signaling Decisions as Hoare Triples



To Signal or Not:

Given CCR C , do we need any signaling for threads blocked on predicate $P \in \text{Preds}(M)$?

If this triple is valid, **no need** to **signal** threads blocked on P!

$$\{Inv \wedge \neg P\} C \{ \neg P\}$$

Example: Signal or Not?

CCR: atomic **void** enterReader() {
 waituntil(!writerIn){
 readers++;
 }
 }
 }

Need to signal
writers after
enterReader?

Example: Signal or Not?

Predicate: $P_w \equiv readers = 0 \wedge \neg writerIn$ $\neg P_w \equiv readers \neq 0 \vee writerIn$

CCR: atomic **void** enterReader() {
 waituntil(!writerIn) {
 readers++;
 }
}

Need to signal
writers after
enterReader?

Example: Signal or Not?

Predicate: $P_w \equiv \text{readers} = 0 \wedge \neg \text{writerIn}$ $\neg P_w \equiv \text{readers} \neq 0 \vee \text{writerIn}$ **Invariant:** $\text{readers} \geq 0$

CCR:

```
atomic void enterReader() {
    waituntil(!writerIn) {
        readers++;
    }
}
```

Need to signal
writers after
enterReader?

Example: Signal or Not?

Predicate: $P_w \equiv \text{readers} = 0 \wedge \neg \text{writerIn}$ $\neg P_w \equiv \text{readers} \neq 0 \vee \text{writerIn}$ **Invariant:** $\text{readers} \geq 0$

CCR:

```
atomic void enterReader() {
    {readers ≥ 0 ∧ ¬Pw}
    waituntil(!writerIn) {
        readers++;
    }
    {¬Pw}
}
```

Need to signal
writers after
enterReader?

Example: Signal or Not?

Predicate: $P_w \equiv \text{readers} = 0 \wedge \neg \text{writerIn}$ $\neg P_w \equiv \text{readers} \neq 0 \vee \text{writerIn}$ **Invariant:** $\text{readers} \geq 0$

CCR:

```
atomic void enterReader() {
    {readers ≥ 0 ∧ ¬Pw}
    waituntil(!writerIn) {
        readers++;
    }
    {¬Pw}
}
```

Need to signal
writers after
enterReader?

Triple is valid, no need to signal writers!

Summary of Signaling



Key Idea: Reduce signaling decisions to validity of Hoare triples

- Once the signaling decisions are made, straightforward to instrument code with operations on condition vars
- To ensure atomicity and data-race-freedom, use one global lock acquired/released before/after each operation

Implementation & Evaluation



- Evaluated on several monitors from popular GitHub repos.
 - Real-world applications: *Spring*, *Gradle*, etc.
 - Efficient explicit-signal implementations.
- Converted explicit monitor to implicit and pass it to Espresso



Evaluation, cont

Compared **EXPRESSO** against:

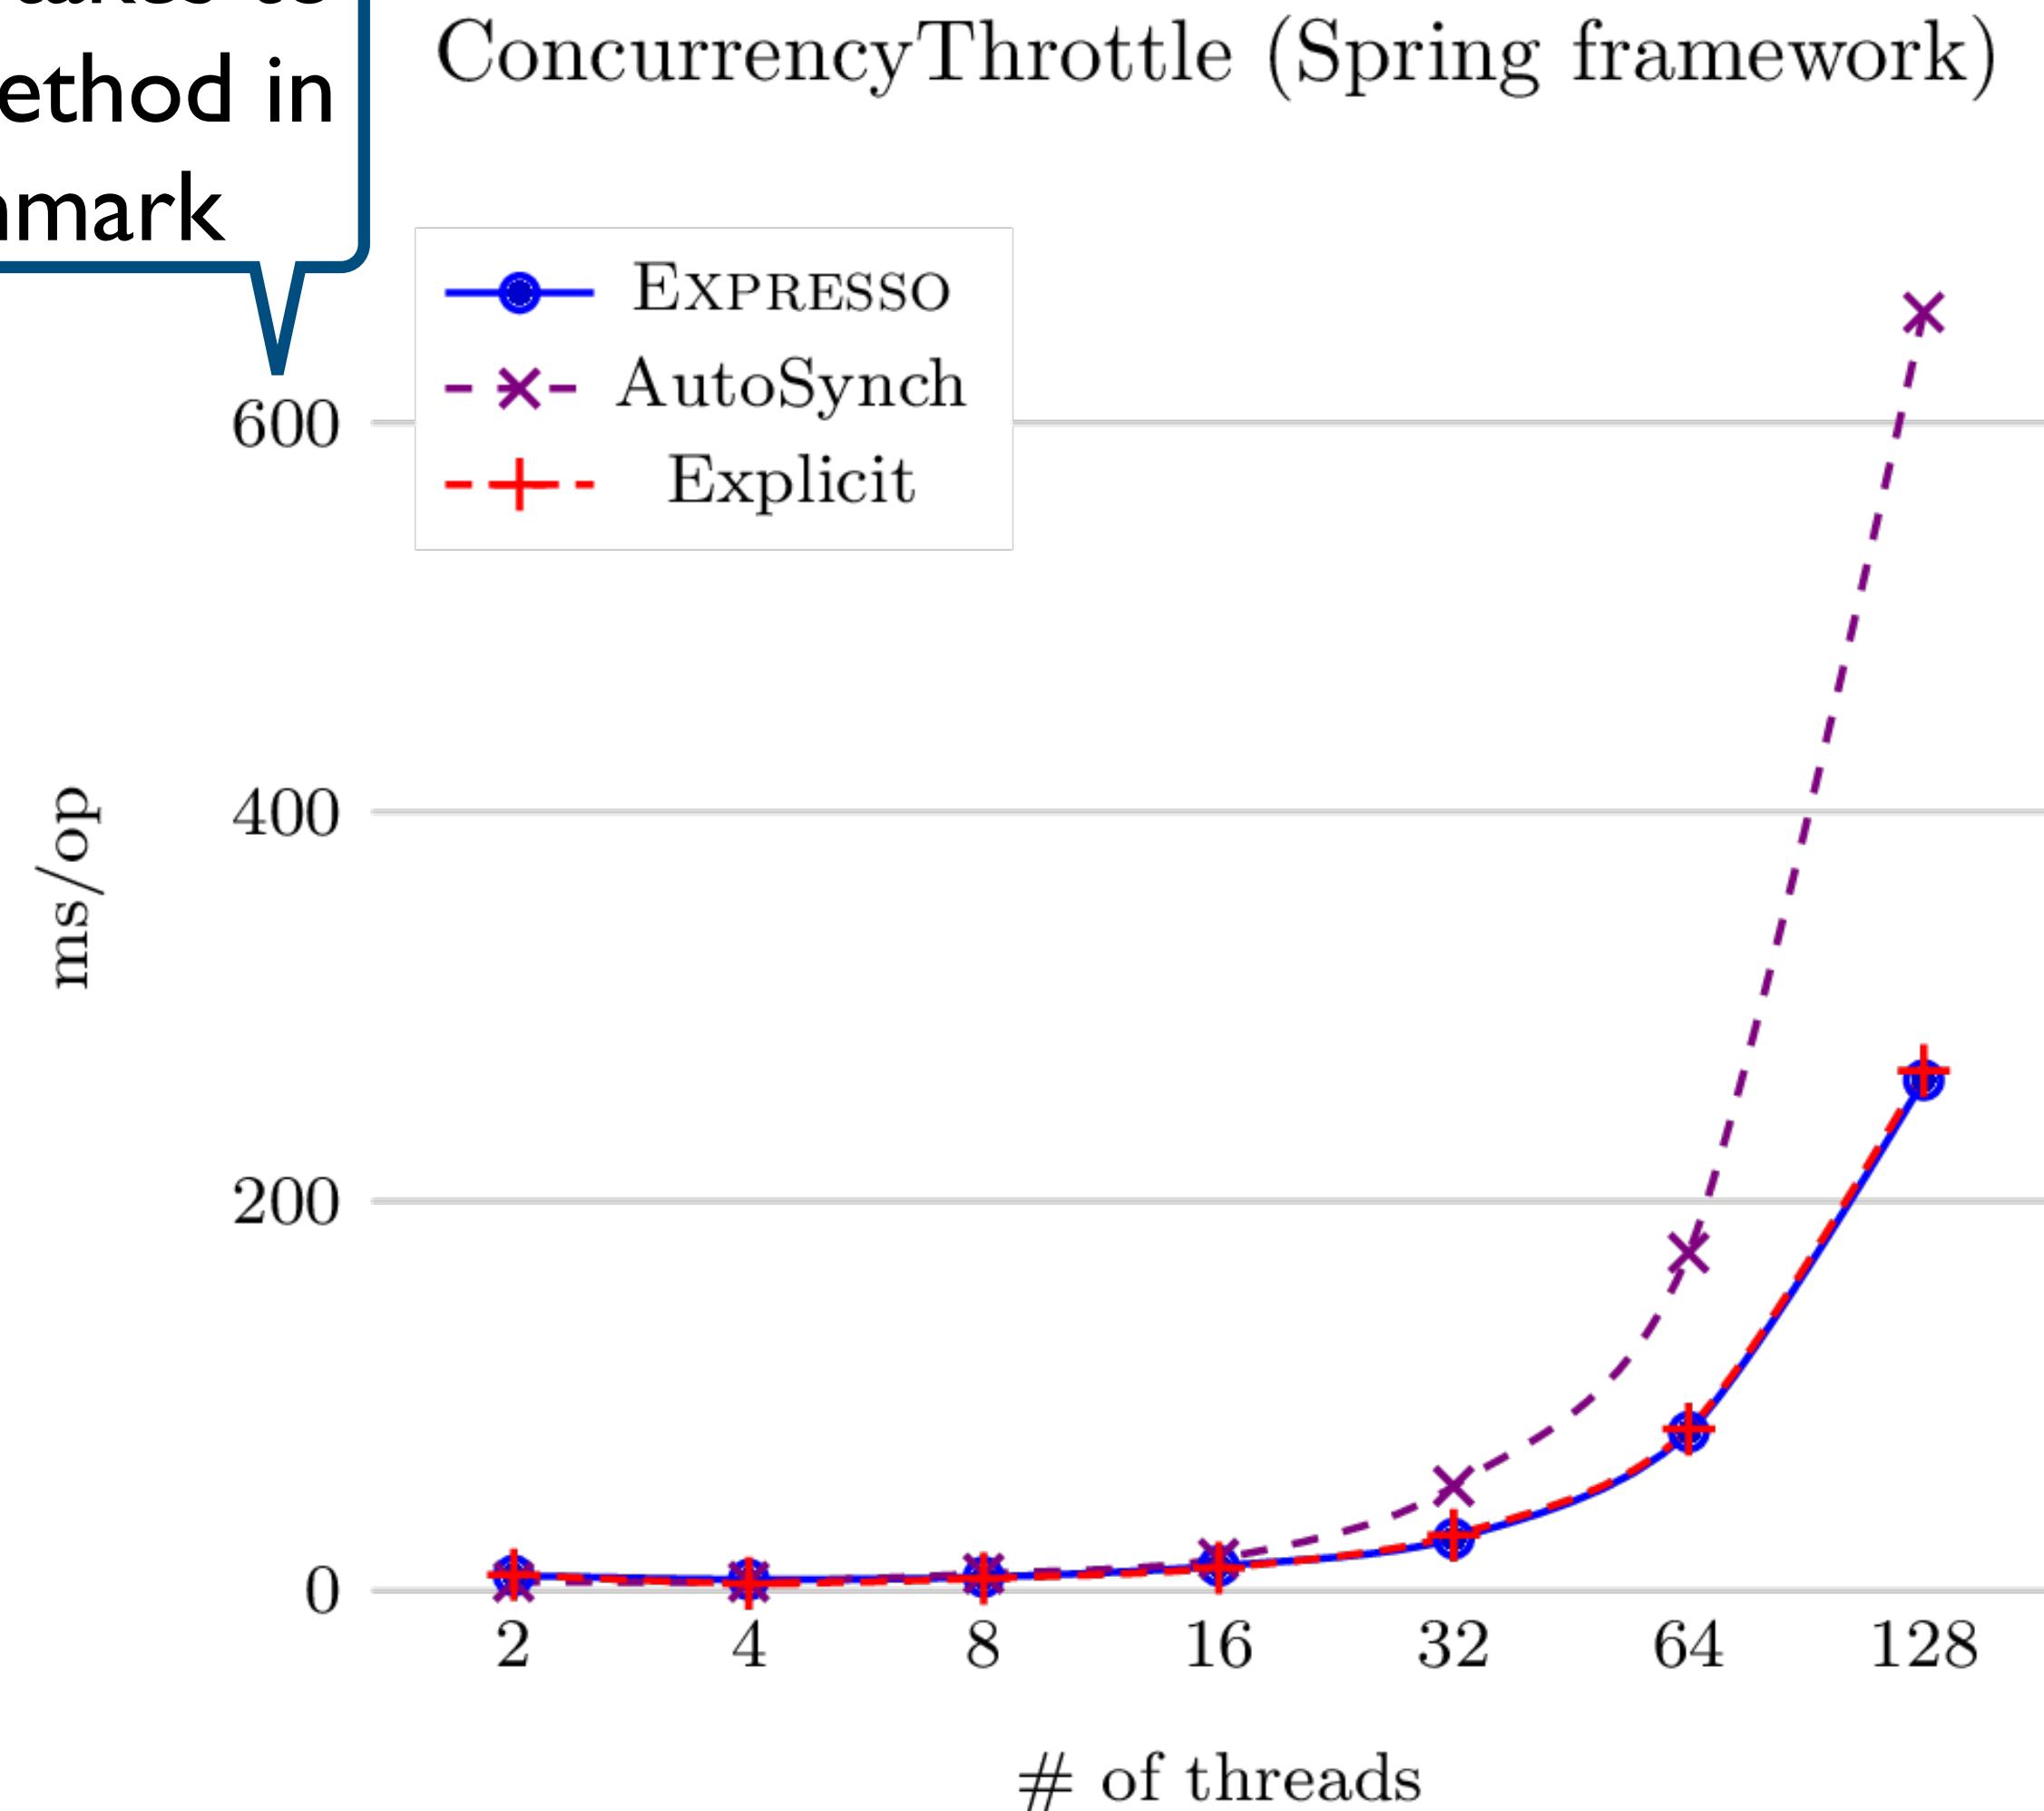
- AutoSynch [I]: **state-of-the-art** runtime system for implicit monitors.
- Explicit: original implementation by expert programmer.



[I] Wei-Lun Hung et al. **AutoSynch**: An Automatic signal Monitor Based on Predicate Tagging **PLDI'13**

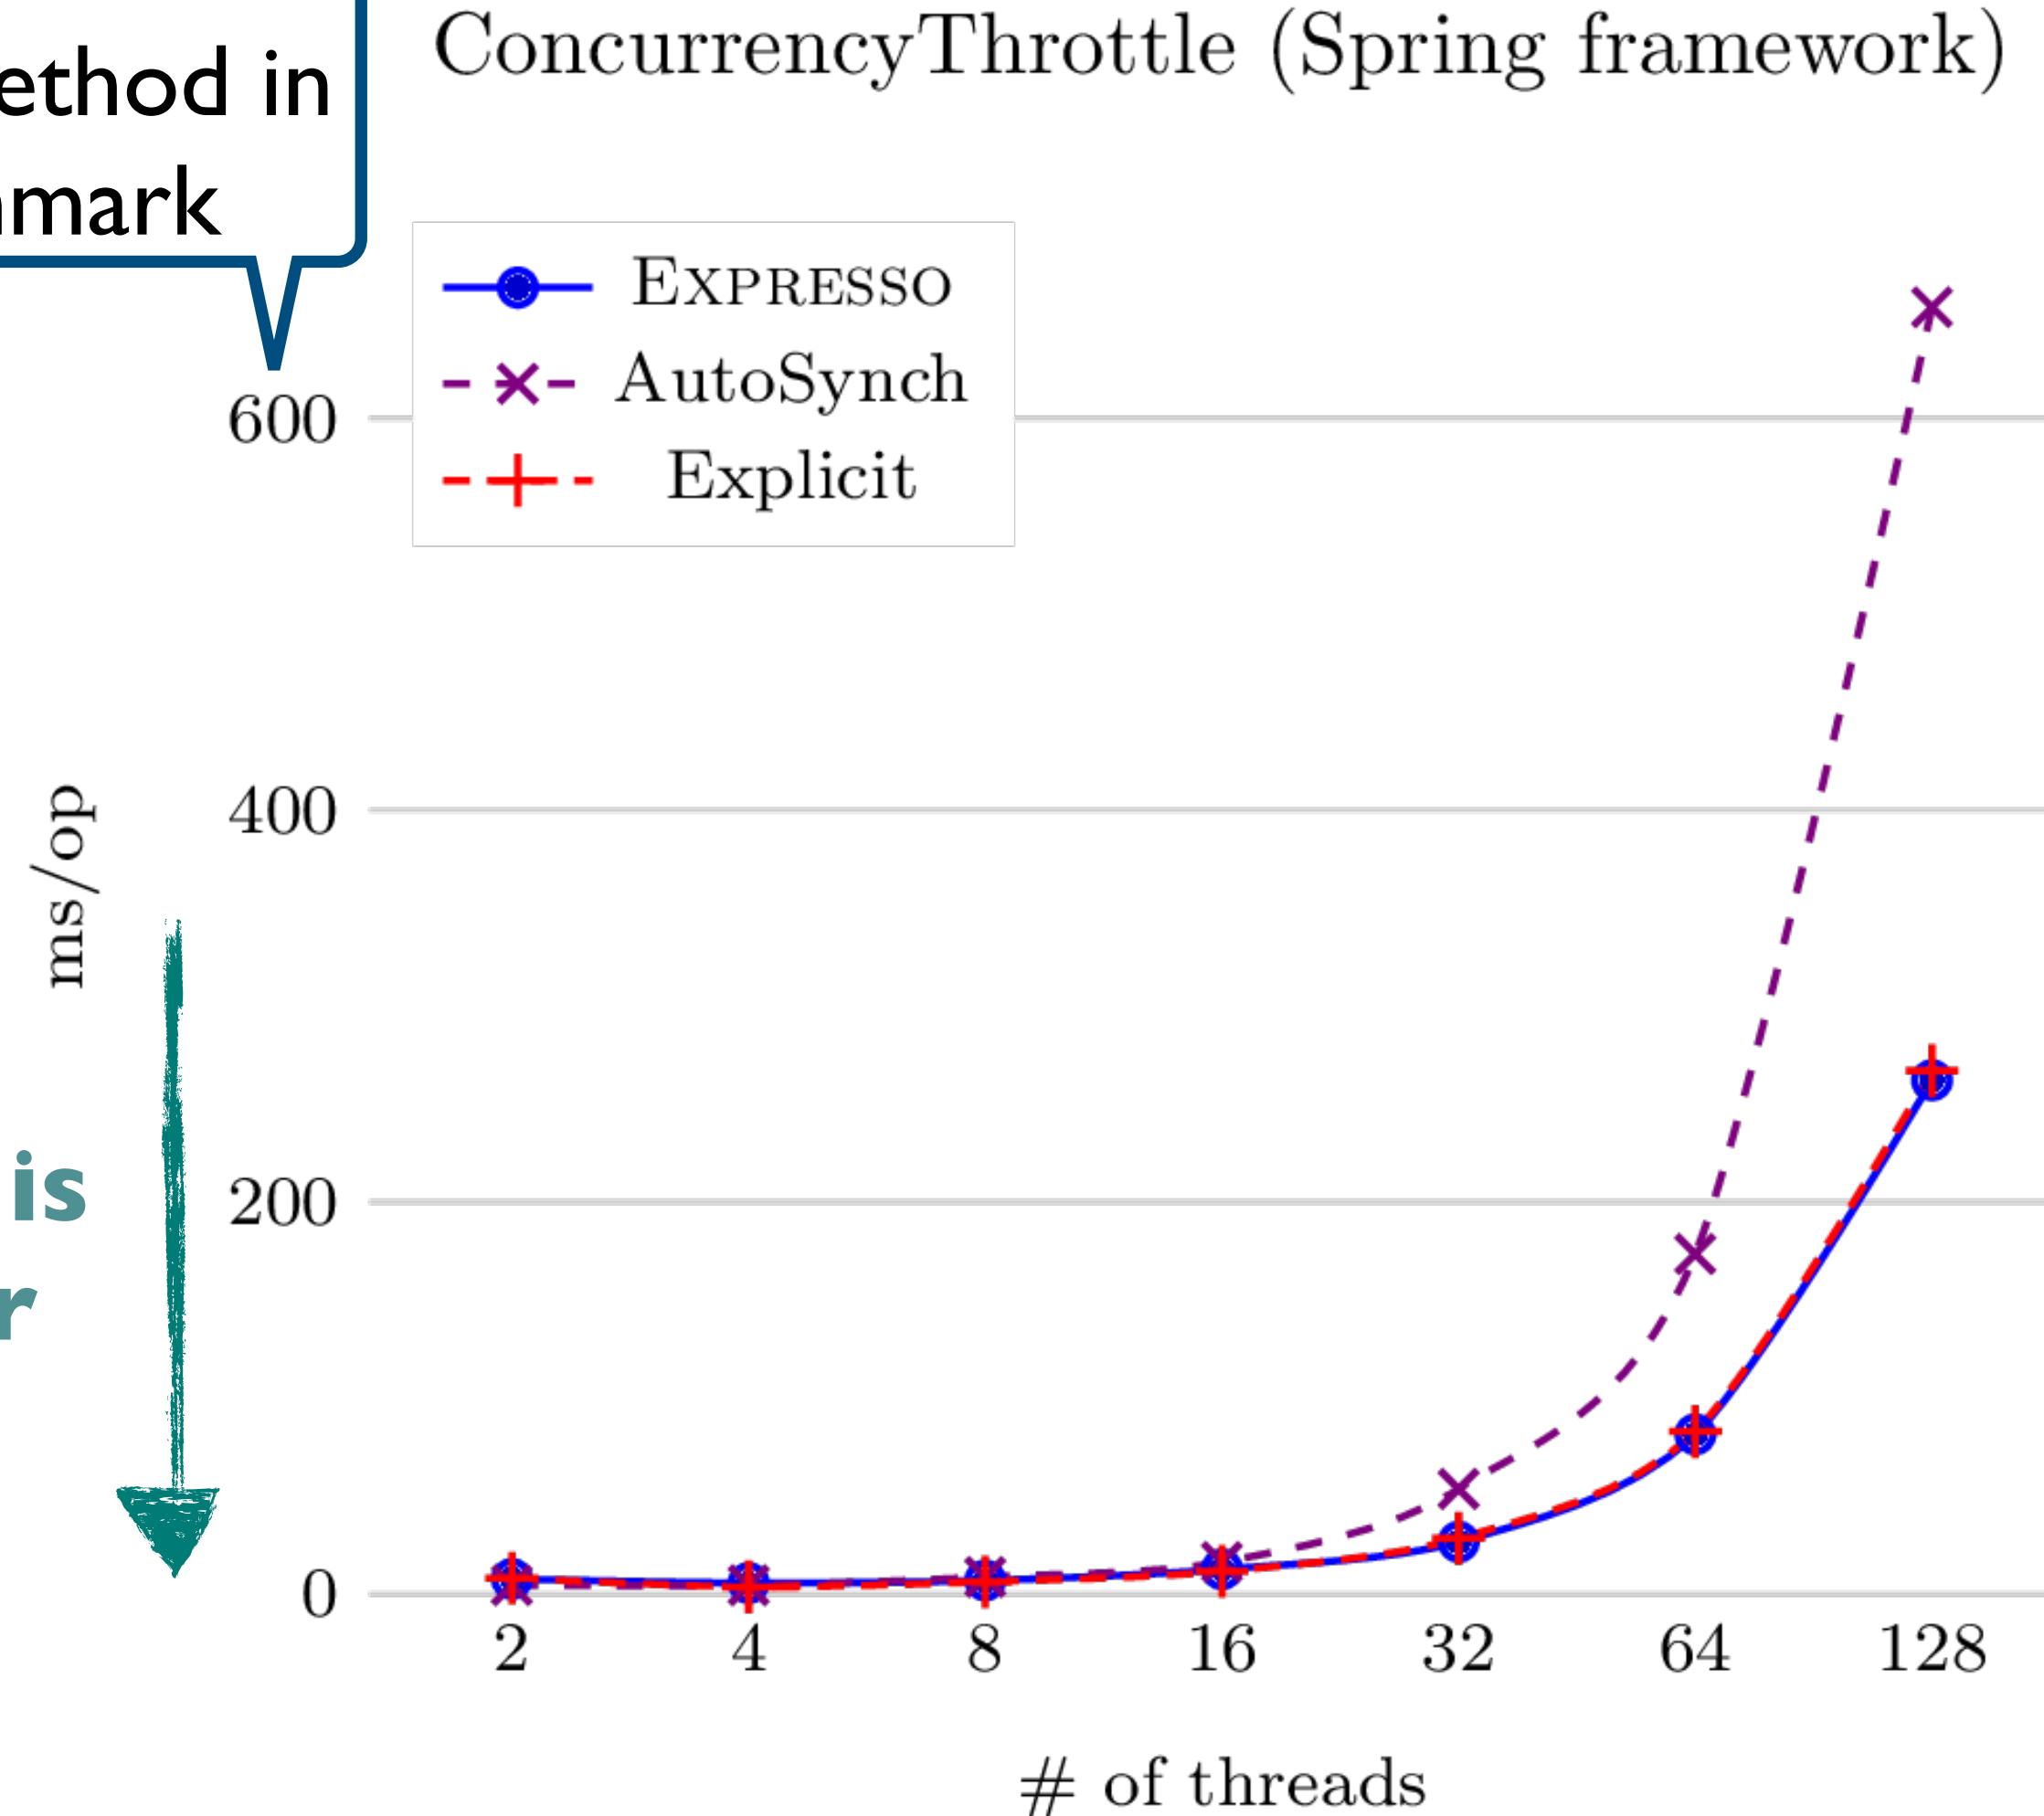
Results

How long it takes to execute a method in the benchmark



Results

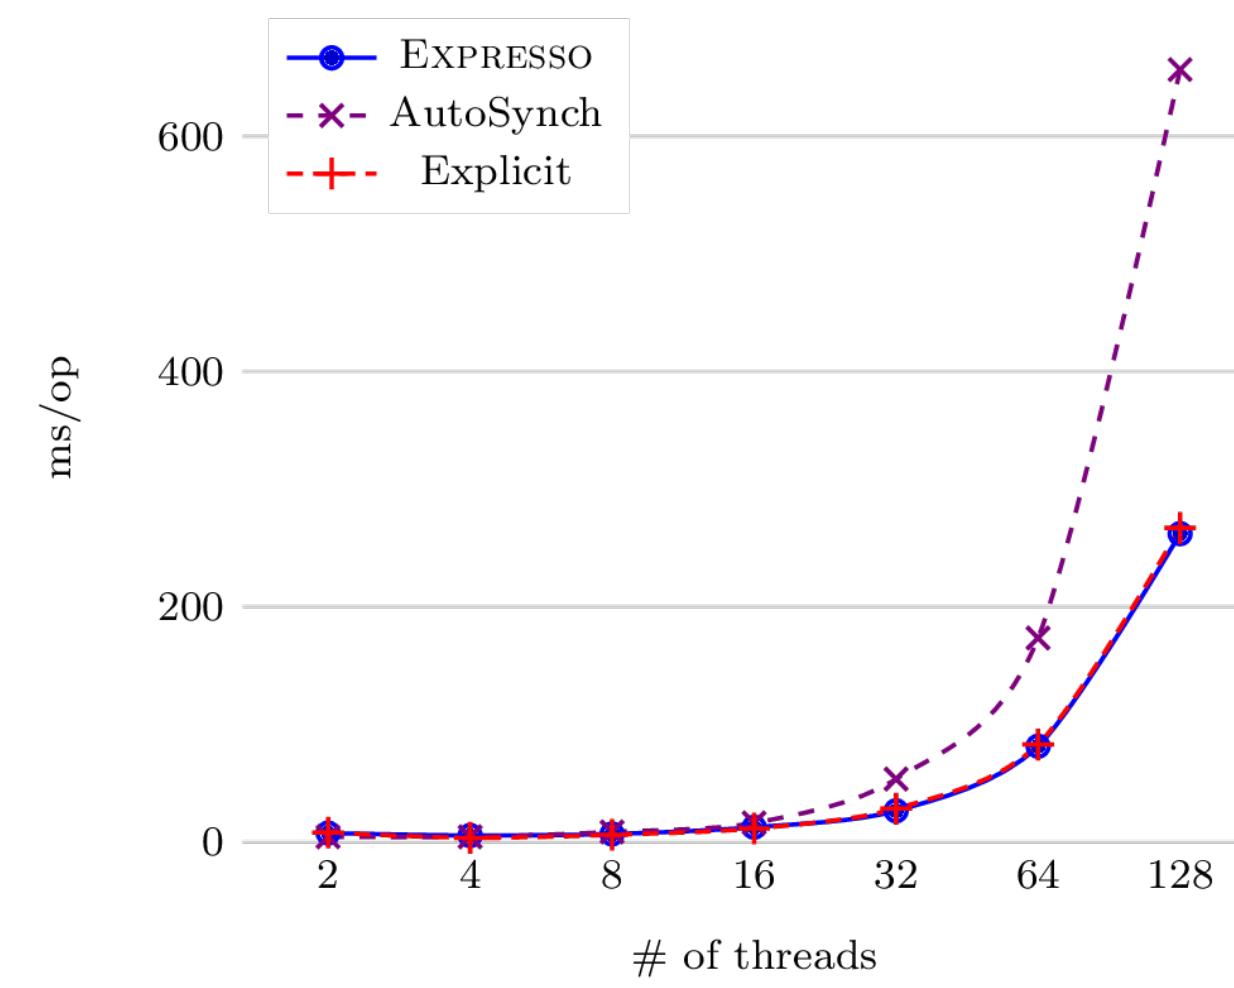
How long it takes to execute a method in the benchmark



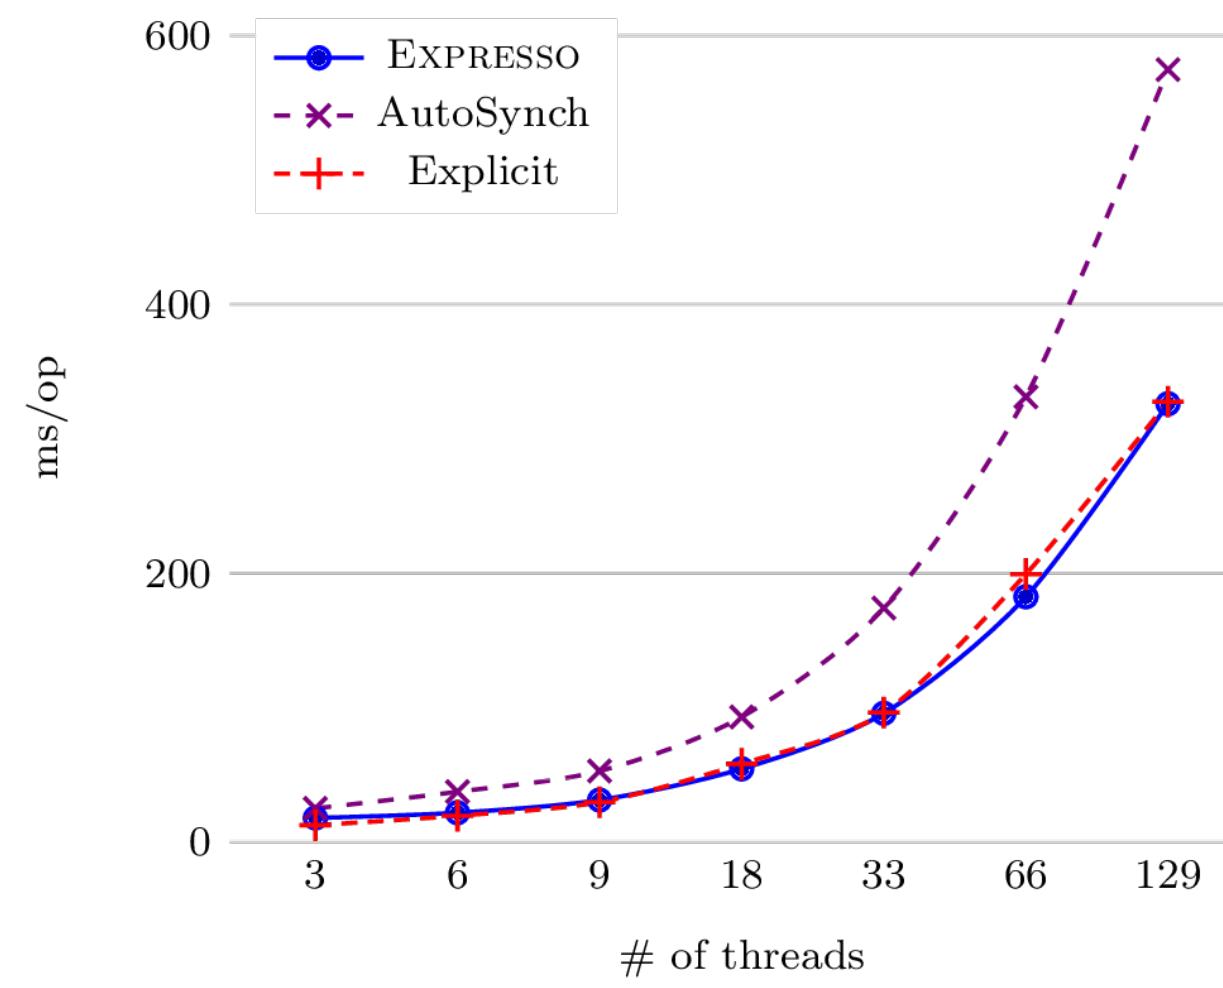
Lower is better

Results, cont

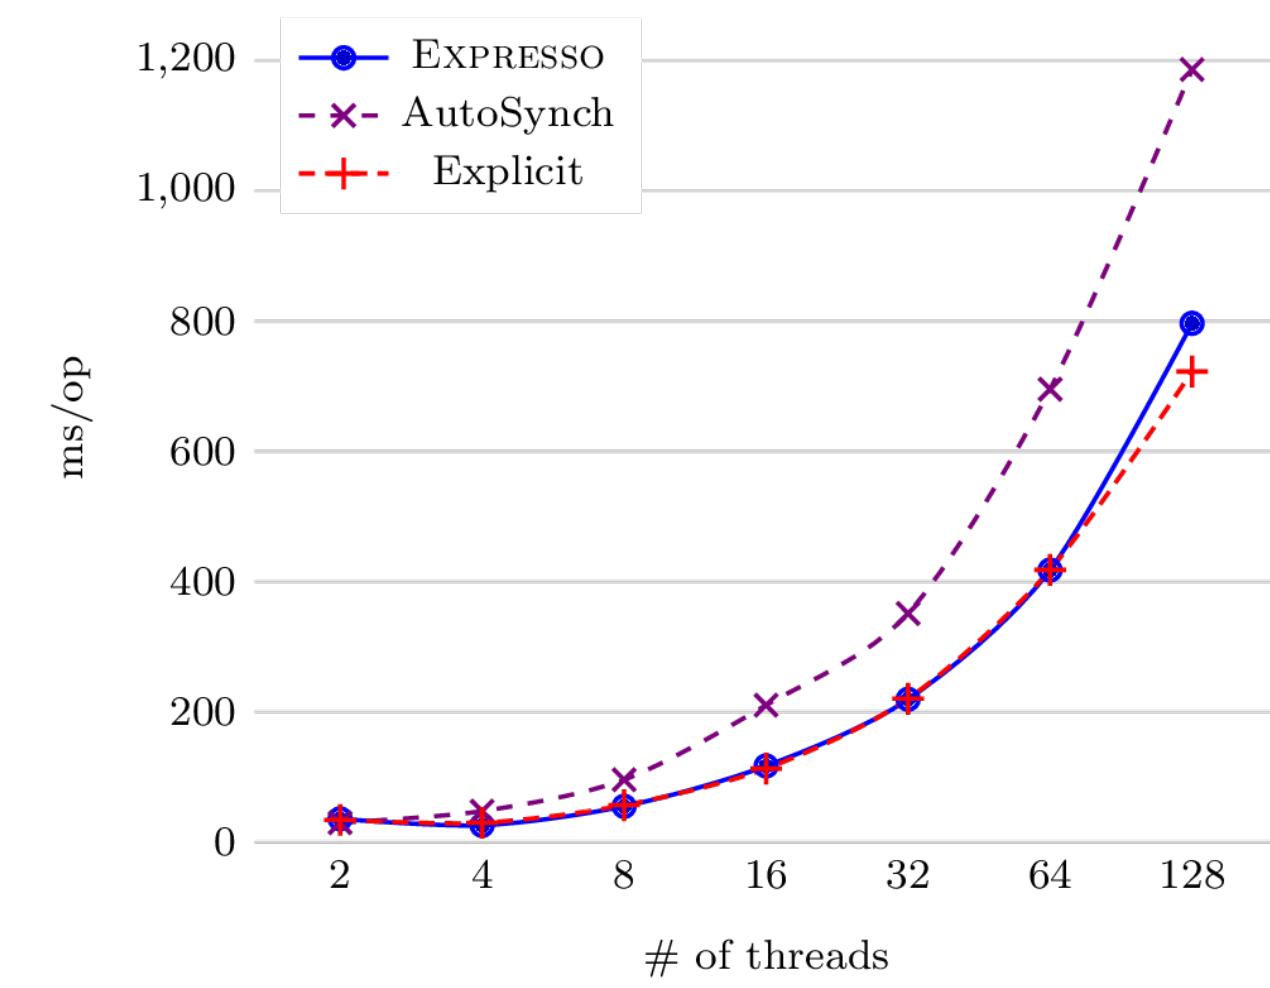
ConcurrencyThrottle (Spring framework)



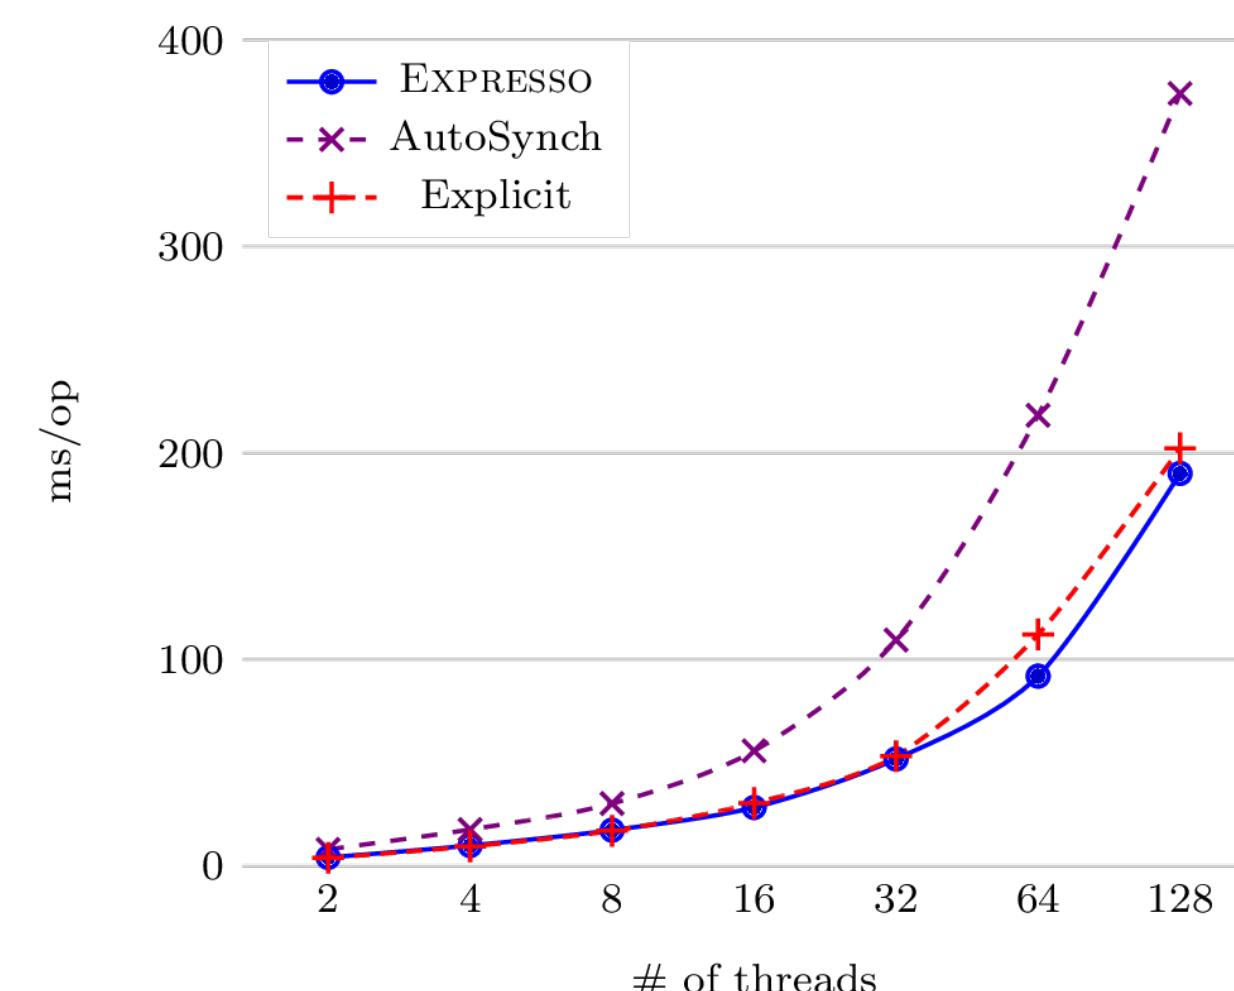
PendingPostQueue (EventBus)



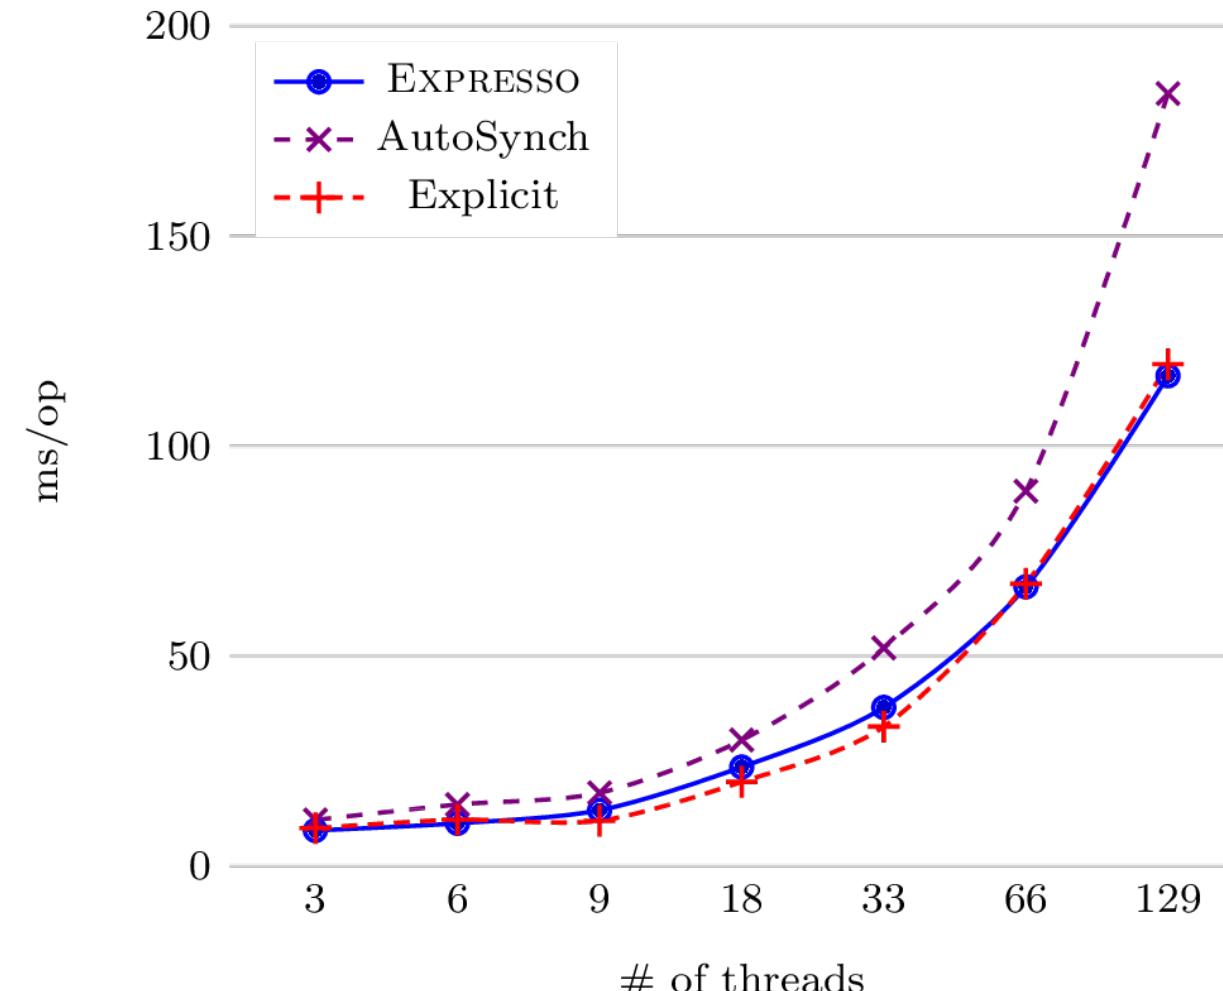
AsyncDispatch (Gradle)



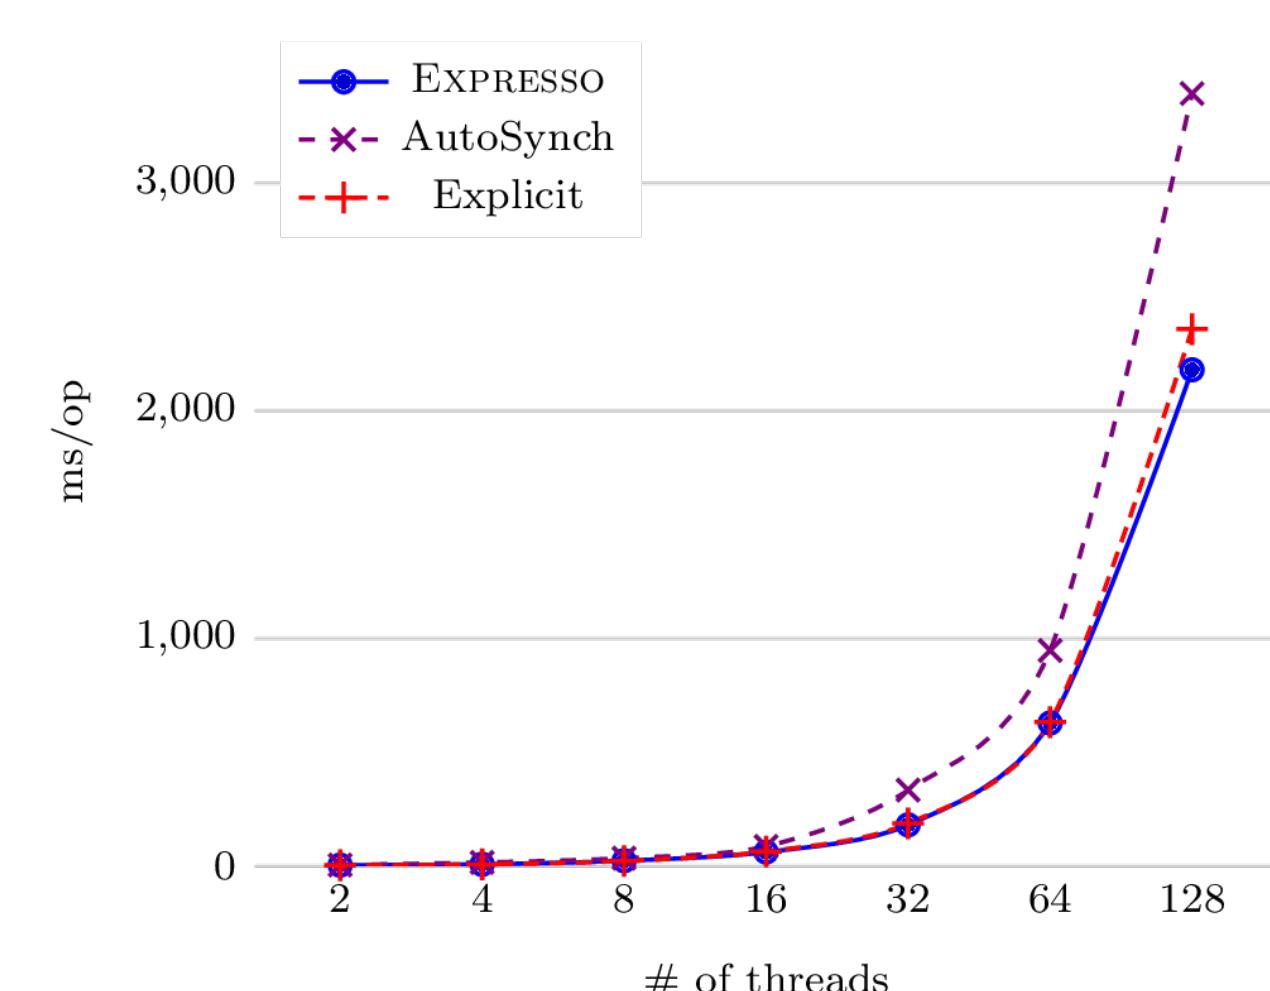
SimpleBlockingDeployment (Gradle)



SimpleDecoder (ExoPlayer)

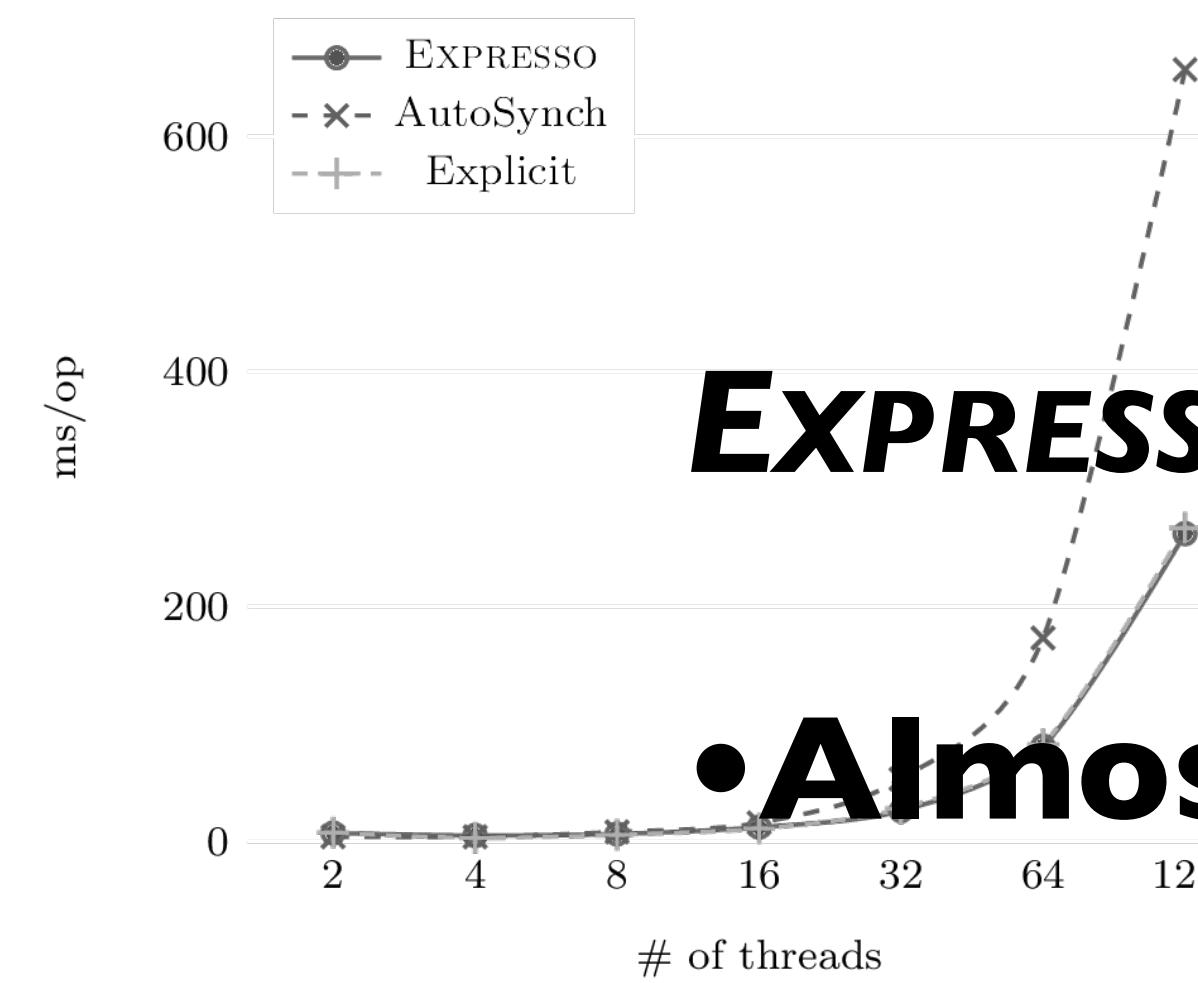


AsyncOperationExecutor (greenDAO)

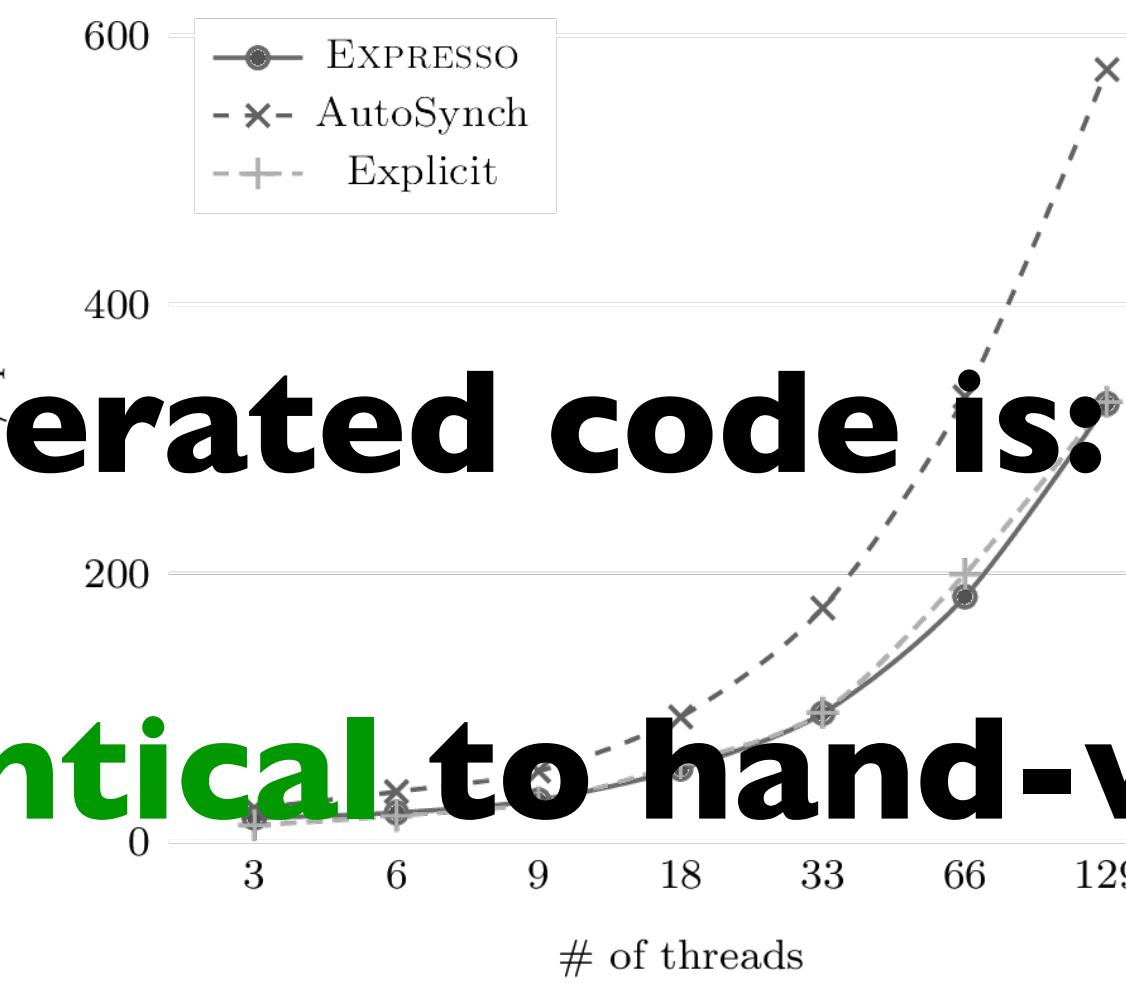


Results, cont

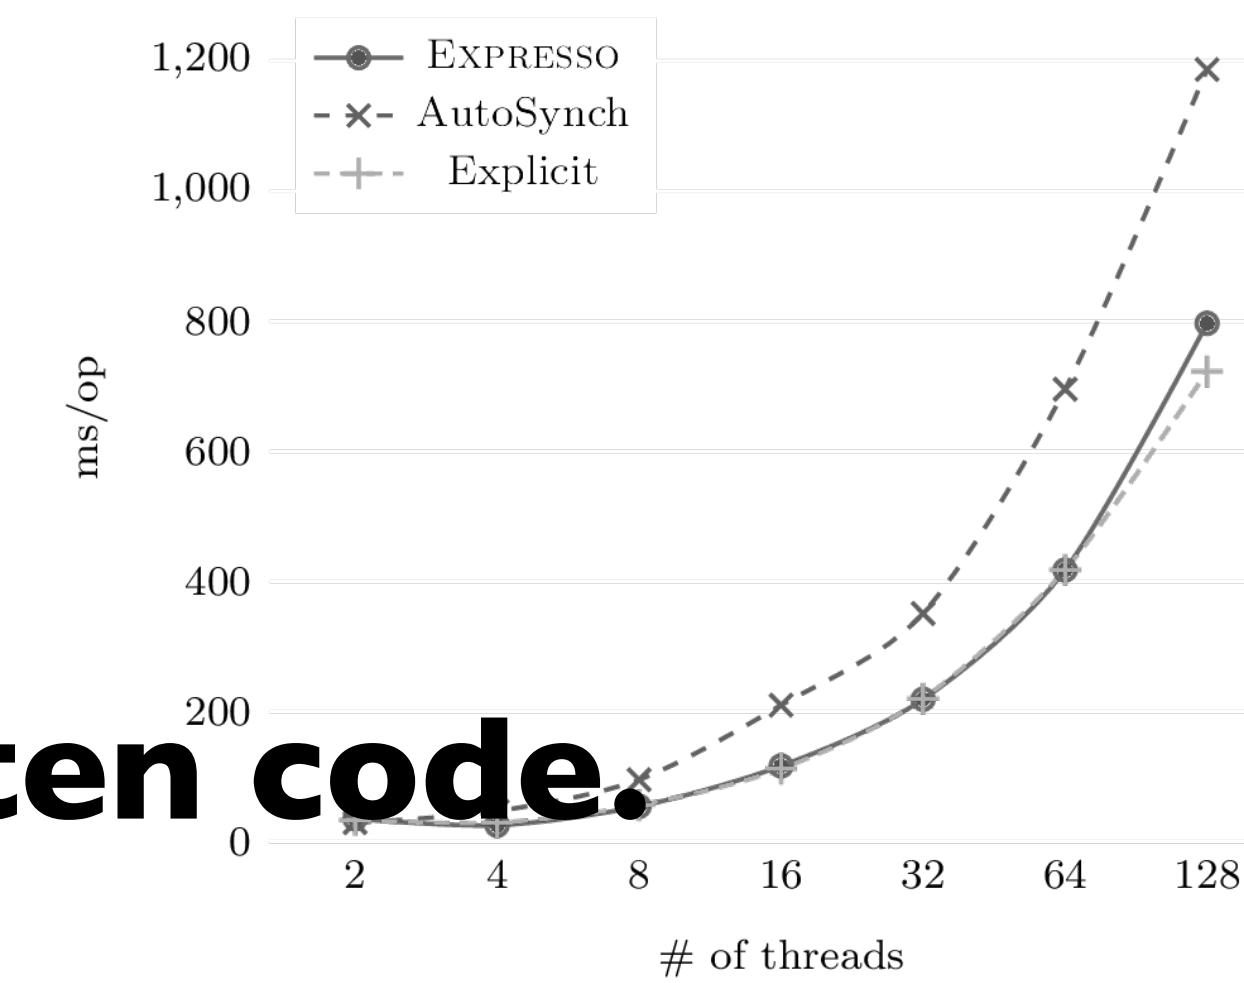
ConcurrencyThrottle (Spring framework)



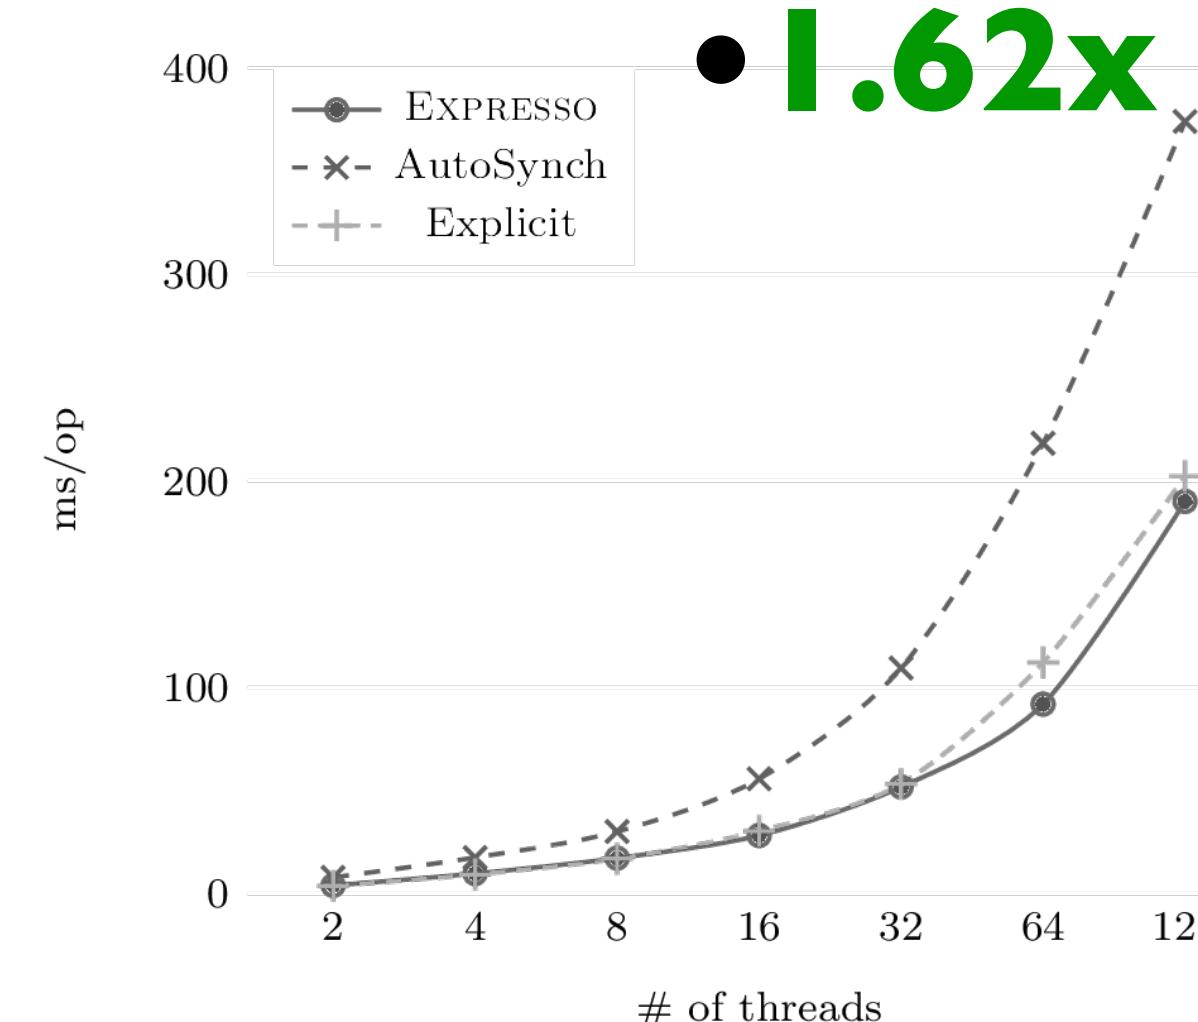
PendingPostQueue (EventBus)



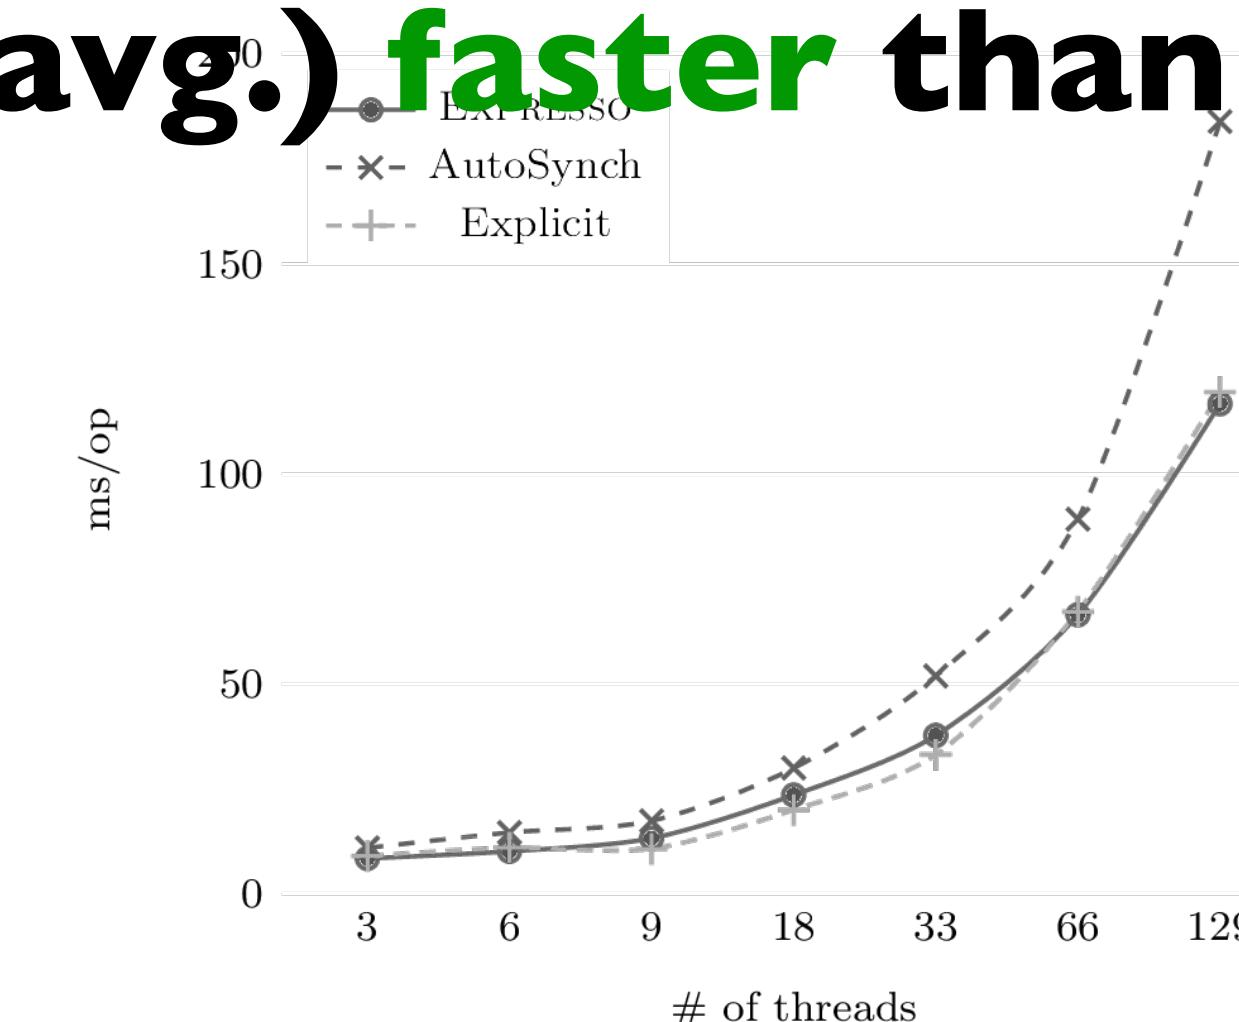
AsyncDispatch (Gradle)



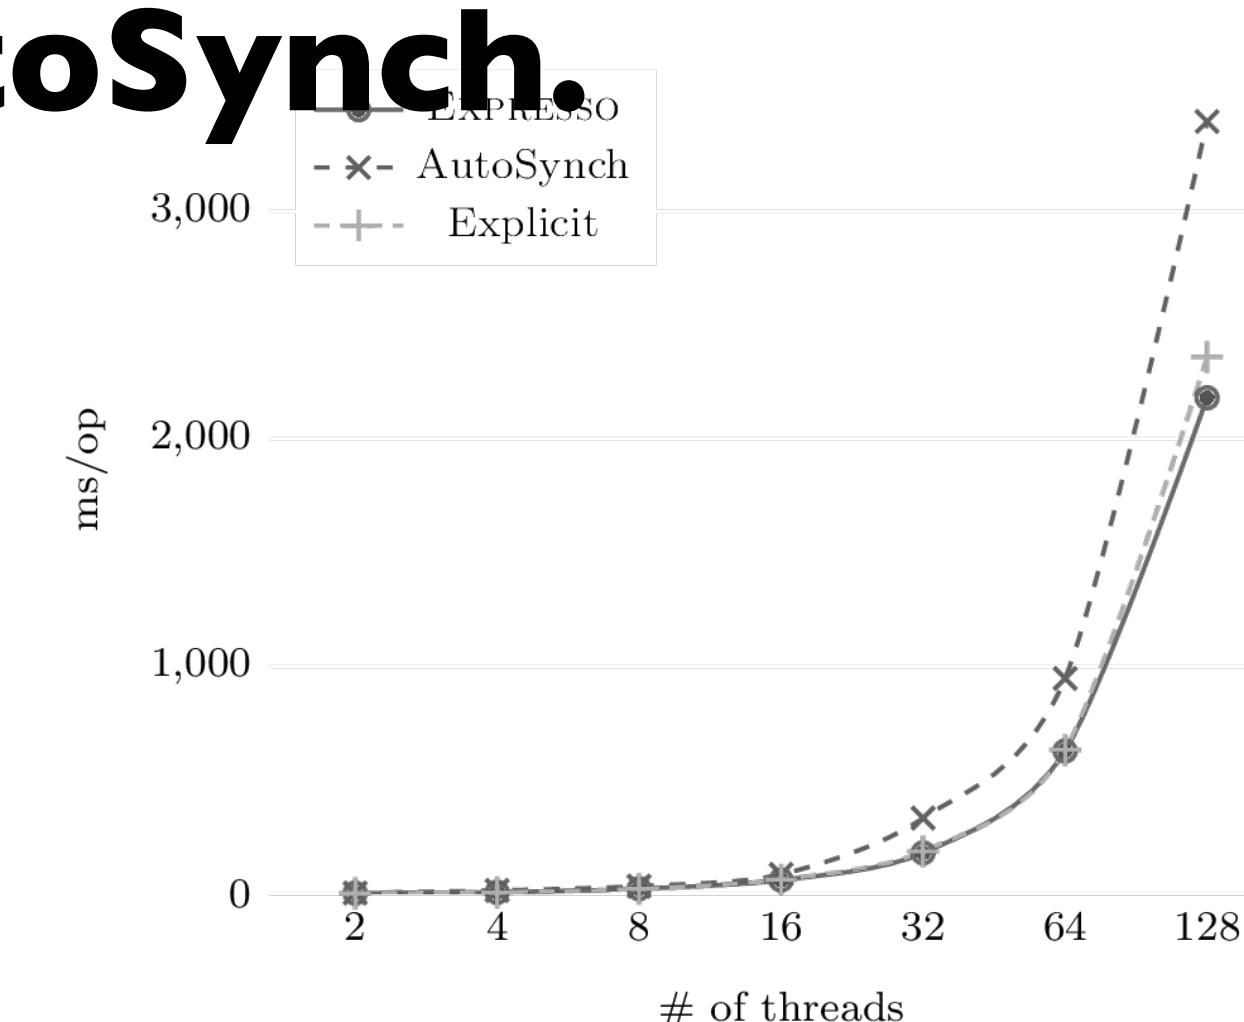
SimpleBlockingDeployment (Gradle)



SimpleDecoder (ExoPlayer)



AsyncOperationExecutor (greenDAO)



Road Map

Part I: How to automatically synthesize necessary **signaling** code



Symbolic Reasoning for Automatic Signal Placement.
Ferles et al, [PLDI'18](#).



Part II: How to automatically synthesize efficient **locking**



Synthesizing Fine-grained Synchronization Protocols for Implicit Monitors.
Ferles et al, [OOPSLA'22](#).

Road Map

Part I: How to automatically synthesize necessary **signaling** code



Symbolic Reasoning for Automatic Signal Placement.
Ferles et al, [PLDI'18](#).



Part II: How to automatically synthesize efficient **locking**



Synthesizing Fine-grained Synchronization Protocols for Implicit Monitors.
Ferles et al, [OOPSLA'22](#).

Recall



Expresso uses a *single global lock* acquired by all CCRs to ensure atomicity



This can be suboptimal in many cases!

Part II: Synthesizing Fine-Grained Locking

In some cases, multiple threads can run **concurrently** in a monitor without violating the object's atomicity

Part II: Synthesizing Fine-Grained Locking

In some cases, multiple threads can run **concurrently** in a monitor without violating the object's atomicity

Each method *appears* to execute atomically even though execution of methods can be interleaved

Part II: Synthesizing Fine-Grained Locking

In some cases, multiple threads can run **concurrently** in a monitor without violating the object's atomicity

Each method *appears* to execute atomically even though execution of methods can be interleaved



Idea: Improve performance by synthesizing **fine-grained locking** in a way that will provide an *illusion* of atomicity

Example

```
class ArrayBlockingQueue {
    int first = 0, last = 0, count = 0;
    Object[] queue;

    ArrayBlockingQueue(int capacity) {
        assert capacity < 1;
        this.queue = new Object[capacity];
    }

    atomic void put(Object o) {
        waituntil(count < queue.len){
            queue[last] = o;
            last = (last + 1) % queue.len;
            count++;
        }
    }

    atomic Object take() {
        waituntil(count > 0){
            Object r = queue[first];
            queue[first] = null;
            first = (first + 1) % queue.len;
            count--;
            return r;
        }
    }
}
```

Blocking FIFO Queue:

- **put** blocks if queue is full
- **take** blocks if queue is empty

Example

```
class ArrayBlockingQueue {
    int first = 0, last = 0, count = 0;
    Object[] queue;

    ArrayBlockingQueue(int capacity) {
        assert capacity < 1;
        this.queue = new Object[capacity];
    }

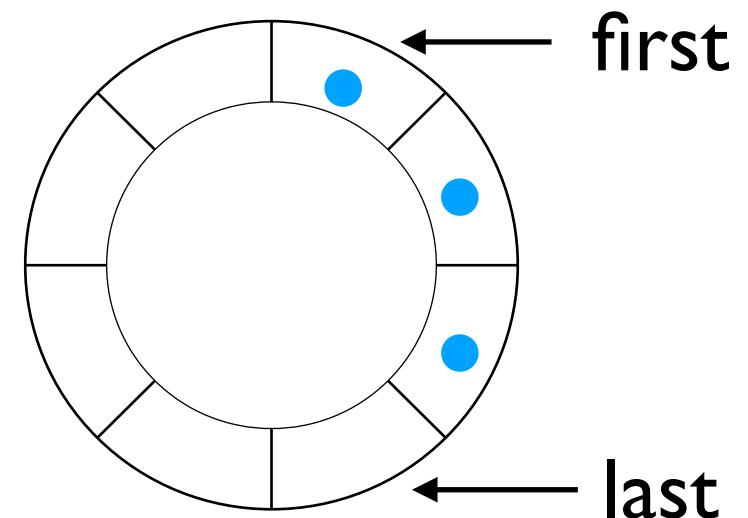
    atomic void put(Object o) {
        waituntil(count < queue.len){
            queue[last] = o;
            last = (last + 1) % queue.len;
            count++;
        }
    }

    atomic Object take() {
        waituntil(count > 0){
            Object r = queue[first];
            queue[first] = null;
            first = (first + 1) % queue.len;
            count--;
            return r;
        }
    }
}
```

Blocking FIFO Queue:

- **put** blocks if queue is full
- **take** blocks if queue is empty

Note: **put** and **take** can run **concurrently**



Example

```
class ArrayBlockingQueue {
    int first = 0, last = 0, count = 0;
    Object[] queue;

    ArrayBlockingQueue(int capacity) {
        assert capacity < 1;
        this.queue = new Object[capacity];
    }

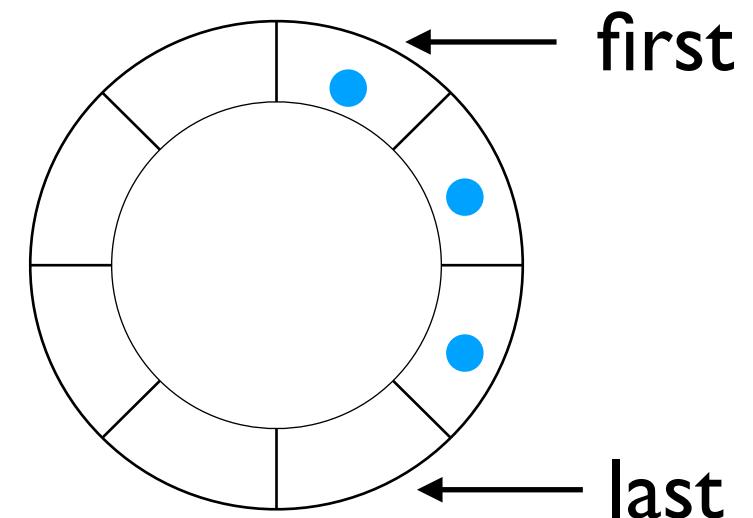
    atomic void put(Object o) {
        waituntil(count < queue.len){
            queue[last] = o;
            last = (last + 1) % queue.len;
            count++;
        }
    }

    atomic Object take() {
        waituntil(count > 0){
            Object r = queue[first];
            queue[first] = null;
            first = (first + 1) % queue.len;
            count--;
            return r;
        }
    }
}
```

Blocking FIFO Queue:

- **put** blocks if queue is full
- **take** blocks if queue is empty

Note: **put** and **take** can run **concurrently**



! A global lock would prevent such parallelization opportunities

Example, cont

```
class ArrayBlockingQueue {  
    Lock putLock = new Lock(),  
        takeLock = new Lock();  
    Condition notFull = putLock.newCondition();  
    Condition notEmpty = takeLock.newCondition();  
  
    int first = 0, last = 0; Object[] queue;  
    AtomicInteger count = new AtomicInteger(0);  
  
    void put(Object o) {  
        putLock.lock();  
        while (count.get() == queue.length)  
            notFull.await();  
        queue[last] = o;  
        last = (last + 1) % queue.length;  
        count.getAndIncrement();  
        putLock.unlock();  
        // Signalling code ...  
    }  
  
    Object take() {  
        takeLock.lock();  
        while (count.get() == 0)  
            notEmpty.await();  
        Object r = queue[first];  
        queue[first] = null;  
        first = (first + 1) % queue.length;  
        count.getAndDecrement();  
        takeLock.unlock();  
        // Signalling code ...  
        return r;  
    }  
}
```

Optimal:
maximizes parallelism

Example, cont

```
class ArrayBlockingQueue {  
    Lock putLock = new Lock(),  
        takeLock = new Lock(); ← two locks  
    Condition notFull = putLock.newCondition();  
    Condition notEmpty = takeLock.newCondition();  
  
    int first = 0, last = 0; Object[] queue;  
    AtomicInteger count = new AtomicInteger(0);  
  
    void put(Object o) {  
        putLock.lock();  
        while (count.get() == queue.length)  
            notFull.await();  
        queue[last] = o;  
        last = (last + 1) % queue.length;  
        count.getAndIncrement();  
        putLock.unlock();  
        // Signalling code ...  
    }  
  
    Object take() {  
        takeLock.lock();  
        while (count.get() == 0)  
            notEmpty.await();  
        Object r = queue[first];  
        queue[first] = null;  
        first = (first + 1) % queue.length;  
        count.getAndDecrement();  
        takeLock.unlock();  
        // Signalling code ...  
        return r;  
    }  
}
```

Optimal:
maximizes parallelism

Example, cont

```
class ArrayBlockingQueue {  
    Lock putLock = new Lock(),  
        takeLock = new Lock(); ← two locks  
    Condition notFull = putLock.newCondition();  
    Condition notEmpty = takeLock.newCondition();  
  
    int first = 0, last = 0; Object[] queue;  
    AtomicInteger count = new AtomicInteger(0);  
  
    void put(Object o) { ← put uses putLock  
        putLock.lock();  
        while (count.get() == queue.length)  
            notFull.await();  
        queue[last] = o;  
        last = (last + 1) % queue.length;  
        count.getAndIncrement();  
        putLock.unlock();  
        // Signalling code ...  
    }  
  
    Object take() {  
        takeLock.lock();  
        while (count.get() == 0)  
            notEmpty.await();  
        Object r = queue[first];  
        queue[first] = null;  
        first = (first + 1) % queue.length;  
        count.getAndDecrement();  
        takeLock.unlock();  
        // Signalling code ...  
        return r;  
    }  
}
```

Optimal:
maximizes parallelism

Example, cont

```
class ArrayBlockingQueue {  
    Lock putLock = new Lock(),  
        takeLock = new Lock(); ← two locks  
    Condition notFull = putLock.newCondition();  
    Condition notEmpty = takeLock.newCondition();  
  
    int first = 0, last = 0; Object[] queue;  
    AtomicInteger count = new AtomicInteger(0);  
  
    void put(Object o) { ← put uses putLock  
        putLock.lock();  
        while (count.get() == queue.length)  
            notFull.await();  
        queue[last] = o;  
        last = (last + 1) % queue.length;  
        count.getAndIncrement();  
        putLock.unlock();  
        // Signalling code ...  
    }  
  
    Object take() { ← take uses takeLock  
        takeLock.lock();  
        while (count.get() == 0)  
            notEmpty.await();  
        Object r = queue[first];  
        queue[first] = null;  
        first = (first + 1) % queue.length;  
        count.getAndDecrement();  
        takeLock.unlock();  
        // Signalling code ...  
        return r;  
    }  
}
```

Optimal:
maximizes parallelism

Example, cont

```
class ArrayBlockingQueue {  
    Lock putLock = new Lock(),  
        takeLock = new Lock(); ← two locks  
    Condition notFull = putLock.newCondition();  
    Condition notEmpty = takeLock.newCondition();  
  
    int first = 0, last = 0; Object[] queue;  
    AtomicInteger count = new AtomicInteger(0); ← put uses putLock  
  
    void put(Object o) {  
        putLock.lock();  
        while (count.get() == queue.length)  
            notFull.await();  
        queue[last] = o;  
        last = (last + 1) % queue.length;  
        count.getAndIncrement(); ← count is an AtomicInteger  
        putLock.unlock();  
        // Signalling code ...  
    }  
  
    Object take() {  
        takeLock.lock(); ← take uses takeLock  
        while (count.get() == 0)  
            notEmpty.await();  
        Object r = queue[first];  
        queue[first] = null;  
        first = (first + 1) % queue.length;  
        count.getAndDecrement();  
        takeLock.unlock();  
        // Signalling code ...  
        return r;  
    }  
}
```

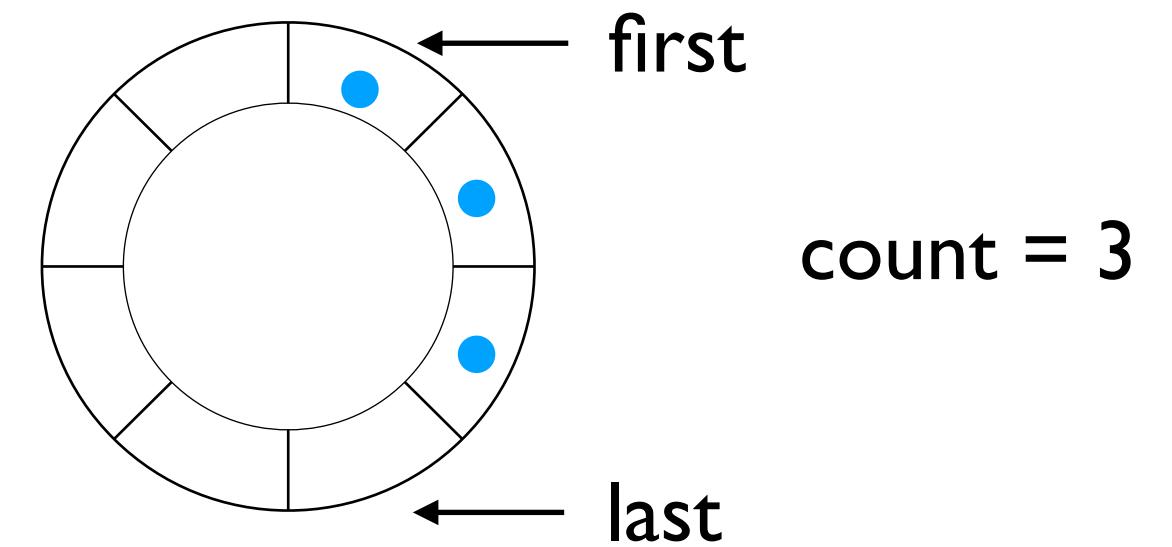
Optimal:
maximizes parallelism

count is an AtomicInteger

Why is this correct?

(I) Does not contain any races

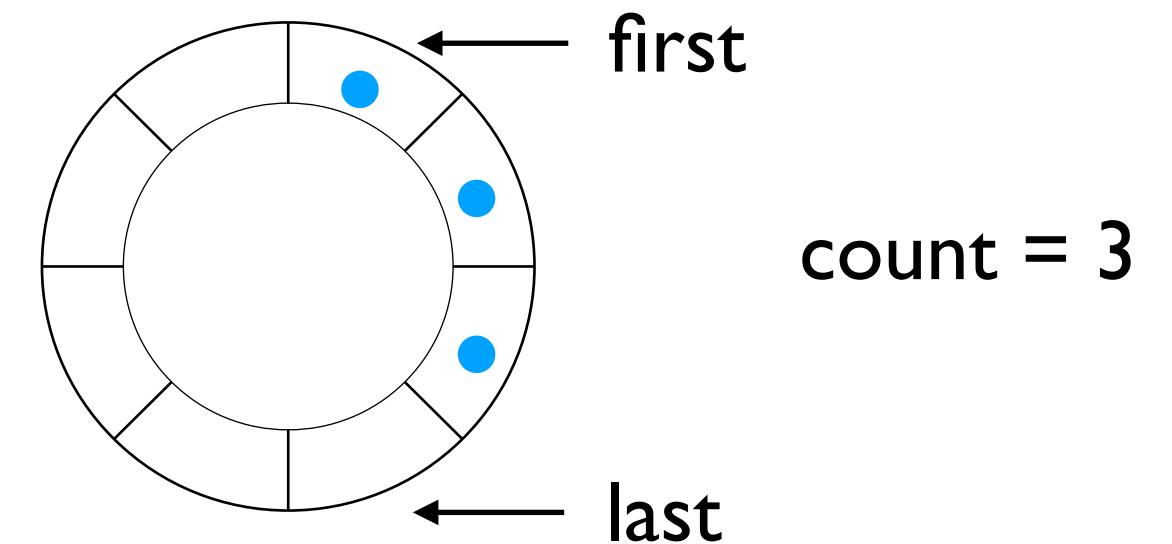
- No race between queue[first] and queue[last]
- Races on count are handled by AtomicInteger



Why is this correct?

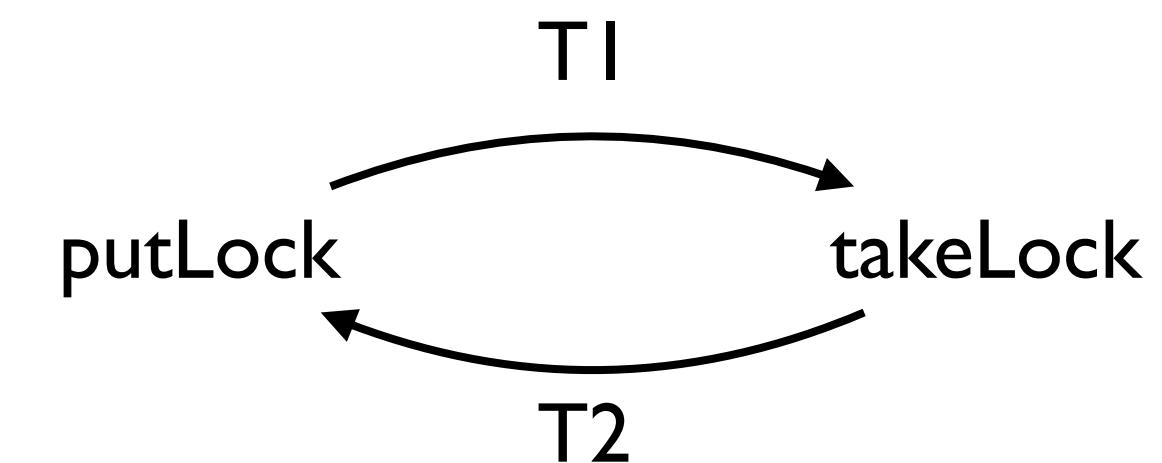
(1) Does not contain any races

- No race between `queue[first]` and `queue[last]`
- Races on `count` are handled by `AtomicInteger`



(2) Monitor is deadlock-free

- No cycles in the lock acquisition order



Why is this correct? (cont.)

(3) All operations *appear* to execute as one indivisible unit

- Executions where threads interleave are equivalent to sequential executions.

T1

```
wait(count < len)
```

```
queue[last] = o;...
```

```
count++;
```

T2

```
wait(count > 0)
```

```
r = queue[first];...
```

```
count--;
```

Why is this correct? (cont.)

(3) All operations *appear* to execute as one indivisible unit

- Executions where threads interleave are equivalent to sequential executions.

T1

```
wait(count < len)  
queue[last] = o;...  
count++;
```

T2

```
wait(count > 0)  
r = queue[first];...  
count--;
```

Why is this correct? (cont.)

(3) All operations *appear* to execute as one indivisible unit

Safe interleaving

- Executions where threads interleave are equivalent to sequential executions.

T1

```
wait(count < len)  
queue[last] = o;...  
count++;
```

T2

```
wait(count > 0)  
r = queue[first];...  
count--;
```

Why is this correct? (cont.)

(3) All operations *appear* to execute as one indivisible unit

Safe interleaving

- Executions where threads interleave are equivalent to sequential executions.

T1

```
wait(count < len)    queue[last] = o;...    count++;
```

T2

```
wait(count > 0)    r = queue[first];...    count--;
```

- The locking scheme only allows such *safe interleavings*



Gives an **illusion** of atomicity!

Take-Away from Example

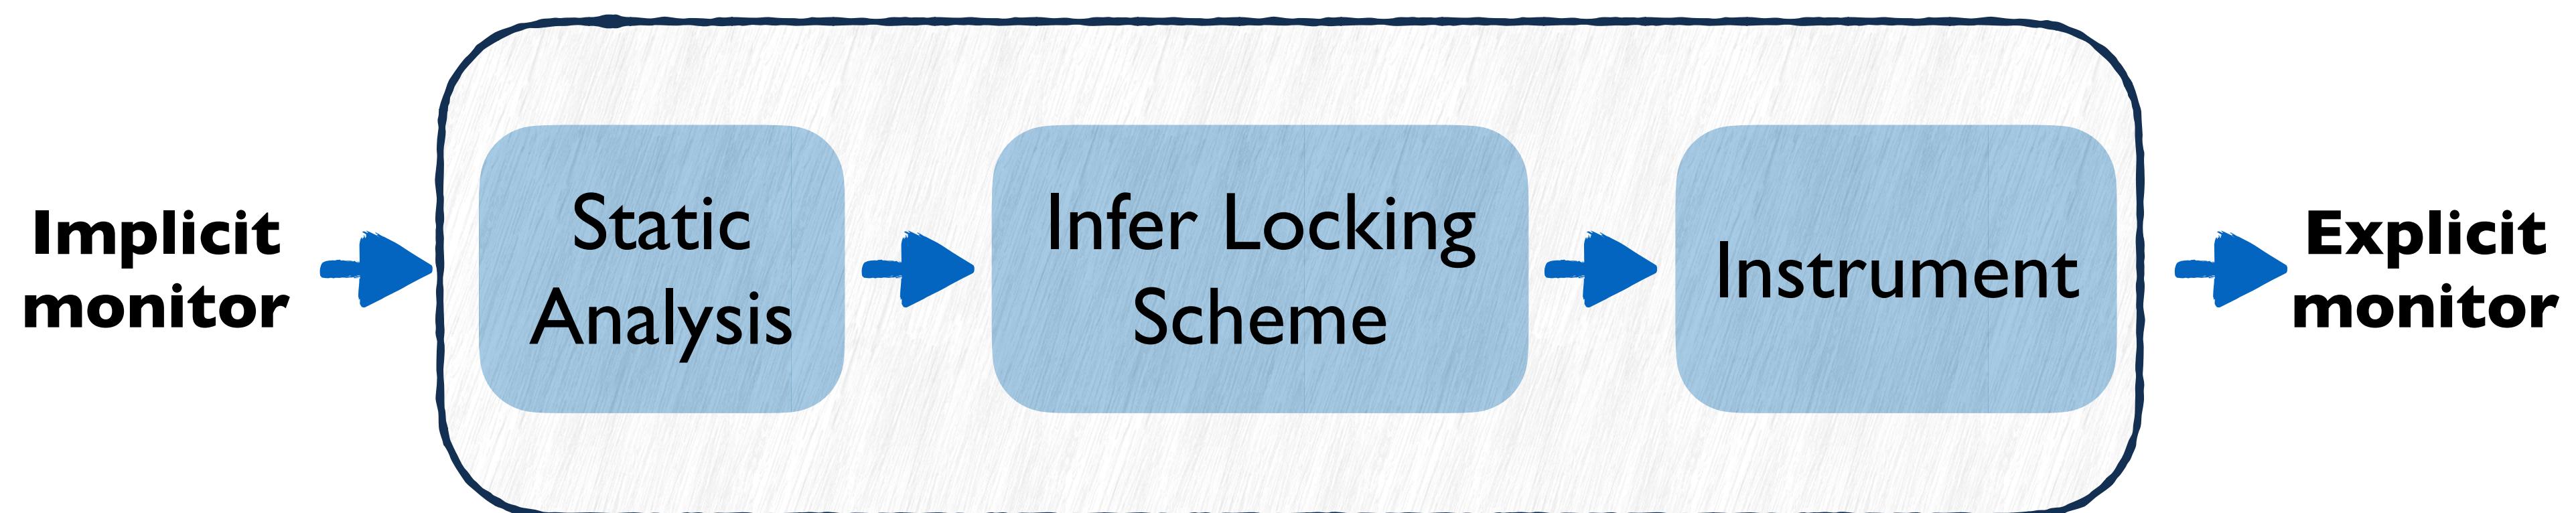


Goal: Maximize parallelization opportunities & obey correctness requirements:
data-race- and deadlock-freedom + atomicity

Take-Away from Example



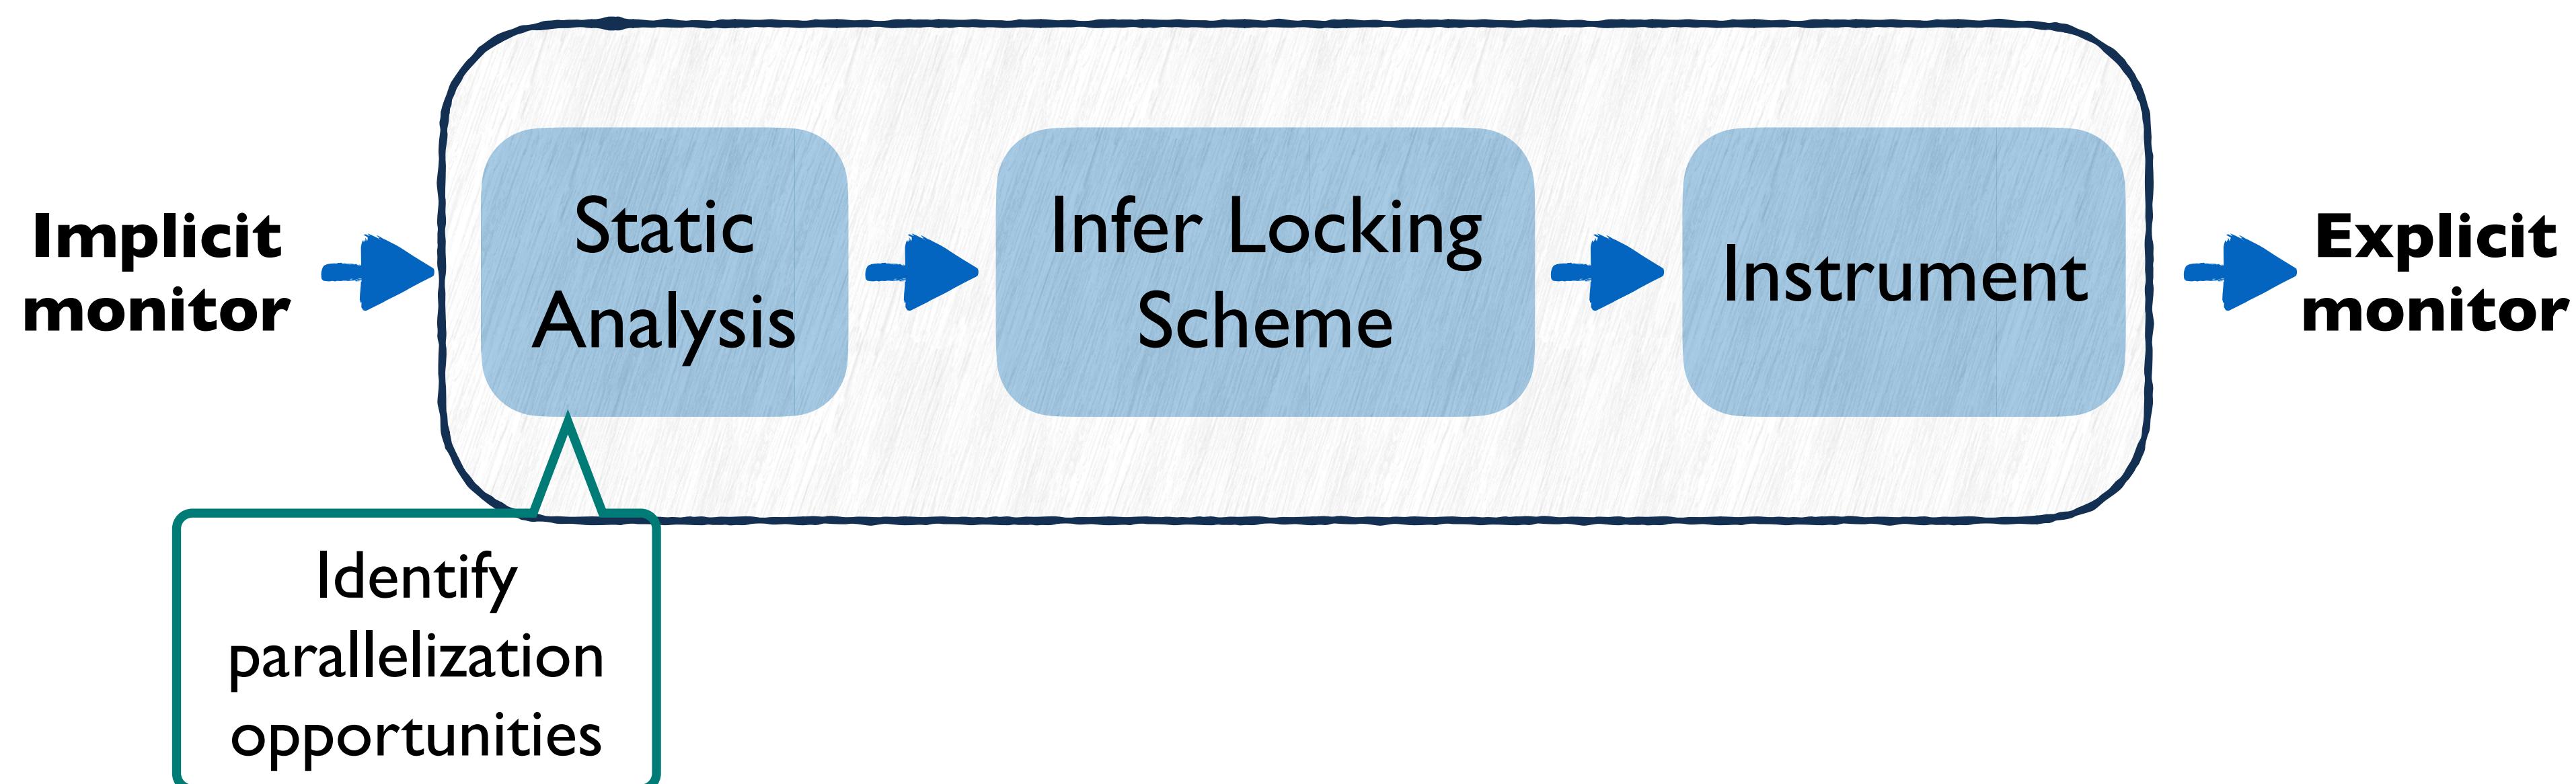
Goal: Maximize parallelization opportunities & obey correctness requirements:
data-race- and deadlock-freedom + atomicity



Take-Away from Example



Goal: Maximize parallelization opportunities & obey correctness requirements:
data-race- and deadlock-freedom + atomicity

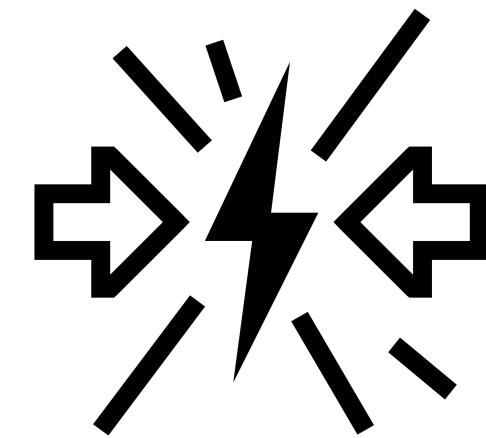


Static Analysis

First splits input monitor to disjoint code snippets

Static Analysis

First splits input monitor to disjoint code snippets

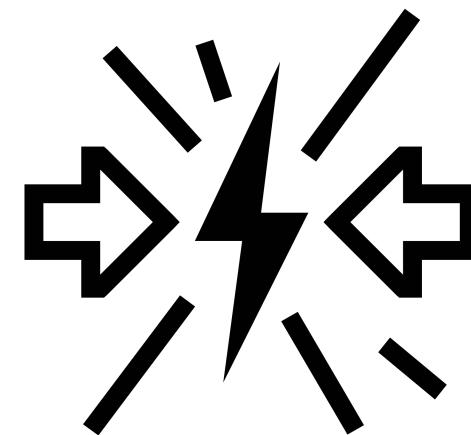


Data races

$\{(S, S') \mid \text{Racy}(S, S')\}$

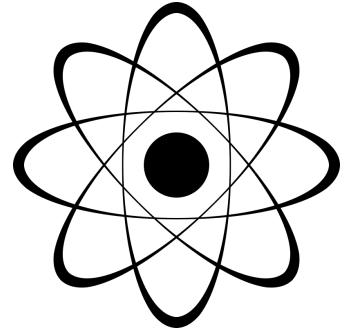
Static Analysis

First splits input monitor to disjoint code snippets



Data races

$\{(S, S') \mid \text{Racy}(S, S')\}$

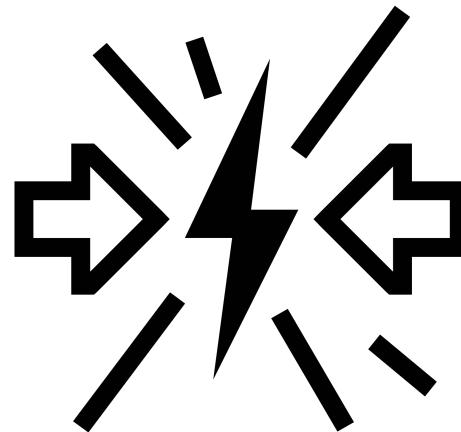


Fields that can be implemented using atomic type

$\{f \mid \text{AtomicTypeExists}(\text{Type}(f), \text{Updates}(f))\}$

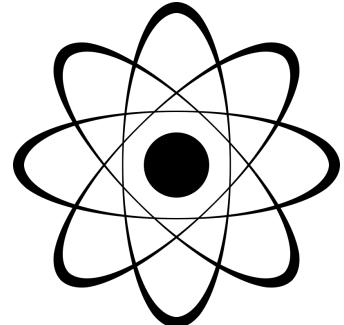
Static Analysis

First splits input monitor to disjoint code snippets



Data races

$\{(S, S') \mid \text{Racy}(S, S')\}$



Fields that can be implemented using atomic type

$\{f \mid \text{AtomicTypeExists}(\text{Type}(f), \text{Updates}(f))\}$



Safe interleaving opportunities

$\{(S, S1, S2) \mid \text{SafeInterleaving}(S, S1, S2)\}$

Detecting Safe Interleavings



Theorem: S can be interleaved in between S1 and S2 if it *left-commutes* with every *predecessor* of S1 and *right-commutes* with every *successor* of S2

Detecting Safe Interleavings



Theorem: S can be interleaved in between S1 and S2 if it *left-commutes* with every *predecessor* of S1 and *right-commutes* with every *successor* of S2

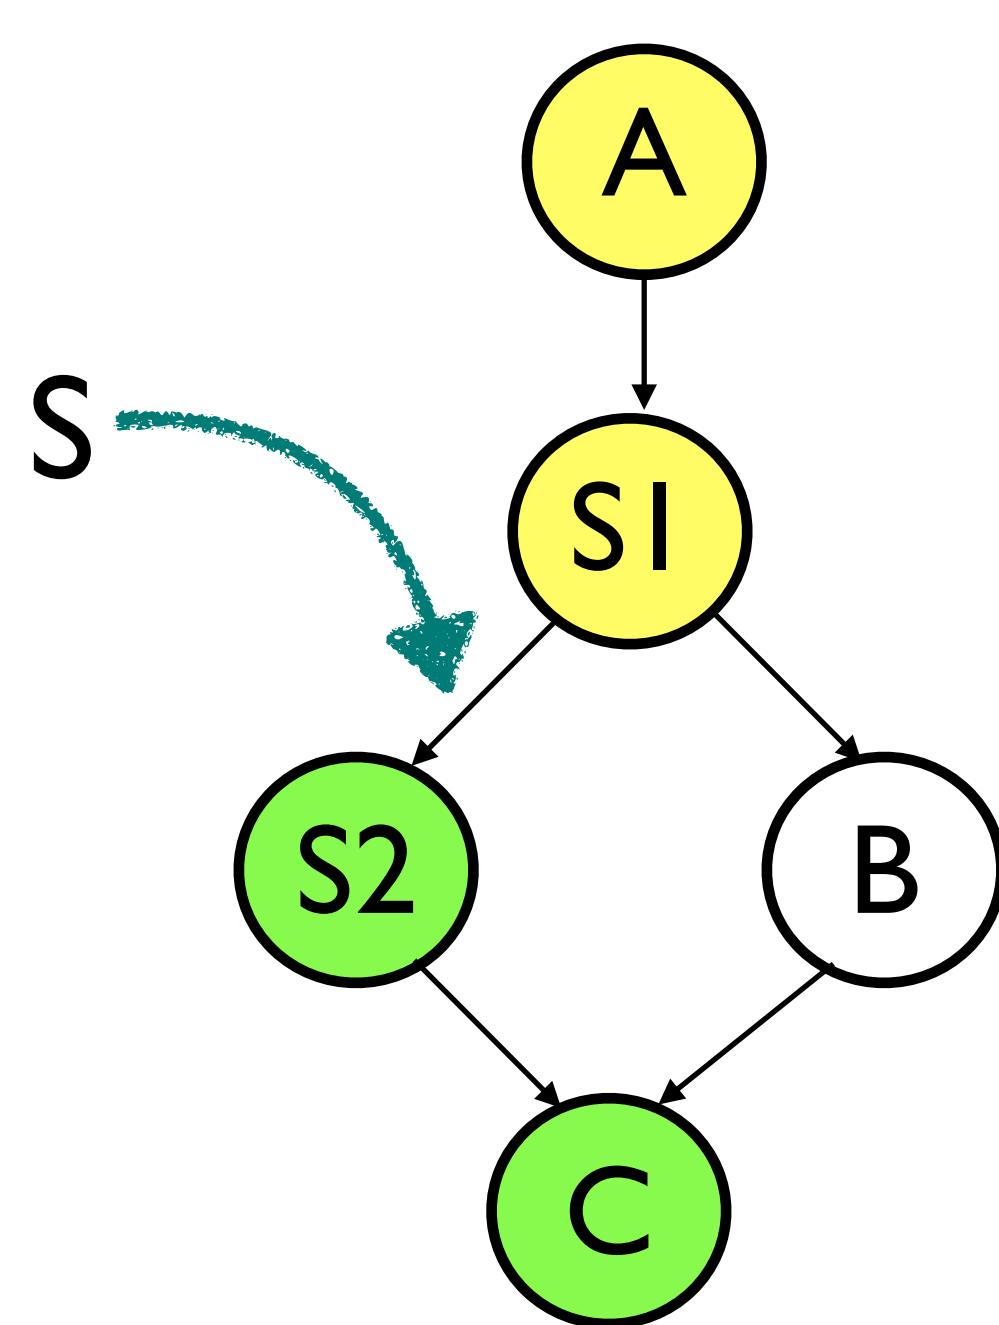
X left-commutes with Y, if whenever X executes right after Y, state is same as when X executes before Y

$Y; X \equiv X; Y$

Detecting Safe Interleavings



Theorem: S can be interleaved in between S1 and S2 if it *left-commutes* with every *predecessor* of S1 and *right-commutes* with every *successor* of S2



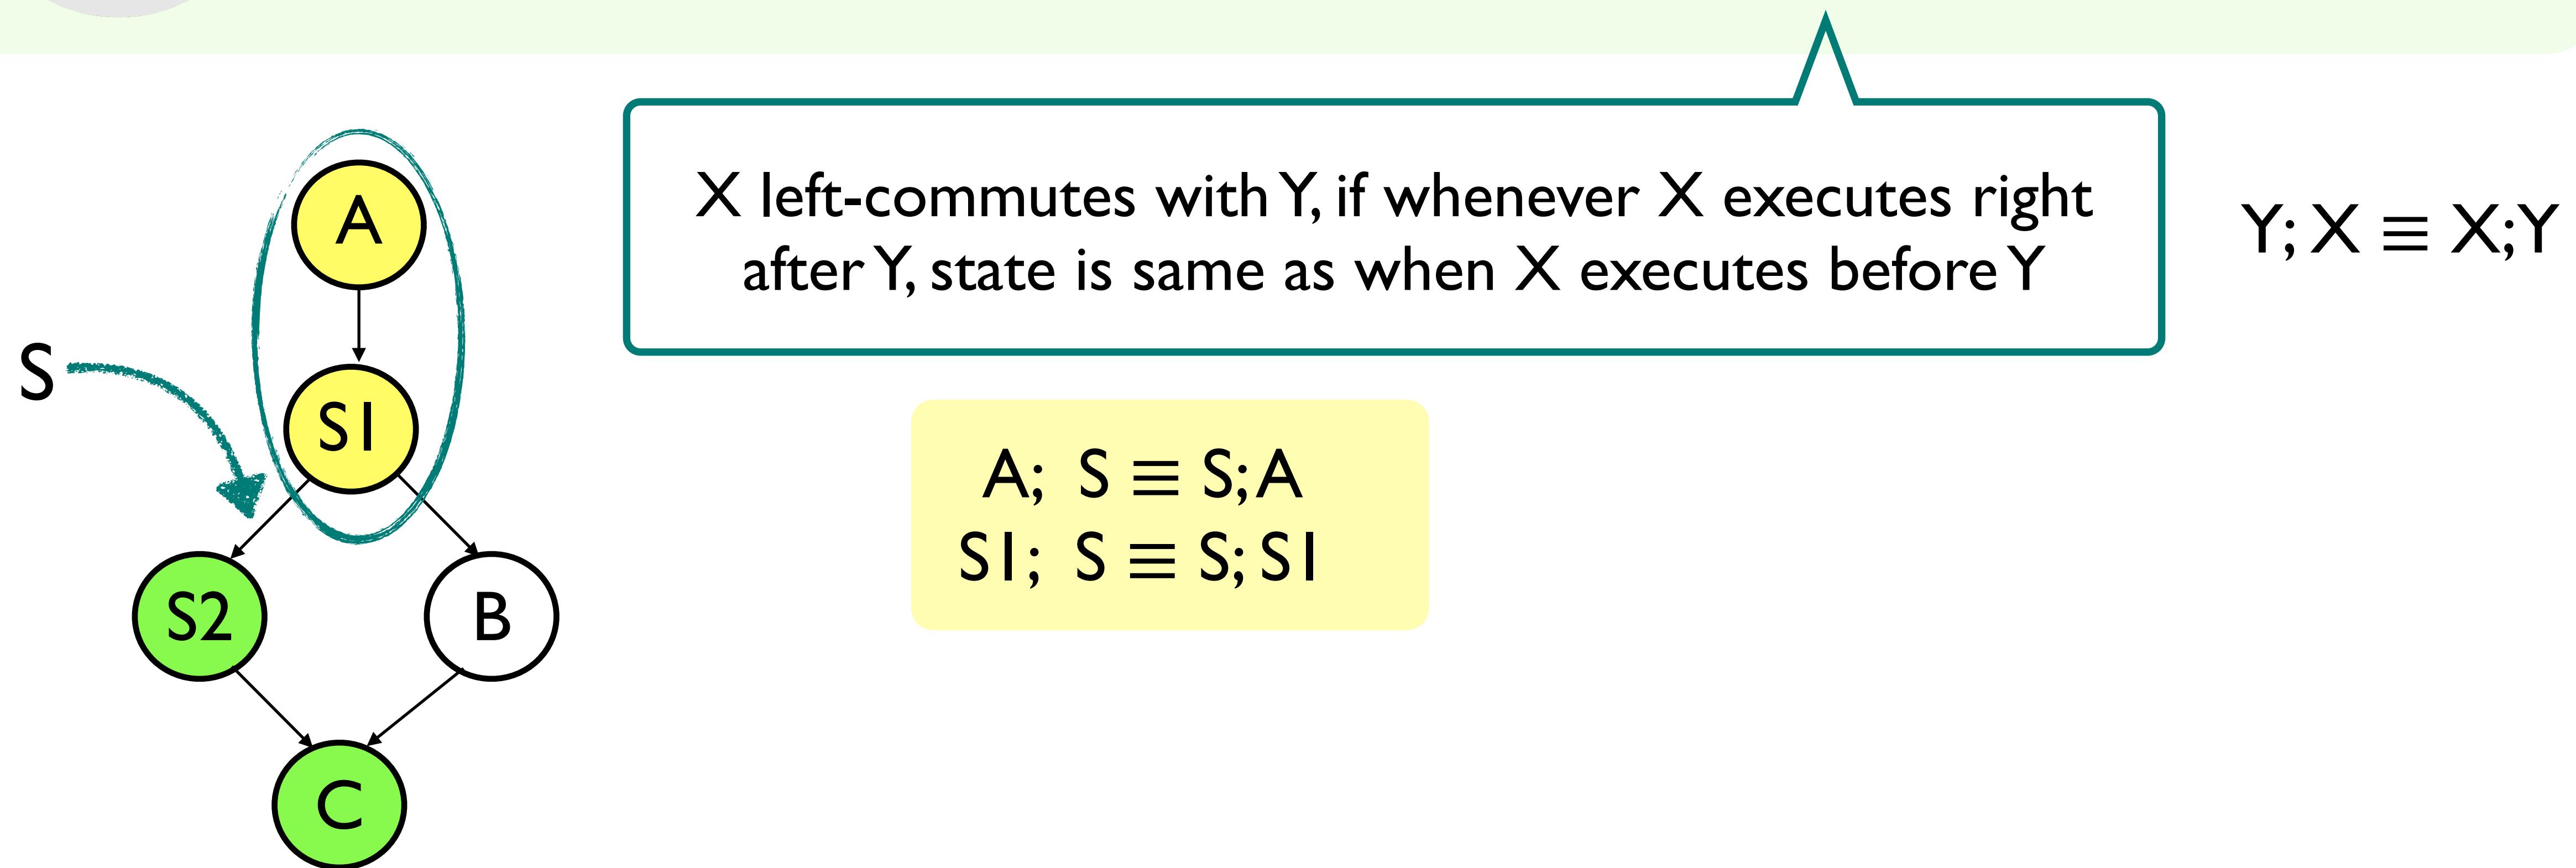
X left-commutes with Y, if whenever X executes right after Y, state is same as when X executes before Y

$Y; X \equiv X; Y$

Detecting Safe Interleavings



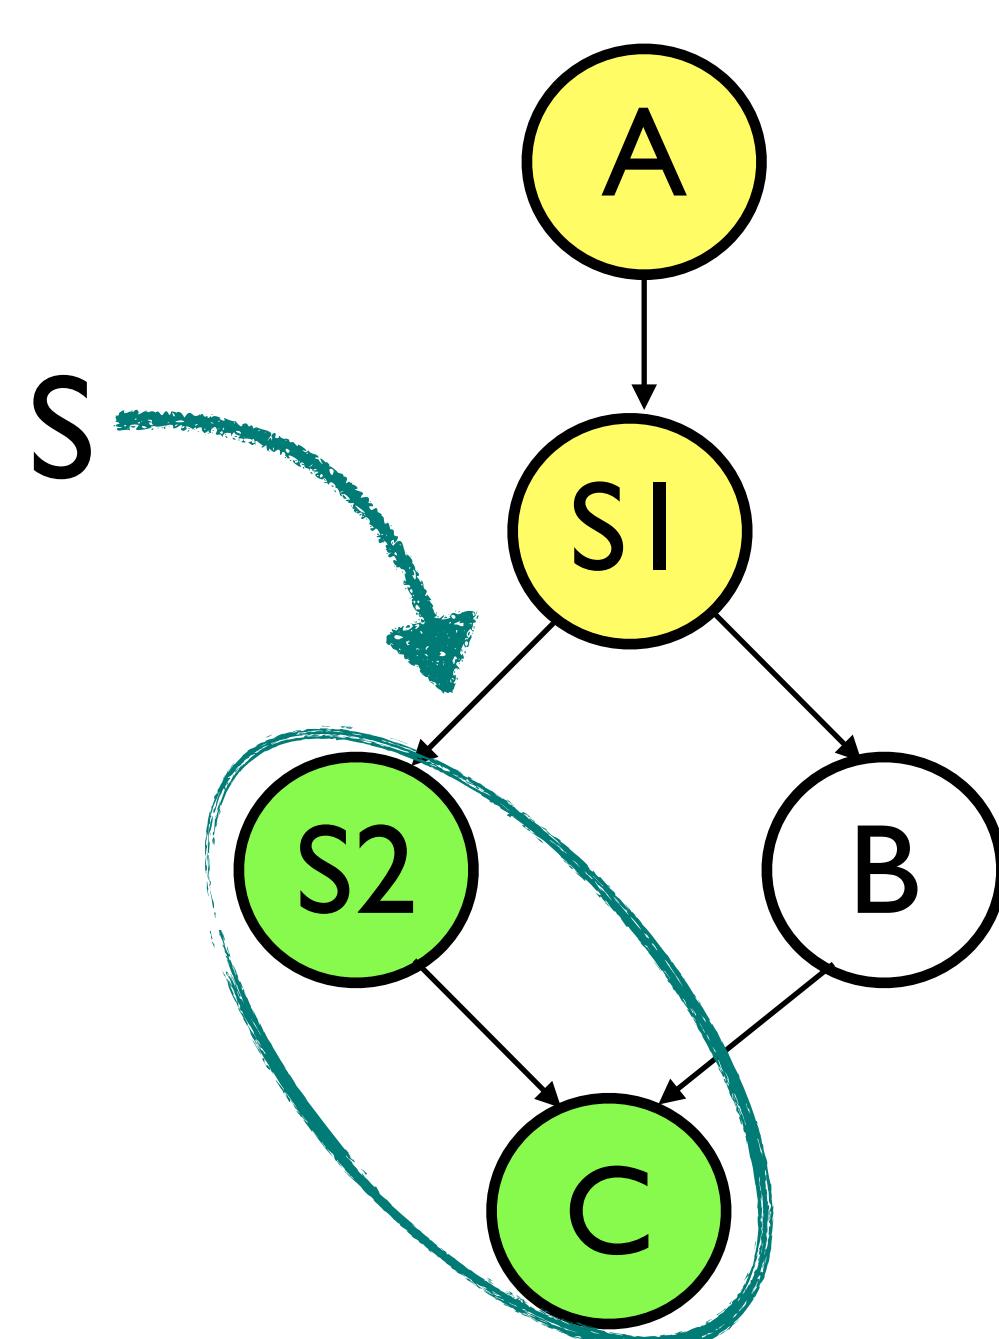
Theorem: S can be interleaved in between S1 and S2 if it *left-commutes* with every *predecessor* of S1 and *right-commutes* with every *successor* of S2



Detecting Safe Interleavings



Theorem: S can be interleaved in between S1 and S2 if it *left-commutes* with every *predecessor* of S1 and *right-commutes* with every *successor* of S2



X left-commutes with Y, if whenever X executes right after Y, state is same as when X executes before Y

$Y; X \equiv X; Y$

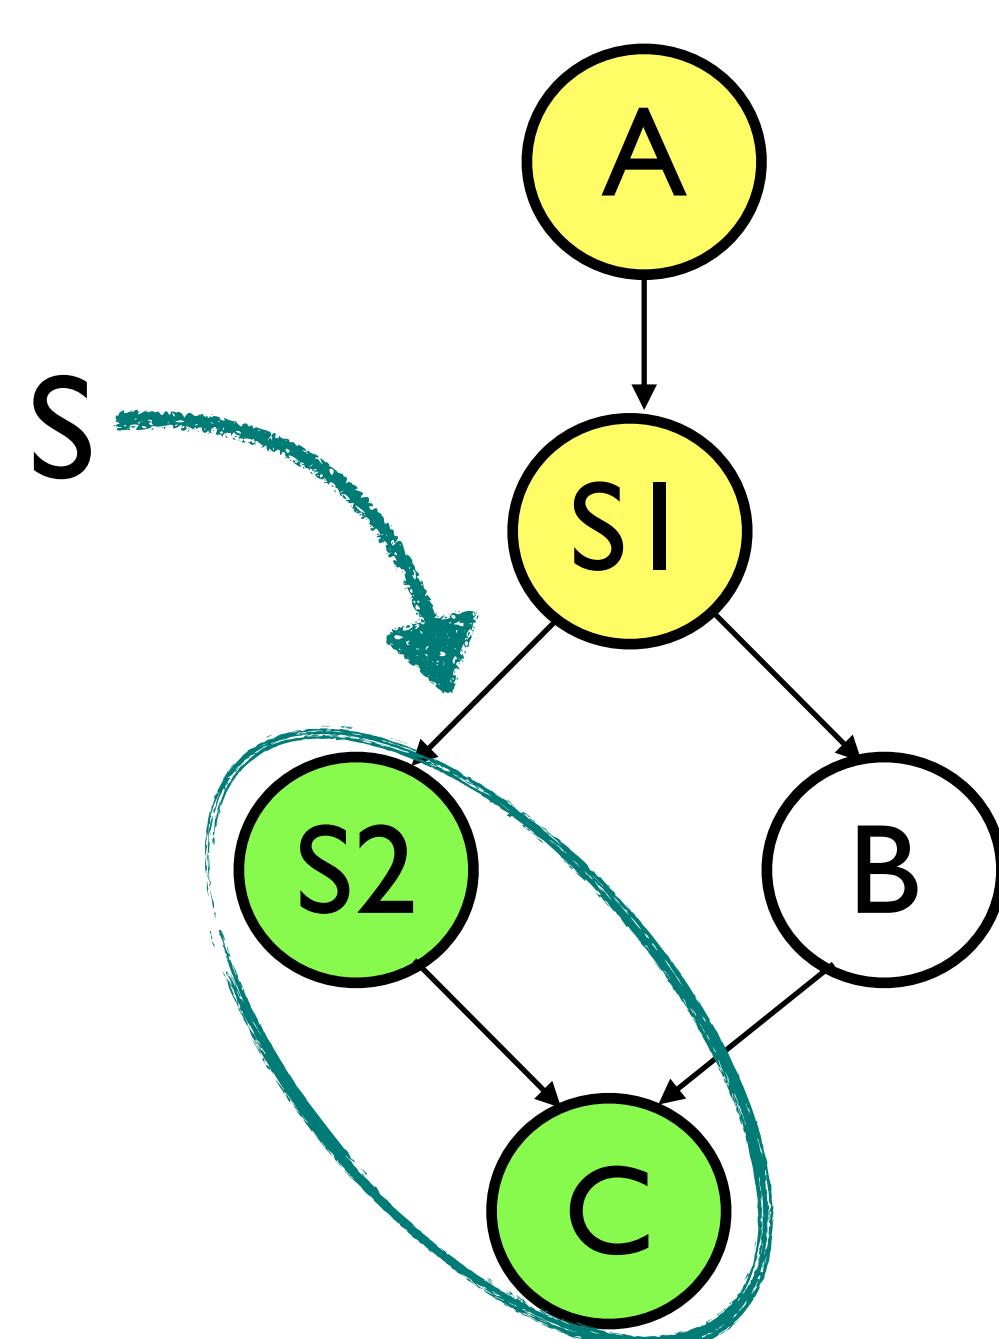
$A; S \equiv S; A$
 $SI; S \equiv S; SI$

$S; S2 \equiv S2; S$
 $S; C \equiv C; S$

Detecting Safe Interleavings



Theorem: S can be interleaved in between S1 and S2 if it *left-commutes* with every *predecessor* of S1 and *right-commutes* with every *successor* of S2



X left-commutes with Y, if whenever X executes right after Y, state is same as when X executes before Y

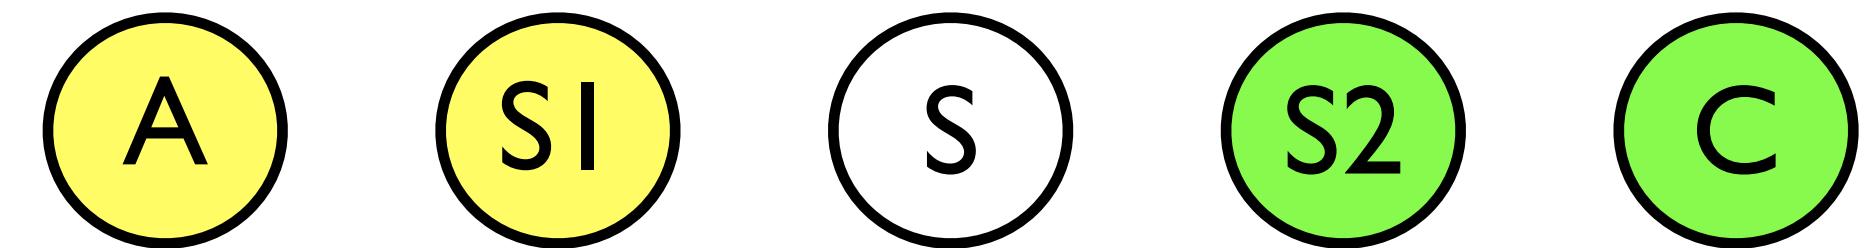
$Y; X \equiv X; Y$

$A; S \equiv S; A$
 $SI; S \equiv S; SI$

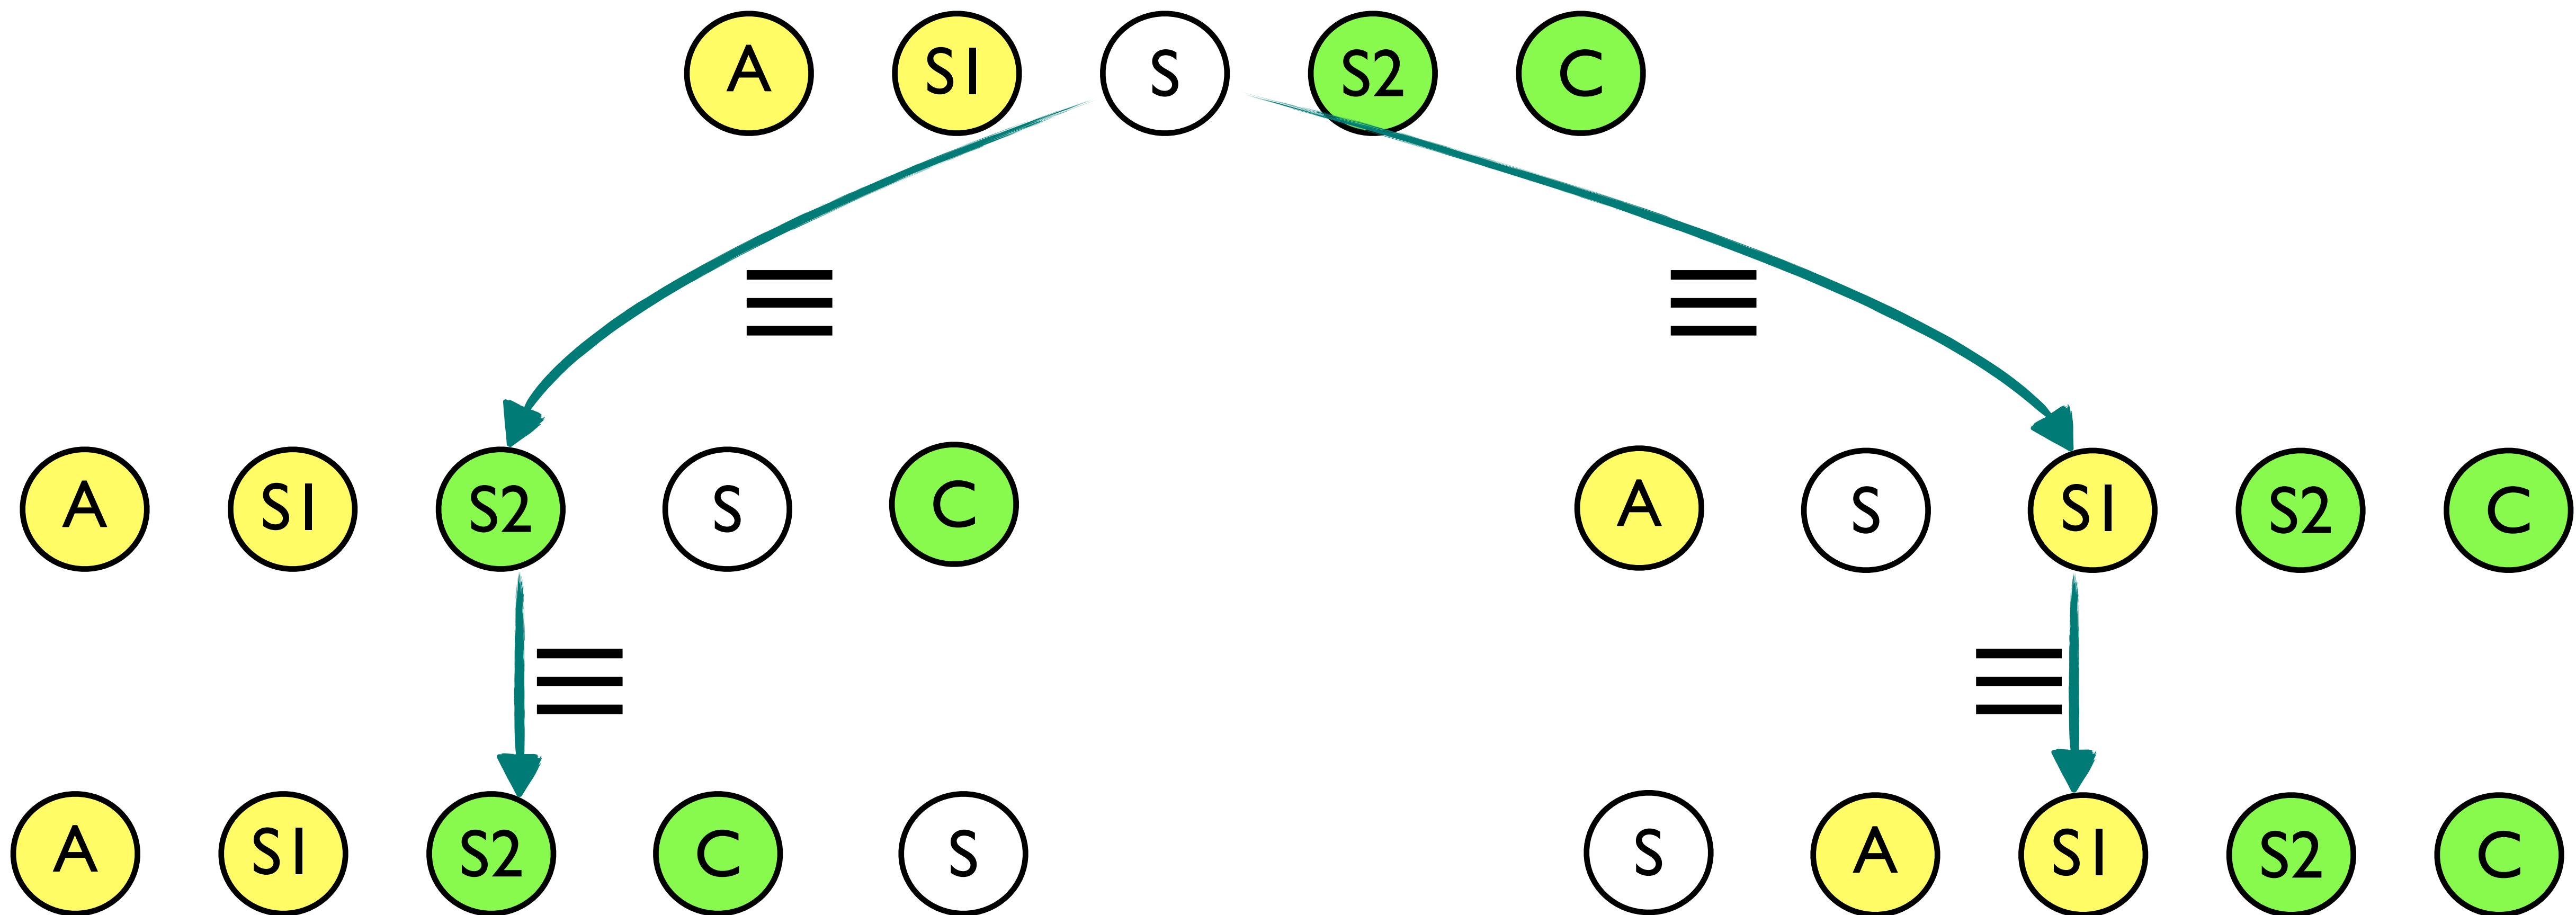
$S; S2 \equiv S2; S$
 $S; C \equiv C; S$

Program equivalence check

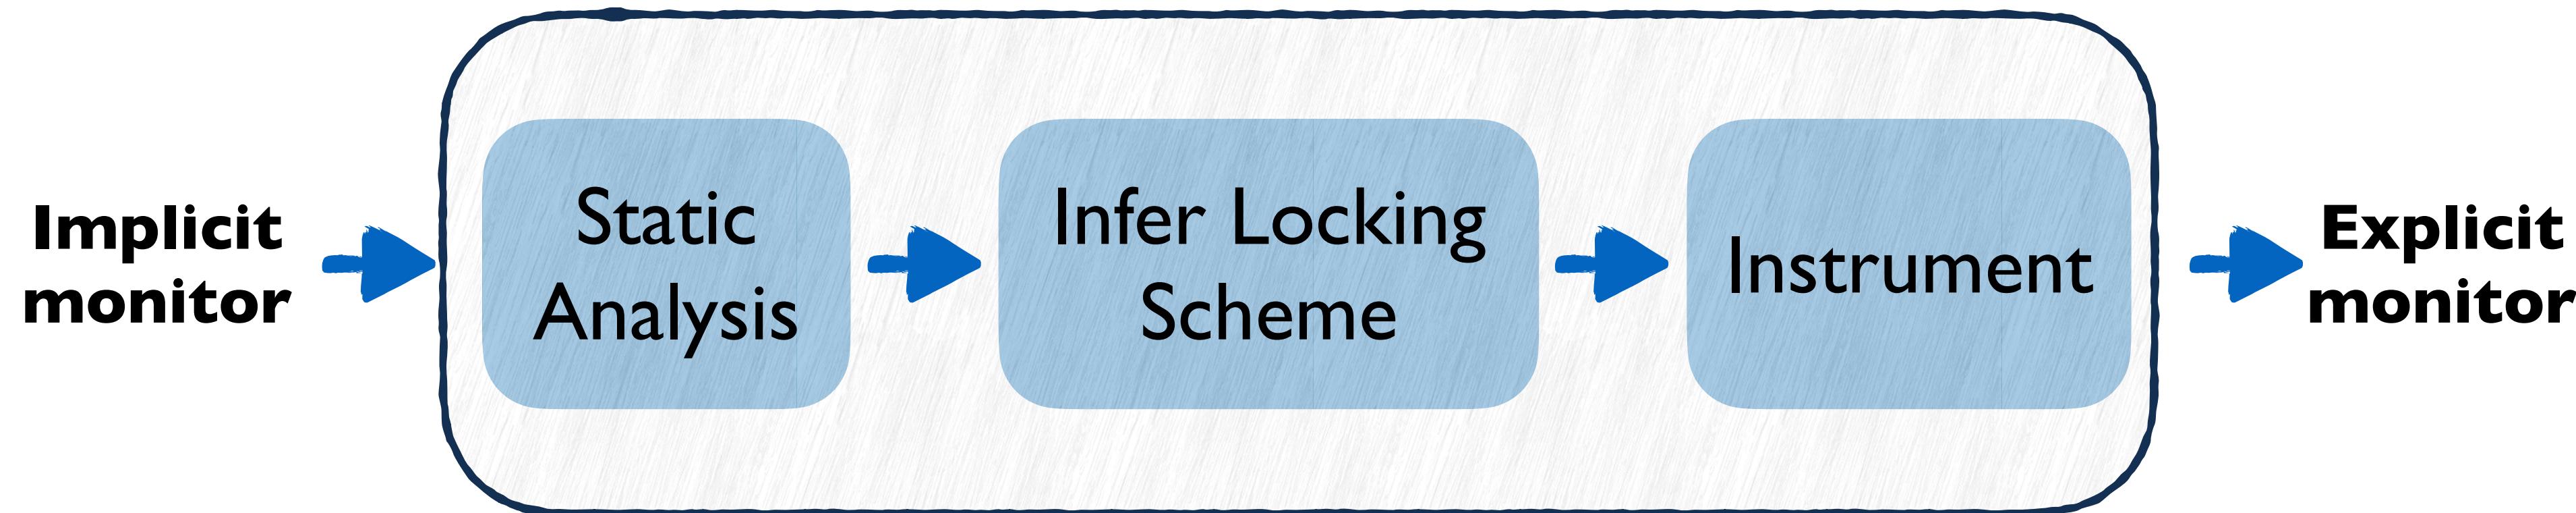
Detecting Safe Interleavings



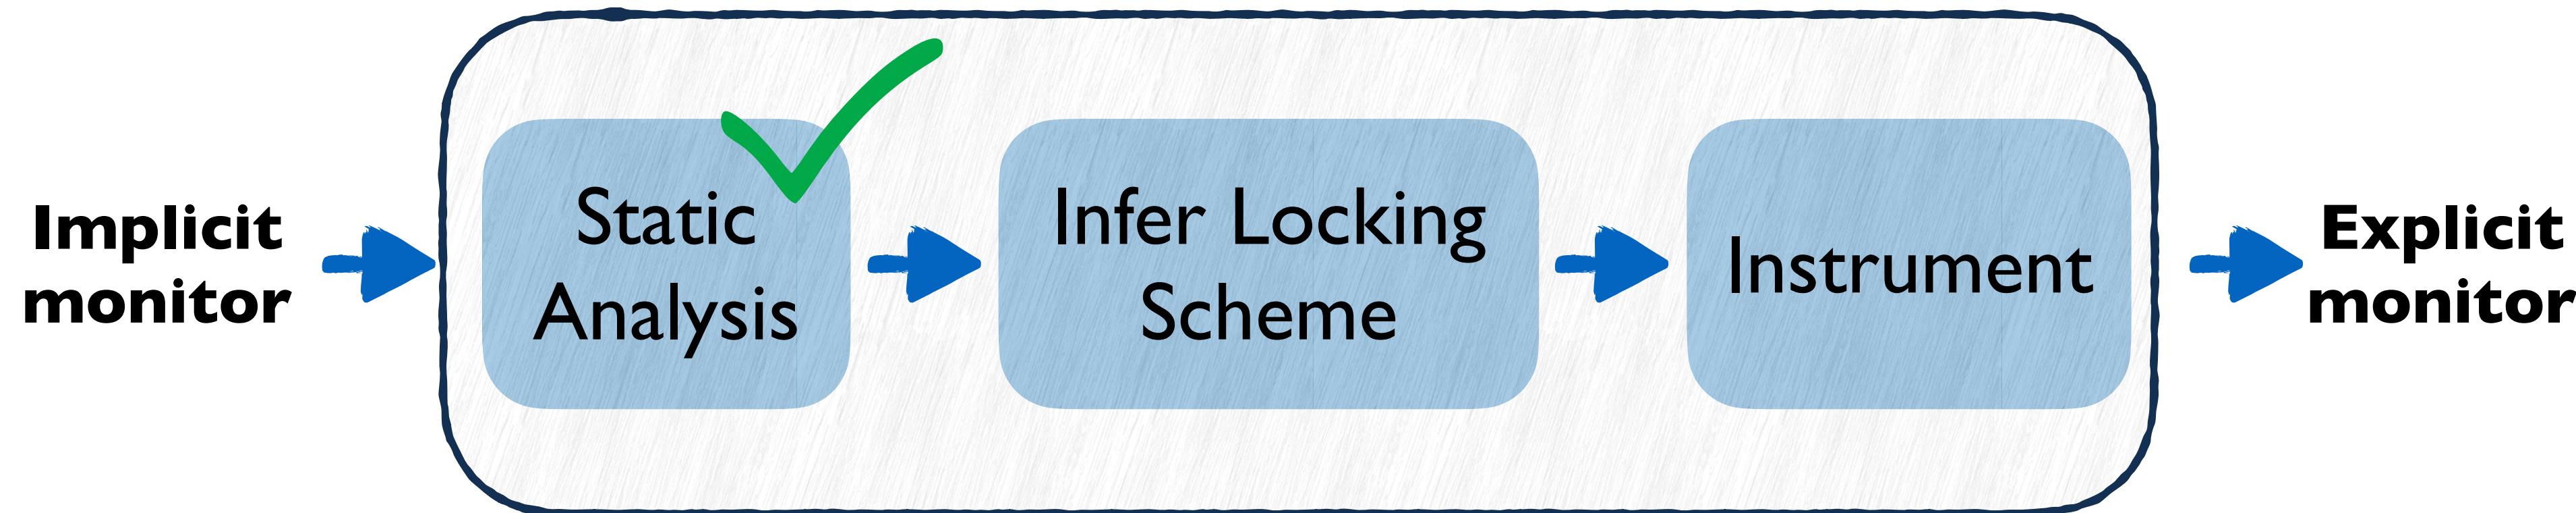
Detecting Safe Interleavings



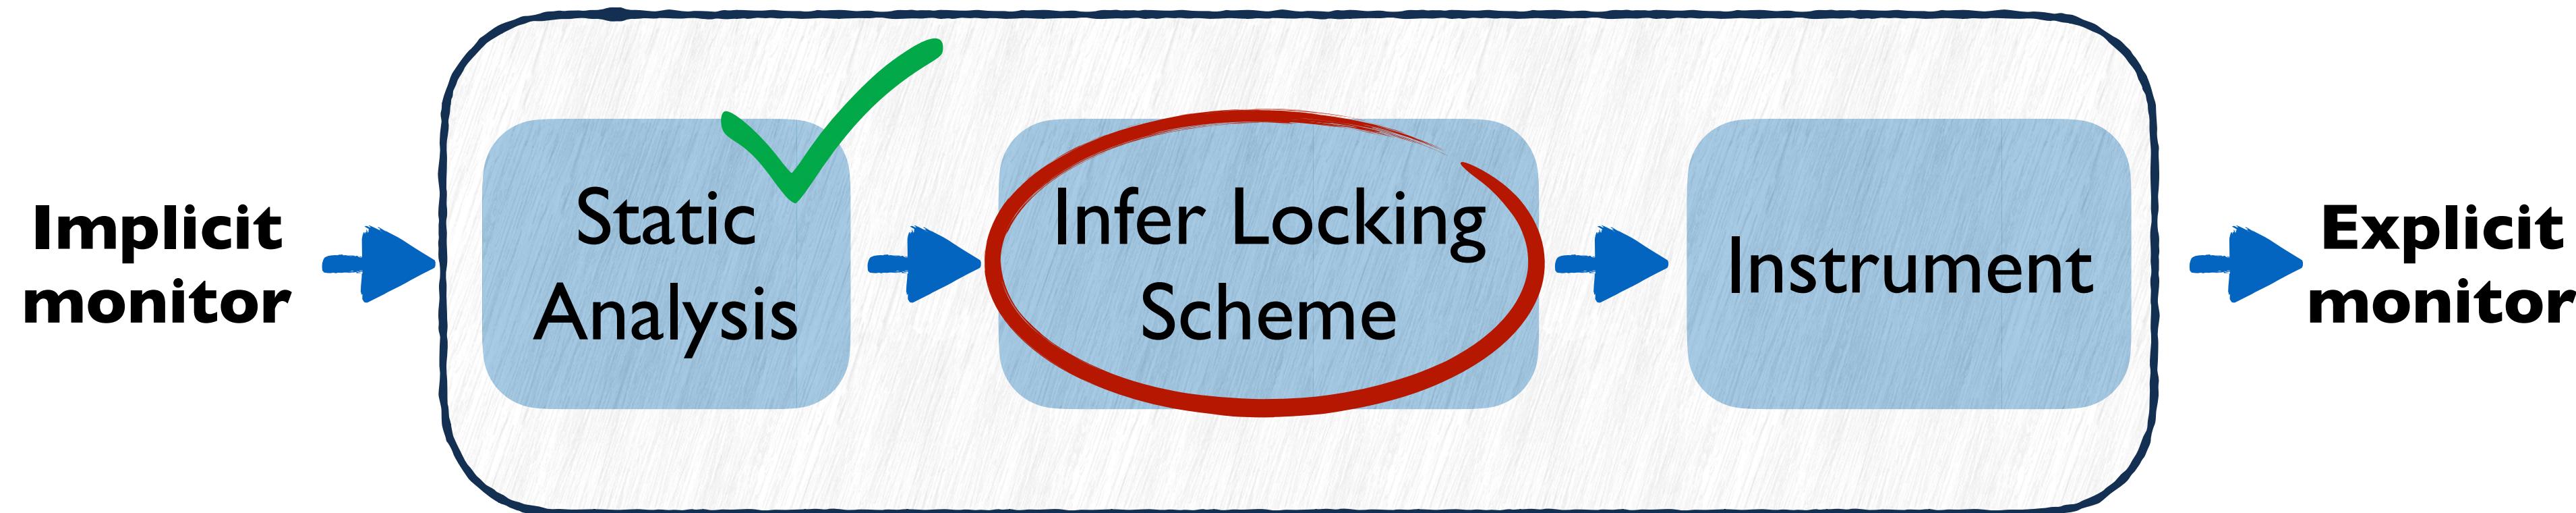
Synthesis Framework



Synthesis Framework



Synthesis Framework



Synthesizing Optimal Locking Scheme

Locking Scheme Requirement

Synthesizing Optimal Locking Scheme

Locking Scheme Requirement

Maximize parallelization
among monitor methods

Synthesizing Optimal Locking Scheme

Locking Scheme Requirement

Satisfy correctness requirements
(data race/deadlock freedom & atomicity)



Maximize parallelization
among monitor methods

Synthesizing Optimal Locking Scheme

Locking Scheme Requirement

Satisfy correctness requirements
(data race/deadlock freedom & atomicity)



Maximize parallelization
among monitor methods

Optimization problem

Synthesizing Optimal Locking Scheme

Locking Scheme Requirement

Satisfy correctness requirements
(data race/deadlock freedom & atomicity)



Maximize parallelization
among monitor methods

Optimization problem



MaxSAT Encoding

Synthesizing Optimal Locking Scheme

Locking Scheme Requirement

Satisfy correctness requirements
(data race/deadlock freedom & atomicity)



Maximize parallelization
among monitor methods

Optimization problem



MaxSAT Encoding

Set of Boolean
Variables

Synthesizing Optimal Locking Scheme

Locking Scheme Requirement

Satisfy correctness requirements
(data race/deadlock freedom & atomicity)



Maximize parallelization
among monitor methods

Optimization problem



MaxSAT Encoding

Set of Boolean
Variables



Set of
Hard Constraints

Synthesizing Optimal Locking Scheme

Locking Scheme Requirement

Satisfy correctness requirements
(data race/deadlock freedom & atomicity)



Maximize parallelization
among monitor methods

Optimization problem



MaxSAT Encoding

Set of Boolean
Variables



Set of
Hard Constraints

All Must be
Satisfied

Synthesizing Optimal Locking Scheme

Locking Scheme Requirement

Satisfy correctness requirements
(data race/deadlock freedom & atomicity)



Maximize parallelization
among monitor methods

Optimization problem



MaxSAT Encoding

Set of Boolean
Variables



Set of
Hard Constraints

All Must be
Satisfied



Set of
Soft Constraints

Synthesizing Optimal Locking Scheme

Locking Scheme Requirement

Satisfy correctness requirements
(data race/deadlock freedom & atomicity)



Maximize parallelization
among monitor methods

Optimization problem



MaxSAT Encoding

Set of Boolean
Variables



Set of
Hard Constraints

All Must be
Satisfied



Set of
Soft Constraints

Max set that **can**
be satisfied

Synthesizing Optimal Locking Scheme

Locking Scheme Requirement

Satisfy correctness requirements
(data race/deadlock freedom & atomicity)



Maximize parallelization
among monitor methods

Optimization problem



MaxSAT Encoding

Map Scheme to
Boolean variables

Set of Boolean
Variables



Set of
Hard Constraints

All Must be
Satisfied



Set of
Soft Constraints

Max set that **can**
be satisfied

Synthesizing Optimal Locking Scheme

Locking Scheme Requirement

Satisfy correctness requirements
(data race/deadlock freedom & atomicity)



Maximize parallelization
among monitor methods

Optimization problem

Map Scheme to
Boolean variables

Set of Boolean
Variables



MaxSAT Encoding

Set of
Hard Constraints

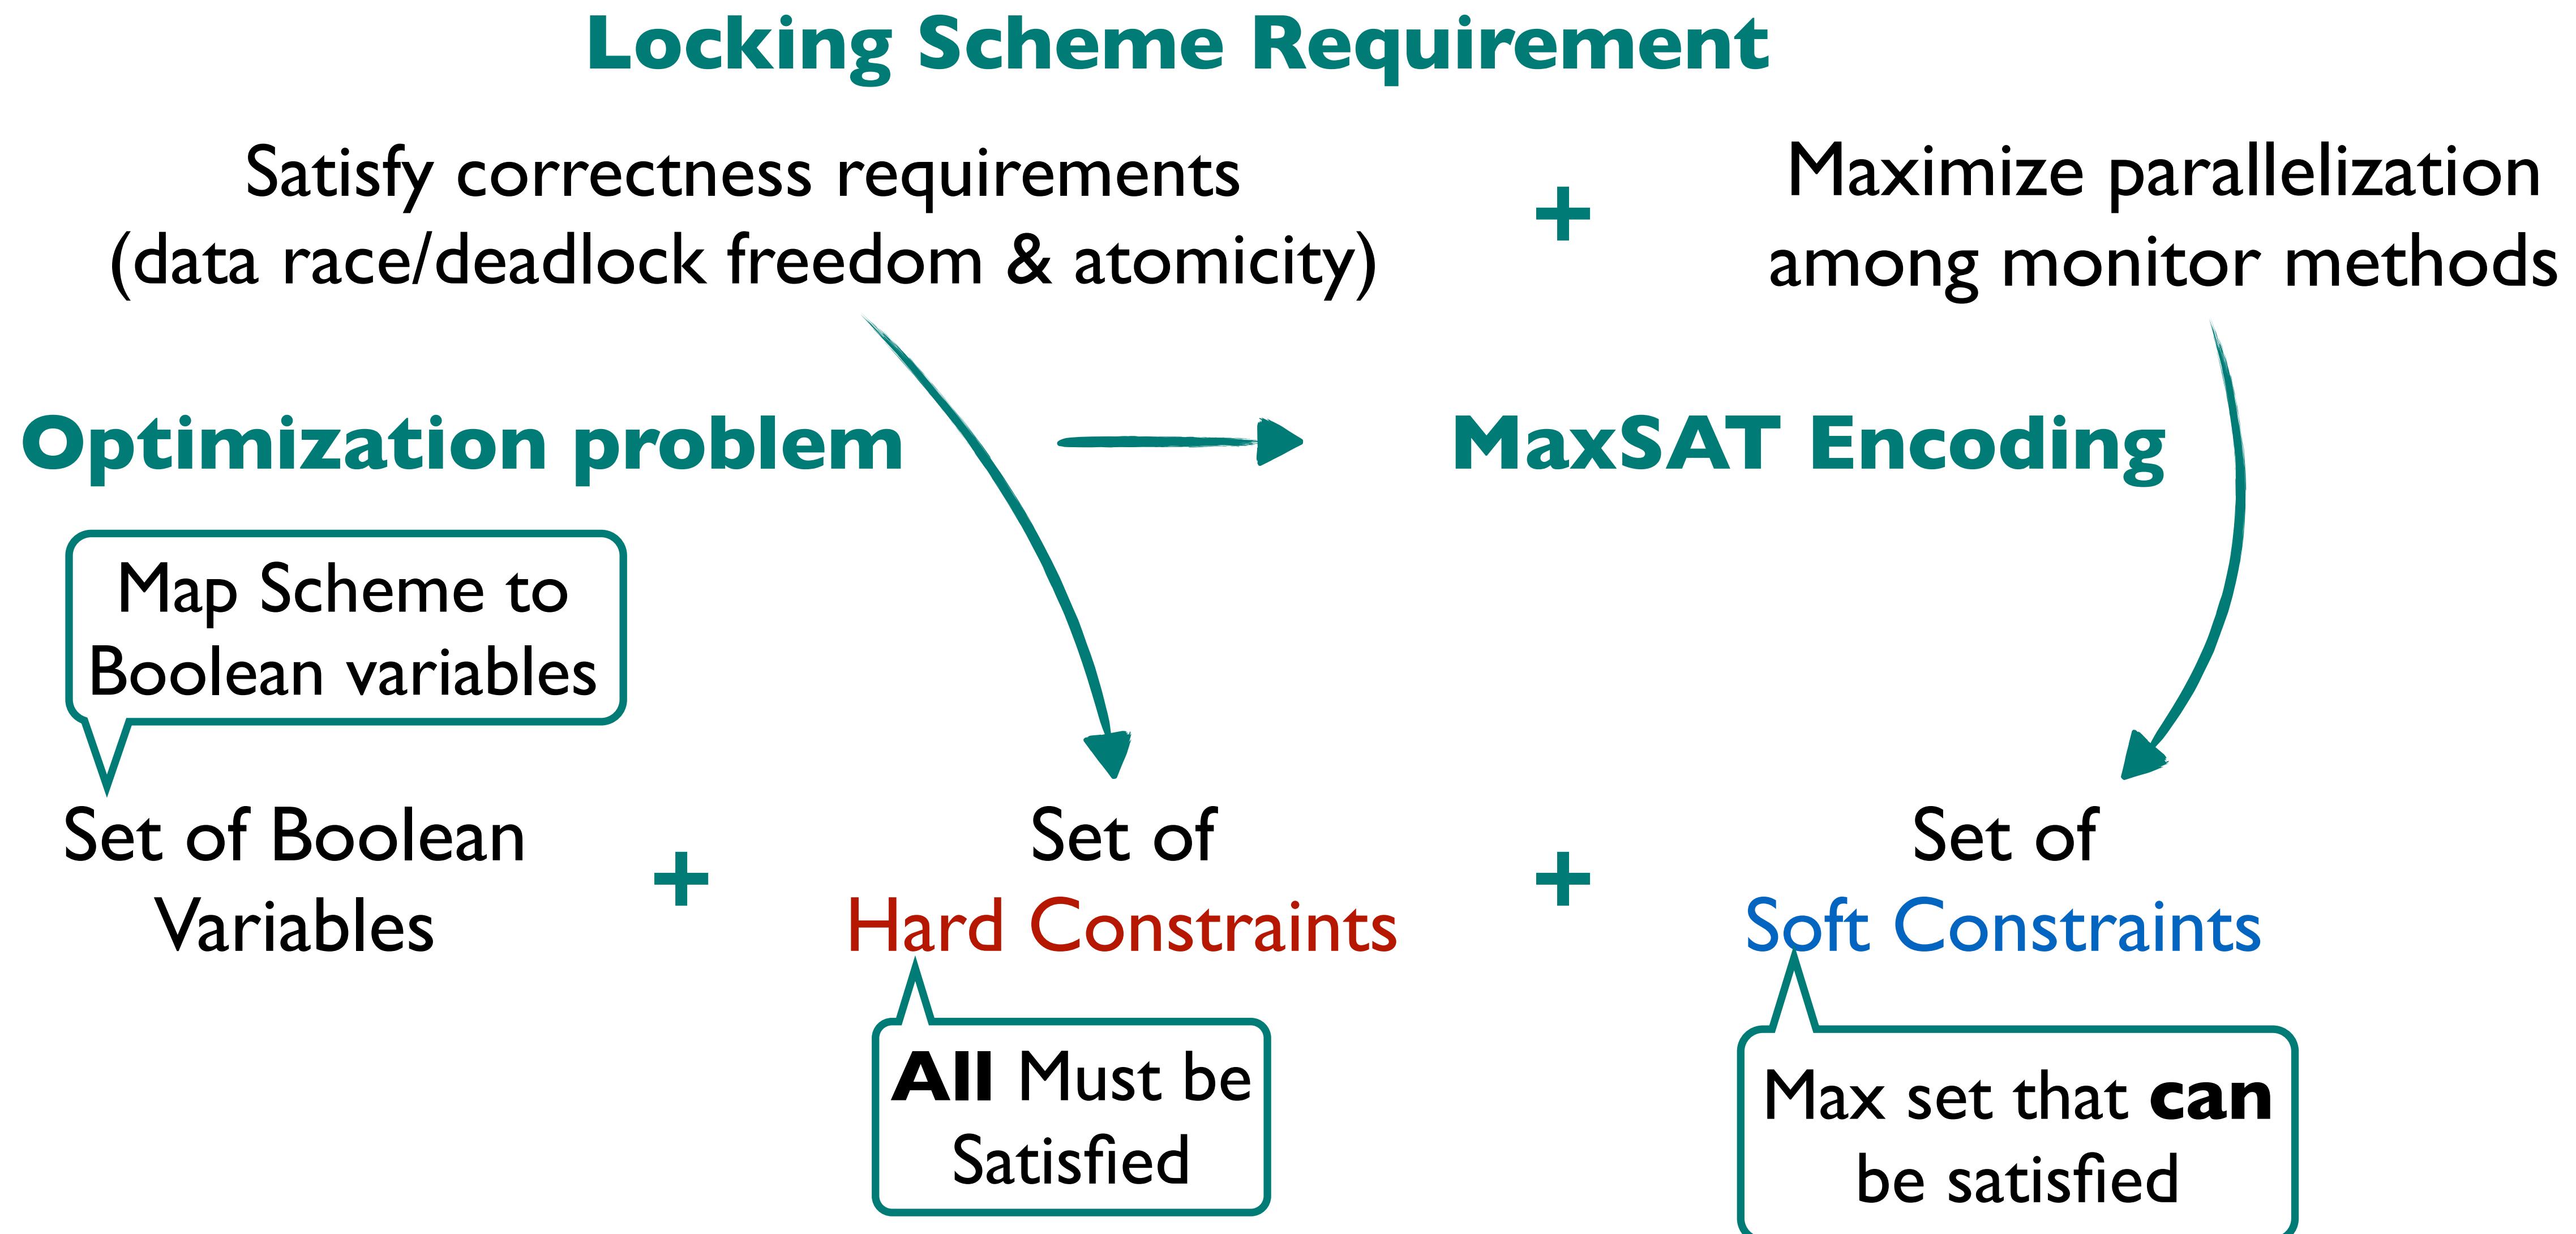
All Must be
Satisfied



Set of
Soft Constraints

Max set that **can**
be satisfied

Synthesizing Optimal Locking Scheme



MaxSAT Encoding Variables



Introduce two types of boolean variables that represent the locking scheme.

MaxSAT Encoding Variables



Introduce two types of boolean variables that represent the locking scheme.

h_s^l

Indicates that code snippet S must hold lock l

$$\mathcal{L}(S) = \{l \mid h_s^l\}$$

Lockset for snippet S

MaxSAT Encoding Variables



Introduce two types of boolean variables that represent the locking scheme.

h_s^l

Indicates that code snippet S must hold lock l

$$\mathcal{L}(S) = \{l \mid h_s^l\}$$

Lockset for
snippet S

α_f

Indicates that field f must be implemented as atomic

$$f \in PossibleAtomicFields(M)$$

Correctness Requirements (Hard constraints)

Data-race freedom

Atomicity

Correctness Requirements (Hard constraints)

Data-race freedom

Atomicity

Resolve data races with locks

If S1 and S2 have a data race,
then generate constraint:

$$\mathcal{L}(S1) \cap \mathcal{L}(S2) \neq \emptyset$$

Correctness Requirements (Hard constraints)

Data-race freedom

Resolve data races with locks

If S1 and S2 have a data race,
then generate constraint:

$$\mathcal{L}(S1) \cap \mathcal{L}(S2) \neq \emptyset$$

Atomicity

Only permit safe interleavings

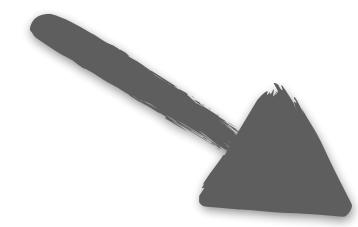
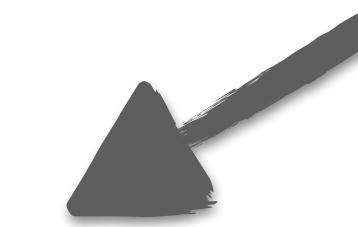
If (S, S1, S2) is not a safe interleaving,
then generate the following constraint

$$\mathcal{L}(S) \cap \mathcal{L}(S_1) \cap \mathcal{L}(S_2) \neq \emptyset$$

Performance Objectives (Soft constraints)



Encode performance objective
using soft constraints



Performance Objectives (Soft constraints)



Encode performance objective
using soft constraints

Maximize
number of code
snippets that can
run in parallel

Performance Objectives (Soft constraints)



Encode performance objective
using soft constraints

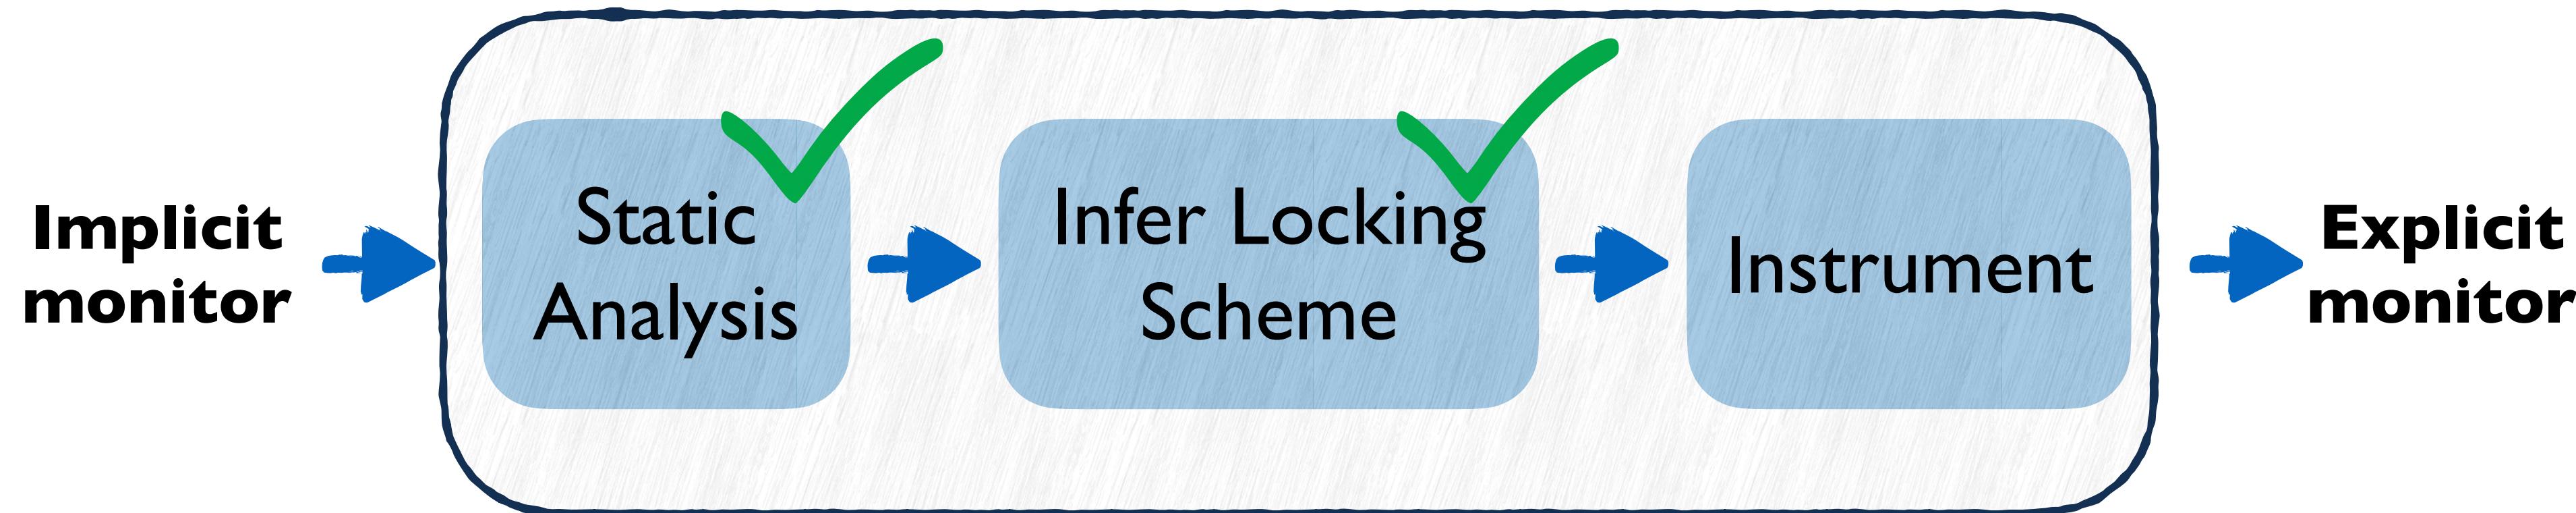
Maximize
number of code
snippets that can
run in parallel

Minimize number
of locks and
atomic fields

Synthesis Framework



Synthesis Framework



Synthesis Framework



Key Idea: Statically detect parallelism and reduce locking scheme inference to MaxSAT

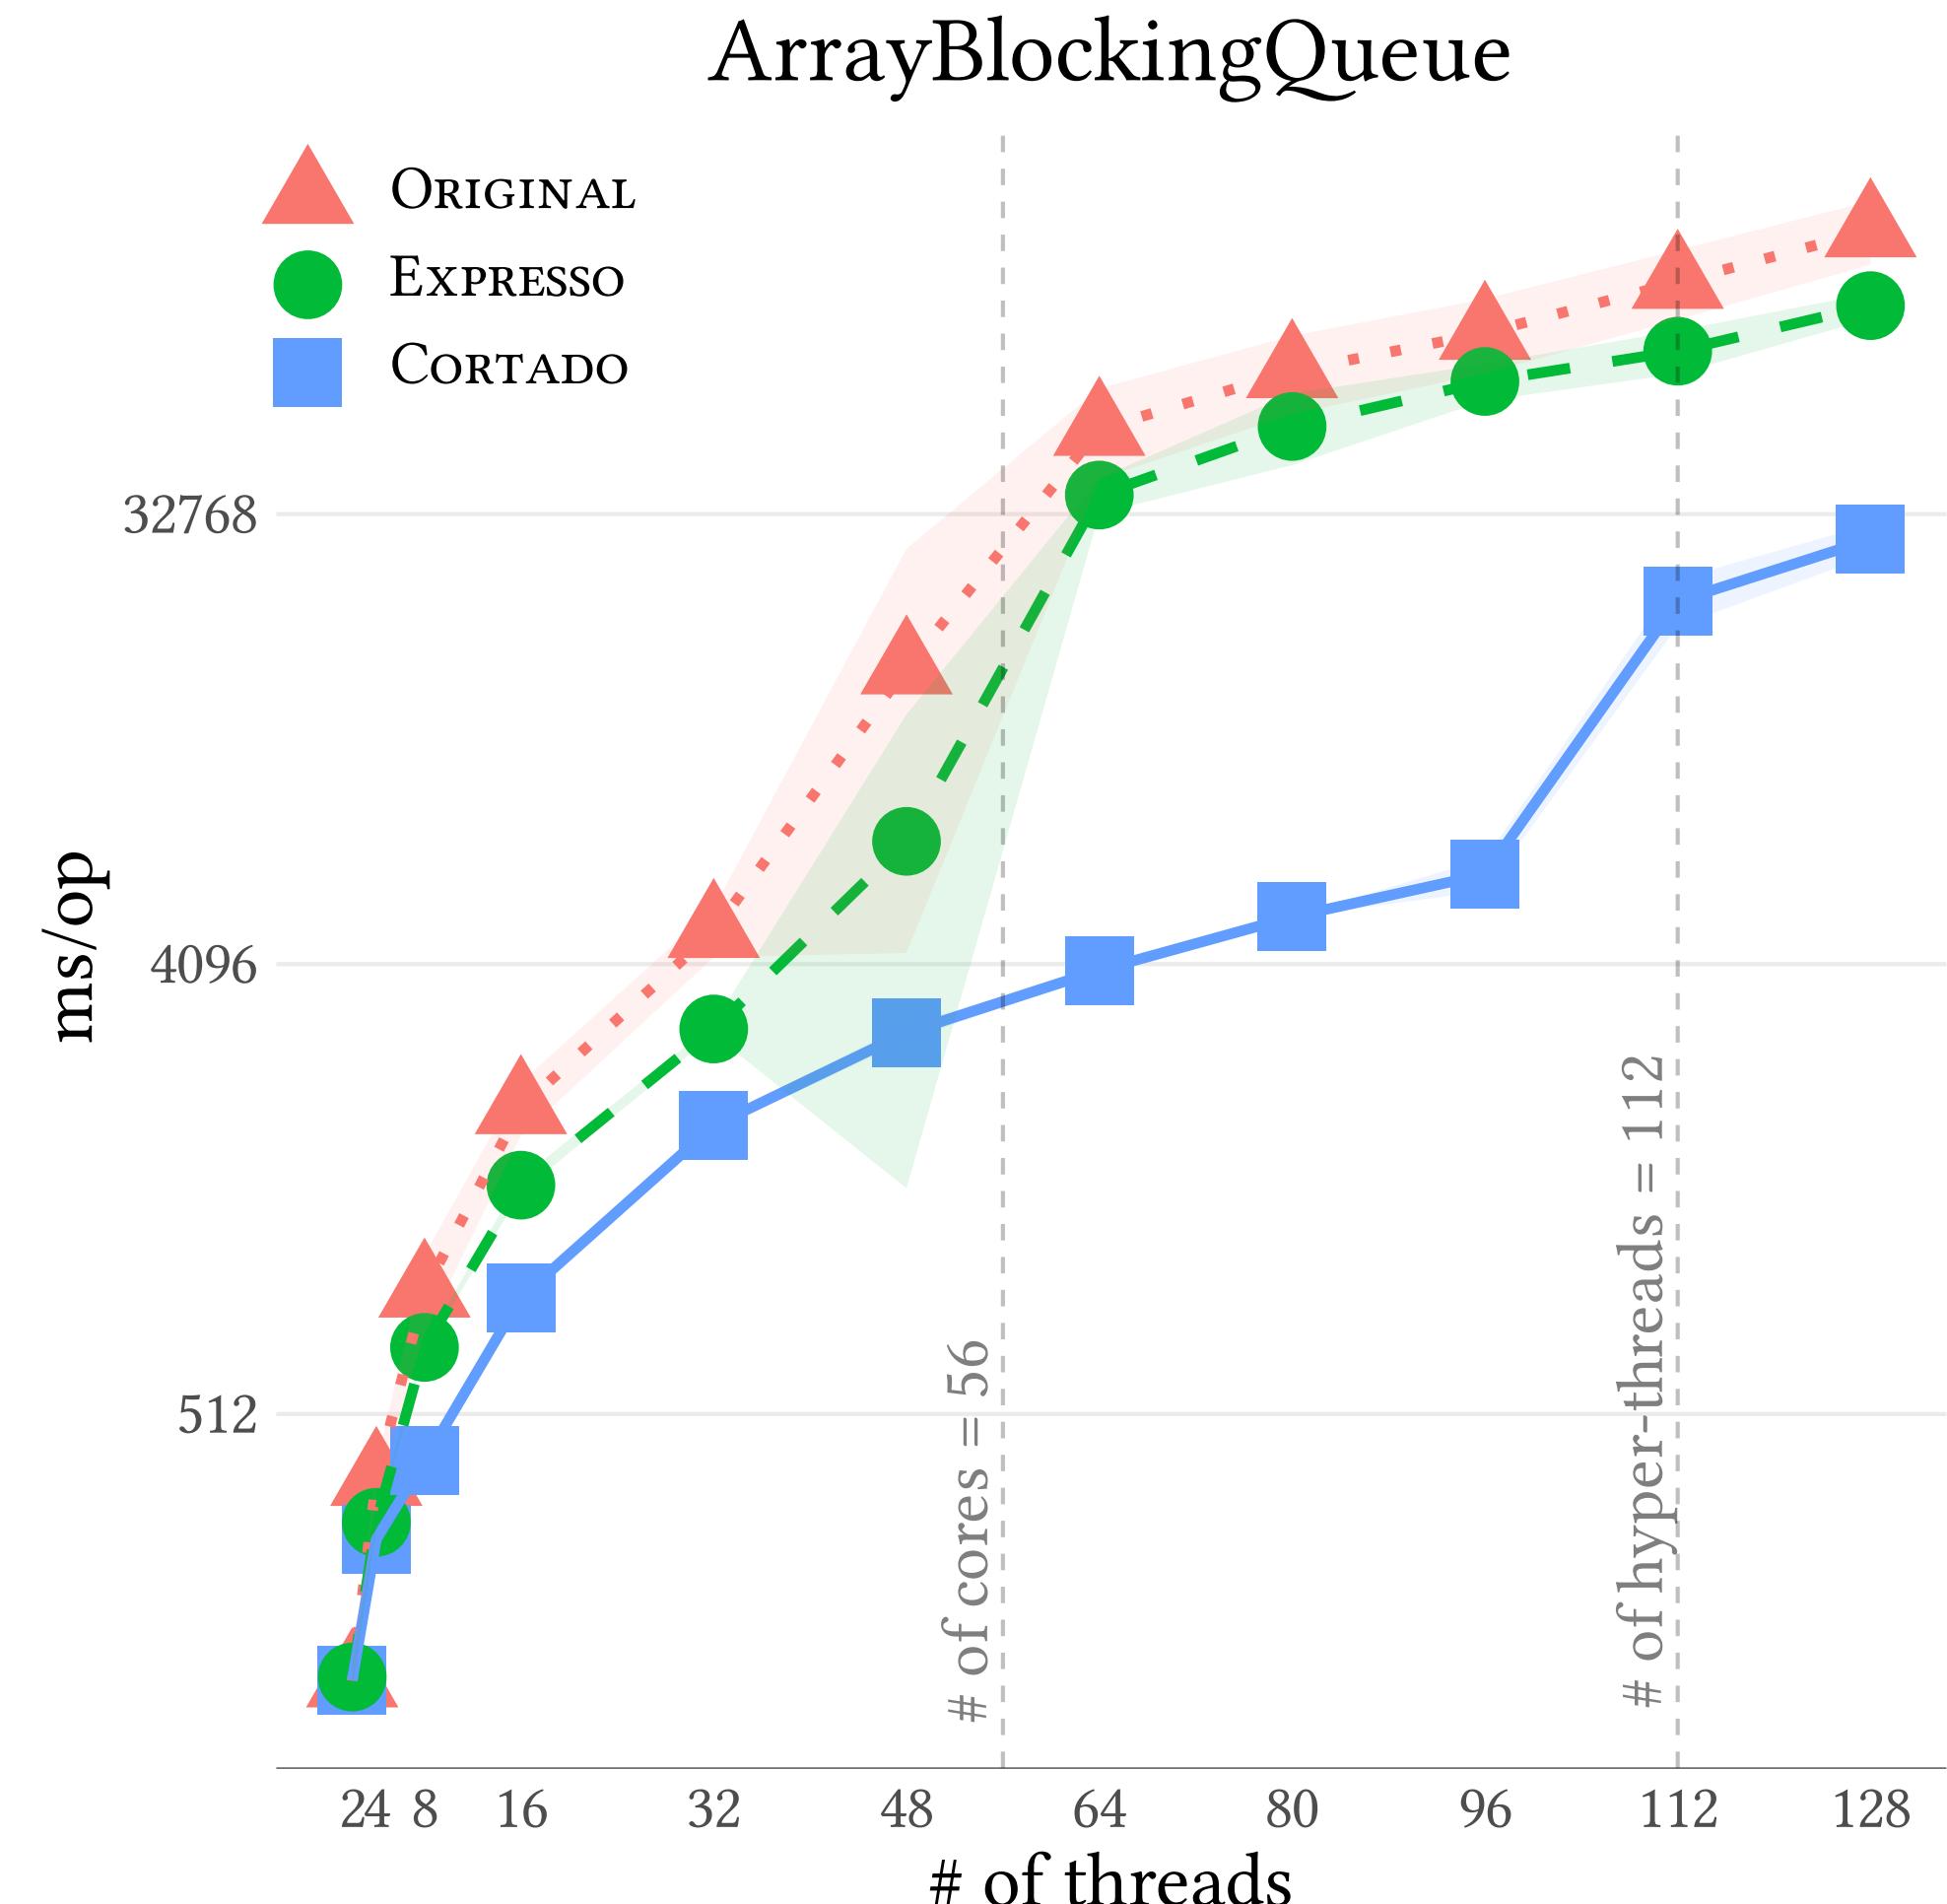


Implementation & Evaluation

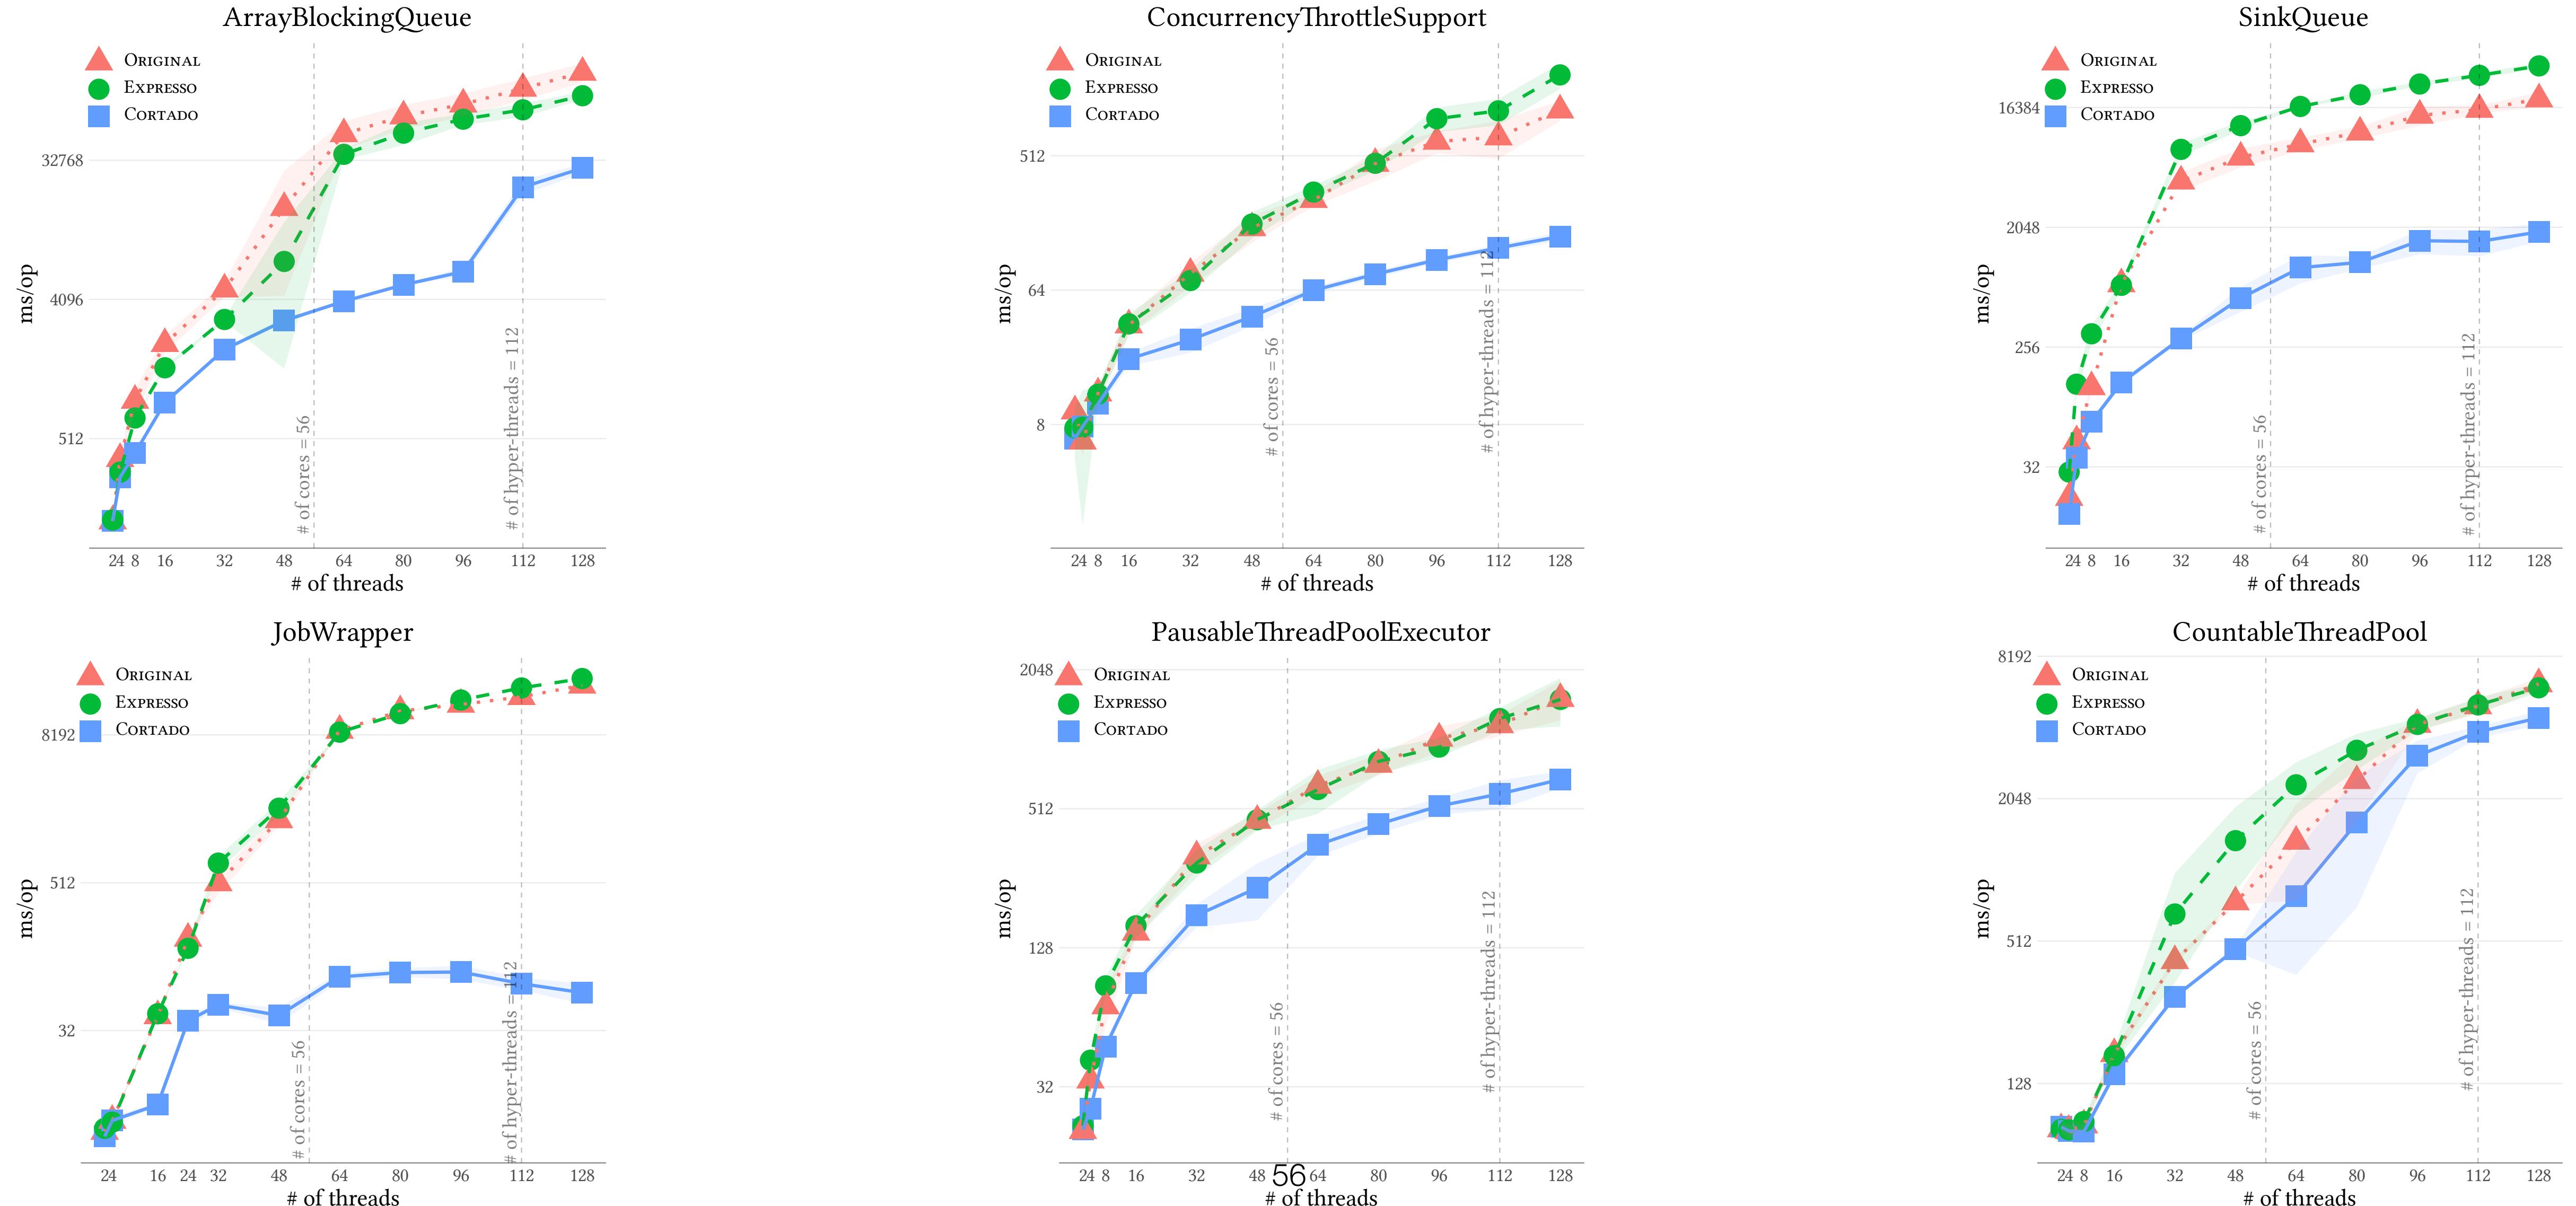


- Evaluated Cortado on monitors with parallelization opportunities
- Manually converted explicit monitor to implicit version
- Compared performance of Cortado-synthesized code against:
 - Manual version as collected from the project
 - The version generated by Espresso (uses a global lock)

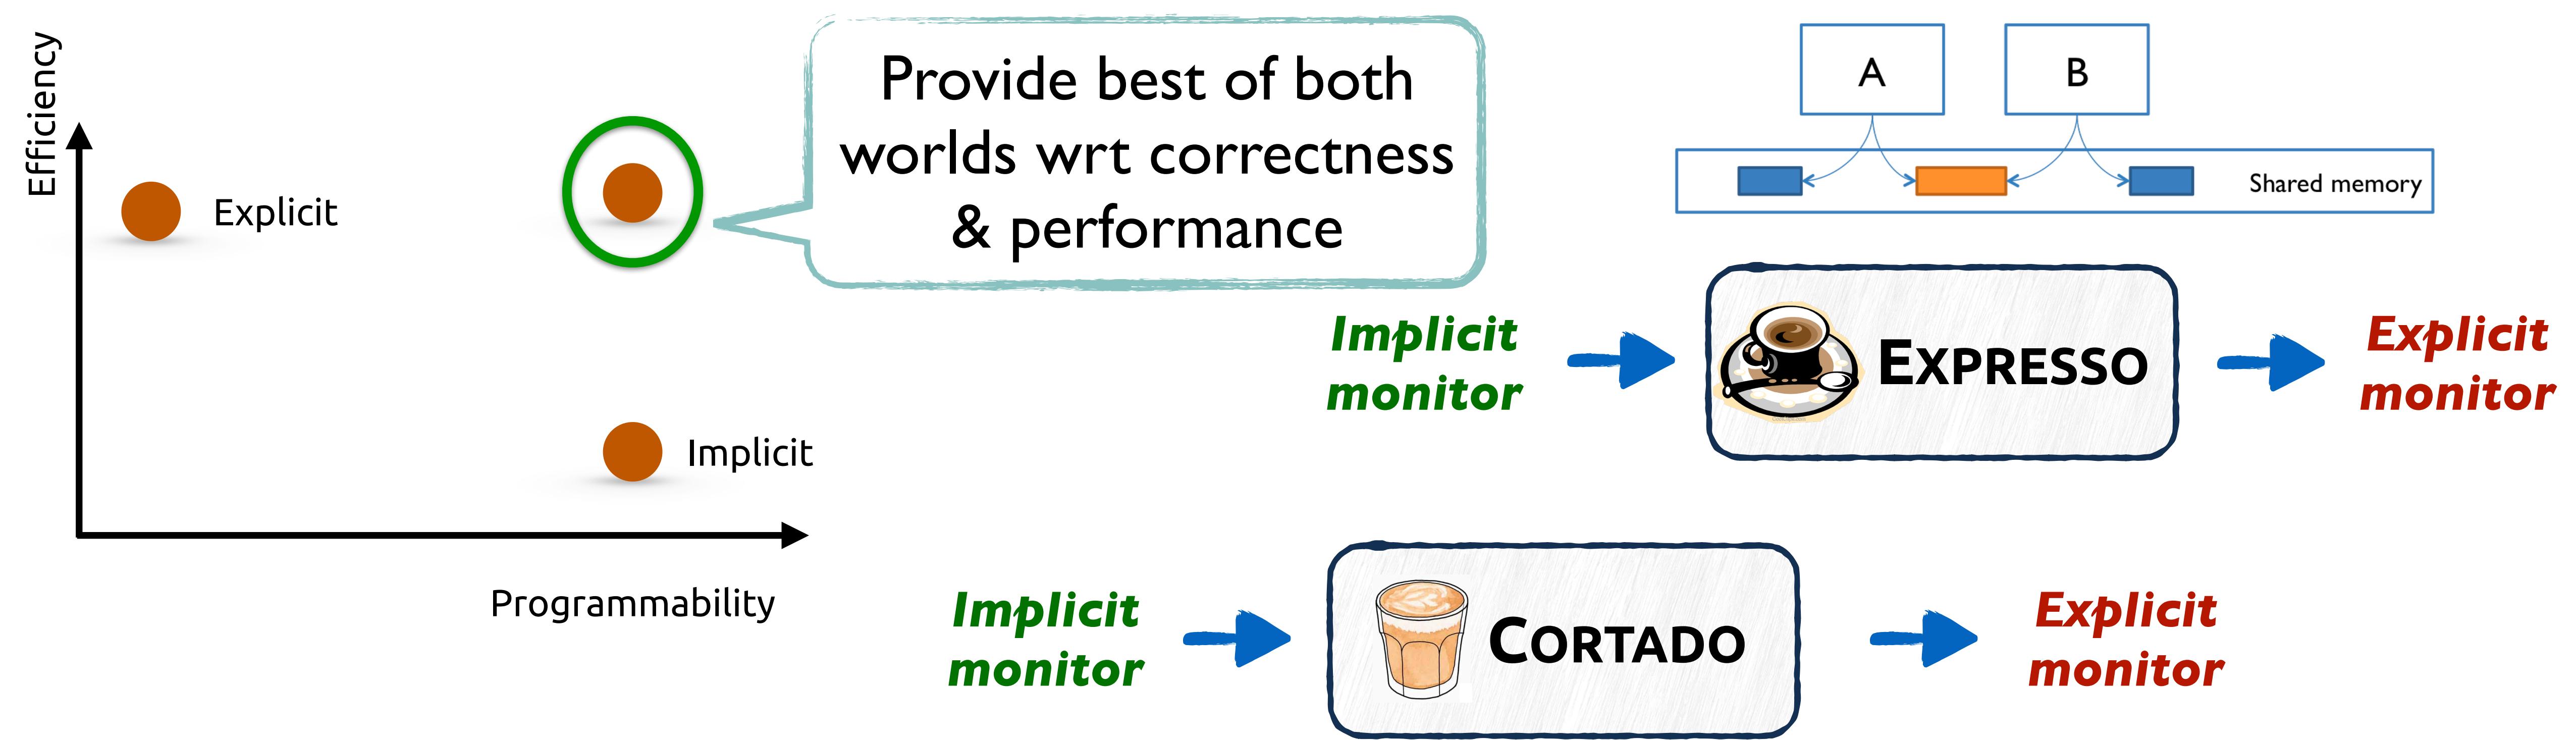
Results



Results



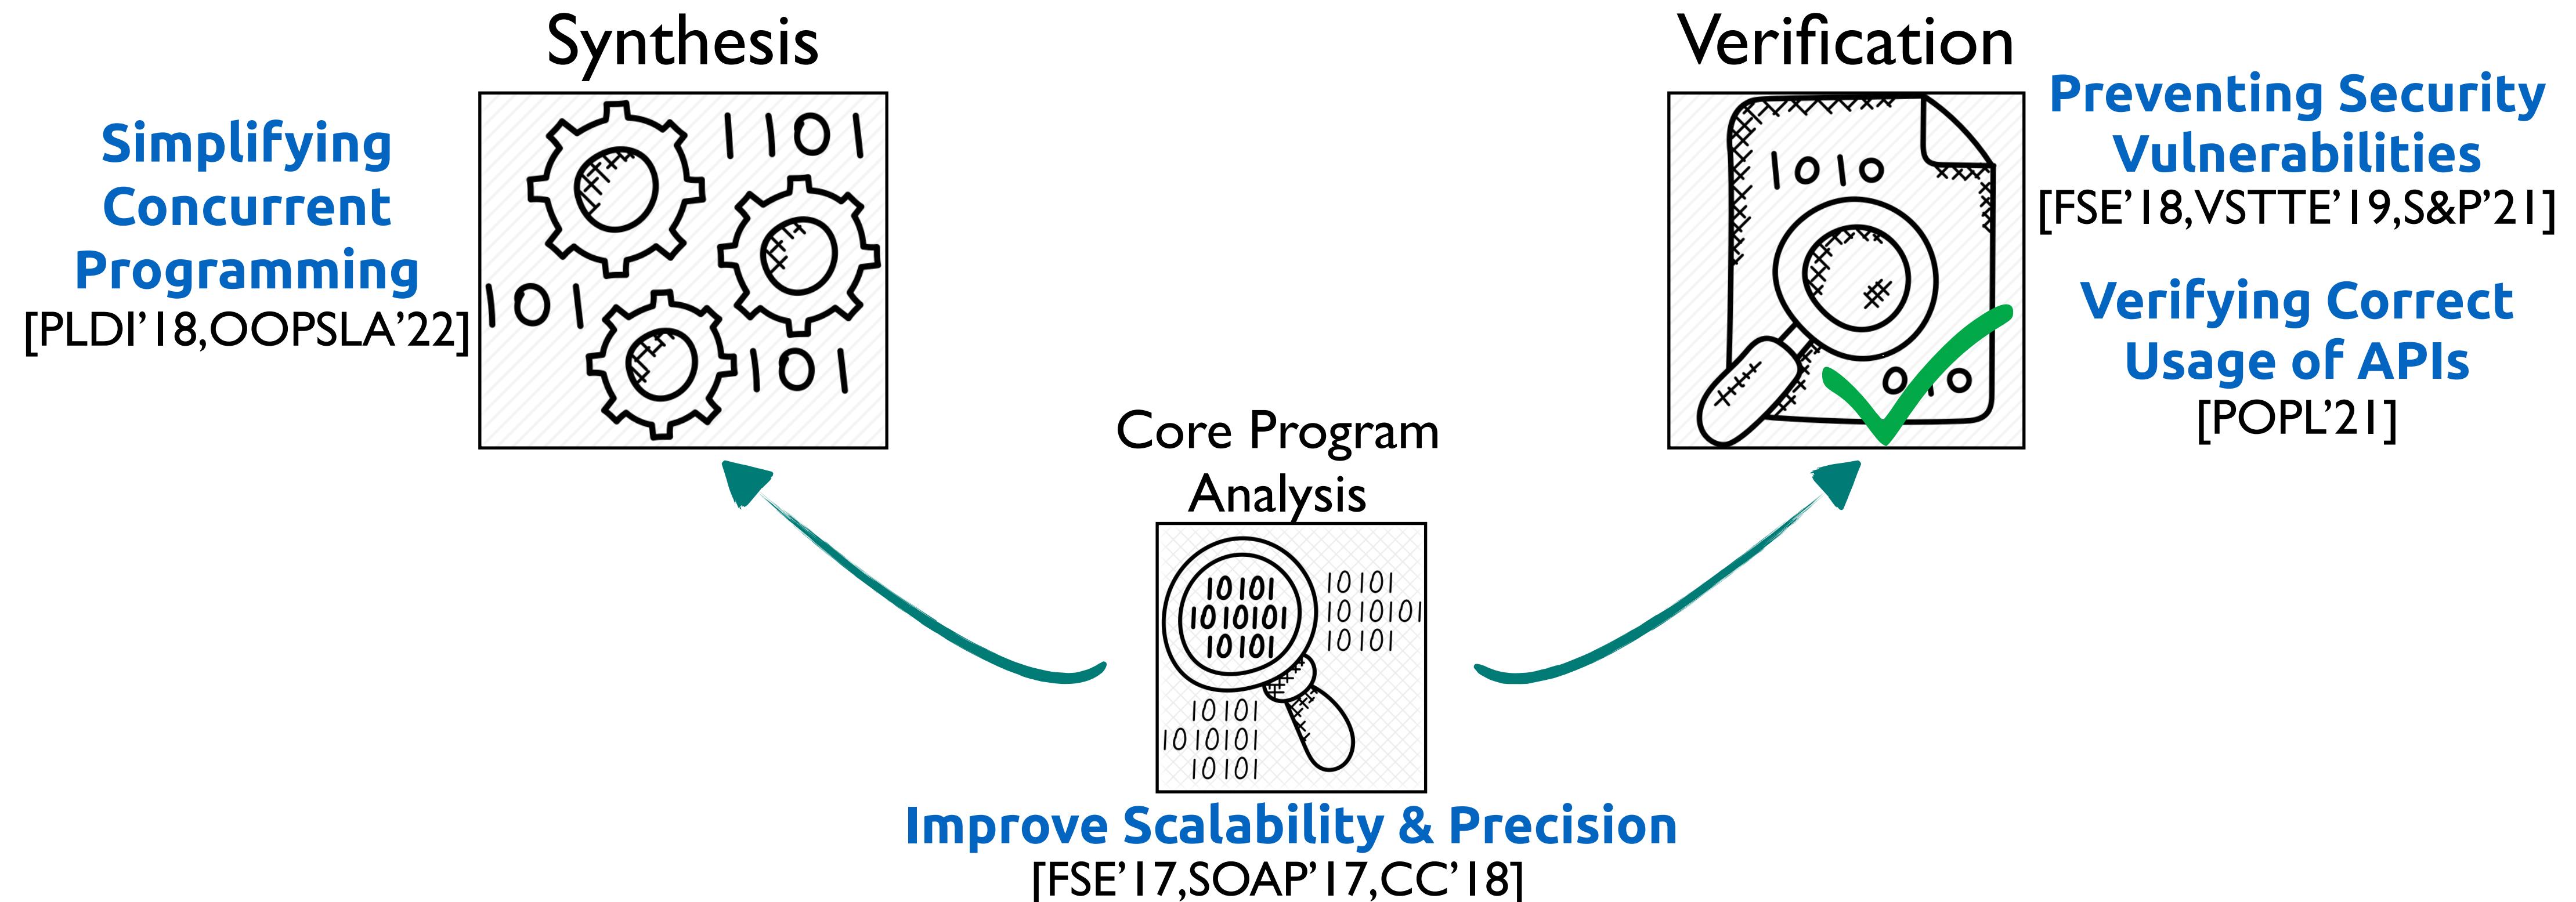
Summary



A set of formal methods techniques for simplifying concurrent programming

Research Highlights

Research Goal: *Formal methods for correct, efficient, and secure code!*

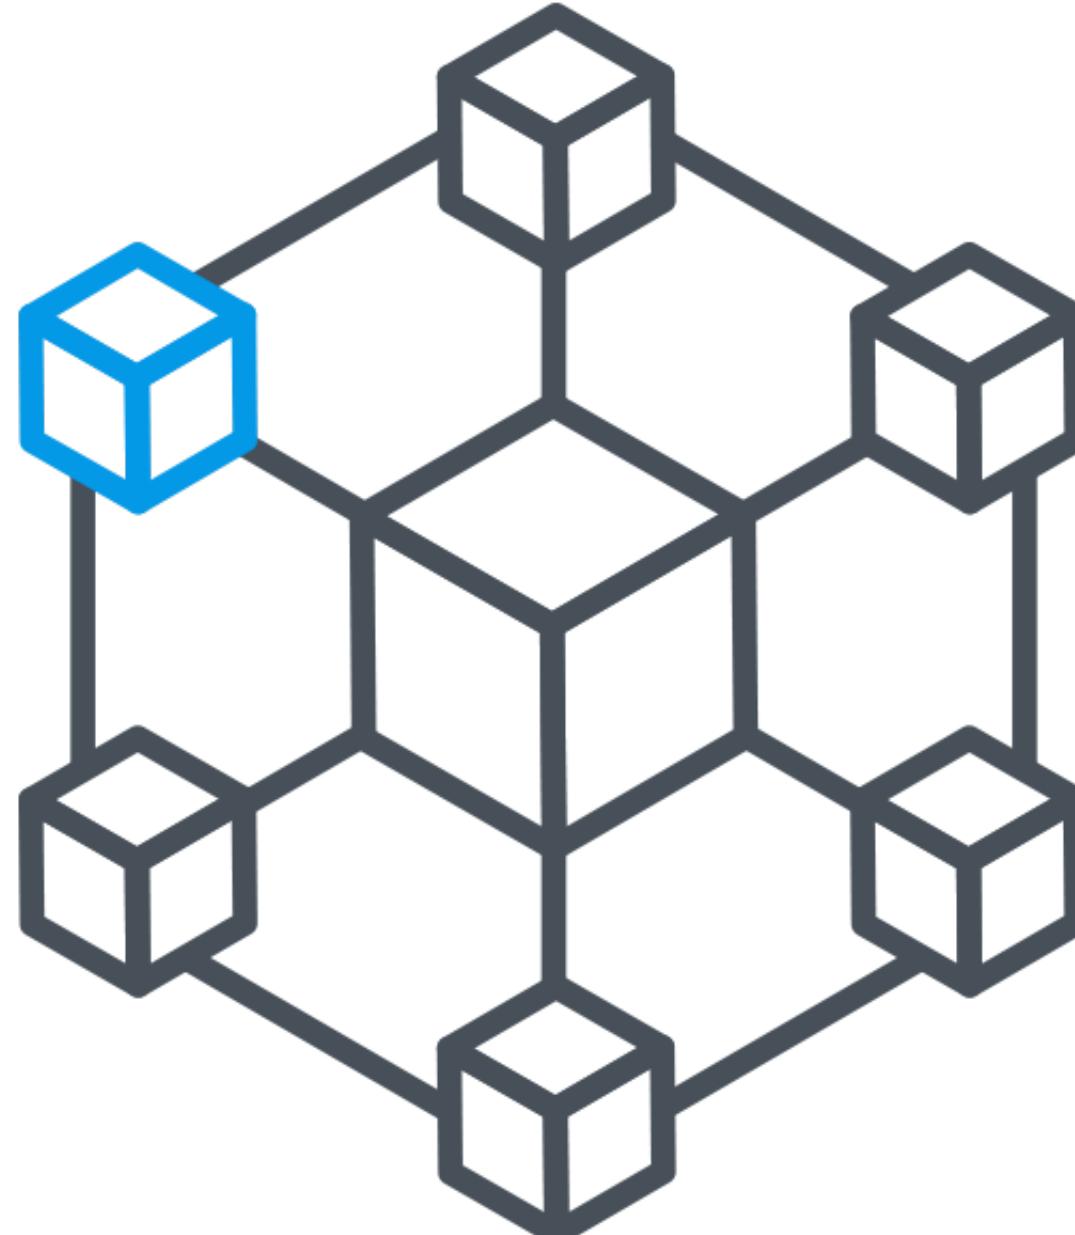


Verification Techniques

Preventing Security Vulnerabilities In Smart Contracts



**Programs running on
a decentralized ledger**



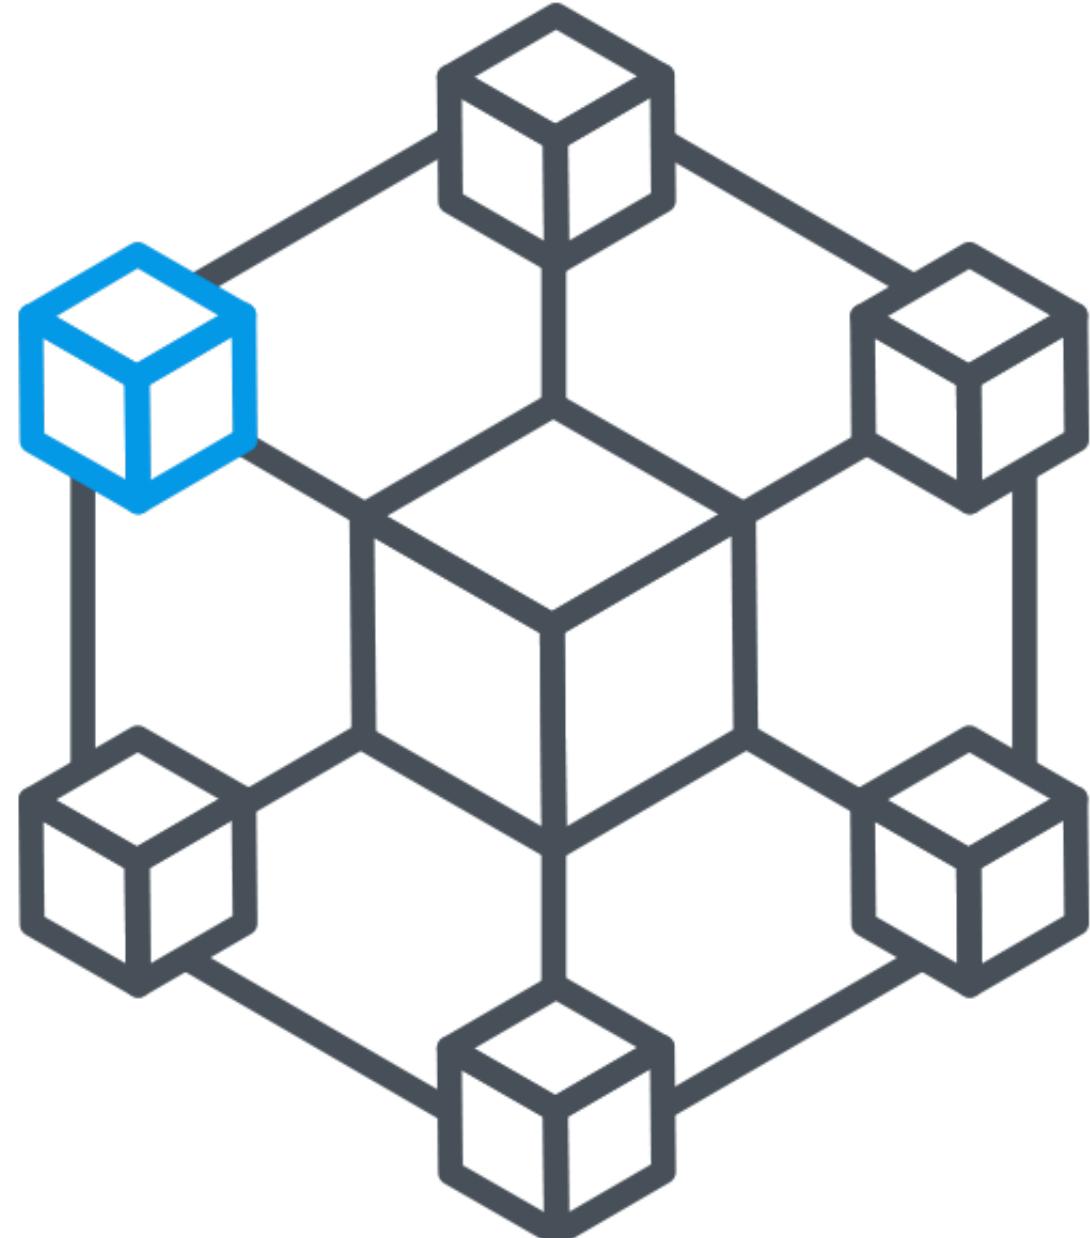
- *Motivation:* They manage a huge mount of user monetary funds
- Vulnerabilities in a contract can severely harm stakeholders

Verification Techniques

Preventing Security Vulnerabilities In Smart Contracts



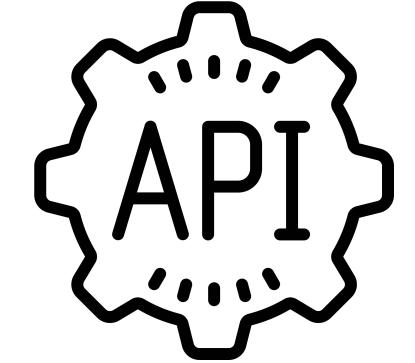
Programs running on
a decentralized ledger



- Motivation: They manage a huge mount of user monetary funds
- Vulnerabilities in a contract can severely harm stakeholders
- Verification to the rescue!
 - **Verisol[VSTTE'19]**: contract won't reach an undesired state
 - **SmartPulse[S&P'21]**: contract will reach a desired state

Verification Techniques

Verifying Correct API Usage

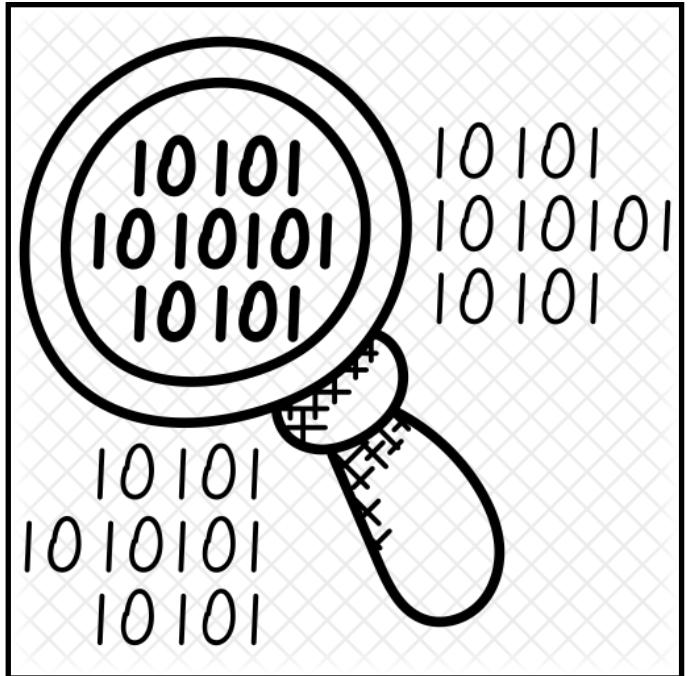


- *Motivation:* Developers tend to misuse APIs written by others
- Prior work cannot handle important classes of APIs

Reentrant lock API: `l.acquire()n l.release()n`

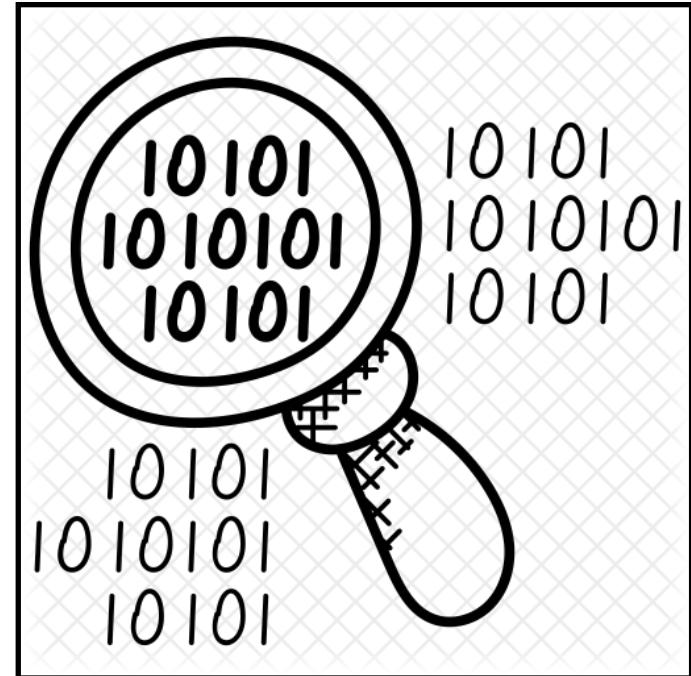
- **CFPChecker[POPL'21]:** closes the theoretical gap for such APIs

Core Program Analysis

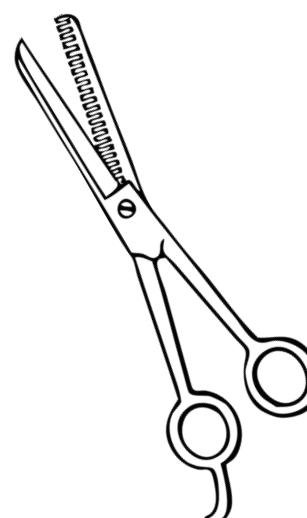


- *Motivation:* Verifiers and Synthesizers need to understand your code!
- Better program analysis tools imply better verifiers and synthesizers

Core Program Analysis



- *Motivation:* Verifiers and Synthesizers need to understand your code!
- Better program analysis tools imply better verifiers and synthesizers
- **Timmer[FSE'17]:** Program transformation that eliminates safe paths
- Transparently increases precision and scalability of other tools



Future Directions



Can we push at the boundaries of formal method applications and further boost developers' confidence in their code?

Direction I

Verification & Synthesis Techniques
in More Contexts

Direction II

Make Formal Method Techniques
More Accessible

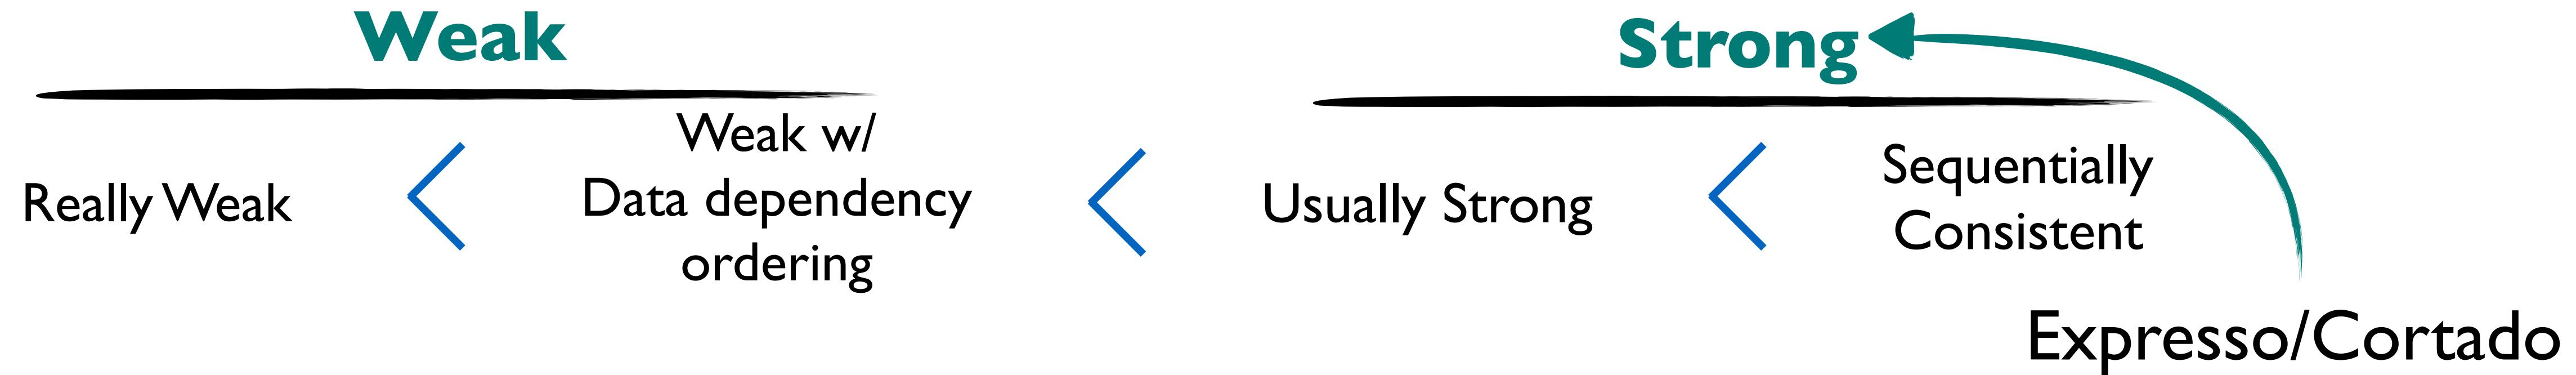
Simplify More Aspects of Shared Memory Concurrency

Several Concurrent Programming Setups Based on Guarantees Provided by Underlying Memory System



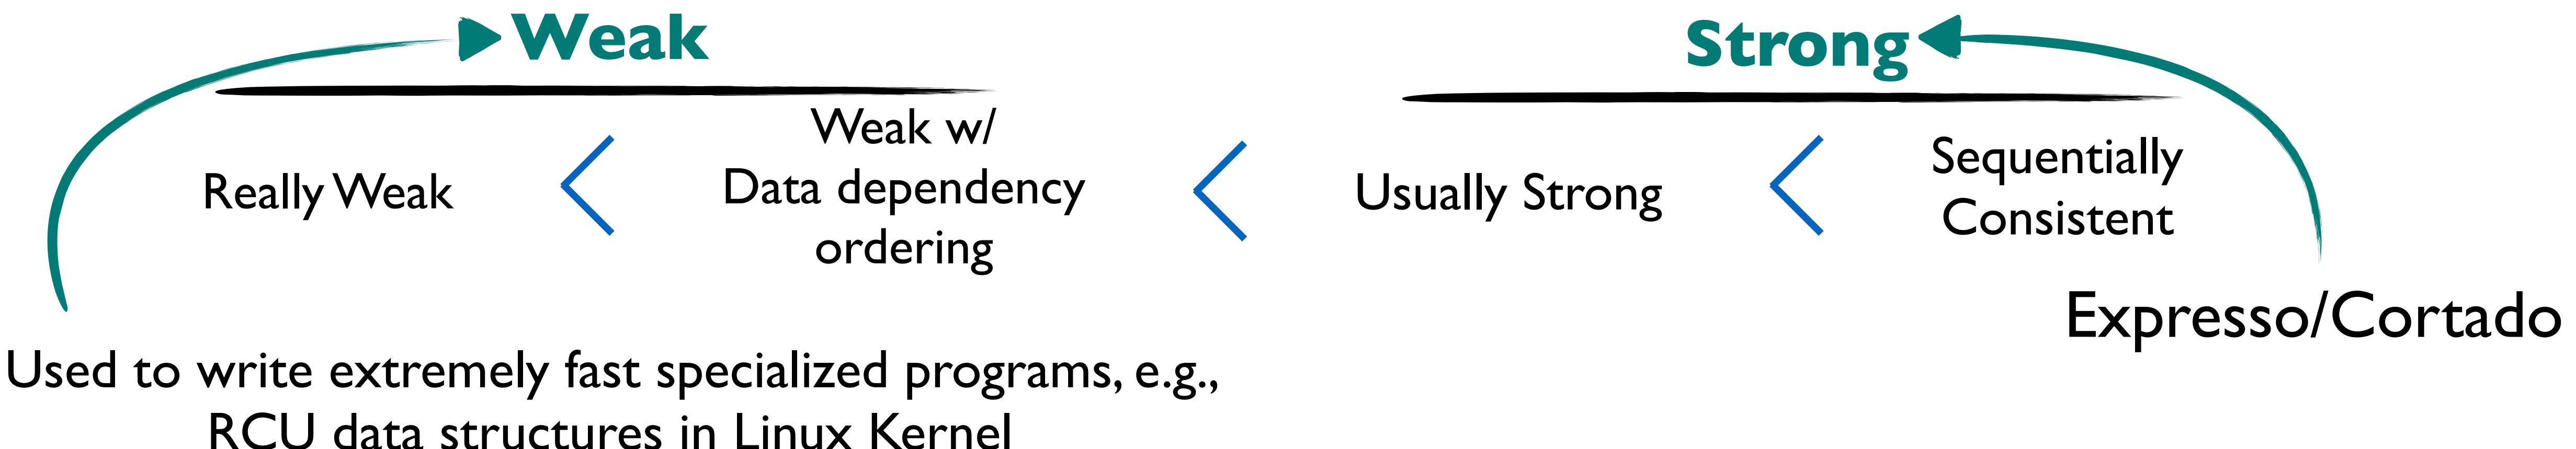
Simplify More Aspects of Shared Memory Concurrency

Several Concurrent Programming Setups Based on Guarantees Provided by Underlying Memory System



Simplify More Aspects of Shared Memory Concurrency

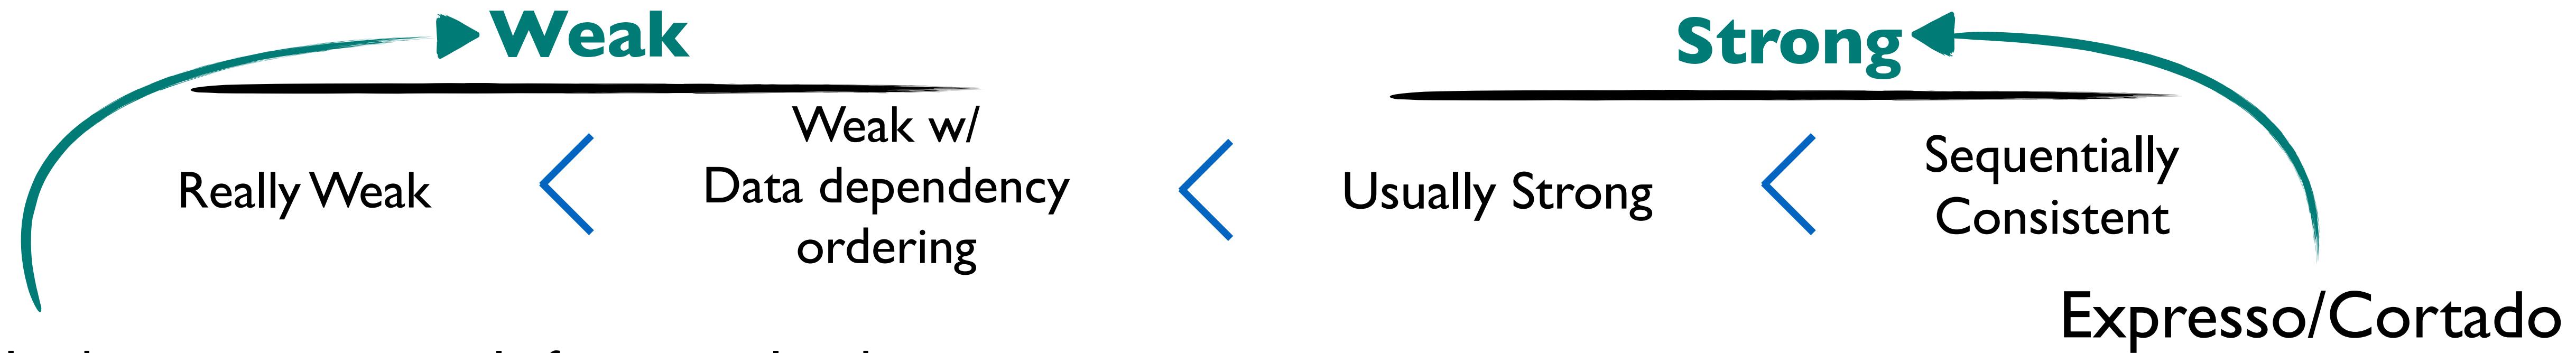
Several Concurrent Programming Setups Based on Guarantees Provided by Underlying Memory System



Used to write extremely fast specialized programs, e.g.,
RCU data structures in Linux Kernel

Simplify More Aspects of Shared Memory Concurrency

Several Concurrent Programming Setups Based on Guarantees Provided by Underlying Memory System



Used to write extremely fast specialized programs, e.g.,
RCU data structures in Linux Kernel

- Challenging to prove correctness of programs written with weak memory primitives
 - Many implementations have manually-written proofs

Simplify More Aspects of Shared Memory Concurrency

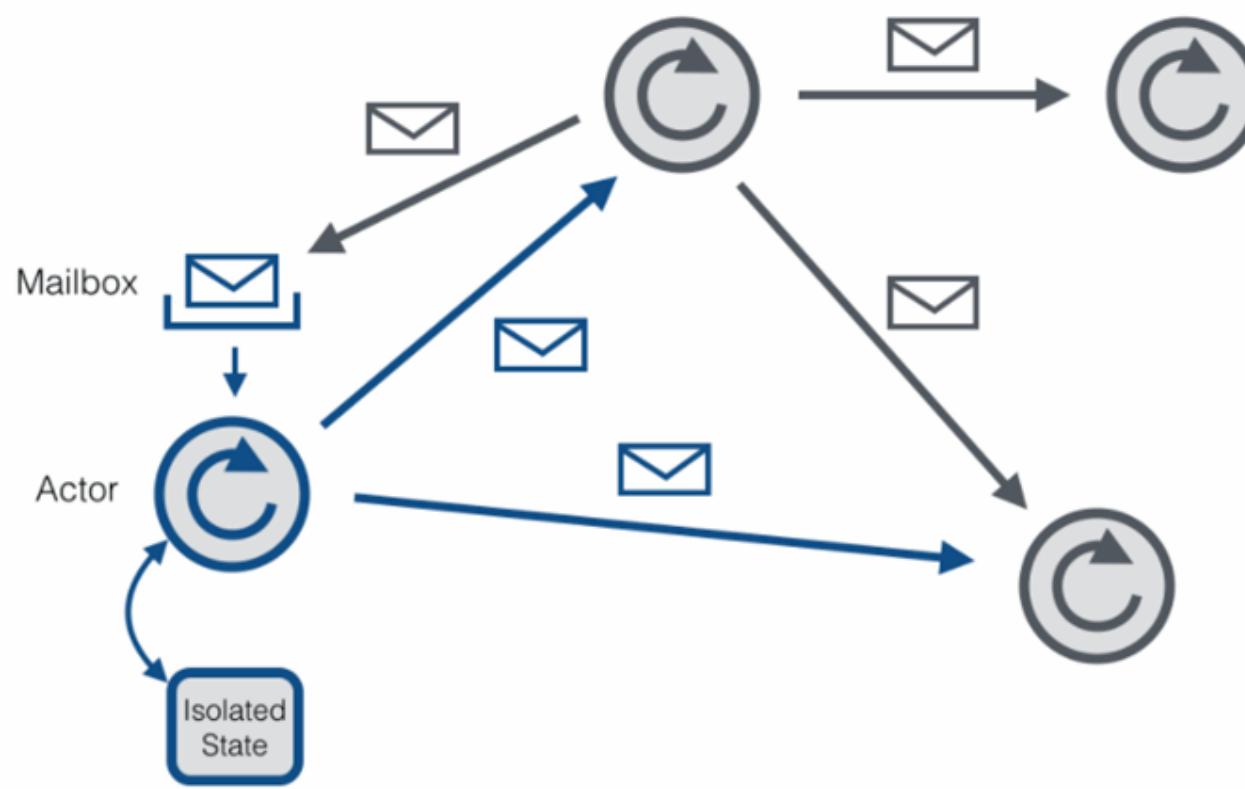
Several Concurrent Programming Setups Based on Guarantees Provided by Underlying Memory System



Used to write extremely fast specialized programs, e.g.,
RCU data structures in Linux Kernel

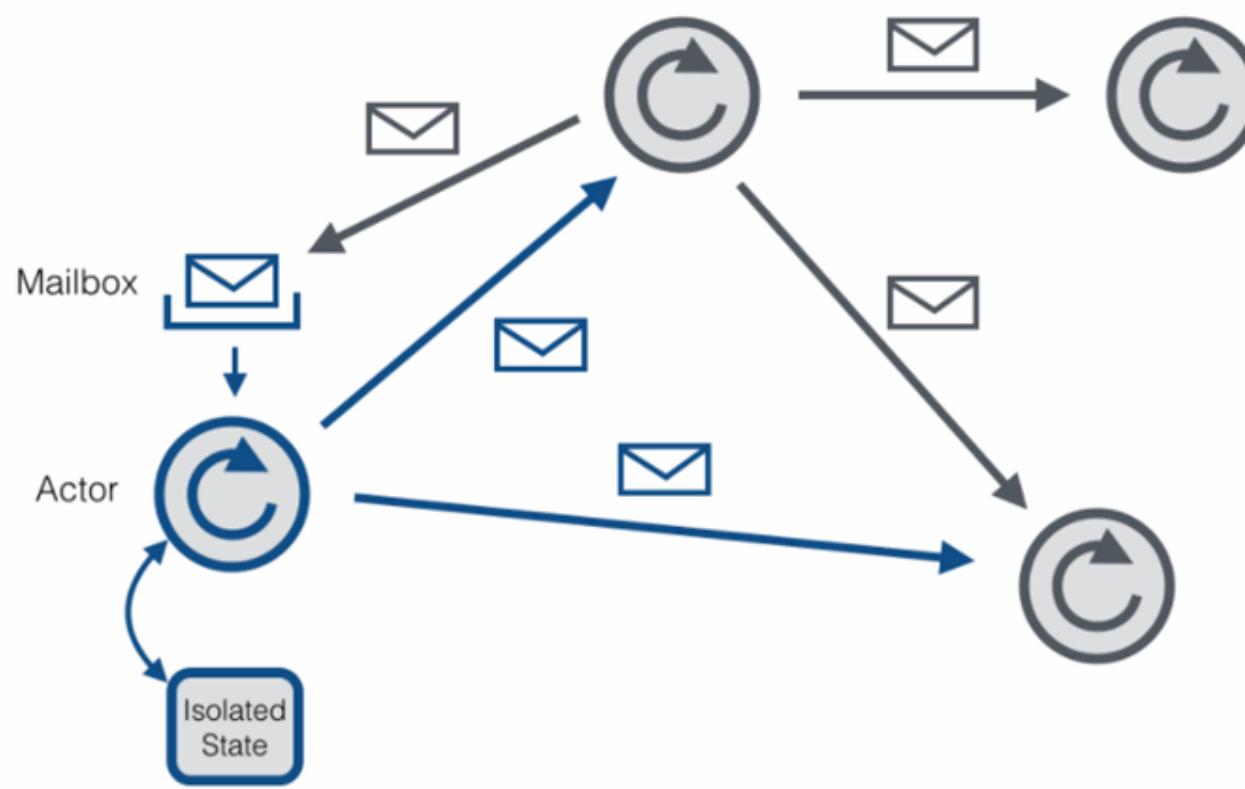
- Challenging to prove correctness of programs written with weak memory primitives
 - Many implementations have manually-written proofs
- **Solutions:**
 1. More automation in verification
 2. Synthesize concurrent programs w/ weak memory primitives

Simplify Other Concurrency Models



- Several other concurrency models based on message passing, e.g., MPI, Actor Model, etc.
- Nodes (i.e., threads) must conform to a communication protocol
 - Developers fail to implement the protocol correctly

Simplify Other Concurrency Models

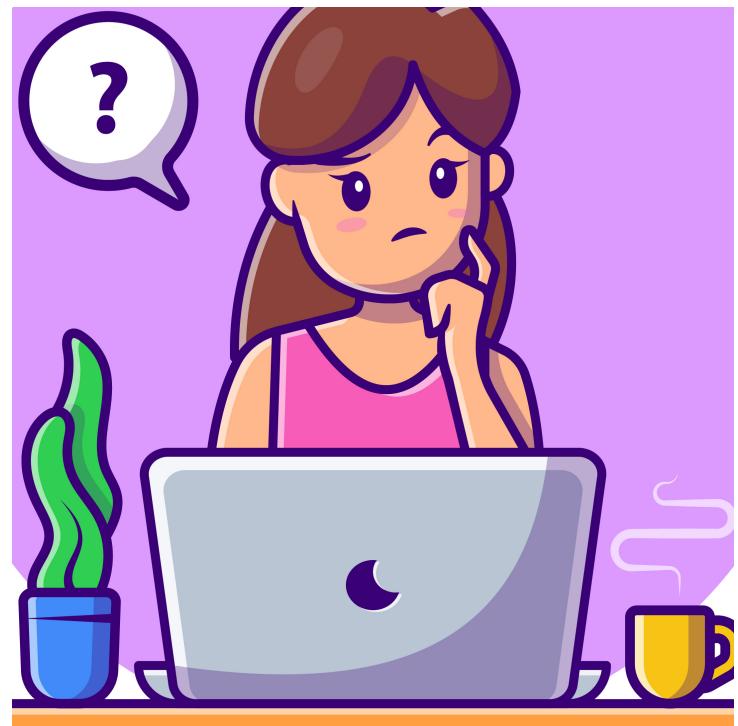


- Several other concurrency models based on message passing, e.g., MPI, Actor Model, etc.
- Nodes (i.e., threads) must conform to a communication protocol
 - Developers fail to implement the protocol correctly



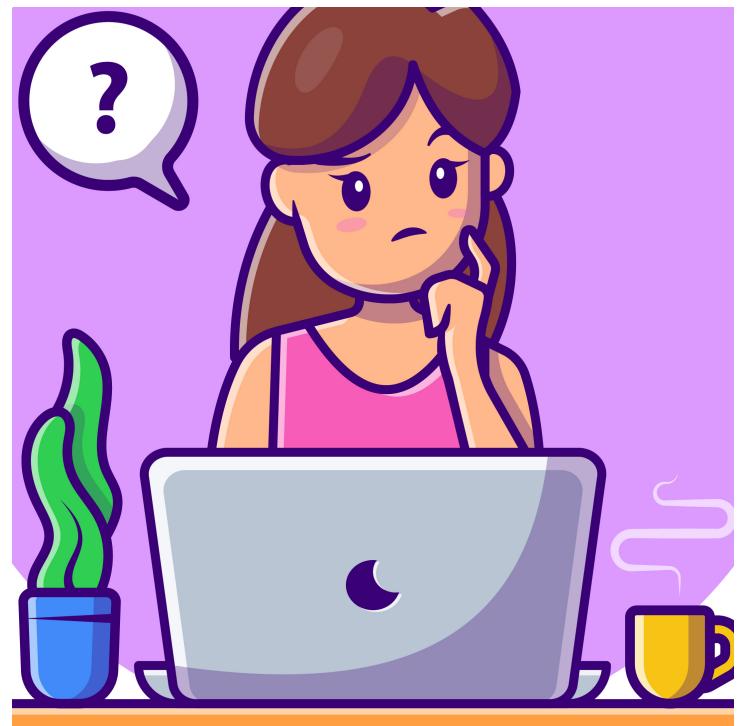
- Automatically verify conformance to protocol given its formal description
- Automatically synthesize message passing programs from their specification

Make Formal Methods More Accessible



- Users must be familiar with complex concepts and write specifications
- Even experts can write buggy specifications

Make Formal Methods More Accessible



- Users must be familiar with complex concepts and write specifications
- Even experts can write buggy specifications



- Develop techniques that help developers fix wrong specifications
- Synthesize specifications from codebase artifacts (e.g., test cases)

Acknowledgements



Isil Dillig



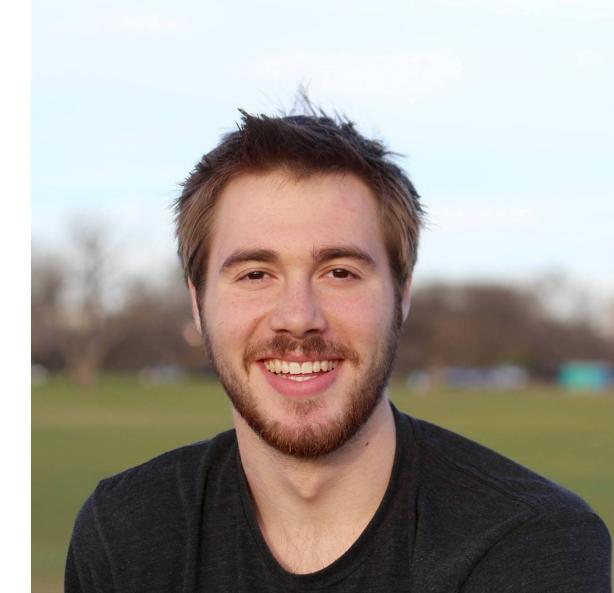
Yannis Smaragdakis



Ben Sepanski



Jon Stephens



Jacob Van Geffen



Yuepeng Wang



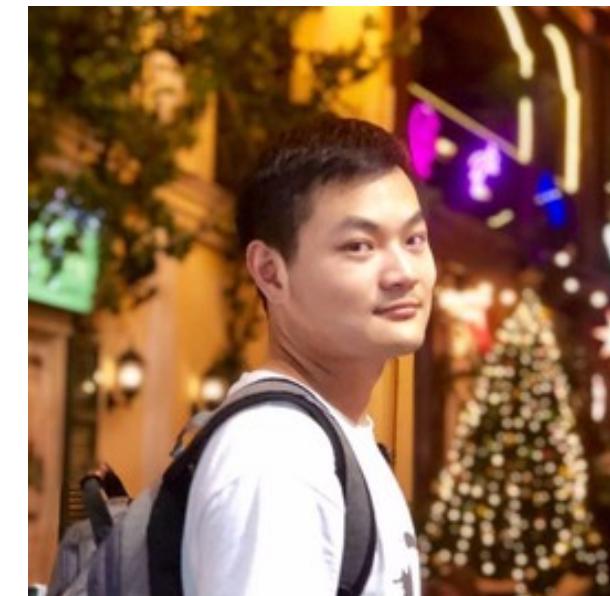
James Bornholt



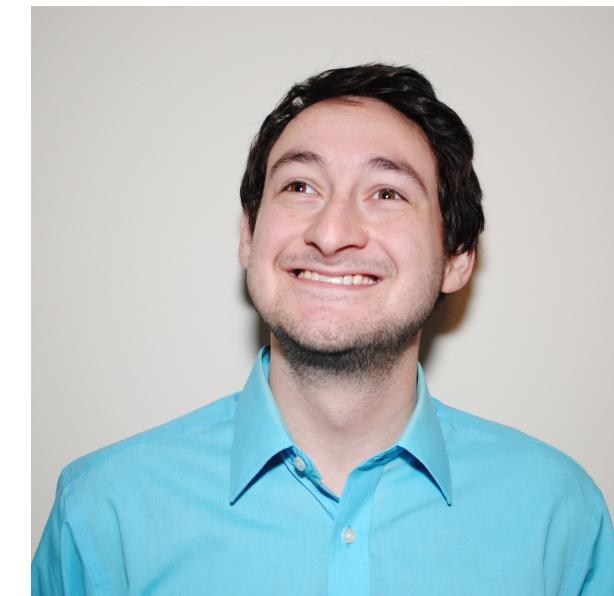
Shuvendu Lahiri



Yu Feng



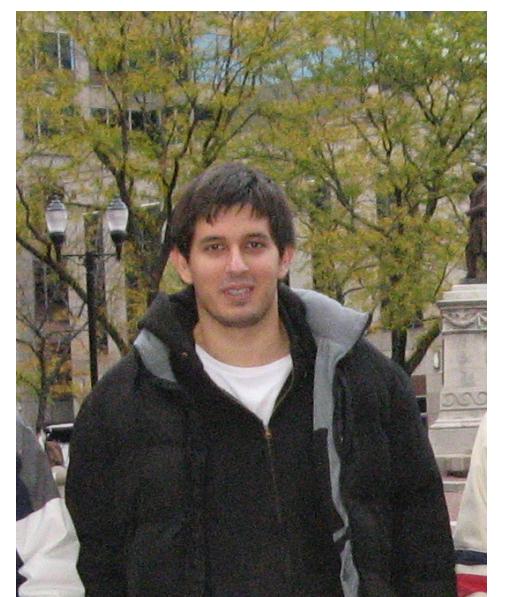
Jiayi Wei



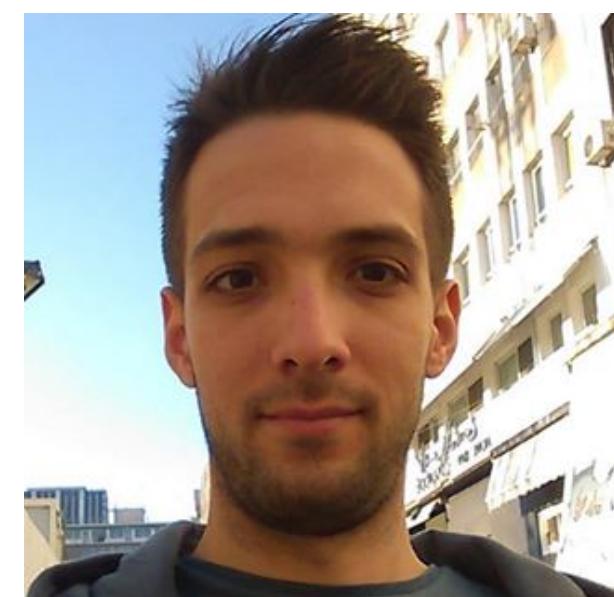
Ben Mariano



Rahul Krishnan



George Balatshouras



George Kastrinis



Maria Christakis

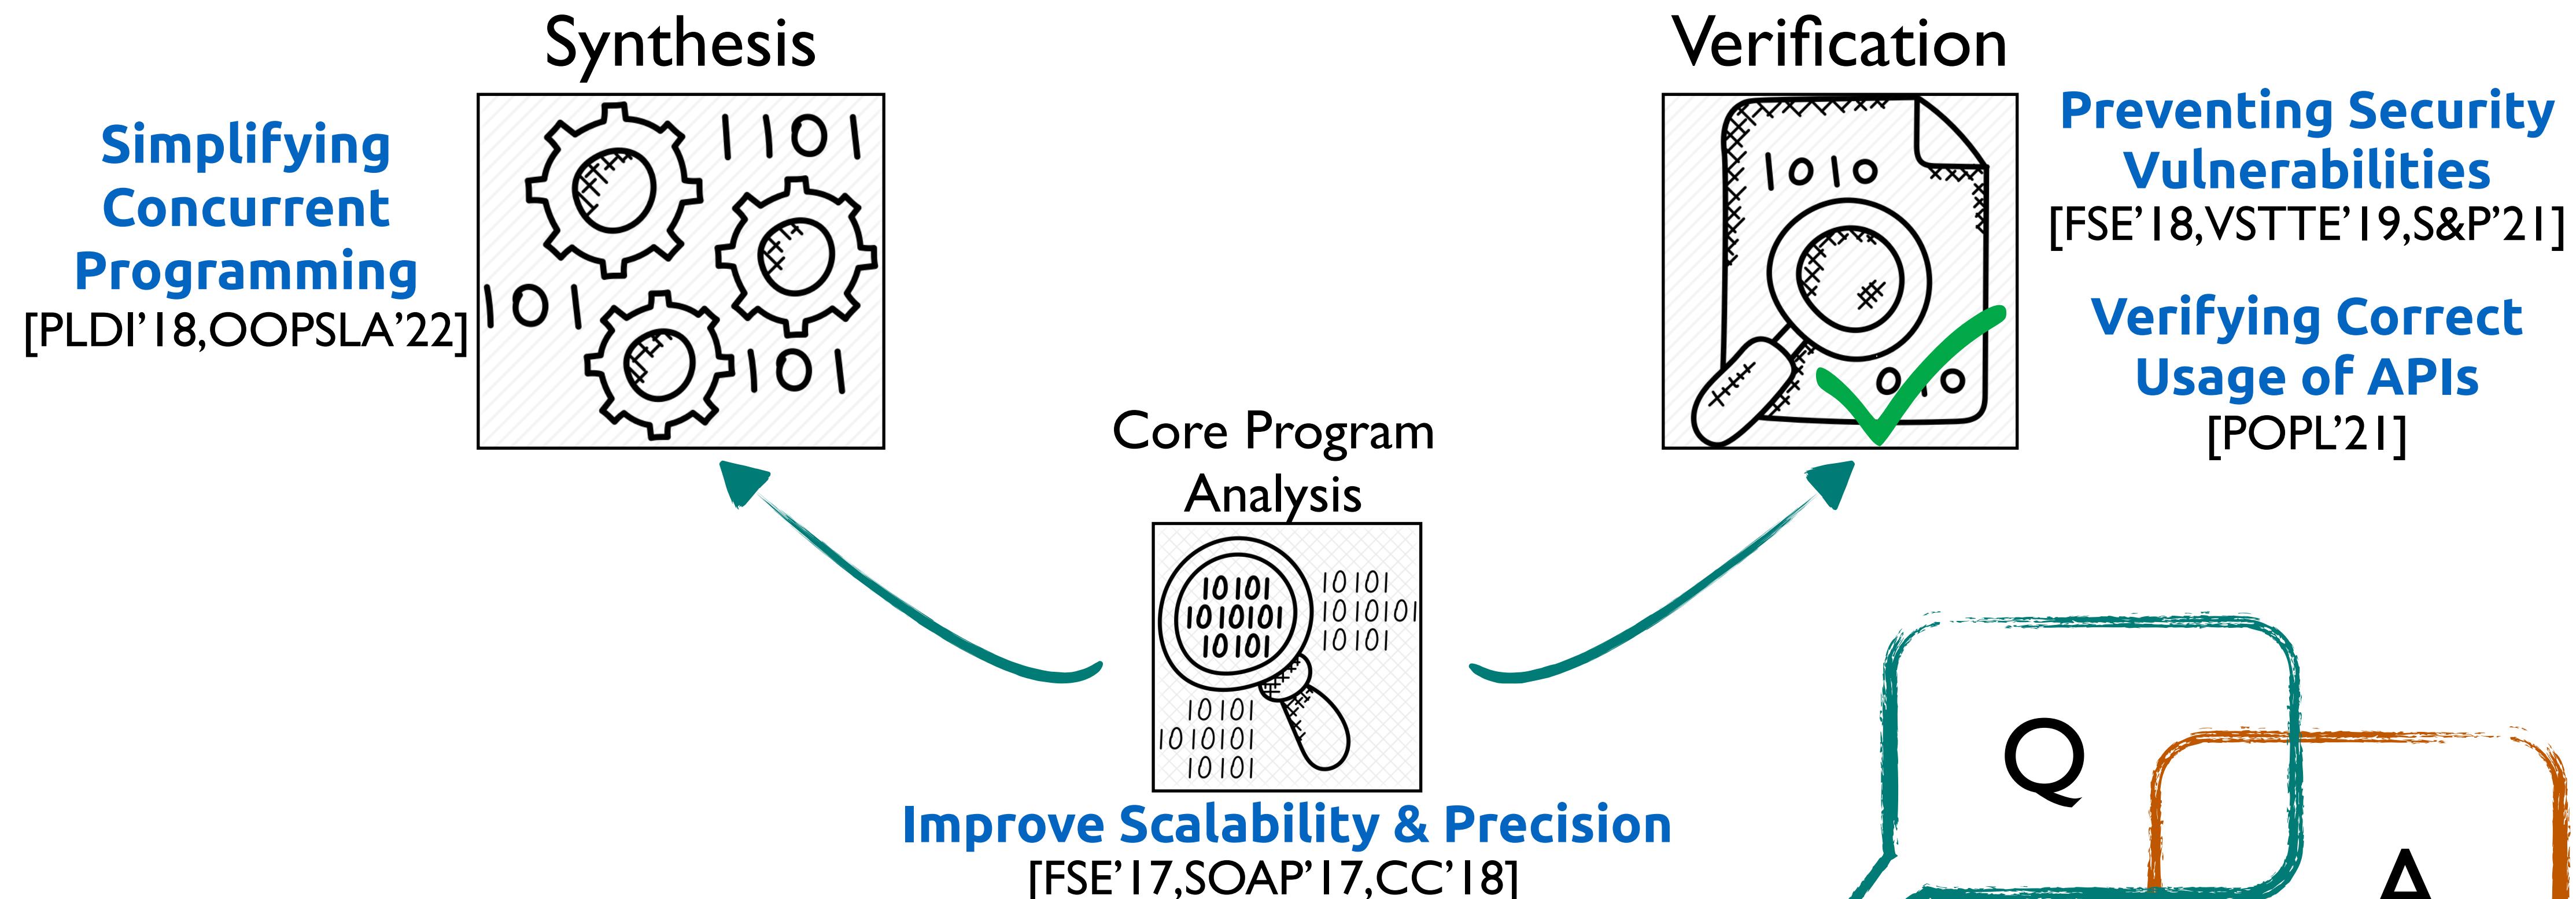


66 Valentin Wüstholtz

• • •

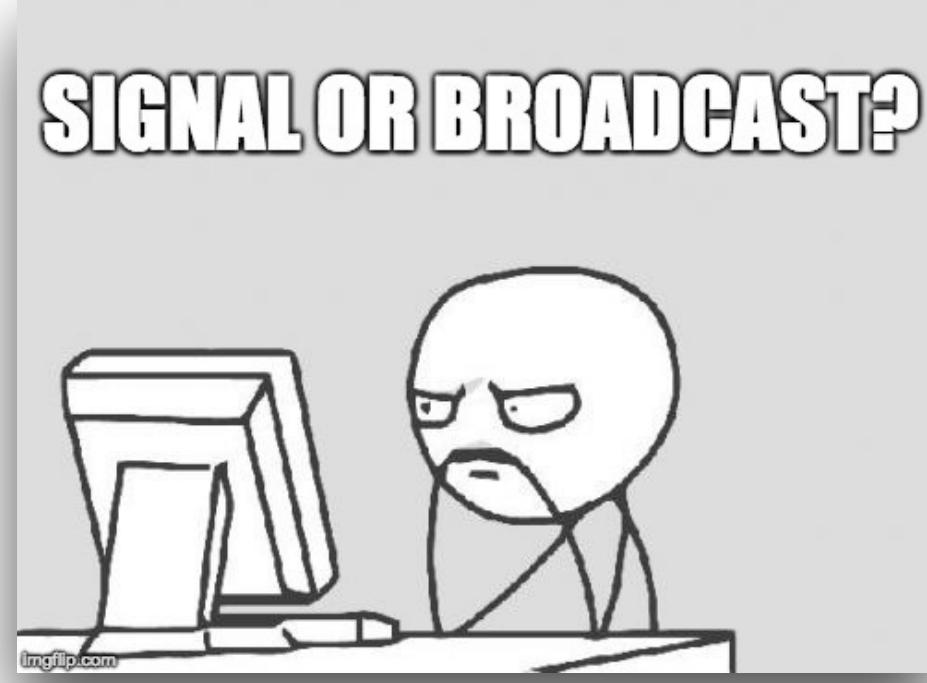
Questions?

Research Goal: Formal methods for correct, efficient, and secure code!



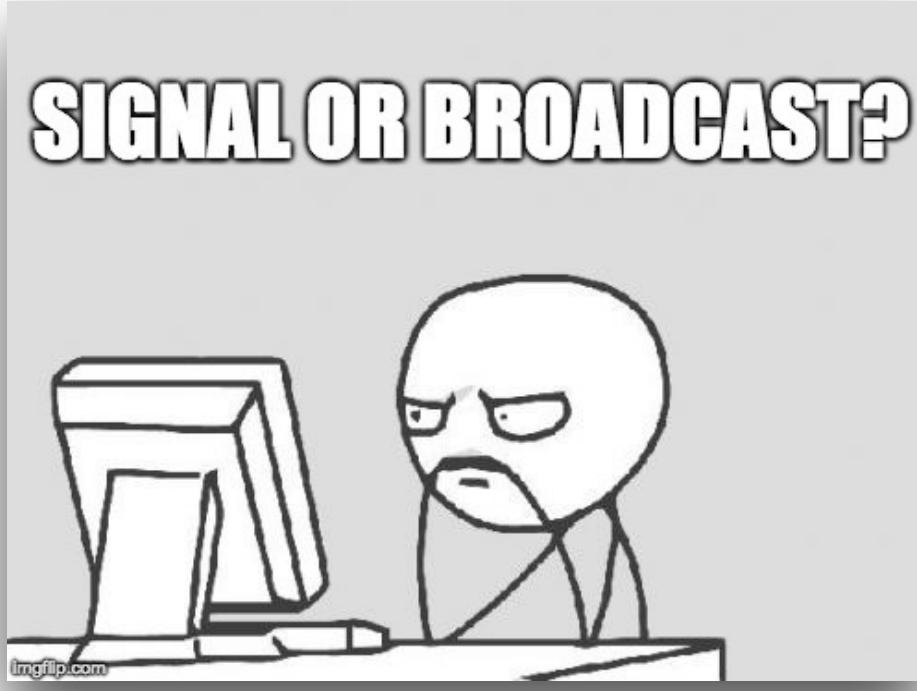
Back Up Slides

Signal or broadcast?



Do we need to wake up all threads block on $P \in Preds(M)$?

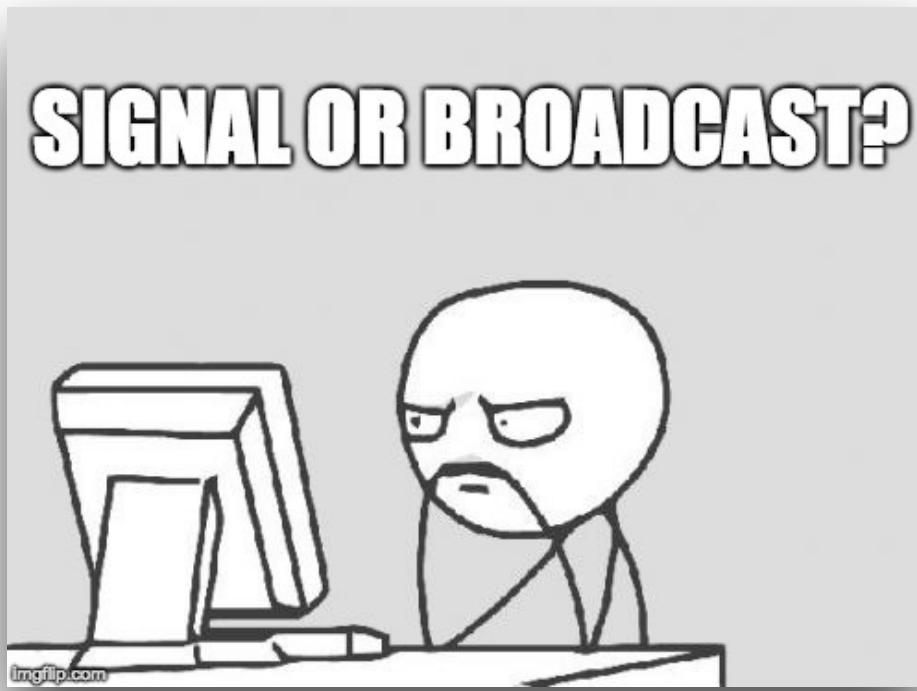
Signal or broadcast?



Do we need to wake up all threads block on $P \in Preds(M)$?

Simulate what happens when a thread gets woken up

Signal or broadcast?



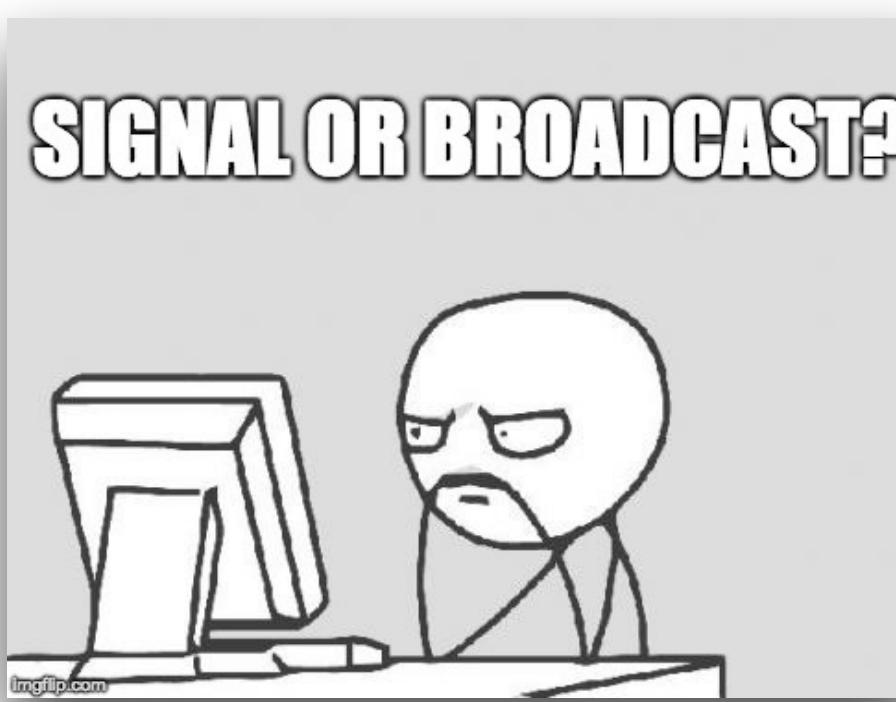
Do we need to wake up all threads block on $P \in \text{Preds}(M)$?

Simulate what happens when a thread gets woken up

For each CCR C of the form `waituntil(P) S`, check validity of Hoare triple:

$$\{\text{Inv}\} \ C \ \{\neg P\}$$

Signal or broadcast?



Do we need to wake up all threads block on $P \in \text{Preds}(M)$?

Simulate what happens when a thread gets woken up

For each CCR C of the form `waituntil(P) S`, check validity of Hoare triple:

$$\{\text{Inv}\} \ C \ \{\neg P\}$$

If **all** of them are valid, no need to broadcast!

Example: Signal or broadcast?

Predicate: $P_w \equiv readers = 0 \wedge \neg writerIn$ $\neg P_w \equiv readers \neq 0 \vee writerIn$ **Invariant:** $readers \geq 0$

CCR:

```
atomic void enterWriter() {  
    waituntil(readers == 0 && !writerIn);  
    writerIn = true;  
}
```

Example: Signal or broadcast?

Predicate: $P_w \equiv \text{readers} = 0 \wedge \neg \text{writerIn}$ $\neg P_w \equiv \text{readers} \neq 0 \vee \text{writerIn}$ **Invariant:** $\text{readers} \geq 0$

CCR: atomic **void** enterWriter() {
 { *readers* ≥ 0 }
 waituntil(*readers* == 0 $\&\&$!*writerIn*);
 writerIn = true;
 { $\neg P_w$ }
}

Example: Signal or broadcast?

Predicate: $P_w \equiv \text{readers} = 0 \wedge \neg \text{writerIn}$ $\neg P_w \equiv \text{readers} \neq 0 \vee \text{writerIn}$ **Invariant:** $\text{readers} \geq 0$

CCR:

```
atomic void enterWriter() {
    { readers ≥ 0 }
    waituntil(readers == 0 && !writerIn);
    writerIn = true;
}
{ ¬Pw }
```

Triple is valid, safe to notify a single writer!

Conditional or unconditional signal?

Decision 2: Given CCR
`waituntil(C) S` and predicate P,
do we need to check P before
signaling threads blocked on it?

Needed to avoid
unnecessary
context-switch

Conditional or unconditional signal?

Decision 2: Given CCR
`waituntil(C) S` and predicate P,
do we need to check P before
signaling threads blocked on it?

Needed to avoid
unnecessary
context-switch

$$\{Inv \wedge \neg P\} \ S \ \{P\}$$

Conditional or unconditional signal?

Decision 2: Given CCR
`waituntil(C) S` and predicate P,
do we need to check P before
signaling threads blocked on it?

Needed to avoid
unnecessary
context-switch

$$\{Inv \wedge \neg P\} \ S \ \{P\}$$

If this triple is valid, no need to test P at run-time!

Example: Conditional Signal?

Can we signal
readers
unconditionally?

CCR: atomic **void** exitWriter() {
 writerIn = **false**;
 }
 }

Example: Conditional Signal?

Invariant: $readers \geq 0$

Predicate: $P_r \equiv \neg writerIn$

Can we signal
readers
unconditionally?

CCR: atomic void exitWriter() {

 writerIn = **false**;

}

Example: Conditional Signal?

Invariant: $readers \geq 0$

Predicate: $P_r \equiv \neg writerIn$

Can we signal
readers
unconditionally?

CCR: atomic **void** exitWriter() {
 $\{readers \geq 0 \wedge \neg P_r\}$
 writerIn = **false**;
}

Example: Conditional Signal?

Invariant: $readers \geq 0$

Predicate: $P_r \equiv \neg writerIn$

Can we signal
readers
unconditionally?

CCR: atomic **void** exitWriter() {
 $\{readers \geq 0 \wedge \neg P_r\}$
 writerIn = **false**;
 $\{P_r\}$
}

Example: Conditional Signal?

Invariant:

$$readers \geq 0$$

Predicate: $P_r \equiv \neg writerIn$

Can we signal
readers
unconditionally?

CCR:

```
atomic void exitWriter() {
    {readers ≥ 0 ∧ ¬Pr}
    writerIn = false;
    {Pr}
}
```

Triple is valid, signal readers unconditionally!

Hard constraints

Data-race freedom

Atomicity

Deadlock-freedom

Hard constraints

Data-race freedom

Resolve race through
locks or atomic fields

Racy(S_1, S_2)

\Rightarrow

$\mathcal{L}(S_1) \cap \mathcal{L}(S_2) \neq \emptyset$

\vee

$\exists f. (Races(S_1, S_2) = \{f\} \wedge a_f)$

Atomicity

Deadlock-freedom

Hard constraints

Data-race freedom

Resolve race through locks or atomic fields

$$\text{Racy}(S_1, S_2)$$

\Rightarrow

$$\mathcal{L}(S_1) \cap \mathcal{L}(S_2) \neq \emptyset$$

\vee

$$\exists f. (\text{Races}(S_1, S_2) = \{f\} \wedge a_f)$$

Atomicity

Only permit safe interleavings

$$\neg \text{Safe}(S, S_1, S_2)$$

\Rightarrow

$$(\mathcal{L}(S) \cap \mathcal{L}(S_1) \cap \mathcal{L}(S_2) \neq \emptyset)$$

Deadlock-freedom

Hard constraints

Data-race freedom

Resolve race through locks or atomic fields

$\text{Racy}(S1, S2)$

\Rightarrow

$\mathcal{L}(S1) \cap \mathcal{L}(S2) \neq \emptyset$

\vee

$\exists f. (\text{Races}(S1, S2) = \{f\} \wedge a_f)$

Atomicity

Only permit safe interleavings

$\neg \text{Safe}(S, S1, S2)$

\Rightarrow

$(\mathcal{L}(S) \cap \mathcal{L}(S_1) \cap \mathcal{L}(S_2) \neq \emptyset)$

Deadlock-freedom

Impose a static lock acquisition order

$\text{Edge}(S1, S2)$

\Rightarrow

$(l_1 \in \mathcal{L}(S1) \cap \mathcal{L}(S2) \wedge l_2 \in \mathcal{L}(S2) \setminus \mathcal{L}(S1) \Rightarrow l_1 < l_2)$

Soft constraints

Maximize
parallelism

Minimize
locks

Minimize
atomics

Soft constraints

Maximize
parallelism

Minimize
locks

Minimize
atomics

Prefer: Non-racy snippet
pairs **not** sharing a lock.

$$\neg \text{Racy}(S1, S2)$$

\Rightarrow

$$\mathcal{L}(s1) \cap \mathcal{L}(s2) = \emptyset$$

Soft constraints

Maximize
parallelism

Prefer: Non-racy snippet
pairs **not** sharing a lock.

$$\neg \text{Racy}(S1, S2)$$

\Rightarrow

$$\mathcal{L}(s1) \cap \mathcal{L}(s2) = \emptyset$$

Minimize
locks

Prefer: Lock not being used

$$\bigwedge_s \neg h_s^l$$

Minimize
atomics

Soft constraints

Maximize parallelism

Prefer: Non-racy snippet pairs **not** sharing a lock.

$$\neg \text{Racy}(S1, S2)$$

\Rightarrow

$$\mathcal{L}(s1) \cap \mathcal{L}(s2) = \emptyset$$

Minimize locks

Prefer: Lock not being used

$$\bigwedge_s \neg h_s^l$$

Minimize atomics

Prefer: field not converted to atomic

$$\neg a_f$$